

THE EXPERT'S VOICE® IN WEB DEVELOPMENT

THIRD EDITION

Beginning Rails 4

*LEARN TO BUILD WEB APPLICATIONS
WITH RUBY ON RAILS*

Adam Gamble, Cloves Carneiro Jr., and Rida Al Barazi

Apress®

For your convenience Apress has placed some of the front matter material after the index. Please use the Bookmarks and Contents at a Glance links to access them.



Contents at a Glance

About the Authors.....	xvii
About the Technical Reviewer	xix
Introduction	xxi
■ Chapter 1: Introducing the Rails Framework	1
■ Chapter 2: Getting Started	13
■ Chapter 3: Getting Something Running	27
■ Chapter 4: Introduction to the Ruby Language.....	45
■ Chapter 5: Working with a Database: Active Record.....	61
■ Chapter 6: Advanced Active Record: Enhancing Your Models.....	81
■ Chapter 7: Action Pack: Working with the View and the Controller	121
■ Chapter 8: Advanced Action Pack	151
■ Chapter 9: JavaScript and CSS	191
■ Chapter 10: Sending and Receiving E-Mail	203
■ Chapter 11: Testing Your Application	219
■ Chapter 12: Internationalization.....	251
■ Chapter 13: Deploying Your Rails Applications	267
■ Appendix A: Databases 101	273
■ Appendix B: The Rails Community.....	281
■ Appendix C: Git	285
Index.....	297

Introduction

What Is This Book About?

In the past several years, the Web has exploded to include information on every facet of our lives. It touches everything we do; even some refrigerators have included access to the Internet. Ruby on Rails has played a part in fueling that explosion. This book will equip you with the knowledge you need to build real production web applications. It leads you through installing the required prerequisites on Windows, OS X, or Linux and then jumps straight into building applications. It is meant for the novice programmer who has some command line experience but little or no programming experience. At the end of the book, you should have a firm grasp on the Ruby language and the Rails framework.

Chapter 1 introduces you to the current web landscape and then goes over some of the ideals and principles that the Rails framework is built on. It teaches you about the MVC paradigm and shows how Rails implements each piece of that paradigm (model, view, and controller).

Chapter 2 walks you through installing Ruby, Rails, and the SQLite database. It is broken down by operating system, and when finished, will give a level platform among all three. You should be able to follow along with the book no matter which platform you choose. It also will show you how to build a quick “Hello World” application to make sure everything is working correctly.

Chapter 3 dives right in and starts the blog application that we’ll use throughout the rest of the book. We’ll continually build on this application, enhancing and refactoring as we go along. You’ll create your first model in this chapter, the article model. We’ll cover how migrations work and even get Rails to construct our first scaffold. At the end of this chapter, you’ll have a working blog application, although it will be lacking features. We’ll add those in the following chapters.

Chapter 4 slows down a little bit from the previous chapter and takes you on a tour of the Ruby language. If you’ve used Ruby for a while and feel comfortable with it, feel free to skim over this. If you’re new to Ruby, this chapter will teach you everything you need to know to use Rails. Ruby is an easy to pick up language, and the syntax is very inviting and easy to read. Although we won’t add any code to our blog application here, you will get to use the Ruby language inside the Ruby console.

Chapter 5 shows you how Rails uses Active Record to let you interact with any number of databases. Rails abstracts away the difficult bits (unless you need them) and lets you interact with databases in an object-oriented way. You’ll learn how to create new records, find records, and even update and delete them. We’ll also apply some basic validations so we can be sure our data are just the way they should be.

Chapter 6 expounds on the previous chapter. You’ll dive deeper into Active Record and your models. You will build more complex validations and custom instance methods. A major component of this chapter is the relation between your models and how Rails lets you define those relations. Your models for the blog application will have complex relations and validations.

In Chapter 7 we’ll cover the view and controller parts of MVC. We will flesh out the blog application and walk through the code that Rails generated for the scaffold of our controllers and views.

Chapter 8 modifies the controller and views in more advanced ways, and at this point the features of our blog application have come together. You’ll learn about controller callbacks and strong parameters that were added in Rails 4. We’ll also give our application a fresh coat of paint with some Cascading Style Sheets (CSS).

■ INTRODUCTION

Chapter 9 goes over the Asset Pipeline that was added in Rails 3.2 and how to add JavaScript and CSS. We'll enhance our application with JavaScript dabbling in Ajax and animation. This chapter covers CoffeeScript and SASS and how they integrate into the Rails landscape. At the end of this chapter, your application will have a nice layer of spit and polish.

Chapter 10 adds e-mail capability to our application. You will be able to suggest articles to friends and even be notified when your article has new comments. It will also discuss methods to receive e-mail into your application.

Chapter 11 covers one of the most important topics in Rails applications: testing. You can be sure that after this chapter you'll be able to add new features without breaking old ones. You'll test whether your application behaves exactly the way you think it should.

Chapter 12 covers internationalization. After all, it is the World Wide Web, and not everyone speaks the same language. We'll translate our web application into another language, and along the way you'll learn how to translate the application into as many languages as you like.

Chapter 13 will show you how to deploy your web application to Heroku, one of the leading Platform As A Service (PAAS) providers. This will allow you to present your application to the world quickly and easily so you can start building a user base.

The three appendices cover the Git version control system, SQL, and where to find help in the Rails community.

CHAPTER 1



Introducing the Rails Framework

Rails is a web application framework for the Ruby programming language. Rails is well thought out and practical: it will help you build powerful web sites quickly, with code that's clean and easy to maintain.

The goal of this book is to give you a thorough and complete understanding of how to build dynamic web applications with Rails. This means more than just showing you how to use the specific features and facilities of the framework, and more than just giving you a working knowledge of the Ruby language. Rails is quite a bit more than just another tool: it represents a way of thinking. To completely understand Rails, it's essential that you know about its underpinnings, its culture and aesthetics, and its philosophy of web development.

If you haven't heard it already, you're sure to notice the phrase "the Rails way" cropping up every now and again. It echoes a familiar phrase that has been floating around the Ruby community for a number of years: "the Ruby way." The Rails way is usually the easiest way—the path of least resistance, if you will. This isn't to say that you can't do things your way, nor is it meant to suggest that the framework is constraining. It simply means that if you choose to go off the beaten path, you shouldn't expect Rails to make it easy for you. If you've been around the UNIX circle for any length of time, you may think this idea bears some resemblance to the UNIX mantra: "Do the simplest thing that could possibly work." You're right. This chapter's aim is to introduce you to the Rails way.

The Rise and Rise of the Web Application

Web applications are extremely important in today's world. Almost everything we do today involves web applications. We check our e-mail on the Web, and we do our banking on the Web. We even use our phones to access the web more than we actually make phone calls! As connections get faster, and as broadband adoption grows, web-based software, and similarly networked client or server applications, are poised to displace software distributed by more traditional (read, outdated) means.

For consumers, web-based software affords greater convenience, allowing us to do more from more places. Web-based software works on every platform that supports a web browser (which is to say, all of them), and there's nothing to install or download. And if Google's stock value is any indication, web applications are really taking off. All over the world, people are waking up to the new Web and the beauty of being web based. From e-mail and calendars, photos and videos, to bookmarking, banking, and bidding, we're living increasingly inside the browser.

Due to the ease of distribution, the pace of change in the web-based software market is fast. Unlike traditional software, which must be installed on each individual computer, changes in web applications can be delivered quickly, and features can be added incrementally. There's no need to spend months or years perfecting the final version or getting in all the features before the launch date. Instead of spending months on research and development, you can go into production early and refine in the wild, even without all the features in place.

Can you imagine having a million CDs pressed and shipped, only to find a bug in your software as the FedEx truck is disappearing into the sunset? That would be an expensive mistake! Software distributed this way takes notoriously long to get out the door because before a company ships a product, it needs to be sure the software is bug free. Of course, there's no such thing as bug-free software, and web applications aren't immune to these unintended features. But with a web application, bug fixes are easy to deploy.

When a fix is pushed to the server hosting the web application, all users get the benefit of the update at the same time, usually without any interruption in service. That's a level of quality assurance you can't offer with store-bought software. There are no service packs to tirelessly distribute and no critical updates to install. A fix is often only a browser refresh away. And as a side benefit, instead of spending large amounts of money and resources on packaging and distribution, software developers are free to spend more time on quality and innovation.

Web-based software has the following advantages; it is:

- Easier to distribute
- Easier to deploy
- Easier to maintain
- Platform-independent
- Accessible from anywhere

The Web Isn't Perfect

As great a platform as the Web is, it's also fraught with constraints. One of the biggest problems is the browser itself. When it comes to browsers, there are several contenders, each of which has a slightly different take on how to display the contents of a web page. Although there has been movement toward unification and the state of standards compliance among browsers is steadily improving, there is still much to be desired. Even today, it's nearly impossible to achieve 100% cross-browser compatibility. Something that works in Internet Explorer doesn't necessarily work in Firefox, and vice versa. This lack of uniformity makes it difficult for developers to create truly cross-platform applications, as well as harder for users to work in their browser of choice.

Browser issues aside, perhaps the biggest constraint facing web development is its inherent complexity. A typical web application has dozens of moving parts: protocols and ports, the HTML and cascading style sheets (CSS), the database and the server, the designer and the developer, and a multitude of other players, all conspiring toward complexity.

Despite these problems, the new focus on the Web as a platform means the field of web development is evolving rapidly and quickly overcoming obstacles. As it continues to mature, the tools and processes that have long been commonplace in traditional, client-side software development are beginning to make their way into the world of web development.

The Good Web Framework

Among the tools making their way into the world of web development is the framework. A *framework* is a collection of libraries and tools intended to facilitate development. Designed with productivity in mind, a good framework provides a basic but complete infrastructure on top of which to build an application.

Having a good framework is a lot like having a chunk of your application already written for you. Instead of having to start from scratch, you begin with the foundation in place. If a community of developers uses the same framework, you have a community of support when you need it. You also have greater assurance that the foundation you're building on is less prone to pesky bugs and vulnerabilities, which can slow the development process.

A good web framework can be described as follows:

- *Full stack:* Everything you need for building complete applications should be included in the box. Having to install various libraries or configure multiple components is a drag. The different layers should fit together seamlessly.
- *Open source:* A framework should be open source, preferably licensed under a liberal, free-as-in-free license like the Berkeley Software Distribution (BSD) or that of the Massachusetts Institute of Technology (MIT).
- *Cross-platform:* A good framework is platform independent. The platform on which you decide to work is a personal choice. Your framework should remain as neutral as possible.

A good web framework provides you with the following:

- *A place for everything:* Structure and convention drive a good framework. In other words, unless a framework offers a good structure and a practical set of conventions, it's not a very good framework. Everything should have a proper place within the system; this eliminates guesswork and increases productivity.
- *A database abstraction layer:* You shouldn't have to deal with the low-level details of database access, nor should you be constrained to a particular database engine. A good framework takes care of most of the database grunt work for you, and it works with almost any database.
- *A culture and aesthetic to help inform programming decisions:* Rather than seeing the structure imposed by a framework as constraining, see it as liberating. A good framework encodes its opinions, gently guiding you. Often, difficult decisions are made for you by virtue of convention. The culture of the framework helps you make fewer menial decisions and helps you focus on what matters most.

Enter Rails

Rails is a best-of-breed framework for building web applications. It's complete, open source, and cross-platform. It provides a powerful database abstraction layer called *Active Record*, which works with all popular database systems. It ships with a sensible set of defaults and provides a well-proven, multilayer system for organizing program files and concerns.

Above all, Rails is opinionated software. It has a philosophy of the art of web development that it takes very seriously. Fortunately, this philosophy is centered on beauty and productivity. You'll find that as you learn Rails, it actually makes writing web applications pleasurable.

Originally created by David Heinemeier Hansson, Rails first took shape in the form of a wiki application called Instiki. The first version, released in July 2004, of what is now the Rails framework, was extracted from a real-world, working application: Basecamp, by 37signals. The Rails creators took away all the Basecamp-specific parts and what remained was Rails.

Because it was extracted from a real application and not built as an ivory tower exercise, Rails is practical and free of needless features. Its goal as a framework is to solve 80% of the problems that occur in web development, assuming that the remaining 20% are problems that are unique to the application's domain. It may be surprising that as much as 80% of the code in an application is infrastructure, but it's not as far-fetched as it sounds. Consider all the work involved in application construction, from directory structure and naming conventions, to the database abstraction layer and the maintenance of state.

Rails has specific ideas about directory structure, file naming, data structures, method arguments, and, well, nearly everything. When you write a Rails application, you're expected to follow the conventions that have been laid out for you. Instead of focusing on the details of knitting the application together, you get to focus on the 20% that really matters.

Since 2004, Rails has come a long way. The Rails team continues to update the framework to support the latest technologies and methodologies available. You'll find that as you use Rails, it's obvious that the core team has kept the project at the forefront of web technology. The Rails 4 release is better than ever. Speed, security, and ease of use were major focuses of this release and it shows.

Rails Is Ruby

There are a lot of programming languages out there. You've probably heard of many of them. C, C#, Lisp, Java, Smalltalk, PHP, and Python are popular choices. And then there are others you've probably never heard of: Haskell, IO, and maybe even Ruby. Like the others, Ruby is a programming language. You use it to write computer programs, including, but certainly not limited to, web applications.

Before Rails came along, not many people were writing web applications with Ruby. Other languages like PHP and ASP were the dominant players in the field, and a large part of the Web is powered by them. The fact that Rails uses Ruby is significant because Ruby is considerably more expressive and flexible than either PHP or ASP. This makes developing web applications not only easy, but also a lot of fun. Ruby has all the power of other languages, but it was built with the main goal of developer happiness.

Ruby is a key part of the success of Rails. Rails uses Ruby to create what's called a *domain-specific language* (DSL). Here, the domain is that of web development; when you're working in Rails, it's almost as if you're writing in a language that was specifically designed to construct web applications—a language with its own set of rules and grammar. Rails does this so well that it's sometimes easy to forget that you're writing Ruby code. This is a testimony to Ruby's power, and Rails takes full advantage of Ruby's expressiveness to create a truly beautiful environment.

For many developers, Rails is their introduction to Ruby—a language with a following before Rails that was admittedly small at best, at least in the West. Although Ruby had been steadily coming to the attention of programmers outside Japan, the Rails framework brought Ruby to the mainstream.

Invented by Yukihiro Matsumoto in 1994, it's a wonder Ruby remained shrouded in obscurity as long as it did. As far as programming languages go, Ruby is among the most beautiful. Interpreted and object oriented, elegant and expressive, Ruby is truly a joy to work with. A large part of Rails' grace is due to Ruby and to the culture and aesthetics that permeate the Ruby community. As you begin to work with the framework, you'll quickly learn that Ruby, like Rails, is rich with idioms and conventions, all of which make for an enjoyable, productive programming environment.

In summary, Ruby can be described as follows:

- An interpreted, object-oriented scripting language
- Elegant, concise syntax
- Powerful metaprogramming features
- Well suited as a host language for creating DSLs

This book includes a complete Ruby primer. If you want to get a feel for what Ruby looks like now, skip to Chapter 3 and take a look. Don't worry if Ruby seems a little unconventional at first. You'll find it quite readable, even if you're not a programmer. It's safe to follow along in this book learning it as you go and referencing Chapter 3 when you need clarification. If you're looking for a more in-depth guide, Peter Cooper has written a fabulous book titled *Beginning Ruby: From Novice to Professional*, Second Edition (Apress, 2009). You'll also find the Ruby community more than helpful in your pursuit of the language. Be sure to visit <http://ruby-lang.org> for a wealth of Ruby-related resources.

Rails Encourages Agility

Web applications aren't traditionally known for agility. They have a reputation of being difficult to work with and a nightmare to maintain. It's perhaps in response to this diagnosis that Rails came onto the scene, helping to usher in a movement toward agile programming methodologies in web development. Rails advocates and assists in the achievement of the following basic principles of software development:

- Individuals and interactions over processes and tools
- Working software over comprehensive documentation
- Customer collaboration over contract negotiation
- Responding to change over following a plan

So reads the Agile Manifesto,¹ which was the result of a discussion among 17 prominent figures (including Dave Thomas, Andy Hunt, and Martin Fowler) in the field of what was then called "lightweight methodologies" for software development. Today, the Agile Manifesto is widely regarded as the canonical definition of agile development.

¹<http://agilemanifesto.org>

Rails was designed with agility in mind, and it takes each of the agile principles to heart almost obsessively. With Rails, you can respond to the needs of customers quickly and easily, and Rails works well during collaborative development. Rails accomplishes this by adhering to its own set of principles, all of which help make agile development possible.

Dave Thomas and Andy Hunt's seminal work on the craft of programming, *The Pragmatic Programmer* (Addison-Wesley, 1999), reads almost like a roadmap for Rails. Rails follows the *don't repeat yourself* (DRY) principle, the concepts of rapid prototyping, and the *you ain't gonna need it* (YAGNI) philosophy. Keeping important data in plain text, using convention over configuration, bridging the gap between customer and programmer, and, above all, postponing decisions in anticipation of change are institutionalized in Rails. These are some of the reasons that Rails is such an apt tool for agile development, and it's no wonder that one of the earliest supporters of Rails was Dave Thomas himself.

The sections that follow take you on a tour through some of Rails mantras and, in doing so, demonstrate how well suited Rails is for agile development. Although we want to avoid getting too philosophical, some of these points are essential to grasp what makes Rails so important.

Less Software

One of the central tenets of Rails' philosophy is the notion of *less software*. What does less software mean? It means using convention over configuration, writing less code, and doing away with things that needlessly add to the complexity of a system. In short, less software means less code, less complexity, and fewer bugs.

Convention Over Configuration

Convention over configuration means that you need to define only configuration that is unconventional.

Programming is all about making decisions. If you were to write a system from scratch, without the aid of Rails, you'd have to make a lot of decisions: how to organize your files, what naming conventions to adopt, and how to handle database access are only a few. If you decided to use a database abstraction layer, you would need to sit down and write it or find an open source implementation that suited your needs. You'd need to do all this before you even got down to the business of modeling your domain.

Rails lets you start right away by encompassing a set of intelligent decisions about how your program should work and alleviating the amount of low-level decision making you need to do up front. As a result, you can focus on the problems you're trying to solve and get the job done more quickly.

Rails ships with almost no configuration files. If you're used to other frameworks, this fact may surprise you. If you've never used a framework before, you should be surprised. In some cases, configuring a framework is nearly half the work.

Instead of configuration, Rails relies on common structures and naming conventions, all of which employ the often-cited *principle of least surprise* (POLS). Things behave in a predictable, easy-to-decipher way. There are intelligent defaults for nearly every aspect of the framework, relieving you from having to explicitly tell the framework how to behave. This isn't to say that you can't tell Rails how to behave: most behaviors can be customized to your liking and to suit your particular needs. But you'll get the most mileage and productivity out of the defaults, and Rails is all too willing to encourage you to accept the defaults and move on to solving more interesting problems.

Although you can manipulate most things in the Rails setup and environment, the more you accept the defaults, the faster you can develop applications and predict how they will work. The speed with which you can develop without having to do any explicit configuration is one of the key reasons why Rails works so well. If you put your files in the right place and name them according to the right conventions, things *just work*. If you're willing to agree to the defaults, you generally have less code to write.

The reason Rails does this comes back to the idea of less software. Less software means making fewer low-level decisions, which makes your life as a web developer a lot easier. And easier is a good thing.

Don't Repeat Yourself

Rails is big on the DRY principle, which states that information in a system should be expressed in only one place.

For example, consider database configuration parameters. When you connect to a database, you generally need credentials, such as a username, a password, and the name of the database you want to work with. It may seem acceptable to include this connection information with each database query, and that approach holds up fine if you're making only one or two connections. But as soon as you need to make more than a few connections, you end up with a lot of instances of that username and password littered throughout your code. Then, if your username and password for the database change, you have to do a lot of finding and replacing. It's a much better idea to keep the connection information in a single file, referencing it as necessary. That way, if the credentials change, you need to modify only a single file. That's what the DRY principle is all about.

The more duplication exists in a system, the more room bugs have to hide. The more places the same information resides, the more there is to be modified when a change is required, and the harder it becomes to track these changes.

Rails is organized so it remains as DRY as possible. You generally specify information in a single place and move on to better things.

Rails Is Opinionated Software

Frameworks encode opinions. It should come as no surprise then that Rails has strong opinions about how your application should be constructed. When you're working on a Rails application, those opinions are imposed on you, whether you're aware of it or not. One of the ways that Rails makes its voice heard is by gently (sometimes, forcefully) nudging you in the right direction. We mentioned this form of encouragement when we talked about convention over configuration. You're invited to do the right thing by virtue of the fact that doing the wrong thing is often more difficult.

Ruby is known for making certain programmatic constructs look more natural by way of what's called *syntactic sugar*. Syntactic sugar means the syntax for something is altered to make it appear more natural, even though it behaves the same way. Things that are syntactically correct but otherwise look awkward when typed are often treated to syntactic sugar.

Rails has popularized the term *syntactic vinegar*. Syntactic vinegar is the exact opposite of syntactic sugar: awkward programmatic constructs are discouraged by making their syntax look sour. When you write a snippet of code that looks bad, chances are it *is* bad. Rails is good at making the right thing obvious by virtue of its beauty and the wrong thing equally obvious by virtue of ugliness.

You can see Rails' opinion in the things it does automatically, the ways it encourages you to do the right thing, and the conventions it asks you to accept. You'll find that Rails has an opinion about nearly everything related to web application construction: how you should name your database tables, how you should name your fields, which database and server software to use, how to scale your application, what you need, and what is a vestige of web development's past. If you subscribe to its worldview, you'll get along with Rails quite well.

Like a programming language, a framework needs to be something you're comfortable with—something that reflects your personal style and mode of working. It's often said in the Rails community that if you're getting pushback from Rails, it's probably because you haven't experienced enough pain from doing web development the old-school way. This isn't meant to deter developers; rather, it means that in order to truly appreciate Rails, you may need a history lesson in the technologies from whose ashes Rails has risen. Sometimes, until you've experienced the hurt, you can't appreciate the cure.

Rails Is Open Source

The Rails culture is steeped in open source tradition. The Rails source code is, of course, open. And it's significant that Rails is licensed under the MIT license, arguably one of the most "free" software licenses in existence.

Rails also advocates the use of open source tools and encourages the collaborative spirit of open source. The code that makes up Rails is 100% free and can be downloaded, modified, and redistributed by anyone at any time.

Moreover, anyone is free to submit patches for bugs or features, and hundreds of people from all over the world have contributed to the project over the past nine years.

You'll probably notice that a lot of Rails developers use Macs. The Mac is clearly the preferred platform of many core Rails team developers, and most Rails developers are using UNIX variants (of which Mac OS X is one). Although there is a marked bias toward UNIX variants when it comes to Rails developers, make no mistake, Rails is truly cross-platform. With a growing number of developers using Rails in a Windows environment, Rails has become easy to work with in all environments. It doesn't matter which operating system you choose: you'll be able to use Rails on it. Rails doesn't require any special editor or Integrated Development Environment (IDE) to write code. Any text editor is fine, as long as it can save files in plain text. The Rails package even includes a built-in, stand-alone web server called WEBrick, so you don't need to worry about installing and configuring a web server for your platform. When you want to run your Rails application in development mode, simply start up the built-in server and open your web browser. Why should it be more difficult than that?

The next chapter takes you step by step through the relatively painless procedure of installing Rails and getting it running on your system. But before you go there, and before you start writing your first application, let's talk about how the Rails framework is architected. This is important because, as you will see in a minute, it has a lot to do with how you organize your files and where you put them. Rails is a subset of a category of frameworks named for the way in which they divide the concerns of program design: the model-view-controller (MVC) pattern. Not surprisingly, the MVC pattern is the topic of our next section.

The MVC Pattern

Rails employs a time-honored and well-established architectural pattern that advocates dividing application logic and labor into three distinct categories: the model, view, and controller. In the MVC pattern, the model represents the data, the view represents the user interface, and the controller directs all the action. The real power lies in the combination of the MVC layers, which Rails handles for you. Place your code in the right place and follow the naming conventions, and everything should fall into place.

Each part of the MVC—the model, view, and controller—is a separate entity, capable of being engineered and tested in isolation. A change to a model need not affect the views; likewise, a change to a view should have no effect on the model. This means changes in an MVC application tend to be localized and low impact, easing the pain of maintenance considerably while increasing the level of reusability among components.

Contrast this to the situation that occurs in a highly coupled application that mixes data access, business logic, and presentation code (PHP, we're looking at you). Some folks call this *spaghetti code* because of its striking resemblance to a tangled mess. In such systems, duplication is common, and even small changes can produce large ripple effects. MVC was designed to help solve this problem.

MVC isn't the only design pattern for web applications, but it's the one Rails has chosen to implement. And it turns out that it works great for web development. By separating concerns into different layers, changes to one don't have an impact on the others, resulting in faster development cycles and easier maintenance.

The MVC Cycle

Although MVC comes in different flavors, control flow generally works as follows (Figure 1-1):

- The user interacts with the interface and triggers an event (e.g., submits a registration form).
- The controller receives the input from the interface (e.g., the submitted form data).
- The controller accesses the model, often updating it in some way (e.g., by creating a new user with the form data).
- The controller invokes a view that renders an updated interface (e.g., a welcome screen).
- The interface waits for further interaction from the user, and the cycle repeats.

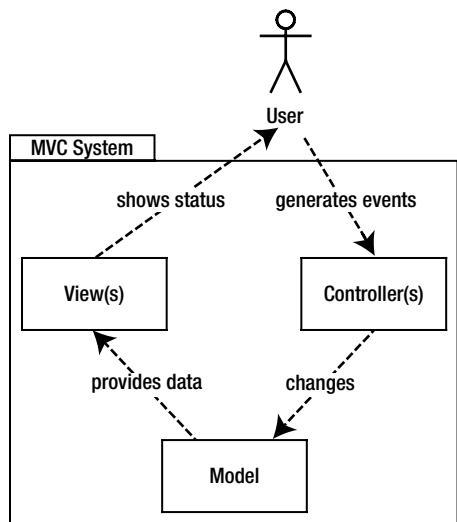


Figure 1-1. The MVC cycle

If the MVC concept sounds a little involved, don't worry. Although entire books have been written on this pattern, and people will argue over its purest implementation for all time, it's easy to grasp, especially the way Rails does MVC.

Next, we'll take a quick tour through each letter in the MVC and then learn how Rails handles it.

The Layers of MVC

The three layers of the MVC pattern work together as follows:

- *Model:* The information the application works with
- *View:* The visual representation of the user interface
- *Controller:* The director of interaction between the model and the view

Models

In Rails, the model layer represents the database. Although we call the entire layer the model, Rails applications are usually made up of several individual models, each of which (usually) maps to a database table. For example, a model called `User` may map to a table called `users`. The `User` model assumes responsibility for all access to the `users` table in the database, including creating, reading, updating, and deleting rows. So, if you want to work with the table and, say, search for someone by name, you do so through the model, like this:

```
User.where(:name => 'Linus').first
```

This snippet, although very basic, searches the `users` table for the first row with the value `Linus` in the `name` column and returns the results. To achieve this, Rails uses its built-in database abstraction layer, Active Record. Active Record is a powerful library; needless to say, this is only a small portion of what you can do with it.

Chapters 5 and 6 will give you an in-depth understanding of Active Record and what you can expect from it. For the time being, the important thing to remember is that models represent data. All rules for data access, associations, validations, calculations, and routines that should be executed before and after save, update, or destroy operations are

neatly encapsulated in the model. Your application's world is populated with Active Record objects: single ones, lists of them, new ones, and old ones. And Active Record lets you use Ruby language constructs to manipulate all of them, meaning you get to stick to one language for your entire application.

Controllers

For the discussion here, let's rearrange the MVC acronym and put the *C* before the *V*. As you'll see in a minute, in Rails, controllers are responsible for rendering views, so it makes sense to introduce them first.

Controllers are the conductors of an MVC application. In Rails, controllers accept requests from the outside world, perform the necessary processing, and then pass control to the view layer to display the results. It's the controller's job to field web requests, like processing server variables and forming data, asking the model for information, and sending information back to the model to be saved in the database. It may be a gross oversimplification, but controllers generally perform a request from the user to create, read, update, or delete a model object. You see these words a lot in the context of Rails, most often abbreviated as **CRUD**. In response to a request, the controller typically performs a **CRUD** operation on the model, sets up variables to be used in the view, and then proceeds to render or redirect to another action after processing is complete.

Controllers typically manage a single area of an application. For example, in a recipe application, you probably have a controller just for managing recipes. Inside the recipes controller, you can define what are called *actions*. Actions describe what a controller can do. If you want to be able to create, read, update, and delete recipes, you create appropriately named actions in the recipes controller. A simple recipes controller would look something like this:

```
class RecipesController < ApplicationController
  def index
    # logic to list all recipes
  end

  def show
    # logic to show a particular recipe
  end

  def create
    # logic to create a new recipe
  end

  def update
    # logic to update a particular recipe
  end

  def destroy
    # logic to delete a particular recipe
  end
end
```

Of course, if you want this controller to do anything, you need to put some instructions inside each action. When a request comes into your controller, it uses a URL parameter to identify the action to execute; and when it's done, it sends a response to the browser. The response is what you look at next.

Views

The view layer in the MVC forms the visible part of the application. In Rails, views are the templates that (most of the time) contain HTML markup to be rendered in a browser. It's important to note that views are meant to be free of all but the simplest programming logic. Any direct interaction with the model layer should be delegated to the controller layer, to keep the view clean and decoupled from the application's business logic.

Generally, views have the responsibility of formatting and presenting model objects for output on the screen, as well as providing the forms and input boxes that accept model data, such as a login box with a username and password or a registration form. Rails also provides the convenience of a comprehensive set of helpers that make connecting models and views easier, such as being able to prepopulate a form with information from the database or the ability to display error messages if a record fails any validation rules, such as required fields.

You're sure to hear this eventually if you hang out in Rails circles: a lot of folks consider the interface to *be* the software. We agree with them. Because the interface is all the user sees, it's the most important part. Whatever the software is doing behind the scenes, the only parts that an end user can relate to are the parts they see and interact with. The MVC pattern helps by keeping programming logic out of the view. With this strategy in place, programmers get to deal with code, and designers get to deal with templates called erb (Embedded Ruby). These templates take plain HTML and use Ruby to inject the data and view specific logic as needed. Designers will feel right at home if they are familiar with HTML. Having a clean environment in which to design the HTML means better interfaces and better software.

The Libraries That Make Up Rails

Rails is a collection of libraries, each with a specialized task. Assembled together, these individual libraries make up the Rails framework. Of the several libraries that compose Rails, three map directly to the MVC pattern:

- *Active Record*: A library that handles database abstraction and interaction.
- *Action View*: A templating system that generates the HTML documents the visitor gets back as a result of a request to a Rails application.
- *Action Controller*: A library for manipulating both application flow and the data coming from the database on its way to being displayed in a view.

These libraries can be used independently of Rails and of one another. Together, they form the Rails MVC development stack. Because Rails is a full-stack framework, all the components are integrated, so you don't need to set up bridges among them manually.

Rails Is Modular

One of the great features of Rails is that it was built with modularity in mind from the ground up. Although many developers appreciate the fact that they get a full stack, you may have your own preferences in libraries, either for database access, template manipulation, or JavaScript libraries. As we describe Rails features, we mention alternatives to the default libraries that you may want to pursue as you become more familiar with Rails' inner workings.

Rails Is No Silver Bullet

There is no question that Rails offers web developers a lot of benefits. After using Rails, it's hard to imagine going back to web development without it. Fortunately, it looks like Rails will be around for a long time, so there's no need to worry. But it brings us to an important point.

As much as we've touted the benefits of Rails, it's important for you to realize that there are no silver bullets in software design. No matter how good Rails gets, it will never be all things to all people, and it will never solve all problems. Most important, Rails will never replace the role of the developer. Its purpose is to assist developers in getting their job done. Impressive as it is, Rails is merely a tool, which when used well can yield amazing results. It's our hope that as you continue to read this book and learn how to use Rails, you'll be able to leverage its strength to deliver creative and high-quality web-based software.

Summary

This chapter provided an introductory overview of the Rails landscape, from the growing importance of web applications to the history, philosophy, evolution, and architecture of the framework. You learned about the features of Rails that make it ideally suited for agile development, including the concepts of less software, convention over configuration, and DRY. Finally, you learned the basics of the MVC pattern and received a primer on how Rails does MVC.

With all this information under your belt, it's safe to say you're ready to get up and running with Rails. The next chapter walks you through the Rails installation so you can try it for yourself and see what all the fuss is about. You'll be up and running with Rails in no time.



Getting Started

For various reasons, Rails has gained an undeserved reputation of being difficult to install. This chapter dispels this myth. The truth is that installing Rails is relatively easy and straightforward, provided you have all the right ingredients. The chapter begins with an overview of what you need to get Rails up and running and then provides step-by-step instructions for the installation. Finally, you'll start your first Rails application.

An Overview of Rails Installation

The main ingredient you need for Rails is, of course, Ruby. Some systems, such as OS X come with Ruby preinstalled, but it's often outdated. To make sure you have the best experience with this book it's best if you start from a clean slate, so you'll install it. After you have Ruby installed, you can install a *package manager* (a program designed to help you install and maintain software on your system) called RubyGems. You use that to install Rails.

If you're a Ruby hacker and already have Ruby and RubyGems installed on your computer, Rails is ridiculously easy to get up and running. Because it's packaged as a gem, you can install it with a single command:

```
$ gem install rails
```

That's all it comes down to—installing Rails is a mere one-liner. The key is in having a working installation of Ruby and RubyGems. Before you get there, though, you need one other ingredient to use Rails: a database server.

As you're well aware by now, Rails is specifically meant for building web applications. Well, it's a rare web application that isn't backed by a database. Rails is so sure you're using a database for your application that it's downright stubborn about working nicely without one. Although Rails works with nearly every database out there, in this chapter you use one called SQLite. SQLite is open source, easy to install, and incredibly easy to develop with. Perhaps that's why it's the default database for Rails.

You start by installing Ruby and RubyGems, and you use the magical one-liner to install Rails. Then, you install SQLite and make sure it's working properly. Here are the steps in order:

1. Install Ruby 2.0.
2. Install Rails.
3. Install SQLite.

Before you begin, note that the “many ways to skin a cat” adage applies to Rails installation. Just as the Rails stack runs on many platforms, there are as many ways to install it. This chapter describes what we feel is the easiest and most reliable way to install Rails for your platform. You go about the process differently for OS X, Linux, and Windows, but they all amount to the same thing.

No matter which platform you're using, you need to get familiar with the command line. This likely isn't a problem for the Linux crowd, but it's possible that some OS X users and certainly many Windows users don't have much experience with it. If you're using OS X, you can find a terminal emulator in /Applications/Utilities/Terminal.app. If you're on Windows, you can open a command prompt by choosing Start ▶ Run, typing cmd, and clicking OK. Note that you'll use the command line extensively in your travels with Rails. A growing number of IDEs make developing applications with Rails even simpler, and they completely abstract the use of a command-line tool; but stick to the command line to make sure you grasp all the concepts behind many commands. If you later decide to use an IDE such as Aptana's RadRails, JetBrains' RubyMine, or Oracle's NetBeans, you'll have a great understanding of Rails and will understand even better where the IDE is speeding up your work.

Also, a quick note for OS X users: if you're using a Mac and would prefer to use a package manager such as Fink or MacPorts, the Linux instructions will prove useful.

Go ahead and flip to the section that describes your platform (OS X, Windows, or Linux), and let's begin.

Installing on Mac OS X 10.8 Mountain Lion

You'd think that given the prevalence of OS X among Rails developers (the entire core team uses OS X), installing Rails on OS X would be easy. And you'd be correct. First, we need to install Apple's Developer Tools so that we can compile packages. Note that SQLite is preinstalled on OS X, so that's one thing we don't need to worry about.

Installing the Apple Developer Tools (Xcode)

You need the Apple Developer Tools installed to be able to compile some of the Ruby gems you may need in the future. Before you can compile source code on your Mac, you need to install a compiler. Apple's Developer Tools package, Xcode Tools, includes a compiler and provides the easiest way to set up a development environment on your Mac. The easiest way to download Xcode is from the Mac App Store. Click the Apple menu ▶ App Store and search for Xcode.

Note Xcode is a large download, over 1 GB, so be patient. It may take a while.

Installing Command Line Tools

Once you have Xcode installed you now need to install the command line tools. Inside of Xcode click Xcode ▶ Preferences ▶ Downloads. You should see a listing for the command line tools with a download button next to it. Once that is downloaded, you can continue to the next step.

Installing Homebrew

For this next piece you'll need to dig into the terminal a bit, but don't worry we'll guide you through it. Homebrew is a great package manager for OS X that is written in Ruby no less. It will help you to install the other pieces you'll need as you go. To install Homebrew, enter the following command into the terminal (Applications ▶ Utilities ▶ Terminal):

```
ruby -e "$(curl -fsSL https://raw.github.com/mxcl/homebrew/go)"
```

Installing RVM

Now you'll need to install a common Ruby tool called the Ruby Version Manager (RVM). It helps you manage versions of Ruby and various gems that you may use for projects. Its install is just as simple as Homebrew. It's just one command:

```
\curl -L https://get.rvm.io | bash -s stable
```

You can test to see if Ruby is installed correctly by asking Ruby for its version number:

```
ruby --version
```

```
ruby 2.0.0p0 (2013-02-24 revision 39474) [x86_64-darwin12.3.0]
```

If your output isn't exactly like this, don't panic. Ruby is often updated with new security patches and features, but it should at least say Ruby 2.x.xpx.

Installing Rails

To install Rails, use the command line:

```
gem install rails
```

That's it! Rails is installed, let's check the version to make sure everything went well:

```
rails -v
```

```
Rails 4.0
```

Great! Ruby, SQLite, and Rails are installed and working correctly.

Installing on Windows

Installation on Windows is easy thanks to installer packages. You start by installing Ruby 2.0.

Installing Ruby

Installing Ruby on Windows is marvelously easy thanks largely to the one-click installer for Ruby. You can read more and download the installer from its web site: <http://rubyinstaller.org/>.

The latest version of the installer at the time of this writing is Ruby 2.0.0-p195 for Ruby 2.0, which you can download using this URL: <http://rubyforge.org/frs/download.php/76955/rubyinstaller-2.0.0-p195.exe>.

After you've downloaded the installer, start the installation by double-clicking its icon. What follows is standard installer fare, and the defaults are sufficient for your purposes. When you select the location where you want to put Ruby (usually C:\Ruby20), as shown in Figure 2-1, select the Add Ruby Executables to Your PATH check box; the installer takes care of the rest. You have a fully functioning Ruby installation in minutes.

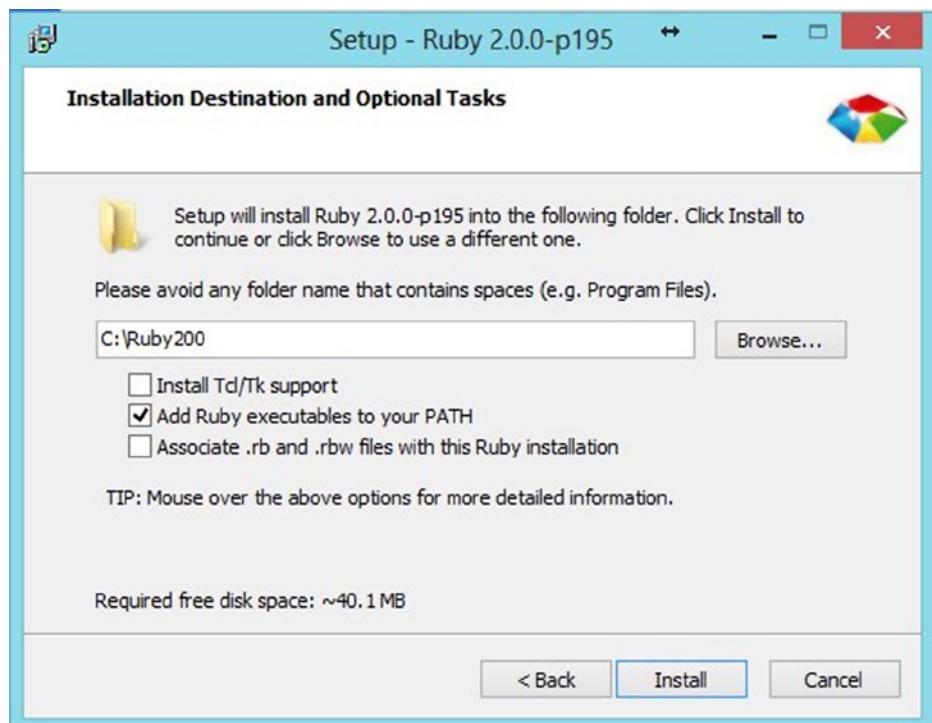


Figure 2-1. Ruby installer for Windows

When the installer is finished, you can test to see if Ruby is working and that your environment is correctly configured by opening your command prompt and asking Ruby its version number:

```
> ruby --version
```

```
ruby 2.0.0p195 (2013-02-24 revision 26273) [i386-mingw32]
```

Installing Rails

You'll be pleased to know that Ruby 2.0 comes bundled with RubyGems, a package-management system for Ruby (<http://rubygems.org>), which makes installing Ruby libraries, utilities, and programs a breeze. This includes Rails installation.

First, let's update RubyGems and its sources list. Open your command prompt and issue the following gem command:

```
> gem update --system
```

Now, to install Rails, issue the following gem command in your command prompt:

```
> gem install rails
```

Be forewarned that the `gem` command can take some time. Don't be discouraged if it seems to be sitting there doing nothing for a few minutes; it's probably updating its index file. RubyGems searches for gems in its remote repository (<http://rubygems.org>), so you need to be connected to the Internet for this command to work.

After spitting out some text to the screen and generally chugging away for a few minutes, the `gem` program should exit with something like the following before dumping you back at the command prompt:

```
Successfully installed rails-4.0
```

That's all there is to it! The one-click installer takes care of most of the work by installing and configuring Ruby; and because Rails is distributed as a RubyGem, installing it is a simple one-liner.

You can double check that Rails was installed successfully by issuing the `rails -v` command at the command prompt:

```
> rails -v
```

```
Rails 4.0
```

Installing SQLite

To install SQLite on Windows, download the following files from the SQLite web site (www.sqlite.org/download.html):

`sqlite-3_6_23_1.zip` - http://www.sqlite.org/sqlite-3_6_23_1.zip
`sqlitedll-3_6_23_1.zip` - http://www.sqlite.org/sqlitedll-3_6_23_1.zip

Note that the version number may be different by the time you read this. Unzip both files, and move their contents to the Ruby bin directory `C:\Ruby20\bin`. When you're done, you can test that you've correctly installed SQLite by issuing the following command from the command prompt:

```
> sqlite3 -version
```

```
3.6.23.1
```

Now that you've installed SQLite, let's install its Ruby binding—a Ruby library that allows you to talk with SQLite. To install the SQLite3 Ruby binding, issue the following `gem` command from the command prompt:

```
> gem install sqlite3-ruby
```

With Ruby, Rails, and SQLite happily installed, it's time to take them for a test drive. Unless you feel like reading the installation instructions for Linux, you're free to skip ahead to the "Creating Your First Rails Application" section.

Installing on Linux

Linux (and UNIX-based systems in general) comes in a variety of different flavors, but they share a lot in common. These instructions use a Debian-based variant called Ubuntu Linux, but they should apply to most UNIX systems with varying mileage.

Note Ubuntu Linux is a top-notch distribution that's rapidly gaining mindshare in the Linux community. At the time of this writing, it's the most popular Linux distribution for general use and is largely responsible for the increased viability of Linux as a desktop platform. It's freely available from <http://ubuntu.org> and highly recommended.

Just about all Linux distributions (including Ubuntu) ship with a package manager. Whether you're installing programs or code libraries, they usually have dependencies; a single program may depend on dozens of other programs in order to run properly, which can be a struggle to deal with yourself. A package manager takes care of these tasks for you, so you can focus on better things.

Ubuntu Linux includes the Debian package manager apt, which is what the examples in this book use. If you're using a different distribution, you likely have a different package manager, but the steps should be reasonably similar.

Before you begin installing Ruby, Rails, and SQLite, update the package library using the `apt-get update` command:

```
$ sudo apt-get update
```

The `apt-get` program keeps a cached index of all the programs and their versions in the repository for faster searching. Running the `update` command ensures that this list is up to date, so you get the most recent versions of the software you need.

Installing Ruby

Before you install Ruby, you need to install a few libraries required by the components you're installing. Enter the following command:

```
$ sudo apt-get install build-essential curl
```

You're going to use the Ruby Version Manager to let Ruby install it for you. This makes everything a snap:

```
$ \curl -L https://get.rvm.io | bash -s stable
```

You can test that this is working by asking Ruby for its version number:

```
$ ruby --version
```

```
ruby 2.0.0p195 (2013-05-14) [i686-linux]
```

Installing Rails

Now you can use RubyGems to install Rails. Enter this command:

```
$ gem install rails
```

After spitting out some text to the screen and generally chugging away for a little while, the `gem` program should exit with a message like the following:

```
Successfully installed rails-4.0
```

You can verify this claim by asking Rails for its version number:

```
$ rails --version
```

```
Rails 4.0
```

With Ruby and Rails happily installed, you're ready to move on to the next step: installing SQLite.

Installing SQLite

To install SQLite with `apt-get`, issue the following command:

```
$ sudo apt-get install sqlite3 libsqlite3-dev
```

If all goes according to plan, you can test your SQLite3 installation by invoking the `sqlite3` program and asking for its version number:

```
$ sqlite3 --version
```

```
3.6.16
```

Now that you've installed SQLite, let's install its Ruby binding—a Ruby library that allows you to talk with SQLite. To install the SQLite3 Ruby binding, issue the following `gem` command from the command prompt:

```
$ sudo gem install sqlite3-ruby
```

With Ruby, Rails, and SQLite happily installed, it's time to take them for a test drive.

Creating Your First Rails Application

You'll start by using the `rails` command to create a new Rails project. Go to the directory where you want your Rails application to be placed; the `rails` command takes the name of the project you want to create as an argument and creates a Rails skeleton in a new directory by the same name. The newly created directory contains a set of files that Rails generates for you to bootstrap your application. To demonstrate, create a new project called (what else?) `hello`:

```
$ rails new hello
```

```
create
create README
create Rakefile
create config.ru
create .gitignore
create Gemfile
create app
create app/controllers/application_controller.rb
create app/helpers/application_helper.rb
create app/views/layouts/application.html.erb
create app/mailers
...
```

If you look closely at the output, you see that the subdirectories of `app/` are named after the MVC pattern introduced in Chapter 1. You also see a name that was mentioned briefly in Chapter 1: *helpers*. Helpers help bridge the gap between controllers and views; Chapter 7 will explain more about them.

Rails generated a new directory called `hello`. If you look at the folder structure, you'll see the following:

<code>Gemfile</code>	<code>app</code>	<code>db</code>	<code>log</code>	<code>test</code>
<code>README</code>	<code>config</code>	<code>doc</code>	<code>public</code>	<code>tmp</code>
<code>Rakefile</code>	<code>config.ru</code>	<code>lib</code>	<code>script</code>	<code>vendor</code>

Starting the Built-In Web Server

Next, let's start up a local web server so you can test your new project in the browser. True, you haven't written any code yet, but Rails has a nice welcome screen that you can use to test whether the project is set up correctly. It even gives you some information about your Ruby environment.

Ruby ships with a built-in, zero-configuration, pure Ruby web server that makes running your application in development mode incredibly easy. You start up the built-in web server using the `rails server` command. To start the server now, make sure you're inside the directory of your Rails application, and then enter the following commands:

```
$ bundle install
```

```
Fetching source index from http://rubygems.org/
```

```
...
```

```
Your bundle is complete! Use 'bundle show [gemname]' to see where a
bundled gem is installed.
```

```
$ rails server
```

```
=> Booting WEBrick
=> Rails 4.0.0 application starting in development on http://0.0.0.0:3000
=> Call with -d to detach
=> Ctrl-C to shutdown server
[2013-05-22 14:26:54] INFO  WEBrick 1.3.1
[2013-05-22 14:26:54] INFO  ruby 2.0.0 (2013-05-14) [i386-darwin10.2.0]
[2013-05-22 14:26:54] INFO  WEBrick::HTTPServer#start: pid=5181 port=3000
```

The message from the `rails server` command tells you that a web server is running at the IP address 0.0.0.0 on port 3000. Don't be alarmed by this all-zeros address—it simply means that the server is running locally on your machine. The hostname `localhost` also resolves to your local machine and is thus interchangeable with the IP address. We prefer to use the hostname variant.

With the server running, if you open <http://localhost:3000> in your browser, you'll see the Rails welcome page, as shown in Figure 2-2. Congratulations! You've put Ruby on Rails.

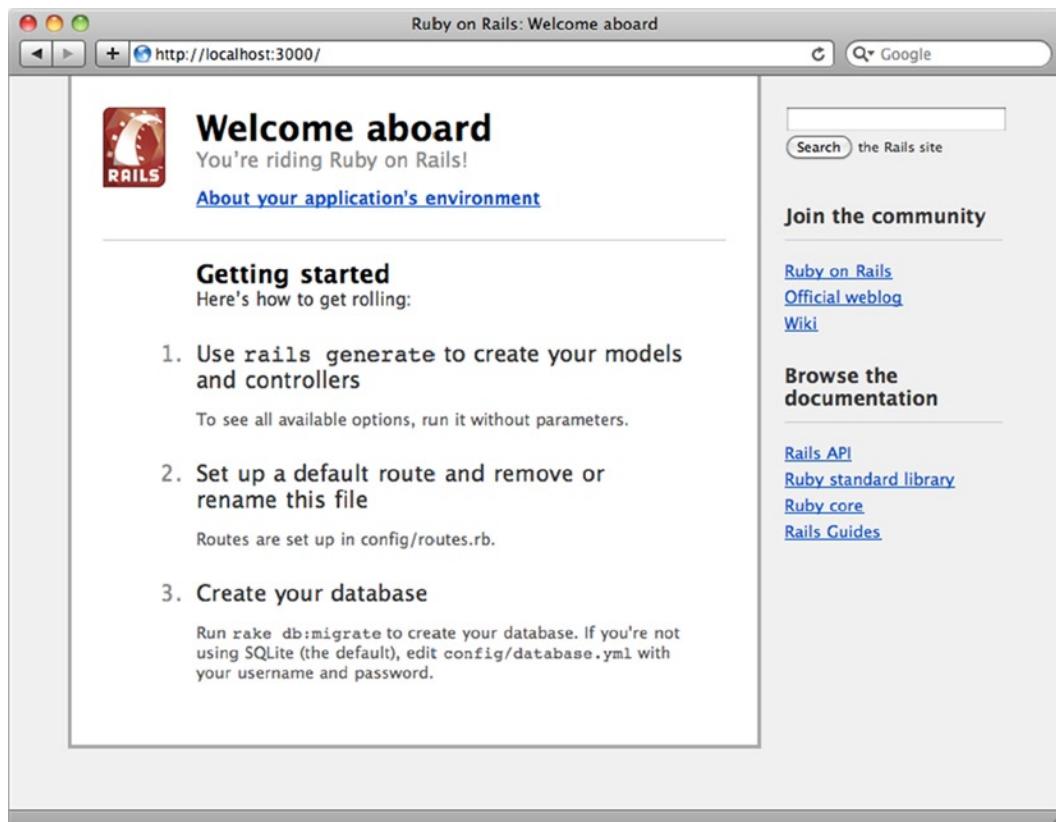


Figure 2-2. Rails welcome page

The welcome page is nice, but it doesn't teach you much. The first step in learning how Rails works is to generate something dynamic. You're about to learn why you called this project "hello"!

We're sure it would be in violation of the law of programming books if we didn't start with the ubiquitous "Hello World" example. And who are we to disobey? In the next few steps, you make your Rails application say hello; and in doing so, you learn a few new concepts. Your goal is to have a request with the URL <http://localhost:3000/salutation/hello> respond with a friendly "Hello World!" message.

First things first: stop the web server by pressing Ctrl+C in the command prompt window. That should bring you back to your prompt.

Note Notice how easy it is to start and stop a local server? That's the whole point of the built-in server in a nutshell. You shouldn't need to be a system administrator to develop a Rails application.

Generating a Controller

You use the `rails` command's `generate` option to create certain files within your project. Because you're dealing with the request and response cycle (you request a URL, and the browser receives a response), you generate a controller that is responsible for handling salutations:

```
$ rails generate controller salutation
```

```
create  app/controllers/salutation_controller.rb
invoke  erb
create    app/views/salutation
invoke  test_unit
create    test/functional/salutation_controller_test.rb
invoke  helper
create    app/helpers/salutation_helper.rb
invoke  test_unit
create    test/unit/helpers/salutation_helper_test.rb
```

Not unlike the `rails` command you used to generate your application, the `rails generate controller` command creates a bunch of new files. These are mostly empty, containing only skeletal code (often called *stubs*). You could easily create these files on your own. The generator merely saves you time and the effort of needing to remember which files to create and where to put them.

The salutation controller was created in the `app/controllers` directory and is sensibly named `salutation_controller.rb`. If you open it with a text editor, you see that there's not much to it, as shown in Listing 2-1.

Listing 2-1. The `app/controllers/salutation_controller.rb` File

```
class SalutationController < ApplicationController
end
```

Creating an Action

If you want `SalutationController` to respond to a request for `hello`, you need to make an action for it. Open `salutation_controller.rb` in your text editor and add the `hello` action, as shown in Listing 2-2.

Listing 2-2. The Updated app/controllers/salutation_controller.rb File: <http://gist.github.com/319866>

```
class SalutationController < ApplicationController
  def hello
    @message = 'Hello World!'
  end
end
```

Actions are implemented as Ruby methods. You can always tell a method definition because of the `def` keyword. Inside the action, you set a Ruby instance variable called `@message`, the value of which you output to the browser.

Creating a Template

With your action successfully defined, your next move is to add some HTML into the mix. Rails makes it easy by separating the files that contain HTML into their own directory as per the MVC pattern. In case you haven't guessed, HTML is the responsibility of the view.

If you look in the `app/views` directory, you see another product of the controller generator: a directory called `salutation`. It's linked to the `salutation` controller, and it's where you put template files that correspond to your `salutation` actions.

Note Because Rails allows you to embed Ruby code in your HTML by using the ERb Templating library, you use the `.html.erb` (HTML + ERb) extension for your templates.

The default way to render a template in response to a request for an action is remarkably simple: name it the same as the action. This is another case of using a predefined Rails convention. Because you want to show a response to the `hello` action, name your file `hello.html.erb`, and Rails renders it automatically. This is easy to grasp in practice. Figure 2-3 gives a visual cue as to how controllers and templates correspond.

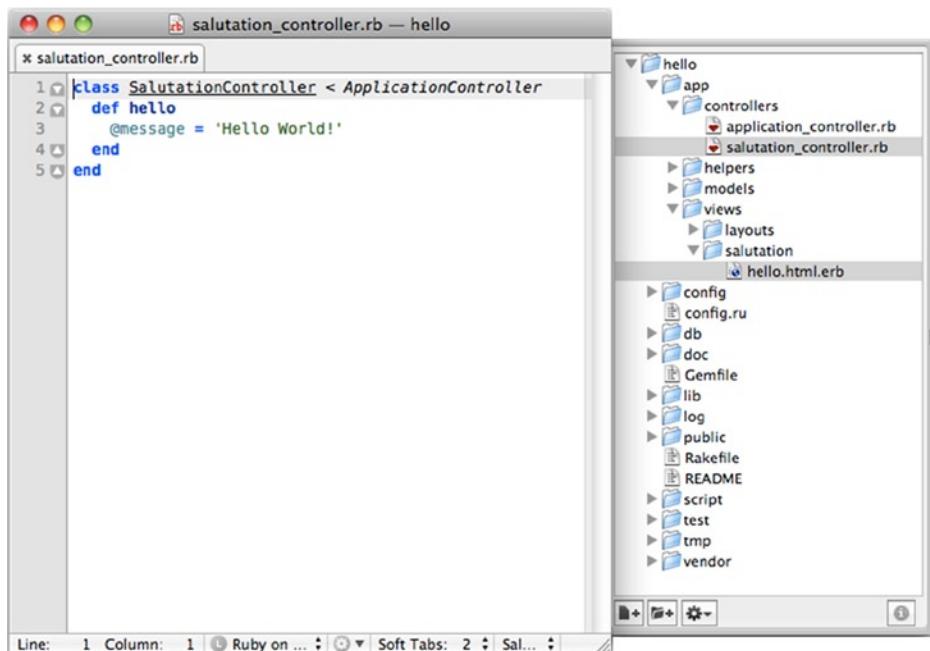


Figure 2-3. Controllers correspond to a directory in app/views

Start by creating a new, blank file in app/views/salutation/. Name it hello.html.erb, and add the code shown in Listing 2-3. Notice the <%= %> syntax that surrounds the @message variable: these are known as Embedded Ruby (ERb) output tags. Chapter 7 explains more about ERb. For now, it's only important to know that whenever you see <%= %> in a template, whatever is between the tags is evaluated as Ruby, and the result is printed out.

Listing 2-3. The app/views/salutation/hello.html.erb File: <http://gist.github.com/319910>

```
<html>
  <body>
    <h1><%= @message %></h1>
  </body>
</html>
```

You now have to tell your Rails application how to respond to a URL. You do that by updating the config/routes.rb file. You don't need to worry about the details of how the routes file work for now, Chapter 8 will cover that. Replace the contents of your config/routes.rb file and make sure it looks like Listing 2-4.

Listing 2-4. The config/routes.rb File: <http://gist.github.com/319933>

```
Hello::Application.routes.draw do
  get ':controller(/:action(/.:format))'
end
```

It looks like you're all set. The salutation controller fields the request for hello and automatically renders the hello.html.erb template. Start up the web server again using the rails server command, and request the URL <http://localhost:3000/salutation/hello> in your browser. You should see the result shown in Figure 2-4.

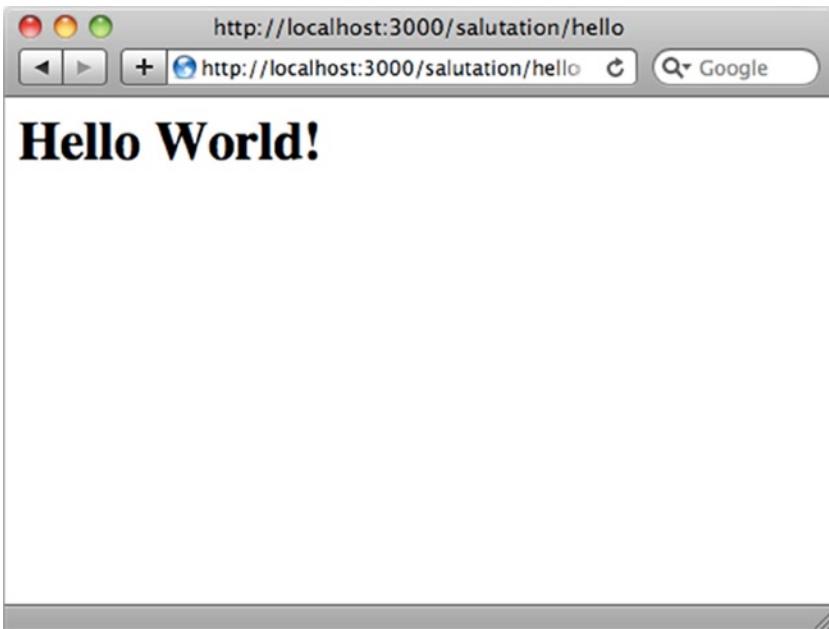


Figure 2-4. The “Hello World” application

Sure enough, there’s your greeting! The `hello` template reads the `@message` variable that you set in the controller and, with a little help from ERb, printed it out to the screen.

In case you didn’t notice, the URL `http://localhost:3000/salutation/hello` maps directly to the controller and action you created because of the change you made to your `config/routes.rb` file. This is achieved by the following URL pattern:

```
:controller(/:action(/:id(.:format)))
```

For now, it’s enough that you can make the mental connection between URL segments and the code in your application. When you read the URL `http://localhost:3000/salutation/hello`, you can translate it into a request for the `hello` action on the `salutation` controller.

Summary

This chapter covered a lot, so you should be proud of yourself. You went from not having Rails installed to getting a basic Rails application up and running. You learned how to install Ruby and how to manage packages with RubyGems (which you used to install Rails). You also learned how to create a new Rails project using the `rails` command and how to use the generator to create a new controller. And you learned how controller actions correspond to templates. The stage is now set for the next chapter, where you begin building a more full-featured project.

CHAPTER 3



Getting Something Running

The best way to learn a programming language or a web framework is to dig in and write some code. After reading the first two chapters, you should have a good understanding of the Rails landscape. Chapter 4 will lead you through the Ruby language, but first let's write a little code to whet your appetite. This chapter builds a foundation and will get you excited by walking you through the construction of a basic application. You will learn how to create a database and how to connect it to Rails, as well as how to use a web interface to get data in and out of the application.

You will receive a lot of information in this chapter, but it shouldn't be more than you can absorb. The goal is to demonstrate, not to overwhelm. Rails makes it incredibly easy to get started, and that's a feature this chapter highlights. There are a few places where Rails really shines, and getting something running is one of them. By the end of this chapter, you'll have a working web application to play with, explore, and learn from. You'll build on this application throughout the rest of the book, adding features and refining functionality.

An Overview of the Project

This chapter will walk you through building a simple blog application that lets you create and publish articles, like WordPress or Blogger. The first iteration focuses on the basics: creating and editing articles.

Before you start coding, let's sketch a brief summary of the goals and flow of the application at a very high level. The idea isn't to focus on the nitty-gritty, but instead to concentrate on the general case.

Your application will have two kinds of users: those who post and publish articles and those who wish to comment on existing articles. In some cases, people will play both roles. Not all users will need to create an account by registering on the site. It will also be nice if people can notify their friends about interesting articles using a feature that sends a friendly e-mail notification to interested parties.

You will add some of these features in later chapters. Other application requirements will likely come up as you continue, but these are enough to get started. In the real world, specifications are seldom correct the first time around, so it's best not to dwell on them. Rails doesn't penalize you for making changes to an application that's under construction, so you can engage in an iterative style of development, adding and incrementing functionality as you go.

You start with what matters most: articles. You may wonder why you don't begin with users. After all, without users, who will post the articles? If you think about it, without articles, what could users do? Articles are the epicenter of the application, so it makes the most sense to start there and work out the details as you go. Ready? Let's get started!

Creating the Blog Application

As you saw in Chapter 2, the first step is to create a new Rails application. You could come up with a fancy name, but let's keep it simple and call the application `blog`. It's not going to win any awards for creativity, but it works.

To begin, from the command line, go to the directory where you want to place your new application; then, issue the `rails` command to generate the application skeleton and base files:

```
$ rails new blog
```

```
create
create  README.rdoc
create  Rakefile
create  config.ru
create  .gitignore
create  Gemfile
...

```

As you recall from the example in Chapter 2, the `rails` command takes as an argument the name of the project you want to create and generates a directory of the same name that contains all the support files. In this case, it creates a subdirectory called `blog` in the current working directory. Change into the `blog` directory and get oriented.

Figure 3-1 shows the directory structure.

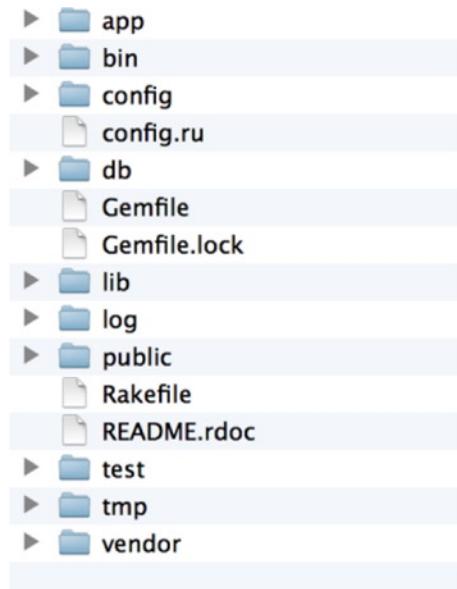


Figure 3-1. The Rails directory structure

You'll quickly get used to the Rails directory structure, because all Rails applications follow this standard. This is another benefit of conventions: you always know where to locate files if you have to work on a Rails project that was developed by someone else. Table 3-1 briefly explains the directory structure.

Table 3-1. Rails Directory Structure

Folder	Description
app	All the components of your application.
bin	Executables to support Rails.
config	Configuration files for all of the components of your application.
config.ru	A file used by rack servers to start the application.
db	Files related to the database you're using, and a folder for migrations.
Gemfile	Used by the bundler gem to keep a list of gems used in your application.
Gemfile.lock	Canonical resource of what gems should be installed.
lib	Libraries that may be used in your application.
log	Log files that your application may require.
public	Static assets served by your application, such as images, JavaScript, and CSS files.
Rakefile	Lists available for tasks used by Rake.
README.rdoc	Human readable file generated to describe an application.
test	Directory containing test unit tests for your application.
tmp	Contains temporary files supporting your application.
vendor	External libraries, such as gems and plug-ins, that your application bundles.

Your first stop is the config directory. Of the little configuration there is to do in a Rails application, most of it takes place in this aptly named location. To get an idea of what Rails expects as far as databases go, open the config/database.yml file in your editor and take a peek. You should see something like the file shown in Listing 3-1 (comments are omitted here).

Listing 3-1. The config/database.yml File

```
# SQLite version 3.x
#   gem install sqlite3
#
#   Ensure the SQLite 3 gem is defined in your Gemfile
#   gem 'sqlite3'
development:
  adapter: sqlite3
  database: db/development.sqlite3
  pool: 5
  timeout: 5000

# Warning: The database defined as "est" will be erased and
# re-generated from your development database when you run "rake".
# Do not set this db to the same as development or production.
test:
  adapter: sqlite3
  database: db/test.sqlite3
  pool: 5
  timeout: 5000
```

```
production:
  adapter: sqlite3
  database: db/production.sqlite3
  pool: 5
  timeout: 5000
```

The first thing you should notice is the different sections: `development`, `test`, and `production`. Rails understands the concept of *environments* and assumes you're using a different database for each environment. Therefore, each has its own database connection settings, and different connection parameters are used automatically. Rails applications run in *development* modes by default, so you really only need to worry about the `development` section at this point. Still, other than the database names (`db/*.sqlite3`), there should be little difference between the connection parameters for each environment.

This example uses the default SQLite database because it's easy to use and set up. However, you can use the database management system of your choice by passing the `-d` or `--database=` option to the `rails` command with one of the following options as per your preference: `mysql`, `oracle`, `postgresql`, `sqlite2`, `sqlite3`, `frontbase`, or `ibm_db`.

If you select a database other than SQLite, the `rails` command may prefill the `database` parameter based on the database server and project name: `blog` in this case. If you give your application a different name (say, a snazzy Web 2.0 name like `blog.ilicio.us *beta`) with a database server such as MySQL, you'll see something different here. It doesn't matter what you name your databases, as long as `database.yml` references the correct one for each environment. Let's stick with the convention and create the databases using the default names.

WHAT IS YAML?

The `.yml` extension refers to a YAML file. YAML (a recursive acronym that stands for "YAML Ain't Markup Language") is a special language for expressing objects in plain text. Rails can work with YAML natively and can turn what looks like plain text into Ruby objects that it can understand.

YAML is whitespace sensitive: it uses spaces (not tabs) to convey structure and meaning. Make sure your editor knows the difference between tabs and spaces, and be sure that when you're editing YAML files, you use only spaces.

Creating the Project Databases

You may think that to create a new database, you'll use your favorite database administration tool. However, because you already told Rails the database connection details, you can now run a Rake task that talks to the database and issues all the necessary commands to set up the databases. Jump to the command prompt and type:

```
$ rake db:create
```

When using SQLite, you aren't forced to create the database, because a new database file is automatically created if one doesn't exist; but it will come in handy when you try a different database engine. You also may see some messages like `db/development.sqlite3 already exists`. Don't be afraid—this is an indication that an SQLite file was found. If you see that message, rest assured that your existing database was left untouched, and no database file has been harmed.

Note Rake is a build language for Ruby. Rails uses Rake to automate several tasks, such as running database migrations, running tests, and updating Rails support files. You can think of Rake tasks as little utility programs. For a list of all available Rake tasks, run `rake -T` from your Rails project directory. For more information about Rake, including complete documentation, see <http://rake.rubyforge.org/>.

Regardless of the database management system you select, you should notice that the databases you want to use are created. This is another case in which Rails removes some complexity from your mind and helps you focus on your application.

Note Depending on how your environment is set up, you may not need to specify the username, password, and other options in your `config/databases.yml` file to create the database.

Although you're only concerned with the development environment at this time, it doesn't hurt to create the other databases while you're at it. Go ahead and create two more databases, one each for the test and production environments:

```
$ rake db:create:all
```

You can confirm the creation of the database by using the `rails dbconsole` program to look at the databases currently on your system:

```
$ rails dbconsole
```

```
SQLite version 3.7.12Enter ".help" for instructions
Enter SQL statements terminated with a ";"
```

sqlite> .exit

At this point, you can issue any number of SQL (Structured Query Language) statements and look at the tables and records that eventually will be in your application. (If you aren't familiar with SQL, you can learn more about it in Appendix B.) When you're finished with the SQLite console, type the `.exit` command to go back to your regular prompt. You can test to see if your connection is working by running the following command:

```
$ rake db:migrate
```

If nothing exceptional is returned, congratulations! Rails can connect to your database. However, if you're using a database engine other than SQLite, and you may see something like this:

```
rake aborted!
Access denied for user 'root'@'localhost' (using password: NO)
```

Then you need to adjust your connection settings. If you're having problems, make sure the database exists and that you've entered the correct username and password in the `config/database.yml` configuration file.

Creating the Article Model

Now that you can connect to the database, this section will explain how you create a model. Remember that models in Rails correspond to database table names. Because you want to model articles, let's create a model named `Article`. By convention, model names are camel-cased singular and correspond to lower-cased plural table names. So, an `Article` model expects a table named `articles`; a `Person` model expects a table named `people`.

Note Camel case means that each word begins with a capital letter and is written without spaces. For instance, a class that described blog images would be written as `BlogImage`. Refer to <http://en.wikipedia.org/wiki/CamelCase> for more information.

Rails is smart enough to use the correct plural name for most common words; it doesn't try to create a `persons` table.

Like most things in Rails, models have their own generator script that makes it easier to get started. The generator automatically creates a new model file in the `app/models` directory and also creates a bunch of other files. Among these are a unit test (for testing your model's functionality, as discussed in Chapter 11) and a database migration. A *database migration* contains instructions for modifying the database table and columns. Whenever you generate a new model, a migration is created along with it.

Note If you want to skip generation of the migration when generating a new model, you can pass the `--skip-migration` argument to the generator. This may be useful if you're creating a model for an existing database or table.

To see the generator's usage information, run it without arguments:

```
$ rails generate model
```

```
Usage:  
  rails generate model NAME [field[:type][:index] field[:type][:index]] [options]  
...
```

As you can see from the usage banner, the generator takes a model name as its argument and an optional list of fields. The model name may be given in camel-cased or snake-cased format, and options can be provided if you want to automatically populate the resulting migration with column information.

Note Snake-cased words are written in all lowercase with underscores replacing spaces, for instance, `blog_image`. For more information visit http://en.wikipedia.org/wiki/Snake_case.

Let's run the generator now to create the first model, Article:

```
$ rails generate model Article
```

```
invoke  active_record
create  db/migrate/20130406203049_create_articles.rb
create  app/models/article.rb
invoke  test_unit
create  test/unit/article_test.rb
create  test/fixtures/articles.yml
```

If you look at the lines that start with `create`, you see that the generator has created an article model, an article test, an articles fixture (which is a textual representation of table data you can use for testing), and a migration named **20130406203049_create_articles.rb**. From that, your model is generated.

Note The first part of the migration file name is the timestamp when the file was generated. So, the file on your computer will have a slightly different name.

Creating a Database Table

You need to create a table in the database. You could do this with a database administration tool or even manually using SQL, but Rails provides a much more efficient facility for table creation and maintenance called a *migration*. It's called a migration because it allows you to evolve, or migrate, your schema over time. (If you're not familiar with databases, tables, and SQL, consult Appendix B for the basics.)

Note *Schema* is the term given to the properties that make up a table: the table's name, its columns, and its column types, as well as any default values a column will have.

What's the best part about migrations? You get to define your schema in pure Ruby. This is all part of the Rails philosophy that you should stick to one language when developing. It helps eliminate context switching and results in higher productivity.

As you can see from the output of the model generator, it created a new file in `db/migrate` called `20130406203049_create_articles.rb`. As mentioned before, migrations are named with a numeric prefix, which is a number that represents the exact moment when the migration file was created. Because multiple developers can create migrations in a development team, this number helps uniquely identify this specific migration in a project.

Let's open this file and take a peek. It's shown in Listing 3-2.

Listing 3-2. The `db/migrate/20130406203049_create_articles.rb` File

```
class CreateArticles < ActiveRecord::Migration
  def change
    create_table :articles do |t|
      t.timestamps
    end
  end
end
```

In its initially generated form, the migration is a blank canvas. But before you go any further, let's note a few important items. First, notice the instance method: `change`. In previous versions of Rails, there would be an up and down class method, but now Rails is smart enough to figure it out based on the modifications you make in this method. You can roll back without ever writing a method that explicitly drops the table. Pretty slick, isn't it?

Listing 3-3 has the details filled in for you. Even without ever having seen a migration before, you should be able to tell exactly what's going on.

Listing 3-3. Completed db/migrate/20130406203049_create_articles.rb File

```
class CreateArticles < ActiveRecord::Migration
  def change
    create_table :articles do |t|
      t.string      :title
      t.text       :body
      t.datetime   :published_at

      t.timestamps
    end
  end
end
```

Let's step through the code. First, you use the `create_table` method, giving it the name of the table you want to create. Inside the code block, the `string`, `text`, and `datetime` methods each create a column of the said type named after the parameter; for example, `t.string :title` creates a field named `title` with the type `string`. The `timestamps` method, in the `t.timestamps` call, is used to create a couple of fields called `created_at` and `updated_at`, which Rails sets to the date when the record is created and updated, respectively. (For a full description of the available method types you can create in your migrations, see <http://api.rubyonrails.org/classes/ActiveRecord/Migration.html>.)

On its own, this migration does nothing. Really, it's just a plain old Ruby class. If you want it to do some work and create a table in the database for you, you need to run it. To run a migration, you use the built-in `db:migrate` Rake task that Rails provides.

From the command line, type the following to run the migration and create the `articles` table. This is the same command you used to test the database connection. You sort of hijack it for this test, knowing that it will attempt to connect to the database and thus prove whether the connection works. Because there were no existing migrations when you first ran it, it didn't do anything. Now that you have your first migration, running it results in a table being created:

```
$ rake db:migrate
```

```
-- CreateArticles: migrating =====
-- create_table(:articles)
 -> 0.0019s
== CreateArticles: migrated (0.0023s) =====
```

Just as the output says, the migration created the `articles` table. If you try to run the migration again (go ahead, try it), nothing happens. That's because Rails keeps track of all the migrations it runs in a database table, and in this case there's nothing left to do. If for some reason you decide you need to roll back the migration, you can use the `db:rollback` task to roll back. Try it and you will notice that it dropped the `articles` table. Remember that we never wrote any code to drop the table, Rails just handled it for us. Imagine if you would have edited the database schema directly with a database management tool; if you wanted to roll back, you'd have to remember what it looked like before and exactly what you changed. This makes your life much easier. Okay, before we move on, don't forget to run migrations again since we rolled back.

Generating a Controller

You've created a model and its supporting database table, so the next step is to work on the controller and view side of the application. Let's create a controller named `articles` (remember controllers are plural and models are singular) to control the operation of the application's articles functionality. Just as with models, Rails provides a generator that you can use to create controllers:

```
$ rails generate controller articles
```

```
create    app/controllers/articles_controller.rb
invoke    erb
create    app/views/articles
invoke    test_unit
create    test/controllers/articles_controller_test.rb
invoke    helper
create    app/helpers/articles_helper.rb
invoke    test_unit
create    test/helpers/articles_helper_test.rb
invoke    assets
invoke    coffee
create    app/assets/javascript/articles.js.coffee
invoke    scss
create    app/assets/stylesheets/articles.css.scss
```

The controller generator creates six files:

- `app/controllers/articles_controller.rb`: The controller that is responsible for handling requests and responses for anything to do with articles.
- `test/controllers/articles_controller_test.rb`: The class that contains all functional tests for the `articles` controller (Chapter 11 covers testing applications).
- `app/helpers/articles_helper.rb`: The `helper` class in which you can add utility methods that can be used in your views (Chapters 7 and 8 cover helpers).
- `test/helpers/articles_helper_test.rb`: The class that contains all helper tests for your `helper` class.
- `app/assets/javascripts/articles.js.coffee`: This is a CoffeeScript file for writing the JavaScript associated with the views for our articles. We'll go in to more depth about JavaScript later.
- `app/assets/stylesheets/articles.css.scss`: This is a SASS (Syntactically Awesome Stylesheets) file where you can put style sheets for the associated views.

Note CoffeeScript and SASS are languages that compile into JavaScript and CSS, respectively. Rails supports these by default by using the asset pipeline to compile them into more standard languages automatically.

The controller generator also creates an empty directory in `app/views` called `articles`. This is where you place the templates for the `articles` controller.

Up and Running with Scaffolding

One of the most talked-about features that has given a lot of exposure to Rails is its scaffolding capabilities. *Scaffolding* allows you to create a boilerplate-style set of actions and templates that makes it easy to manipulate data for a specific model. You generate scaffolding using the scaffold generator. You're probably getting used to generators by now. Rails makes heavy use of them because they help automate repetitive tasks and generally remove the chances for errors when creating new files. Unlike you probably would, the generator won't ever forget how to name a file; nor will it make a typo when creating a class. Let's use the scaffold generator now and solve the mystery of how this works:

```
$ rails generate scaffold Article title:string body:text published_at:datetime  
--skip-migration
```

Because you've already generated a few of these files, Rails prompts you before it tries to overwrite any that already exist. Specifically, it warns about the articles controller, the articles functional test, the articles fixture file, and both the SASS and CoffeeScript files. Because your articles controller was empty, it's safe to overwrite it. The same goes for the test and fixture, so go ahead and answer yes by pressing Y when you're asked. You used the --skip-migration option when calling the generator because a migration creating the articles table already existed, and Rails would complain if you tried to create the same table twice.

The scaffold provides methods and pages that allow you to insert, update, and delete records in your database. That's all you need to generate a working scaffold of the Article model. Let's fire up the web server and test it. Start your local web server from the command line (`rails server`), and browse to the articles controller in your browser:

<http://localhost:3000/articles>

You should see the results displayed in your browser, as shown in Figure 3-2.



Figure 3-2. Articles scaffolding

Click the New article link, and you're taken to a screen where you can enter articles. Notice that the URL is <http://localhost:3000/articles/new>, which means you're invoking the new action on the articles controller. Go ahead and add a few articles and generally play with the application. Figure 3-3 shows an example of an article entered on this screen.

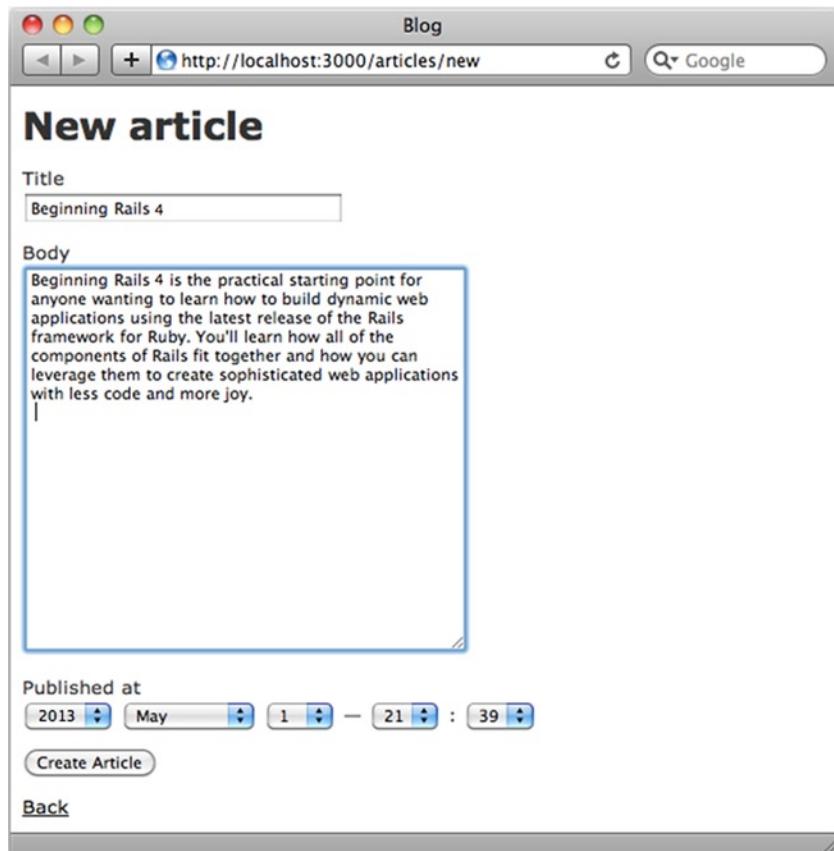


Figure 3-3. Adding an article

Notice that every time you add an article, you're redirected back to the index action, where you see all your articles listed. You can edit them, delete them, or create new ones. If you click the Show link, you're taken to the detail page for that article. You've got to admit, Rails gives you a lot of functionality for free.

Speed is the key benefit here. The scaffold generator allows you to quickly get something running, which is a great way to test your assumptions.

Caution Scaffolding comes with an important disclaimer. You shouldn't use it in production. It exists to help you do exactly what you just did: get something running. By its definition, it's a temporary or unfinished product.

Adding More Fields

Now that you can see the model represented in the browser, let's add some more fields to make it a little more interesting. Whenever you need to add or modify database fields, you should do so using a migration. In this case, let's add the excerpt and location fields to the articles table.

You didn't need to generate the last migration (the one you used to create the articles table), because the model generator took care of that for you. This time around, you can use the migration generator. It works just like the model and controller generators, which you've seen in action. All you need to do is give the migration generator a descriptive name for the transformation:

```
$ rails generate migration add_excerpt_and_location_to_articles
excerpt:string location:string
```

invoke	active_record
create	db/migrate/20130423232337_add_excerpt_and_location_to_articles.rb

As you've already seen, the generator creates a migration class in db/migrate prefixed by a number identifying when the migration was created. If you open the 20130423232337_add_excerpt_and_location_to_articles.rb file, you see the migration class with the code shown in Listing 3-4. As with the model generator, which prefilled the migration to some extent, passing field names and types as options to the migration generator prefills the generated class for you as long as you refer to the correct table name at the end of the migration name—in this case, to_articles.

Listing 3-4. The db/migrate/20100223232337_add_excerpt_and_location_to_articles.rb File

```
class AddExcerptAndLocationToArticles < ActiveRecord::Migration
  def change
    add_column :articles, :excerpt, :string
    add_column :articles, :location, :string
  end
end
```

Looking at the add_column method, the first argument is the table name (`articles`), the second is the field name, and the third is the field type. Remember that the change method knows how to migrate up or down, so if in the unlikely event you want to remove these columns, Rails will know how.

With this new migration in place, use the Rake task to apply it and make the changes to the database:

```
$ rake db:migrate
```

```
--> 0.0236s
--> 0.0013s
==> 0.0164s
```

If all goes according to plan, the articles table now has two new fields. You could edit the view templates in the app/views/articles folder to add form elements for the new fields, but instead let's call the generator again (you'll learn about views in Chapter 7):

```
$ rails generate scaffold Article title:string location:string excerpt:string
body:text published_at:datetime --skip-migration
```

Press Y when asked if you want to overwrite some files, and you're finished, as you can see in Figure 3-4.

The screenshot shows a web browser window with the title "Articles: new". The URL in the address bar is "http://localhost:3000/articles/new". The page content is a form titled "New article". It has four main sections: "Title" (input field containing "Beginning Rails 4"), "Location" (input field containing "Toronto, Canada"), "Excerpt" (input field empty), and "Body" (text area containing a large block of Lorem ipsum text). Below the body section is a "Published at" field, which is a date-time selector showing "2013" and "February" for the year and month respectively, with "23" and "23" for the day and hour, and "40" for the minute. At the bottom of the form are two buttons: "Create Article" and "Back".

Figure 3-4. Additional fields added to the new article form

This exposes one of the issues of this type of scaffolding: when you generate new versions of the scaffold files, you run the risk of overwriting custom changes you may have made. We're doing it this way as an illustration, but you wouldn't normally do this.

Adding Validations

You may wonder what happens if you try to save a new article without giving it any information. Try doing that: Rails doesn't care. Actually, it's the Article model that doesn't care. This is because in Rails, the rules for data integrity (such as required fields) are the responsibility of the model.

To add basic validation for required fields, open the Article model in `app/models/article.rb` and add the validation method shown in Listing 3-5 inside the class body.

Listing 3-5. Validation Added to the `app/models/article.rb` File

```
class Article < ActiveRecord::Base
  validates_presence_of :title, :body
end
```

Save the file, and try creating an empty article again. Instead of saving the record, Rails displays a formatted error message, as shown in Figure 3-5.

The screenshot shows a Mac OS X desktop environment with a web browser window titled "Blog". The address bar indicates the URL is `http://localhost:3000/articles`. The main content area is titled "New article". A red box highlights an error message: "2 errors prohibited this article from being saved:" followed by a bulleted list: "Title can't be blank" and "Body can't be blank". Below this, the "Title" field is empty and highlighted with a red border. The "Location" and "Excerpt" fields are also empty. The "Body" field is empty and highlighted with a red border. A vertical blue scroll bar is visible on the right side of the browser window.

Figure 3-5. Error messages for an article

If you've done any web development before, you know that validating fields is a major nuisance. Thankfully, Rails makes it easy.

Note Notice that you don't need to restart the web server when you make changes to your project files in the app/ directory. This is a convenience provided by Rails when running in development mode.

Chapter 6 goes through all the specifics of model validations. For now, you're using only the most primitive methods of protecting your data. It shouldn't surprise you that Active Record is capable of much more involved validations, such as making sure a numeric value is entered, validating that data are in the correct format using regular expressions, and ensuring unique values, among other checks.

Note *Regular expressions* (regex for short) are expressions that describe patterns in strings. Like most programming languages, Ruby has built-in support for regular expressions.

Generated Files

Now that you've seen the pages in action, let's look at the articles controller again. As you can see in Listing 3-6, the controller is now chock-full of actions. There's one for each of index, show, new, create, edit, update, and destroy—the basic CRUD actions.

Listing 3-6. The app/controllers/articles_controller.rb

```
class ArticlesController < ApplicationController
  before_action :set_article, only: [:show, :edit, :update, :destroy]

  # GET /articles
  # GET /articles.json
  def index
    @articles = Article.all
  end

  # GET /articles/1
  # GET /articles/1.json
  def show
  end

  # GET /articles/new
  def new
    @article = Article.new
  end

  # GET /articles/1/edit
  def edit
  end

  # POST /articles
  # POST /articles.json
  def create
    @article = Article.new(article_params)
```

```

respond_to do |format|
  if @article.save
    format.html { redirect_to @article, notice: 'Article was successfully created.' }
    format.json { render action: 'show', status: :created, location: @article }
  else
    format.html { render action: "new" }
    format.json { render json: @article.errors, status: :unprocessable_entity }
  end
end
end

# PATCH/PUT /articles/1
# PATCH/PUT /articles/1.json
def update
  respond_to do |format|
    if @article.update(article_params)
      format.html { redirect_to @article, notice: 'Article was successfully updated.' }
      format.json { head :no_content }
    else
      format.html { render action: "edit" }
      format.json { render json: @article.errors, status: :unprocessable_entity }
    end
  end
end # DELETE /articles/1
# DELETE /articles/1.xml
def destroy
  @article.destroy

  respond_to do |format|
    format.html { redirect_to articles_url }
    format.json { head :no_content }
  end
end

private
# Use callbacks to share common setup or constraints between actions.
def set_article
  @article = Article.find(params[:id])
end

# Never trust parameters from the scary internet, only allow the white list through.
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at)
end
end

```

As you did in this chapter, after you've generated scaffolding, if you change your model, you have to regenerate it if you want your application to follow suit. Most of the time, however, you make the changes by hand and have a variation of the default scaffold.

It's important to realize why scaffolding exists and to be aware of its limitations. As you've just seen, scaffolding helps when you need to get something running quickly to test your assumptions. It doesn't take you very far in the real world, and eventually you end up replacing most (if not all) of it.

Explore the generated code and see if you can figure out how it hangs together. Don't worry if you can't understand all of it—the chapters that follow will discuss it in depth. With everything you know about Rails already, you should be able to piece together most of it.

Try changing a few things to see what happens. If you inadvertently break something, you can always run the scaffolding generator again to revert to the original. Can you see how the views in `app/views/articles` are related to the actions? What about the response messages, like `Article was successfully created`? What happens when you change them? See if you can find where the error messages for failed validations are rendered. If you remove the message, does the record still get saved? You can learn a lot by exploring, so take as much time as you need.

Summary

This chapter started by outlining the basics of the sample application. Then, you rolled up your sleeves and created a database and configuration files. Based on the goals of the application, you began by creating the tables necessary to run the core of your `Article` model and got a first look at the simplicity and flexibility that migrations give the development process. The scaffolding allowed you to test your assumptions about the model and table you created by getting a firsthand look at it in action. You also took a first crack at adding in validations that ensure you maintain the integrity of your data. The chapters that follow investigate these concepts in depth, starting with the first part of the MVC principle: models.

CHAPTER 4



Introduction to the Ruby Language

Rails is a great framework for the development of web-based applications. One of its greatest advantages over other web frameworks is that it's written in Ruby, a very consistent and elegant object-oriented programming language. In order to increase your productivity as a Rails developer, it's important that you master Ruby. If you're new to programming, don't worry: we explain the concepts in a way you can understand.

Ruby was made to make developers happy. This should be exciting to you because you're a developer, and you want to be happy! Some languages feel like the creator was in a bad mood and hated you. Ruby tries its best to make you feel at ease and in control. As you grow as a developer, you'll understand the importance of this fact more and more, especially if you do this for a living.

This chapter gives you an overview of the features of the Ruby language. It explains how the language is organized and presents its fundamentals. After reading this chapter, you should better understand how the Ruby language that Rails is built on works, and you should be able to create classes and methods and use control-flow statements in your code. The best way to learn is to explore the language using this chapter as a guide. It's important that you run the examples given yourself, and also to try things on your own.

Ruby has far more features than we can mention in this short introduction. We encourage you to investigate more of the complex features of Ruby as you continue using Rails.

Instant Interaction

A lot of languages require that you write some code, compile, and then run the program to see the results. However, Ruby is *dynamic*, which means you can work with the language live. You will get instant feedback from your commands.

Ruby comes with a great little tool: an interactive interpreter called `irb` (for Interactive Ruby). You can start up an `irb` session whenever you want by typing `irb` at the command prompt. Using `irb`, you can play around with code and make sure it works as you expect before you write it into your programs.

You can execute any arbitrary Ruby code in `irb` and do anything you would otherwise do inside your Ruby programs: set variables, evaluate conditions, and inspect objects. The only essential difference between an interactive session and a regular old Ruby program is that `irb` echoes the return value of everything it executes. This saves you from having to explicitly print the results of an evaluation. Just run the code, and `irb` prints the result.

You can tell when you're in an `irb` session by looking for the `irb` prompt, which looks like `irb(main):001:0>`, and the arrow symbol (`=>`), which indicates the response.

To start an `irb` session, go to the command prompt and type `irb`. You should see the `irb` prompt waiting for your input:

```
$ irb
irb(main):001:0>
```

Look at that. You're inside Ruby! If you press Enter, Ruby ignores the line and gives you another prompt, but it ends with an asterisk instead of the greater-than sign to indicate that Ruby is expecting something from you to execute. It can only get more exciting from here.

Note Your `irb` prompt might look slightly different depending on your version of Ruby and your computer environment. This is perfectly okay.

When learning a new programming language, traditionally, the first thing you do is make the language print the string "Hello, World!" Let's go ahead and do that. Type the following after the `irb` prompt:

```
irb(main):001:0> "Hello, World!"  
=> "Hello, World!"
```

Excellent. You just wrote your first bit of Ruby! Some languages require many more lines of code to write the Hello, World! Application, but in Ruby it only took one. One of the ways Ruby makes developers happy is by being concise. One line is certainly concise, wouldn't you say?

So what exactly happened here? Well first, you created a string with the content "Hello, World!" The `irb` command always outputs the value of the last command to the screen, thus you have "Hello, World!" written to the screen. You will notice as you type valid Ruby commands and press Enter that `irb` will continue to output the value of those commands. Try adding two numbers together:

```
irb(main):001:0> 1 + 1  
=> 2
```

Now let's try something a little more difficult. Let's ask Ruby for the current time:

```
irb(main):001:0> Time.now  
=> 2013-04-02 14:14:55 -0500
```

So, Ruby dutifully reported the current time to us, including the date no less. What if you just wanted the current year?

```
irb(main):001:0> Time.now.year  
=> 2013
```

You can see how easy and concise Ruby is. The code is simple and almost reads like an English sentence. If you're wanting a description of exactly what you did in the last two examples here it is: You called a method (`now`) on a Class (`Time`). In the second example, you chained another method call onto the previous one. We'll cover this in depth later, but first let's talk about data types.

Ruby Data Types

A *data type* is a constraint placed on the interpretation of data. Numbers and strings are just two of the data types the Ruby interpreter distinguishes among, and the way Ruby adds numbers is different from the way in which it adds strings. For example, `2 + 3` evaluates to `5`, but, `"2" + "3"` evaluates to `"23"`. The second example may seem surprising at first, but it's simple: anything, including numbers, surrounded by quotes is interpreted as a string. Read on to find out more.

Strings

A *string* is a sequence of characters that usually represents a word or some other form of text. In Ruby, you can create *String* objects by putting the characters inside single or double quotation marks:

```
irb(main):001:0> 'Ruby is a great language'
=> "Ruby is a great language"

irb(main):002:0> "Rails is a great framework"
=> "Rails is a great framework"
```

The main difference between strings delimited by single and double quotes is that the latter are subject to substitutions. Those substitutions are identified by Ruby code inside the `#{} construct`, which is evaluated and replaced by its result in the final *String* object. The technical term for this technique is *string interpolation*.

```
irb(main):003:0> "Now is #{Time.now}"
=> "Now is 2013-04-02 20:54:31 +0000"

irb(main):004:0> 'Now is #{Time.now}'
=> "Now is \#{Time.now}"
```

Note In general, most developers only use double quotes when using string interpolation or if the actual string includes single quotes. This is technically faster, if only slightly.

When you use the hash symbol (#) with the curly braces, Ruby notices and tries to evaluate whatever is between the braces. To *evaluate* means to process it like any other code. So, inside the braces, you say `Time.now`, which returns the current time. However, when you use single quotes, Ruby doesn't check the string for substitutions before sending it through.

The *String* class has a large number of methods you need when doing string manipulation, like concatenation and case-changing operations. The following examples list a few of those methods:

```
irb(main):005:0> "Toronto - Canada".downcase
=> "toronto - canada"

irb(main):006:0> "New York, USA".upcase
=> "NEW YORK, USA"

irb(main):007:0> "a " + "few " + "strings " + "together"
=> "a few strings together"

irb(main):008:0> "HELLO".capitalize
=> "Hello"
```

Tip To get a list of methods available for any object, call the "methods" method using an instance of the object you want to inspect. Type "**a string".methods**" in irb to see all the methods you can call on the String object. If you want to find a certain method, try using grep on that method, too. For example, typing "**a string".methods.grep /case/**" shows all string methods containing the word *case*. Other examples would be:

4.methods

```
["some", "array", "elements"].methods
```

Numbers

Ruby has a few classes to represent numbers: Fixnum, Bignum, and Float. As the names of the classes suggest, Fixnum and Bignum represent whole numbers and are both subclasses of Integer. Float objects represent real numbers, meaning numbers with a fractional part. As in most programming languages, you can perform basic arithmetic operations in Ruby as you would using a calculator:

```
irb(main):001:0> 1 + 2
=> 3

irb(main):002:0> 2323 + 34545
=> 36868

irb(main):003:0> 9093 - 23236
=> -14143

irb(main):004:0> 343 / 4564
=> 0

irb(main):005:0> 3434 / 53
=> 64

irb(main):006:0> 99 * 345
=> 34155

irb(main):007:0> 34545.6 / 3434.1
=> 10.059578928977
```

Note Notice that when whole numbers are divided, the result is always a whole number even if there is a remainder. If one of the numbers is a decimal, then a decimal will always be returned.

```
irb(main):001:0> 6/4
=> 1

irb(main):002:0> 6/4.0
=> 1.5
```

Symbols

Symbols aren't a common feature in most languages. However, as you'll learn when reading this book, they're extremely useful. `Symbol` is a data type that starts with a colon, like `:controller`. Symbols are objects that work just like any other object in Ruby. They're used to point to some data that isn't a traditional `String` object, in a human-readable format. In fact, they're almost like strings, except you can't modify them:

```
irb(main):001:0> :my_symbol
=> :my_symbol

irb(main):002:0> :my_symbol + :second
NoMethodError: undefined method `+' for :my_symbol:Symbol
    from (irb):2
    from /usr/bin/irb:15:in `<main>'

irb(main):003:0> "my_string" + "second"
=> "my_stringsecond"
```

Fancy computer science types refer to this condition as being *immutable*, which really just means you can't modify something. Use symbols when you want to name something nicely and you don't want it changed at all—for example, by having something appended to the end of it. There are also memory advantages to using symbols, but that is out of the scope of this book. The importance of symbols will become clear as you use Ruby more.

Arrays and Hashes

Sometimes you have a lot of data that you need to keep track of—maybe a list of students, users, or anything that you may keep in a collection. Ruby has two different types of *container* objects for storing collections: arrays and hashes.

Arrays are part of almost every modern language. They keep information in order. You can ask for the first item or the last item or put items in a certain order. You can think of an `Array` object as a long series of boxes in which you can put things. You define arrays by using the `[]` notation. Note that in most programming languages, including Ruby, arrays are 0 indexed. This means you always refer to the first element in an array as 0. Read carefully what happens here:

```
irb(main):001:0> city_array = ['Toronto', 'Miami', 'Paris']
=> ["Toronto", "Miami", "Paris"]

irb(main):002:0> city_array[0]
=> "Toronto"

irb(main):003:0> city_array[1] = 'New York'
=> "New York"

irb(main):004:0> city_array << 'London'
=> ["Toronto", "New York", "Paris", "London"]

irb(main):004:0> city_array + ["Los Angeles"]
=> ["Toronto", "New York", "Paris", "London", "Los Angeles"]
```

In the first example we created the array of cities and assigned it to the variable named `city_array`. In the second example we referenced the city array and asked for the object at the index position 0 (remember, with arrays, 0 is the first index). "Toronto" is returned. In the third example, we are replacing the object at index 1 with the string "New York." Notice in the next example when the array is printed to the screen that Miami is no longer in the list but has

been replaced. The fourth example uses what is commonly called the *shovel* operator. Simply put, this just adds the object to the end of the array. So we added the string “London” to the end of our array. Finally, in the last array we add the array that contains “Los Angeles” to our previous array. This returns a new single dimensional array with the contents of both arrays. Arrays are extremely common and useful in Ruby.

The Hash object offers another way to keep a collection. Hashes are different from arrays, because they store items using a *key*. *Hash objects* preserve order, just like arrays, which enables you to call certain methods on them—for example, `hash.first` to get the first key, value pair. In Ruby, you often use symbols for hash keys, but in reality, *any* object can function as a key.

You define hashes with curly braces, `{}`. You can create a Hash object by defining it with `{:key => "value", :other_key => "other value"}`. Then, you can pull out data by using square brackets on the end of the list. For instance, you retrieve a value by typing `my_hash[:key]` from the `my_hash` variable. Here are some examples:

```
irb(main):005:0> my_hash = {:canada => 'Toronto', :france => 'Paris', :uk => 'London'}
=> {:canada=>"Toronto", :france=>"Paris", :uk=>"London"}
```

We've created a hash and assigned it to the “`my_hash`” variable. In this example, the keys of our array are countries, and the values are cities. To reference a specific value of a hash, you pass the hash a key and it will return the value to you:

```
irb(main):006:0> my_hash[:uk]
=> "London"
```

We've passed the hash a key of `:uk`, and it return the value of “London.”

```
irb(main):007:0> my_hash[:canada] = 'Calgary'
=> "Calgary"
```

This is the same idea, but here we're changing the value out for the key Canada. So the value of “Toronto” goes away and is replaced by “Calgary.”

```
irb(main):008:0> my_hash.first
=> [:canada, "Calgary"]
```

In this example, we use the `first` method, which returns the first key value pair. Notice in this case the return value is an array. The first element in the array is the key, and the second is the value. The `keys` method will return an array of all the keys contained in the hash. Here is an example:

```
irb(main):010:0> my_hash.keys
=> [:canada, :france, :uk]
```

It is important to note that in all of our examples, we have assigned strings to different positions to both our hashes and arrays, but any object could be stored in an array or hash. For instance, you might want to store numbers or even another array or hash. The possibilities are unlimited.

```
irb(main):001:0> numbers_array = [1, 2, 3, 4, 5]
=> [1,2,3,4,5]
```

```
irb(main):002:0> numbers_hash = {:one => 1, :two => 2, :three => 3}
=> {:one => 1, :two => 2, :three => 3}
```

Language Basics

Like other programming languages, Ruby includes variables, operators, control-flow statements, and methods. This section shows you how to use them.

Variables

Variables are used to hold values you want to keep for later processing. When you perform a calculation, you probably want to use the result of that calculation somewhere else in your application code, and that's when you need a variable. In Ruby, variables are easily created. You just need to give a variable a name and assign a value to it; there's no need to specify a data type for the variable or define it in your code before you use it.

Let's create a few variables to hold some values you may need later. Notice that you can reuse a variable name by reassigning a value:

```
irb(main):001:0> test_variable = 'This is a string'
=> "This is a string"

irb(main):002:0> test_variable = 2010
=> 2010

irb(main):003:0> test_variable = 232.3
=> 232.3
```

You've created a variable named `test_variable` and assigned a few different values to it. Because everything in Ruby is an object, the `test_variable` variable holds a reference to the object you assigned.

Variable names can be any sequence of numbers and letters, as long as they start with a letter or an underscore; however, the first character of a variable indicates the type of the variable. Variables also have a *scope*, which is the context in which the variable is defined. Some variables are used in a small snippet of code and need to exist for only a short period of time; those are called *local variables*. Table 4-1 lists the different types of variables supported by Ruby and shows how to recognize them when you're coding. Type some variable names in `irb`, and you'll get results similar to those shown here.

Table 4-1. Ruby Variables

Example	Description
<code>@@count</code>	<i>Class variables</i> start with <code>@@</code> . Class variables exist in the scope of a class, so all instances of a specific class have a single value for the class variable.
<code>@name</code>	<i>Instance variables</i> start with <code>@</code> . Instance variables are unique to a given instance of a class.
<code>SERVER_IP</code>	You can create a <i>constant</i> in Ruby by capitalizing the first letter of a variable, but it's a convention that constants are written in all uppercase characters. Constants are variables that don't change throughout the execution of a program. In Ruby, constants can be reassigned; however, you get a warning from the interpreter if you do so.
<code>my_string</code>	<i>Local variables</i> start with a lowercase letter, and they live for only a short period of time. They usually exist only inside the method or block of code where they're first assigned.

In Ruby, it's considered best practice to use long and descriptive variable names. For example, in Java, you may have a variable named phi; but in Ruby, you write out `place_holder_variable` for clarity. The basic idea is that code is much more readable if the person looking at it (probably you) doesn't have to guess what phi stands for. This is extremely important when you come back to a piece of code after a year or so.

Operators

You can combine Ruby code using operators. Many classes implement operators as methods. Table 4-2 lists the most common operators and their functions.

Table 4-2. Ruby Operators

Operator	Description
<code>[] []=</code>	Assignment
<code>* / % + **</code>	Arithmetic
<code><= >= < ></code>	Comparison
<code>... ...</code>	Range
<code>& ^ </code>	AND, exclusive OR, regular OR (bitwise)
<code> && not or and</code>	Logical operators

Ruby contains a ternary operator that you can use as a short notation for `if-else-end`. The ternary operator uses the form expression `? value_if_true : value_if_false`:

```
a = 10
b = 20
a > b ? a : b
# => 20
```

In plain English, we're saying if `a` is greater than `b` then return `a`, otherwise return `b`. The ternary operator is very concise but still easy to read.

Blocks and Iterators

Any method in Ruby can accept a *code block*—a fragment of code between curly braces or `do..end` constructs; it determines whether the method in question calls the given block. The block always appears immediately after the method call, with the start of the block coming on the same line as the method invocation.

Here's an example using the `times` method; `times` executes the given code block once for each iteration. In this case, "Hello" is printed five times:

```
5.times { puts "Hello" }
Hello
Hello
Hello
Hello
Hello
```

If a method yields arguments to a block, the arguments are named between two pipe characters (|) on the same line as the method call. In the next example, the block receives one argument, *item*.

```
[1,2,3,4,5].each { |item| puts item }
```

```
1
2
3
4
5
```

Here, each number is yielded to the block in succession. You store the number in the block variable *item* and use *puts* to print it on its own line.

The convention is to use braces for single-line blocks and *do..end* for multiline blocks. Here's an example similar to the previous one; it uses *each_with_index*, which yields the item and its index in the array:

```
["a", "b", "c"].each_with_index do |item, index|
  puts "Item: #{item}"
  puts "Index: #{index}"
  puts "---"
end
```

```
Item: a
Index: 0
---
Item: b
Index: 1
---
Item: c
Index: 2
---
```

Control Structures

In all of the previous examples, the Ruby interpreter executed the code from top to bottom. However, in the majority of cases, you want to control which methods are to be executed and when they should be executed. The statements you want to be executed may depend on many variables, such as the state of some computation or the user input. For that purpose, programming languages have *control-flow statements*, which allow you to execute code based on conditions. Here are a few examples of how to use *if*, *else*, *elsif*, *unless*, *while*, and *end*. Notice that control structures in Ruby are terminated using the *end* keyword:

```
now = Time.now
# => 2010-04-05 20:55:03 +0000

if now == Time.now
  puts "now is in the past"
elsif now > Time.now
  puts "nonsense"
else
  puts "time has passed"
end
# => time has passed
```

The first `if` statement will never trigger because there is a slight bit of time that passes between when you set the `now` variable and when you test it against `Time.now`. The second conditional won't trigger because the `now` variable will obviously be in the past, if only slightly. The third conditional "`else`" will always trigger, because neither of the first two conditionals triggered.

A trick that makes simple conditionals easy to read is to place `if` and `unless` conditional statements at the end of a code line so they act as *modifiers*. Here's how it looks:

```
a = 5
b = 10
puts "b is greater than a" if a < b
```

```
b is greater than a
```

```
puts "a is greater than b" unless a < b
```

```
nil
```

The `unless` structure was confusing for us at first. Once we started reading it as "if not" it made sense. In the previous example, reading the statement as "puts 'a is greater than b' if not `a < b`" makes the most sense.

You can also use `while` statements, as in all major programming languages:

```
a = 5
b = 10

while a < b
  puts "a is #{a}"
  a += 1
end
```

```
a is 5
a is 6
a is 7
a is 8
a is 9
```

Methods

Methods are little programmable actions that you can define to help your development. Let's leave `irb` for the moment and talk about pure Ruby code. (All of this also works if you type it into `irb`.)

Suppose that, several times in the application you're writing, you need to get the current time as a string. To save yourself from having to retype `Time.now.to_s` over and over, you can build a method. Every method starts with `def`:

```
def time_as_string
  Time.now.to_s
end
```

Anywhere in the application that you want to get the time, you type `time_as_string`:

```
puts time_as_string
```

```
"2010-04-05 20:55:21 +0000"
```

See how easy that is? Obviously with this code you didn't do much, but methods can be much more complex. Methods can also take in variables:

```
def say_hello_to(name)
  "Hello, #{name}!"
end
```

```
puts say_hello_to("John")
```

```
"Hello, John!"
```

Here you defined a method named `say_hello_to` that accepts one argument `name`. That method uses string interpolation to return a string of "Hello, *name that was passed to the method!*" The `puts` then sends the response of that method to the screen.

Next, let's look at how to put methods together into classes to make them really powerful.

Note You already know that local variables must start with a lowercase letter and can't contain any characters other than letters, numbers, and underscores. Method names are restricted to the same rules, which means they often look like variables. Keywords (like `if`, `or`, `when`, `and`, and others) share the same set of properties. How does the Ruby interpreter know the difference? When Ruby encounters a word, it sees it as a local variable name, a method invocation, or a keyword. If it's a keyword, then Ruby knows it and responds accordingly. If there's an equals sign (=) to the right of the word, Ruby assumes it's a local variable being assigned. If it's neither a keyword nor an assignment, Ruby assumes it's a method being invoked and sends the method to the implied receiver, `self`.

Classes and Objects

You've reviewed all the basic types of items in a Ruby application, so let's start using them.

Objects

Ruby is an *object-oriented* (OO) programming language. If you've never worked in an OO language before, the metaphors used can be confusing the first time you hear them. Basically, *objects* are simple ways to separate your code and the data it contains. Objects are just like objects in the real world. You can move them around, make them do things, destroy them, create them, and so forth. In OO programming, you act on objects by either calling methods on them or passing them to other methods. This will become clearer as we go along.

Let's say you're writing a program to help track the athletic program at a school. You have a list of all the students who are currently participating on a team, along with their student IDs. This example looks at the rowing team. You could keep an array of arrays representing the students on the team:

```
rowing_team = [[1975, "Smith", "John"], [1964, "Brown", "Dan"], ...]
```

Note This is called a *multidimensional array*. It's simply an array that contains more arrays as elements. You could reference the first array in the array like so:

```
rowing_team.first
=> [1975, "Smith", "John"]
```

This is an array of [id, first_name, last_name]. You'd probably need to add a comment to explain that. If you wanted multiple teams, you could wrap this in a hash:

```
teams = { :rowing => [[1975, "Smith", "John"], [1964, "Brown", "Dan"], ...],
          :track  => [[1975, "Smith", "John"], [1900, "Mark", "Twain"], ...]
    }
```

That works for now. But it's kind of ugly, and you could easily get confused, especially if you kept adding teams. This style of coding is referred to as *procedural*, and it's not object oriented. You're keeping track of huge data collections that are made up of simple types. Wouldn't it be nice to keep all these data more organized? You'll need to define what your objects will look like, so you'll need a pattern, called a *class*. Then you will instantiate your class to make an *instance*.

Classes

A *class* is like a blueprint for creating an object. You've been using classes all over the place—*Array*, *String*, *User*, and so on. Now, let's construct a *Student* class and a *Team* class.

Here is the basic blueprint for a *Student* class:

```
class Student
  # Setter method for @first_name
  def first_name=(value)
    @first_name = value
  end

  # Getter method for @first_name
  def first_name
    @first_name
  end

  # Setter method for @last_name
  def last_name=(value)
    @last_name = value
  end
```

```
# Getter method for @last_name
def last_name
  @last_name
end

# Returns full name
def full_name
  last_name + ", " + first_name
end
end
```

Note “Getter” and “Setter” methods are methods that get an instance variable or set an instance variable, respectively. It’s that simple. They are used to expose this functionality both inside and outside your instance. In this case, you have a getter and setter method for `last_name` and `first_name`.

Right now, you’re keeping track of the student’s `first_name` and `last_name` strings. As you can see, you define a method named `first_name=(value)`, and you take `value` and put it into an instance variable named `@first_name`. Let’s try using this class:

```
# Take the Class, and turn it into a real Object instance
@student = Student.new
@student.first_name = "Bob"
@student.last_name = "Jones"
puts @student.full_name
```

"Jones, Bob"

Instead of building a dumb array, you’ve built a smart class. When you call `new` on the class, it builds a version of itself called an *object*, which is then stored in the `@student` variable. In the next two lines, you use the `=` methods to store the student’s first and last names. Then, you use the method `full_name` to give a nicely formatted response.

It turns out that creating getter and setter methods like this is a common practice in OO programming. Fortunately, Ruby saves you the effort of creating them by providing a shortcut called `attr_accessor`:

```
class Student
  attr_accessor :first_name, :last_name, :id_number

  def full_name
    last_name + ", " + first_name
  end
end
```

This behaves in exactly the same way as the first version. The `attr_accessor` bit helps by automatically building the methods you need, such as `first_name=`. Also, this time you add an `@id_number`.

Let's build a Team class now:

```
class Team
  attr_accessor :name, :students

  def initialize(name)
    @name = name
    @students = []
  end

  def add_student(id_number, first_name, last_name)
    student = Student.new
    student.id_number = id_number
    student.first_name = first_name
    student.last_name = last_name
    @students << student
  end

  def print_students
    @students.each do |student|
      puts student.full_name
    end
  end
end
```

You've added something new to this class: the `initialize` method. Now, when you call `new`, you can pass in the name. For example, you can type `Team.new('baseball')`, and the `initialize` method is called. Not only does `initialize` set up the name, but it also sets up an instance variable named `@students` and turns it into an empty array. The method `add_student` fills the array with new `Student` objects.

Let's see how you use this class:

```
team = Team.new("Rowing")
team.add_student(1982, "John", "Smith")
team.add_student(1984, "Bob", "Jones")
team.print_students
```

```
Smith, John
Jones, Bob
```

Containing things in objects cleans up your code. By using classes, you ensure that each object only needs to worry about its own concerns. If you were writing this application without objects, everyone's business would be shared. The variables would all exist around one another, and there would be one *huge* object. Objects let you break things up into small working parts.

By now you should have a general idea of what's going on with some of the Ruby code you've seen floating around Rails. There is a *lot* more to Ruby that we haven't touched on here. Ruby has some amazing metaprogramming features you can read about in a book that specifically focuses on Ruby, such as *Beginning Ruby: From Novice to Professional*, second edition, by Peter Cooper (Apress, 2009).

RUBY STYLE

Style is important when you're programming. Ruby programmers tend to be picky about style, and they generally adhere to a few specific guidelines, summarized here:

- Indentation size is two spaces.
- Spaces are preferred to tabs.
- Variables should be lowercase and underscored: `some_variable`, not `someVariable` or `somevariable`.
- Method definitions should include parentheses and no unnecessary spaces: `MyClass.my_method(my_arg)`, not `my_method(my_arg)` or `my_method my_arg`.

Whatever your personal style, the most important thing is to remain consistent. Nothing is worse than looking at code that switches between tabs and spaces or mixed and lowercase variables.

Ruby Documentation

You can refer to the following documentation for more information about Ruby:

- *Core library*: The Ruby distribution comes with a set of classes known as the Ruby Core library, which includes base classes such as `Object`, `String`, `Array`, and others. In the Ruby Core application programming interface (API) documentation, you can find all the classes and methods included in the Core library. In this short chapter, you've already seen a few classes in action. One of the secrets to effectively using Ruby is to know which classes and methods are available to you. We recommend that you go to the Ruby Core API documentation page at www.ruby-doc.org/core/ and start to learn more about Ruby classes and methods.
- *Standard library*: In addition to the Core library, the Ruby distribution comes bundled with the Ruby Standard library. It includes a set of classes that extends the functionality of the Ruby language by helping developers perform common programming tasks, such as network programming and threading. Make sure you spend some time reading the Standard library documentation at www.ruby-doc.org/stdlib/.
- *Online resources*: The Ruby documentation project home page is located at www.ruby-doc.org. There you can find additional reading resources to help you learn Ruby, such as articles and tutorials, as well as the Core and Standard Ruby API documentation.

Summary

This chapter gave a strong introduction to the Ruby language. You now have the tools to start learning the Rails framework and start building web applications. As you progress, you'll more than likely come to love Ruby, especially if you have a background in other languages. Its power is only matched by its simplicity, and its genuinely fun to program with. The next chapter will dive into Active Record and learn how Rails lets you easily interact with your database.

Working with a Database: Active Record

Earlier, you took a whirlwind tour through creating a basic Rails application using the built-in scaffolding feature. You sketched out a basic model for a blog application and created the project databases. You used the built-in web server to run the application locally and practiced adding and managing articles from the web browser. This chapter will take a more in-depth look at how things work, starting with what is arguably the most important part of Rails: Active Record.

You may recall from Chapter 1 that Active Record is the Ruby object-relational mapping library that handles database abstraction and interaction for Rails. Whether you realized it or not, in Chapter 3 all access to the database—adding, editing, and deleting articles—happened through the magic of Active Record.

If you’re not sure what exactly object-relational mapping is, don’t worry. By the end of this chapter, you’ll know. For now, it’s best if you think of Active Record as being an intermediary that sits between your code and your database, allowing you to work with data effectively and naturally. When you use Active Record, you communicate with your database using pure Ruby code. Active Record translates the Ruby you write into a language that databases can understand.

This chapter teaches you how to use Active Record to talk to your database and perform basic operations. It introduces the concepts you need to know about communicating with databases and object-relational mapping. Then, you will look at Active Record and walk through the techniques you need to know to effectively work with a database from Rails. If you don’t have a lot of database experience under your belt, don’t worry. Working with databases through Active Record is a painless and even enjoyable experience. If you’re an experienced database guru, you’ll find that Active Record is an intelligent and efficient way to perform database operations without the need for low-level database-specific commands.

Note If you need to get the code at the exact point where you finished Chapter 3, download the zip file from GitHub (https://github.com/adamgamble/beginning_rails_4/archive/ch3.zip).

Introducing Active Record: Object-Relational Mapping on Rails

The key feature of Active Record is that it maps tables to classes, table rows to objects, and table columns to object attributes. This practice is commonly known as *object-relational mapping* (ORM). To be sure, Active Record isn’t the only ORM in existence, but it may well be the easiest to use of the bunch.

One of the reasons Active Record is so easy to use is that almost no configuration is required to have it map a table to a class. You just need to create a Ruby class that's named after the table you want to map and extend the Active Record Base class:

```
class Book < ActiveRecord::Base
end
```

Notice the part that reads `< ActiveRecord::Base`. The less-than sign indicates that the `Book` class on the left is a subclass of the one on the right, `ActiveRecord::Base`. In Ruby, when you extend a class like this, you automatically gain access to all the functionality in the parent class. There's a lot of code in the `ActiveRecord::Base` class, but you don't need to look at it. Your class merely inherits it, and your work is finished.

Assuming Active Record knows how to find your database and that you have a table called `books` (note that the table name is plural, whereas the class name is singular), the table is automatically mapped. If you know your `books` table contains the fields `title`, `publisher`, and `published_at`, you can do this in any Ruby context:

```
book = Book.new

book.title = "Beginning Rails 4"
book.publisher = "Apress"
book.published_at = "2013-10-21"

book.save
```

These five lines write a new record to the `books` table. You gain a lot of ability by the simple act of subclassing! And that's why Active Record is easy to use. Notice how the table's fields (`title`, `publisher`, and `published_at`) can be read and written to using methods on the object you created (`book`). And you didn't need to tell Active Record what your fields were named, or even that you had any fields. It figured this out on its own. Of course, Active Record doesn't just let you create new records. It can also read, update, and delete records, plus a lot more.

Active Record is database agnostic, so it doesn't care which database software you use, and it supports nearly every database out there. Because it's a high-level abstraction, the code you write remains the same no matter which database you're using. For the record (no pun intended), in this book you use SQLite. As explained in Chapter 2, SQLite is open source, easy to use, and fast, and it's the default database used for Rails development. (Along with the SQLite site, <http://sqlite.org>, the Wikipedia entry on SQLite is a good resource: <http://en.wikipedia.org/wiki/SQLite>.)

Note Rails is also ORM agnostic: it allows you to hook up your ORM of choice. The main ORM competitor for Active Record is DataMapper (<http://datamapper.org>), which you can use if you think Active Record has some deficiencies; however, we feel that sticking to the default ORM is the best way to learn. We don't cover DataMapper in this book.

What About SQL?

To be sure, you don't need Active Record (or any ORM) to talk to and manipulate your database. Databases have their own language: SQL, which is supported by nearly every relational database in existence. Using SQL, you can view column information, fetch a particular row or a set of rows, and search for rows containing certain criteria. You can also use SQL to create, drop, and modify tables and insert, update, and destroy the information stored in those tables. The problem with SQL is that it's not object oriented. If you want to learn the basic SQL syntax, look at Appendix B.

Object-oriented programming and relational databases are fundamentally different paradigms. The relational paradigm deals with relations and is mathematical by nature. The object-oriented paradigm, however, deals with objects, their attributes, and their associations to one another. As soon as you want to make objects persistent using a relational database, you notice something: there is a rift between these two paradigms—the so-called *object-relational gap*. An ORM library like Active Record helps you bridge that gap.

Note Active Record is based on a design pattern. *Design patterns* are standard solutions to common problems in software design. Well, it turns out that when you’re working in an object-oriented environment, the problem of how to effectively communicate with a database (which isn’t object oriented) is quite common. Therefore, many smart people have wrapped their minds around the problem of how best to bring the object-oriented paradigm together with the relational database. One of those smart people is Martin Fowler, who, in his book *Patterns of Enterprise Application Architecture* (Addison-Wesley, 2002), first described a pattern that he called an Active Record. In the pattern Fowler described, a one-to-one mapping exists between a database record and the object that represents it. When Rails creator David Heinemeier Hansson sought to implement an ORM for his framework, he based it on Fowler’s pattern.

Active Record lets you model real-world things in your code. Rails calls these real-world things *models*—the *M* in MVC. A model might be named Person, Product, or Article, and it has a corresponding table in the database: people, products, or articles. Each model is implemented as a Ruby class and is stored in the app/models directory. Active Record provides the link between these classes and your tables, allowing you to work with what look like regular objects, which, in turn, can be persisted to the database. This frees you from having to write low-level SQL to talk to the database. Instead, you work with your data as if they were an object, and Active Record does all the translation into SQL behind the scenes. This means that in Rails, you get to stick with one language: Ruby.

Note Just because you’re using Active Record to abstract your SQL generation doesn’t mean SQL is evil. Active Record makes it possible to execute SQL directly whenever that’s necessary. The truth is that raw SQL is the native language of databases, and there are some (albeit rare) cases when an ORM won’t cut it.

Active Record Conventions

Active Record achieves its zero-configuration reputation by way of convention. Most of the conventions it uses are easy to grasp. After all, they’re conventions, so they’re already in wide use. Although you can override most of the conventions to suit the particular design of your database, you’ll save a lot of time and energy if you stick to them.

Let’s take a quick look at the two main conventions you need to know:

- Class names are singular; table names are plural.
- Tables contain an identity column named `id`.

Active Record assumes that the name of your table is the plural form of the class name. If your table name contains underscores, then your class name is assumed to be in *CamelCase*. Table 5-1 shows some examples.

Table 5-1. Table and Class Name Conventions

Table	Class
events	Event
people	Person
categories	Category
order_items	OrderItem

All tables are assumed to have a unique identity column named `id`. This column should be the table's *primary key* (a value used to uniquely identify a table's row). This is a fairly common convention in database design. (For more information on primary keys in database design, the Wikipedia entry has a wealth of useful information and links: http://en.wikipedia.org/wiki/Unique_key.)

The belief in convention over configuration is firmly entrenched in the Rails philosophy, so it should come as no surprise that there are more conventions at work than those listed here. You'll likely find that they all make good sense, and you can use them without paying much attention.

Introducing the Console

Ruby comes with a great little tool: an interactive interpreter called `irb` (for Interactive Ruby). Most of the time, you invoke `irb` using the console program that ships with Rails, but you can start up an `irb` session whenever you want by typing `irb` at the command prompt. The advantage of the console is that it enjoys the special privilege of being integrated with your project's environment. This means it has access to and knowledge of your models (and subsequently, your database).

You use the console as a means to get inside the world of your `Article` model and to work with it in the exact same way your Rails application would. As you'll see in a minute, this is a great way to showcase the capabilities of Active Record interactively.

You can execute any arbitrary Ruby code in `irb` and do anything you might otherwise do inside your Ruby programs: set variables, evaluate conditions, and inspect objects. The only essential difference between an interactive session and a regular old Ruby program is that `irb` echoes the return value of everything it executes. This saves you from having to explicitly print the results of an evaluation. Just run the code, and `irb` prints the result.

You can tell whenever you're inside an `irb` session by looking for the double greater-than sign (`>>`)—or a slightly different sign depending on your environment—which indicate the `irb` prompt, and the arrow symbol (`=>`), which indicates the response.

As you continue to progress with both Ruby and Rails, you'll find that `irb` is an essential tool. Using `irb`, you can play around with code and make sure it works as you expect before you write it into your programs.

If you've been following along with the previous chapters, then you should have a model called `Article` (in `app/models/article.rb`), and you've probably already entered some sample data when playing with scaffolding in Chapter 3. If not, make sure you get up to speed by reading Chapters 2 and 3 before moving on.

Let's load `irb` and start to experiment with the `Article` model. Make sure you're inside the blog application directory, and then type `rails console` on your command line. This causes the `irb` console to load with your application's *development* environment and leaves you at a simple prompt, waiting for you to enter some code:

```
$ rails console
Loading development environment.
>>
```

From the console, you can interrogate your Article model for information. For instance, you can ask it for its column names:

```
>> Article.column_names
=> ["id", "title", "body", "published_at", "created_at", "updated_at",
"excerpt", "location"]
```

Look at that! All your columns are presented as a Ruby array (you can tell by the fact that they're surrounded by square brackets). Another quick trick you may use often is to type just the name of your model class in the console to find out not only the column names, but also the data type of each column:

```
>> Article
=> Article(id: integer, title: string, body: text, published_at: datetime,
created_at: datetime, updated_at: datetime, excerpt: string, location: string)
```

You get the `column_names` class method courtesy of the `ActiveRecord::Base` class from which your `Article` class inherits. Actually, you get a lot of methods courtesy of `ActiveRecord::Base`. To see just how many, you can ask:

```
>> Article.methods.size
=> 549
```

Note Depending on the version of Rails you are using and what gems you have installed, the number of methods might be different from 549. This is normal.

That's a lot of methods! You may get a different number of methods depending on your environment. Don't worry—you don't need to memorize all of them. Most of them are used internally so you'll never have to use them directly. Still, it's important, if for no other reason than to get a sense of what you get for free just by subclassing Active Record. Although in this case `ActiveRecord::Base` is considered the *superclass*, it sure makes your lowly `Article` class super, doesn't it? (Sorry, enough bad humor.)

A CRASH COURSE IN RUBY CLASS DESIGN

Object-oriented programming is all about objects. You create a class that encapsulates all the logic required to create an object, along with its properties and attributes, and use the class to produce new objects, each of which is a unique instance, distinct from other objects of the same class. That may sound a little abstract (and with good reason—abstraction, after all, is the name of the game), but if it helps, you can think of a class as being an object factory.

The obvious example is that of a car factory. Contained within a car factory are all the resources, tools, workers, and processes required to produce a shiny new car. Each car that comes off the assembly line is unique. The cars may vary in size, color, and shape, or they may not vary from one another much at all. The point is that even if two cars share the exact same attributes, they aren't the same car. You certainly wouldn't expect a change to the color of one car to affect all the others, would you? Well, in object-oriented programming, it's not much different. The

class is the factory that produces objects, which are called *instances* of a class. From a single factory, an infinite number of objects can be produced:

```
class Car
end

car1 = Car.new
car2 = Car.new
```

`car1` is a `Car` object, which is to say it's an instance of the class `Car`. Each car is a different object, created by the same factory. Each object knows which class it belongs to (which factory created it), so if you're ever in doubt, you can ask it:

```
car2.class #=> Car
```

Your `Car` class doesn't really do anything that useful—it has no attributes. So, let's give it some. You start by giving it a `make`—something like `Toyota` or `Nissan`. Of course, you need to define a way to read and write these attributes. You do this by creating aptly named *reader* and *writer* methods. Some object-oriented languages refer to these as *getters* and *setters*. The two sets of terms are pretty much interchangeable, but Ruby favors the former. Let's add a reader and writer for the `make` attribute:

```
class Car
  # A writer method. Sets the value of the @make attribute
  def make=(text)
    @make = text
  end

  # A reader method. Returns the value of the @make attribute
  def make
    @make
  end
end
```

The methods you just defined (`make()` and `make=()`) are instance methods. This is because they can be used only on instances of the class, which is to say, the individual objects that have been created from the class. To create a new instance of the `Car` class, you use the `new` constructor:

```
my_car = Car.new
```

That's all that's required to create a new instance of the class `Car` in a local variable called `my_car`. The variable `my_car` can now be considered a `Car` object. Although you have a new `Car` object, you haven't yet given it a `make`. If you use the reader method you created to ask your car what its `make` is, you see that it's `nil`:

```
my_car.make #=> nil
```

Apparently, if you want your car to have a `make`, you have to set it. This is where the writer method comes in handy:

```
my_car.make = 'Toyota'
```

This sets the value of the `make` attribute for your car to Toyota. If you had other `Car` objects, their `makes` would remain unchanged. You're setting the attribute only on the `my_car` object. Now, when you use the reader method, it confirms that the `make` attribute has been updated:

```
my_car.make #=> 'Toyota'
```

Of course, you can change the value any time you want:

```
my_car.make = 'Mazda'
```

And again, if you ask your `Car` object its `make`, it will tell you:

```
my_car.make #=> 'Mazda'
```

That's a simple example, but it illustrates a couple of very important points: Classes are used to create objects, and objects have attributes. Every object has a unique set of attributes, different from other objects of the same class.

The reason for this crash course in Ruby class design is to illustrate the point that modeling with Active Record is a lot like modeling with standard Ruby classes. If you decided to think of Active Record as being an extension to standard Ruby classes, you wouldn't be very far off. In practice, this fact makes using Active Record in Ruby quite natural. And because Active Record can reflect on your tables to determine which fields to map automatically, you need to define your attributes in only one place: the database. That's DRY (don't repeat yourself)! See Chapter 4 to learn more about Ruby's syntax, classes, and objects.

Active Record Basics: CRUD

Active Record is a big topic, so let's start with the basics. You've seen the so-called big four earlier, but here they are again: create, read, update, and delete, affectionately known as CRUD. In one way or another, most of what you do with Active Record in particular, and with databases in general, relates to CRUD. Rails has embraced CRUD as a design technique and as a way to simplify the modeling process. It's no surprise then that this chapter takes an in-depth look at how to do CRUD with Active Record.

Let's build on the blog application you started in Chapter 3. Although your application doesn't do much yet, it's at a stage where it's easy to demonstrate these concepts in a more concrete fashion.

This section uses the console, so keep it open as you work, and feel free to experiment as much as you want. The more experimentation you do, the deeper your understanding will be.

Creating New Records

You start by creating a new article in the database so you have something to work with. There are a few different ways to create new model objects, but they're all variations on the same theme. This section shows how each approach works and explains the often subtle differences among them.

Using the *new* Constructor

The most basic way to create a new model object is with the `new` constructor. If you read the crash-course section on Ruby classes earlier, you’re sure to recognize it. If you didn’t, then it’s enough that you know `new` is the usual way to create new objects of any type. Active Record classes are no different. Try it now:

```
>> article = Article.new
=> #<Article id: nil, title: nil, body: nil, published_at: nil, created_at: nil,
updated_at: nil, excerpt: nil, location: nil>
```

All you’re doing here is creating a new `Article` object and storing it in the local variable `article`. True to form, the console responds with the return value of the method, which in this case is a *string* representation of the model object. It may look a little funny, but this is what all Ruby objects look like when you inspect them. The response lists the attributes of the `Article` class. Starting here, you could call a few of the `article` variable methods. For example, the `new_record?` method tells you whether this object has been persisted (saved) to the database, and the `attributes` method returns a hash of the attributes that Active Record garnered by reflecting on the columns in the table. Each key of the hash will be the name of a column (`body`, `created_at`, etc.).

```
>> article.new_record?
=> true
>> article.attributes
=> {"body"=>nil, "created_at"=>nil, "excerpt"=>nil, "location"=>nil,
"published_at"=>nil, "title"=>nil, "updated_at"=>nil}
```

Here, you’re using *reader* methods, which read and return the value of the attribute in question. Because this is a brand-new record and you haven’t given it any information, all your attributes are `nil`, which means they have no values. Let’s remedy that now using (what else?) *writer* methods:

```
>> article.title = 'RailsConf'
=> "RailsConf"

>> article.body = 'RailsConf is the official gathering for Rails developers..'
=> "'RailsConf is the official gathering for Rails developers..'

>> article.published_at = '2013-04-13'
=> "2013-04-13"
```

Note A return of `nil` always represents nothing. It’s a helpful little object that stands in the place of nothingness. If you ask an object for something and it returns `false`, then `false` is *something*, so it’s not a helpful representation. As a nerdy fact, in logics, `false` and `true` are equal and opposite values, but they’re values in the end. The same is true of zero (0). The number 0 isn’t truly nothing—it’s an actual representation of an abstract nothing, but it’s still something. That’s why in programming you have `nil` (or `null` in other languages).

Now, when you inspect your Article object, you can see that it has attributes:

```
>> article
=> #<Article id: nil, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2013-04-13 00:00:00",
created_at: nil, updated_at: nil, excerpt: nil, location: nil>
```

You still haven't written a new record. If you were to look at the articles table in the database, you wouldn't find a record for the object you're working with. (If you squint really hard at the preceding object-inspection string, notice that no id has been assigned yet.) That's because you haven't yet saved the object to the database. Fortunately, saving an Active Record object couldn't be any easier:

```
>> article.save
(0.1ms) begin transaction
SQL (2.2ms) INSERT INTO "articles" ("body", "created_at", "published_at",
"title", "updated_at") VALUES (?, ?, ?, ?, ?) [[{"body": "RailsConf is the official
gathering for Rails developers..."}, {"created_at": Sat, 13 Apr 2013 15:50:29 UTC
+00:00}, {"published_at": Wed, 13 Apr 2013 00:00:00 UTC +00:00}, {"title": "RailsConf"}, {"updated_at": Sat, 13 Apr 2013 15:50:29 UTC +00:00}]]
(2.9ms) commit transaction
=> true
```

When you save a new record, an SQL INSERT statement is constructed behind the scenes; notice that Rails has displayed the generated SQL for you. If the INSERT is successful, the save operation returns true; if it fails, save returns false. You can ask for a count of the number of rows in the table just to be sure that a record was created:

```
>> Article.count
=> 1
```

Sure enough, you have a new article! You've got to admit, that was pretty easy. (You may have created some articles during the scaffolding session. If so, don't be surprised if you have more than one article already.) Additionally, if you ask the article whether it's a new_record?, it responds with false. Because it's saved, it's not "new" anymore:

```
>> article.new_record?
=> false
```

Let's create another article. This time, omit all the chatter from the console so you can get a better sense of how the process plays out. You create a new object and place it in a variable, you set the object's attributes, and finally you save the record. Note that although you're using the local variable article to hold your object, it can be anything you want. Usually, you use a variable that indicates the type of object you're creating, like article or, if you prefer shorthand, just a:

```
>> article = Article.new
>> article.title      = "Introduction to SQL"
>> article.body       = "SQL stands for Structured Query Language, ..."
>> article.published_at = Date.today
>> article.save
```

Note Although writer methods look like assignments, they're really methods in disguise. `article.title = 'something'` is the functional equivalent of `article.title=('something')`, where `title=()` is the method. Ruby provides a little syntactic sugar to make writers look more natural.

Now you're rolling! You've already created a few articles and haven't had to write a lick of SQL. Given how easy this is, you may be surprised that you can do it in even fewer steps, but you can. Instead of setting each attribute on its own line, you can pass all of them to `new` at once. Here's how you can rewrite the preceding process of creating a new record in fewer lines of code:

```
>> article = Article.new(:title => "Introduction to Active Record",
:body => "Active Record is Rails's default ORM..", :published_at => Date.today)
>> article.save
```

Not bad, but you can do even better. The `new` constructor creates a new object, but it's your responsibility to save it. If you forget to save the object, it will never be written to the database. There is another method available that combines the creating and saving steps into one.

Using the *create* Method

When you want to create an object and save it in one fell swoop, you can use the `create` method. Use it now to create another article:

```
>> Article.create(:title => "RubyConf 2013", :body => "The annual RubyConf will
take place in..", :published_at => '2013-04-13')
=> #<Article id: 4, title: "RubyConf 2013", body: "The annual RubyConf will take
place in..", published_at: "2013-04-13 00:00:00", created_at: "2013-04-13
23:17:19", updated_at: "2013-04-13 23:17:19", excerpt: nil, location: nil>
```

Instead of returning `true` or `false`, the `create` method returns the object it created—in this case, an `Article` object. You're actually passing a hash of attributes to the `create` method. Although hashes are normally surrounded by curly braces, *when a hash is the last argument to a Ruby method, the braces are optional*. You can just as easily create the attribute's hash first and then give that to `create`:

```
>> attributes = { :title => "Rails Pub Nite", :body     => "Rails Pub Nite is every
3rd Monday of each month, except in December.", :published_at => "2013-04-13"}
=> {:title=>"Rails Pub Nite", :body=>"Rails Pub Nite is every
3rd Monday of each month, except in December.", :published_at=>" 2013-04-13"}
>> Article.create(attributes)

=> #<Article id: 5, title: "Rails Pub Nite", body: "Rails Pub Nite is every 3rd
Monday of each month, e...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:36:07", updated_at: "2013-04-13 23:36:07",
excerpt: nil, location: nil>
```

Let's see how many articles you've created by doing a count:

```
>> Article.count
=> 5
```

You're getting the hang of this now. To summarize, when you want to create a new object and save it manually, use the `new` constructor; when you want to create and save in one operation, use `create`. You've already created five new records, which are plenty for now, so let's move on to the next step: finding records.

Reading (Finding) Records

Now that you have a few articles to play with, it's time to practice finding them. Every model class understands the `find` method. It's quite versatile and accepts a number of options that modify its behavior.

Let's start with the basics. `find` is a class method. That means you use it on the model class rather than an object of that class, just as you did the `new` and `create` methods. Like `new` and `create`, a `find` operation, if successful, returns a new object.

You can call `find` four ways:

- `find(:id)`: Finds a single record by its unique `id`, or multiple records if `:id` is an array of `ids`
- `all`: Finds all records in the table
- `first`: Finds the first record
- `last`: Finds the last record

The following sections go through the different ways to call `find` and explain how to use each.

Finding a Single Record Using an ID

The `find`, `first`, and `last` methods mostly return a single record. The `:id` option is specific; you use it when you're looking for a specific record and you know its unique `id`. If you give it a single `id`, it either returns the corresponding record (if there is one) or raises an exception (if there isn't one). If you pass an array of `ids`—like `[4, 5]`—as the parameter, the method returns an array with all records that match the passed in `ids`. The `first` method is a little more forgiving; it returns the first record in the table or `nil` if the table is empty, as explained in the next section.

You can find a single record using its unique `id` by using `find(:id)`. Here's how it works:

```
>> Article.find(3)
=> #<Article id: 3, title: "Introduction to Active Record", body: "Active Record
is Rails's default ORM..", published_at: "2013-04-13 04:00:00",
created_at: "2013-04-13 23:15:37", updated_at: "2013-04-13 23:15:37",
excerpt: nil, location: nil>
```

As you can see, you found the article with the `id` of 3. If you wanted to take a closer look at what was returned, you can store the result in a local variable:

```
>> article = Article.find(3)
=> #<Article id: 3 ...
>> article.id
=> 3
>> article.title
=> "Introduction to Active Record"
```

Here, you store the object that `find` returned in the local variable `article`. Then, you can interrogate it and ask for its attributes.

All this works because an article with the `id` 3 *actually exists*. If instead you search for a record that you know doesn't exist (say, 1037), Active Record raises an exception:

```
>> Article.find 1037
ActiveRecord::RecordNotFound: Couldn't find Article with ID=1037
...
```

Active Record raises a `RecordNotFound` exception and tells you it couldn't find any articles with the `id` of 1037. Of course it couldn't. You know that no such record exists. The lesson here is that you use `find(:id)` when you're looking for a specific record that you expect to exist. If the record doesn't exist, it's probably an error you want to know about; therefore, Active Record raises `RecordNotFound`.

RECOVERING FROM RECORDNOTFOUND ERRORS

When you use `find` with a single `id`, you expect the record to exist. Usually we don't want to display Rails error messages directly to the user, but we can make them nicer and customize the verbiage. So how can you recover gracefully from a `RecordNotFound` exception if you need to? You can use Ruby's facility for error handling: `begin` and `rescue`. Here's how this works:

```
begin
  Article.find(1037)
rescue ActiveRecord::RecordNotFound
  puts "We couldn't find that record"
end
```

First, you open a `begin` block. Then, you cause a `RecordNotFound` error by deliberately searching for a record that you know doesn't exist. When the error occurs, Ruby runs the code you put inside the `rescue` part of the body, which prints a friendly message.

You can put anything you like in the `rescue` block—you might want to render a specific view here, log the error, or even redirect to another location. Error handling works the same way with other error messages also. If you need to rescue from any error at all you can just use `rescue` without specifying an error class.

Finding a Single Record Using `first`

You can find the first record that the database returns by using the `first` method. This always returns exactly one item, unless the table is empty, in which case `nil` is returned:

```
>> Article.first
=> #<Article id: 1, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:12:09", updated_at: "2010-04-13 23:12:09",
excerpt: nil, location: nil>
```

Keep in mind that this isn't necessarily the first record in the table. It depends on the database software you're using and the default order in which you want your records to be retrieved. Usually records are ordered by either `created_at` or `updated_at`. If you need to be sure you get the first record you're expecting, you should specify an order. It's the equivalent of saying `SELECT * FROM table LIMIT 1` in SQL. If you need to find a record and don't

particularly care which record it is, `first` can come in handy. Note that `first` doesn't raise an exception if the record can't be found.

The `last` method works exactly the same as `first`; however, records are retrieved in the inverse order of `first`. For example, if records from articles are listed in chronological order for `first`, they're retrieved in inverse chronological order for `last`:

```
>> Article.last
=> #<Article id: 5, title: "Rails Pub Nite", body: "Rails Pub Nite is every 3rd
Monday of each month, e...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:36:07", updated_at: "2013-04-13 23:36:07",
excerpt: nil, location: nil>
```

Finding All Records

So far, you've looked at finding a single record. In each case, `find`, `first`, or `last` returns a single `Article` object. But what if you want to find more than one article? In your application, you want to display all the articles on the home page.

If you run the `all` method, it returns all records for that class:

```
>> articles = Article.all
=> [#<Article id: 1,..>, #<Article id: 2,..>, #<Article id: 3,..>,
#<Article id: 4,..>, #<Article id: 5,..>]
```

The square brackets in the response indicate that `all` has returned an array. You can confirm this by asking the `articles` variable what its class is:

```
>> articles.class
=> Array
```

Sure enough, `articles` tells you it's an `Array`. To be precise, it's an array of `Article` objects. Like all Ruby arrays, you can ask for its `size`:

```
>> articles.size
=> 5
```

Because `articles` is an array, you can access the individual elements it contains by using its `index`, which is numeric, starting at 0:

```
>> articles[0]
=> #<Article id: 1, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:12:09", updated_at: "2013-04-13 23:12:09",
excerpt: nil, location: nil>
```

And after you've isolated a single `Article` object, you can find its attributes:

```
>> articles[0].title
=> "RailsConf"
```

What's happening here is that `all` produces an array, and you access the object at the 0 index and call the `title` method. You can also use the `first` method, which all arrays respond to, and get the same result, but with a little more natural syntax:

```
>> articles.first.title
=> "RailsConf"
```

If you want to iterate over the collection, you can use the `each` method, which, again, works with all arrays. Here, you loop over the array, extract each item into a variable called `article`, and print its `title` attribute using the `puts` command:

```
>> articles.each { |article| puts article.title }
RailsConf
Introduction to SQL
Introduction to Active Record
RubyConf 2010
Rails Pub Nite
=> [<Article id: 1,...> #<Article id: 2,...>, #<Article id: 3,...>,
#<Article id: 4,...> , #<Article id: 5,...>]
```

Sometimes you want your results ordered. For example, if you're listing all your articles, you probably want them listed chronologically. To do so, you can use the `order` method, which accepts as argument the name of the column or columns. For you SQL heroes, it corresponds to the SQL ORDER clause:

```
>> articles = Article.order("published_at")
=> [<Article id: 1,...> #<Article id: 2,...>, #<Article id: 3,...>,
#<Article id: 4,...> , #<Article id: 5,...>]
>> articles.each { |article| puts article.published_at }2013-04-13 00:00:00 UTC2013-04-13 04:00:00
UTC2013-04-13 04:00:00 UTC2013-04-13 00:00:00 UTC2013-04-13 00:00:00 UTC
=> [<Article id: 1,...> #<Article id: 2,...>, #<Article id: 3,...>,
#<Article id: 4,...> , #<Article id: 5,...>]
```

Notice that when you call the `order` method, it returns an array object, as you may have expected. One thing that happens on the background is that Active Record allows you to chain calls to multiple methods before sending the command to the database; so you can call `all`, followed by `order`, and some other methods we'll talk about in Chapter 6, to create more precise database queries. Also, Active Record is smart enough to use *lazy loading*, a practice that only hits the database when necessary—in this example, when you call the `each` method.

By default, any column is ordered in ascending order (for example, 1–10, or *a-z*). If you want to reverse this to get descending order, use the `DESC` modifier (the same way you do in SQL, because the value of the `order` parameter is really just an SQL fragment):

```
>> articles = Article.order ('published_at DESC')
=> [<Article id: 4,...> #<Article id: 5,...>, #<Article id: 2,...>,
#<Article id: 3,...> , #<Article id: 1,...>]
>> articles.each { |article| puts article.published_at }2013-04-13 00:00:00 UTC2013-04-13 00:00:00
UTC2013-04-13 00:00:00 UTC2013-04-13 00:00:00 UTC2013-04-13 00:00:00 UTC
=> [<Article id: 4,...> #<Article id: 5,...>, #<Article id: 2,...>,
#<Article id: 3,...> , #<Article id: 1,...>]
```

Finding with Conditions

Although finding a record by its primary key is useful, it requires that you know the `id` to begin with, which isn't always the case. Sometimes you want to find records based on other criteria. This is where conditions come into play. Conditions correspond to the SQL `WHERE` clause. If you want to find a record by its title, you call the `where` method and pass a value that contains either a hash of conditions or an SQL fragment.

Here, you use a hash of conditions to indicate you want the first article with the title `RailsConf`:

```
>> Article.where(:title => 'RailsConf').first
=> #<Article id: 1, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:12:09", updated_at: "2013-04-13 23:12:09",
excerpt: nil, location: nil>
```

Because you use `first`, you get only one record (the first one in the result set, even if there is more than one result). If you instead use `all`, you get back a collection, even if the collection has only one item in it:

```
>> Article.where(:title => 'RailsConf').all
=> [#<Article id: 1, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:12:09", updated_at: "2013-04-13 23:12:09",
excerpt: nil, location: nil>]
```

Notice the square brackets, and remember that they indicate an array. More often than not, when you're doing an `all` operation, you expect more than one record in return. But `all` always produces an array, even if that array is empty:

```
>> Article.where(:title => 'Unknown').all
=> []
```

Updating Records

Updating a record is a lot like creating a record. You can update attributes one at a time and then save the result, or you can update attributes in one fell swoop. When you update a record, an SQL `UPDATE` statement is constructed behind the scenes. First, you use a `find` operation to retrieve the record you want to update; next, you modify its attributes; and finally, you save it back to the database:

```
>> article = Article.first
>> article.title = "Rails 4 is great"
>> article.published_at = Time.now
>> article.save
=> true
```

This should look pretty familiar by now. The only real difference between this process and the process of creating a new record is that instead of creating a brand-new row, you fetch an existing row. You update the attributes the exact same way, and you save the record the same way. Just as when you create a new record, when `save` operates on an existing record, it returns `true` or `false`, depending on whether the operation was successful.

When you want to update an object's attributes and save it in a single operation, you use the `update_attributes` method. Unlike when you create a new record with `create`, because you're updating a record, you need to fetch that record first. That's where the other subtle difference lies. Unlike `create`, which is a class method (it operates on the

class, not on an object), `update_attributes` is an instance method. Instance methods work on objects or instances of a class. Here's an example:

```
>> article = Article.first
>> article.update_attributes(:title => "RailsConf2013", :published_at => 1.day.ago)
=> true
```

Deleting Records

You're finally at the last component of CRUD: delete. When you're working with databases, you inevitably need to delete records. If a user cancels their order, or if a book goes out of stock, or even if you have an error in a given row, you may want to delete it. Sometimes you need to delete all rows in a table, and sometimes you want to delete only a specific row. Active Record makes deleting rows every bit as easy as creating them.

There are two styles of row deletion: `destroy` and `delete`. The `destroy` style works on the *instance*. It instantiates the object, which means it finds a single row first, and then deletes the row from the database. The `delete` style operates on the *class*, which is to say it operates on the table rather than a given row from that table.

Using `destroy`

The easiest and most common way to remove a record is to use the `destroy` method, which means the first thing you need to do is find the record you want to destroy:

```
>> article = Article.last
>> article.destroy
=> #<Article id: 5, title: "Rails Pub Nite", body: "Rails Pub Nite is every 3rd
Monday of each month, e...", published_at: "2013-04-13 00:00:00",
created_at: "2013-04-13 23:36:07", updated_at: "2013-04-13 23:36:07",
excerpt: nil, location: nil>
```

If you're interested, the SQL that Active Record generates in response to the `destroy` operation is as follows:

```
DELETE FROM articles WHERE id = 5;
```

As a result, the article with the `id` of 5 is permanently deleted. But you still have the object hanging around in the variable `article`, so how can it really be gone? The answer is that although the object remains *hydrated* (retains all its attributes), it's *frozen*. You can still access its attributes, but you can't modify them. Let's see what happens if you try to change the location:

```
>> article.location = 'Toronto, ON'
RuntimeError: can't modify frozen hash
```

It appears that the deleted article is now frozen. The object remains, but it's read-only, so you can't modify it. Given this fact, if you're going to delete the record, you don't really need to create an explicit `Article` object after all. You can do the `destroy` in a one-line operation:

```
>> Article.last.destroy
```

Here, the object instantiation is implicit. You're still calling the `destroy` instance method, but you're not storing an `Article` object in a local variable first.

You can still do better. You can use the class method `destroy`, which does a `find` automatically. As with `find` and `create`, you can use `destroy` directly on the class (that is, you don't create an object first). Because it operates on the table and not the row, you need to help it by telling it which row or rows you want to target. Here's how you delete the article with the id 1:

```
>> Article.destroy(1)
=> [#<Article id: 1, title: "RailsConf", body: "RailsConf is the official
gathering for Rails devel...", published_at: "2010-02-27 00:00:00",
created_at: "2010-05-01 23:12:09", updated_at: "2010-05-01 23:12:09",
excerpt: nil, location: nil>]
```

Sometimes you want to destroy more than one record. Just as with `find`, you can give `destroy` an array of primary keys whose rows you want to remove. Use square brackets ([]) to indicate that you're passing an array:

```
>> Article.destroy([2,3])
=> [#<Article id: 2, ...>, #<Article id: 3, ...>]
```

Although `ActiveRecord::Base.destroy` is a class method, it instantiates each object before destroying it. You can tell this by looking at its source:

```
def destroy(id)
  if id.is_a?(Array)
    id.map { |one_id| destroy(one_id) }
  else
    find(id).destroy
  end
end
```

Note The documentation for the classes and methods rails provides can be viewed at <http://api.rubyonrails.org>. In this case, this specific method is located at <http://api.rubyonrails.org/classes/ActiveRecord/Relation.html#method-i-destroy>.

Here you can see that if the received argument is an array, `destroy` iterates over the array and calls the same `destroy` method once for each item in the array. This effectively causes it to take the `else` path of the conditional, which performs a `find` first (instantiating the object) and then calls the instance version of `destroy`. Neat, huh? That pretty much covers `destroy`.

Using `delete`

The second style of row deletion is `delete`. Every Active Record class has class methods called `delete` and `delete_all`. The `delete` family of methods differs from `destroy` in that they don't instantiate or perform callbacks on the object they're deleting. They remove the row immediately from the database.

Just like `find` and `create`, you use `delete` and `delete_all` directly on the class (that is, you don't create an object first). Because the method operates on the table and not the row, you need to help it by telling it which row or rows you want to target:

```
>> Article.delete(4)
=> 1
```

Here you specify a single primary key for the article you want to delete. The operation responds with the number of records removed. Because a primary key uniquely identifies a single record, only one record is deleted.

Just as with `find`, you can give `delete` an array of primary keys whose rows you want to delete. Use square brackets ([]) to indicate that you're passing an array:

```
>> Article.delete([5, 6])
=> 0
```

The return of the `delete` method in this case is 0, since we didn't have records with `id` 5 and 6 in our database. Zero records were deleted.

Note Unlike `find`, which is capable of collecting any arguments it receives into an array automatically, `delete` must be supplied with an array object explicitly. So, although `Model.find(1,2,3)` works, `Model.delete(1,2,3)` fails with an argument error (because it's really receiving three arguments). To delete multiple rows by primary key, you must pass an actual array object. The following works, because it's a single array (containing three items) and thus a single argument: `Model.delete([1,2,3])`.

Deleting with Conditions

You can delete all rows that match a given condition with the `delete_all` class method. The following deletes all articles before a certain date:

```
>> Article.delete_all("published_at < '2011-01-01'")
=> 0
```

The return value of `delete_all` is the number of records deleted.

Caution If you use `delete_all` without any arguments, it deletes all rows in the table, so be careful! Most of the time, you pass it a string of conditions.

When Good Models Go Bad

So far, you've been nice to your models and have made them happy by providing just the information they need. But in Chapter 3 you provided *validations* that prevented you from saving bad records to the database. Specifically, you told the `Article` model that it should never allow itself to be saved to the database if it isn't given a `title` and `body`. Look at the `Article` model, as shown in Listing 5-1, to recall how validations are specified.

Listing 5-1. The `app/models/article.rb` File

```
class Article < ActiveRecord::Base
  validates_presence_of :title, :body
end
```

You may have noticed in your generated scaffolding that you use a helper method called `errors.full_messages` to print out a helpful error message. That helper isn't black magic; it's a bit of code that asks the model associated with the form for its list of errors (also referred to as the *errors collection*) and returns a nicely formatted block of HTML to show the user.

Note You may have noticed that you call methods in Ruby with a dot (.) For instance, you say `@article.errors` to get the error collection back. However, Ruby documentation uses the # symbol along with the class name to let the reader know that there is a method it can call on for an instance of that class. For example, on the Article class, you can use the method `@article.title` as `Article#title`, because it's something that acts on a particular `@article` but not the Article class itself. You've also seen that you can write the code `Article.count`, because you don't need to know about a particular `@article`, but only Article objects in general. Keep this convention in mind when you're reading Ruby documentation.

The secret to this is that every Active Record object has an automatic attribute added to it called `errors`. To get started, create a fresh Article object:

```
>> article = Article.new
=> #<Article id: nil, title: nil, body: nil, published_at: nil, created_at: nil,
updated_at: nil, excerpt: nil, location: nil>
>> article.errors.any?
=> false
```

This seems odd: you know this new article should have errors, because it's invalid—you didn't give it a title or a body. This is because you haven't triggered the validations yet. You can cause them to occur a couple of ways. The most obvious way is to attempt to save the object:

```
>> article.save
=> false
```

Every time you've used `save` before, the model has happily chirped `true` back to you. But this time, `save` returns `false`. This is because before the model allows itself to be saved, it runs through its gauntlet of validations, and one or more of those validations failed.

You would be right to guess that if you tried `article.errors.any?` again, it would return `true`:

```
>> article.errors.any?
=> true
```

Let's interrogate the `errors` collection a little more closely with the `full_messages` method:

```
>> article.errors.full_messages
=> ["Title can't be blank", "Body can't be blank"]
```

Voilà! Look how helpful the model is being. It's passing back an array of error messages.

If there is only one attribute that you care about, you can also ask the `errors` collection for a particular attribute's errors:

```
>> article.errors.messages(:title)
=> "can't be blank"
```

Notice that because you tell it which attribute you're looking for, the message returns a slightly different result than before. What if you ask for an attribute that doesn't exist or doesn't have errors?

```
>> article.errors.messages(:nonexistent)
=> nil
```

You get back `nil`, which lets you know that you didn't find anything.

Another helpful method is `size`, which, as you saw earlier, works with all arrays:

```
>> article.errors.size
=> 2
```

Saving isn't the only way you can cause validations to run. You can ask a model object if it's `valid?`:

```
>> article.valid?
=> false
```

If you try that on a new object, the `errors` collection magically fills up with your pretty errors.

Summary

In this chapter, you've become familiar with using the console to work with models. You've learned how to create, read, update, and destroy model objects. Also, you've briefly looked into how to see the simple errors caused by the validations you set up on your model in Chapter 3.

The next chapter discusses how to create relationships (called *associations*) among your models, and you begin to see how Active Record helps you work with your data in extremely powerful ways. It also expands on the concept of validations and shows how you can do a lot more with `validates`. You'll see that Rails provides a bevy of prewritten validators and an easy way to write your own customized validators.

CHAPTER 6



Advanced Active Record: Enhancing Your Models

Chapter 5 introduced the basics of Active Record and how to use it. This chapter delves more deeply into Active Record and teaches you how to enhance your models.

Model enhancement is a general term. It refers to endowing your models with attributes and capabilities that go beyond what you get from subclassing `ActiveRecord::Base`. A model contains all the logic that governs its citizenship in the world of your application. In the model, you can define how it interacts with other models, what a model should accept as a minimum amount of information for it to be considered valid, and other abilities and responsibilities.

Models need to relate to one another. In the real world, bank accounts have transactions, books belong to authors, and products have categories. These relationships are referred to as *associations*, and Active Record makes them easy to work with. Models also have requirements. For instance, you can't have a transaction without an amount—it might break your system if someone tried to have an empty transaction. So, Active Record gives you easy ways to tell a model what it should expect in order to be saved to the database.

This chapter will teach you how to programmatically enhance your models so they're more than just simple maps of your tables. To demonstrate the concepts, you build on the blog application you started in Chapter 3, so keep it handy if you want to follow along with the examples.

Adding Methods

Let's begin with a brief review of Active Record basics. At the simplest level, Active Record works by automatically wrapping database tables whose names match the plural, underscored version of any classes that inherit from `ActiveRecord::Base`. For example, if you want to wrap the `users` table, you create a subclass of `ActiveRecord::Base` called `User`, like this:

```
class User < ActiveRecord::Base  
end
```

That's all you really need to have Active Record map the `users` table and get all the basic CRUD functionality described in Chapter 5. But few models are actually this bare.

So far, you've left your model classes unchanged. That's a good thing, and it speaks to the power and simplicity of Active Record. However, it leaves something to be desired. Most of the time, your models need to do a lot more than just wrap a table.

Note If you're familiar with SQL, you're probably feeling that Active Record provides only simple-case solutions and can't handle complicated cases. That's entirely untrue. Although SQL is useful for highly customized database queries, most Rails projects rarely need to touch SQL, thanks to some clever tricks in Active Record.

The primary way in which you enhance models is by adding methods to them. This is referred to as adding *domain logic*. With Active Record, all the logic for a particular table is contained in one place: the model. This is why the model is said to *encapsulate* all the domain logic. This logic includes access rules, validations, relationships, and, well, just about anything else you feel like adding.

In addition to all the column-based reader and writer methods you get by wrapping a table, you're free to define your own methods on the class. An Active Record subclass isn't much different from a regular Ruby class; about the only difference is that you need to make sure you don't unintentionally overwrite any of Active Record's methods (`find`, `save`, or `destroy`, for example). For the most part, though, this isn't a problem.

Let's look at a simple example. You often need to format data, rather than accessing a model attribute in its raw form. In the blog application, you want to be able to produce a formatted, long title that includes the title of the article and its date. To accomplish this, all you need to do is define a new instance method called `long_title` that performs the concatenation of those attributes and produces a formatted string. Add the code shown in Listing 6-1 just before the last `end` statement in the `app/models/article.rb` file.

Listing 6-1. Custom `long_title` Method, in `app/models/article.rb`: <http://gist.github.com/323787>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  def long_title
    "#{title} - #{published_at}"
  end
end
```

You've just created an instance method on the model; that is, you've told the `Article` model that it's now endowed with a new attribute called `long_title`. You can address `long_title` the same way you would any other method on the class. Open an irb session and try this on the console. From the terminal window, make sure you're inside the blog application directory, and then start up the Rails console with the following command:

```
$ rails console
```

This should drop you at a simple irb prompt with two right arrows and a blinking cursor; this may look a bit different based on your environment. From here, you create a new article and use it to call the `long_title` method:

```
>> Article.create :title => 'Advanced Active Record', :published_at => Date.today,
:body => 'Models need to relate to each other. In the real world, ...'
=> #<Article id: 6, title: "Advanced Active Record", ...>
>> Article.last.long_title
=> "Advanced Active Record - 2013-04-22 04:00:00 UTC"
```

There is no difference between the methods Active Record creates and those you define. Here, instead of asking the model for one of the attributes garnered from the database column names, you define your *own* method called `long_title`, which does a bit more than the standard `title` method.

The methods you add to your models can be as simple as returning true or false or as complicated as doing major calculations and formatting on the object. The full power of Ruby is in your hands to do with as you please.

Don't worry if you don't feel comfortable adding your own methods to models just yet. The important part to note from this section is that Active Record models are regular Ruby classes that can be augmented, modified, played with, poked, and turned inside out with sufficient Ruby-fu. Knowing this is extremely helpful in being able to pull back the curtain and understand the advanced features of Active Record.

FAT MODELS

Some might be nervous by the `long_title` method you just used. They may see it as a violation of the MVC paradigm. They might ask, "Isn't formatting code supposed to be in the view?" In general, the answer is yes. However, it often helps to have models that act as intelligent objects. If you ask a model for some information about itself, it's natural to assume that it can give you a decent answer that doesn't require a large amount of work later on to figure out what it means. So, small formatted strings and basic data types that faithfully represent the data in the model are good things to have in your code.

An intelligent model like this is often called *fat*. Instead of performing model-related logic in other places (that is, in controllers or views), you keep it in the model, thus making it fat. This makes your models easier to work with and helps your code stay DRY.

A basic rule of thumb while trying to stay DRY is that if you find yourself copying and pasting a bit of code, it may be worth your time to take a moment and figure out if there is a better way to approach the problem. For instance, if you had kept the `Article#long_title` formatting outside the model, you might have needed to repeat the same basic string-formatting procedure every time you wanted a human-friendly representation of an article's title. Then again, creating that method is a waste of time if you're going to use it in only one place in the application and never again.

This is where programmer experience comes in. As you learn and mature in your Rails programming, you'll find it easier and easier to figure out where stuff is supposed to go. If you're always aiming for a goal of having the most maintainable and beautiful code you can possibly write, your projects will naturally become easier to maintain.

Next, let's look at another common form of model enhancement: associations. Active Record's associations give you the ability to define in simple terms how models relate to and interact with one another.

Using Associations

It's a lowly application that has only one table. Most applications have many tables, and these tables typically need to relate to one another in one way or another. *Associations* are a common model enhancement that let you relate tables to one another.

Associations are natural constructs that you encounter all the time in the real world: articles have comments, stores have products, magazines have subscriptions, and so on. In a relational database system, you relate tables using a *foreign key reference* in one table to the *primary key* of another table.

Note The terms *relationship* and *association* can be used pretty much interchangeably. However, when this book refers to *associations*, it generally means the association on the Active Record side, as opposed to the actual foreign-key relationships at the database level.

Let's take the example of articles and comments. In a situation where a given article can have any number of comments attached to it, each comment *belongs to* a particular article. Figure 6-1 demonstrates the association from the database's point of view.

The example in Figure 6-1 uses a column named `article_id` in the `comments` table to identify the related article in the `articles` table. In database-speak, `comments` holds a *foreign-key reference* to `articles`.

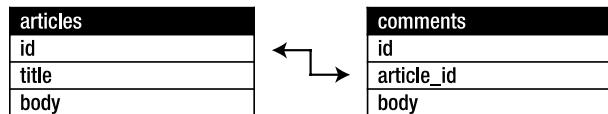


Figure 6-1. The relationship between the `articles` and `comments` tables

By Rails convention, the foreign-key column is the singular, lowercase name of the target class with `_id` appended. So, for products that belong to a particular store, the foreign key is named `store_id`; for subscriptions that belong to magazines, the foreign key is named `magazine_id`; and so on. Here's the pattern:

```
#{singular_name_of_parent_class}_id
```

Table 6-1 shows a few more examples, just to drive this concept home.

Table 6-1. Sample Foreign-Key References

Model	Table	Foreign Key to Reference This Table
Article	articles	article_id
Person	people	person_id
Friend	friends	friend_id
Category	categories	category_id
Book	books	book_id

Whenever you need one table to reference another table, remember to create the foreign-key column in the table doing the referencing. In other words, the model that contains the “`belong_to`” needs to have the foreign key column in it. That's all your table needs before you can put Active Record's associations to work.

Declaring Associations

As you've probably come to expect by now, Active Record makes working with associations easy. You don't need to get down to the bare metal of the database very often. As long as you understand the concept of primary and foreign keys and how to create basic relationships in your tables, Active Record does the proverbial heavy lifting, converting foreign-key relationships into rich object associations. This means you get to access associated objects cleanly and naturally using Ruby:

```
article.comments
store.products
magazine.subscriptions
```

After the relationships are defined in your database tables, you use a set of macro-like class methods in your models to create associations. They look like this:

- `has_one`
- `has_many`
- `belongs_to`
- `has_and_belongs_to_many`

Here's a quick example. The `Message` model declares a `has_many` relationship with `Attachment`; `Attachment` returns the favor by declaring that each of its objects belongs to a particular `Message`:

```
class Message < ActiveRecord::Base
  has_many :attachments
end

class Attachment < ActiveRecord::Base
  belongs_to :message
end
```

Given these instructions, Active Record expects to find a table called `attachments` that has a field in it called `message_id` (the foreign-key reference). It uses this association to let you enter things like `Message.first.attachments` and get an array (or a *collection*) of `Attachment` objects that belongs to the first `Message` in the database. Moreover, you can work with your associations in both directions. So, you can enter `Attachment.first.message` to access the `Message` to which the first `Attachment` belongs. It sounds like a mouthful, but when you get the hang of it, it's quite intuitive.

Whenever you declare an association, Active Record automatically adds a set of methods to your model that makes dealing with the association easier. This is a lot like the way in which Active Record creates methods based on your column names. When it notices you've declared an association, it dynamically creates methods that enable you to work with that association. The following sections go through the different types of associations and describe how to work with them. You also learn about the various options you can use to fine-tune associations.

Creating One-to-One Associations

One-to-one associations describe a pattern where a row in one table is related to exactly one row in another table.

Suppose that in your blog application, you have users and profiles, and each user has exactly one profile. Assume you have `User` and `Profile` models, and the corresponding `users` and `profiles` tables have the appropriate columns. You can tell your `User` model that it *has one* `Profile` and your `Profile` model that it *belongs to* a `User`. Active Record takes care of the rest. The `has_one` and `belongs_to` macros are designed to read like regular English, so they sound natural in conversation and are easy to remember. Each represents a different side of the equation, working in tandem to make the association complete.

Note Part of the Rails philosophy about development is that the gap between programmers and other project stakeholders should be bridged. Using natural language, such as *has one* and *belongs to*, in describing programmatic concepts helps bridge this gap, providing a construct that everyone can understand.

Adding the User and Profile Models

When you started the blog application, you decided to let anyone create new articles. This worked fine when only one person was using the system; but you want this to be a multiple-user application and let different people sign up, sign in, and start writing their own articles separately from one another.

Let's fire up the generator and create the User model:

```
$ rails generate model User email:string password:string
```

Just as you saw in Chapter 5, the model generator creates, among other things, a model file in `app/models` and a migration in `db/migrate`. Open `db/migrate/20130406205458_create_users.rb`, and you should see the now-familiar code in Listing 6-2.

Listing 6-2. Migration to Create the users Table, `db/migrations/20130406205458_create_users.rb`

```
class CreateUsers < ActiveRecord::Migration
  def change
    create_table :users do |t|
      t.string :email
      t.string :password

      t.timestamps
    end
  end
end
```

This is standard migration fare. In the `change` definition, you use the `create_table` method to create a new `users` table. The new table object is *yielded* to the block in the variable, `t`, on which you call the `string` method to create each column. Along with the standard `email` field, you specify a `password` field, which you use for authentication, as explained in the “Reviewing the Updated Models” section later in this chapter. The primary key, `id`, is created automatically, so there's no need to specify it here.

As you probably noticed, the `User` model is extremely simple: it only contains information that allows the user to authenticate into the application. Some users may want to add a lot more detail about themselves and would love the ability to enter personal information such as their birthday, a biography, their favorite color, their Twitter account name, and so on. You can create a `Profile` model to hold such information outside the scope of the `User` model. Just as you did for the `User` model, use the generator again:

```
$ rails generate model Profile user_id:integer name:string birthday:date
  bio:text color:string twitter:string
```

You also have a migration file for the `Profile` model in `db/migrate/20130406210440_create_profiles.rb`—feel free to take a peek. Notice the existence of the foreign key `user_id` in the `profiles` schema. Also recall that you don't need to specify primary keys in migrations because they're created automatically.

Now, all you need to do is run the migrations and create the new tables using the `db:migrate` Rake task. Run the migrations with the following command:

```
$ rake db:migrate
```

```
--> CreateUsers: migrating =====
-- create_table(:users)
-> 0.0019s
--> CreateUsers: migrated (0.0020s) =====
--> CreateProfiles: migrating =====
-- create_table(:profiles)
-> 0.0027s
--> CreateProfiles: migrated (0.0035s) =====
```

With the table and foreign keys in place, Listings 6-3 and 6-4 show how to declare the one-to-one association on the User and Profile models, respectively.

Listing 6-3. The User Model, app/models/user.rb: <http://gist.github.com/323944>

```
class User < ActiveRecord::Base
  has_one :profile
end
```

Listing 6-4. The Profile Model, app/models/profile.rb: <http://gist.github.com/323946>

```
class Profile < ActiveRecord::Base
  belongs_to :user
end
```

The `has_one` declaration on the User model tells Active Record that it can expect to find one record in the `profiles` table that has a `user_id` matching the primary key of a row in the `users` table. The Profile model, in turn, declares that each of its records `belongs_to` a particular User.

Telling the Profile model that it `belongs_to :user` is saying, in effect, that each Profile object references a particular User. You can even go so far as to say that User is the parent and Profile is the child. The child model is dependent on the parent and therefore references it. Figure 6-2 demonstrates the `has_one` relationship.

has_one

```
User has_one :profile
Profile belongs_to :user
```

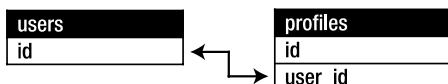


Figure 6-2. The one-to-one relationship between users and profiles

Let's get inside a console session (`rails console`) and see how this comes together. If you have a console session opened, run the `reload!` command in the console session to make sure it loads the newly generated models. Follow along to create objects and relate them to one another. First, create a user and a profile as follows:

```
>> reload!
Reloading...
>> user = User.create(:email => 'user@example.com', :password => 'secret')
=> #<User id: 1, email: "user@example.com", password: "secret",
  created_at: "2013-04-02 15:10:07", updated_at: "2013-04-02 15:10:07">
>> profile = Profile.create(:name => 'John Doe',
:bio => 'Ruby developer trying to learn Rails')
=> #<Profile id: 1, user_id: nil, name: "John Doe", birthday: nil,
bio: "Ruby developer trying to learn Rails", color: nil, twitter: nil,
created_at: "2013-04-02 15:10:55", updated_at: "2013-04-02 15:10:55">
```

Note The `reload!` method reloads the Rails application environment within your console session. You need to call it when you make changes to existing code. It's exactly as if you had restarted your console session—all the variables you may have instantiated are lost.

Although you've successfully created a new user and a new profile, you haven't yet associated them with each other. If you ask the `user` object for its profile, it responds with `nil`:

```
>> user.profile
=> nil
```

To make the association happen, you specify it like any regular assignment on the `user` object and then call `save`, like this:

```
>> user.profile = profile
=> #<Profile id: 1, user_id: 1, name: "John Doe", birthday: nil,
bio: "Ruby developer trying to learn Rails", color: nil, twitter: nil,
created_at: "2013-04-02 15:10:55", updated_at: "2013-04-02 15:10:55">
>> user.save
=> true
```

Assignment is assignment, whether it's a name attribute to which you're assigning the value `Joe` or an association method to which you're assigning an object. Also notice that the profile's `user_id` attribute is updated to the value of `user.id`: This is what bonds both objects together. Now, when you ask the `user` object for its profile, it happily responds with one:

```
>> user.profile
=> #<Profile id: 1, user_id: 1, name: "John Doe", birthday: nil,
bio: "Ruby developer trying to learn Rails", color: nil, twitter: nil,
created_at: "2013-04-02 15:10:55", updated_at: "2013-04-02 15:10:55">
```

That's all there is to it. Although this is pretty good, you can do a bit better. You can create and save the profile in one shot and have it perform the association automatically, like this:

```
>> user.create_profile :name => 'Jane Doe', :color => 'pink'
=> #<Profile id: 2, user_id: 1, name: "Jane Doe", birthday: nil,
bio: nil, color: "pink", twitter: nil, created_at: "2013-04-02 15:18:57",
updated_at: "2013-04-02 15:18:57">
```

Using the `create_profile` method to create a new profile initializes the `Profile` object, sets its foreign key to `user.id`, and saves it to the database; at the same time, it sets the previous profile object's—the one named "John Doe"—`user_id`—field to `nil`. This works for any `has_one` association, no matter what it's named. Active Record automatically generates the `create_#{association_name}` method for you. So if you had an `Employee` model set up with an association like `has_one :address`, you would get the `create_address` method automatically.

These alternatives for doing the same thing may seem confusing, but they're really variations on the same theme. In all cases, you're creating two objects (the parent and the child) and telling each about the other. Whether you choose to do this in a multistep operation or all on one line is entirely up to you.

Earlier, you learned that declaring a `has_one` association causes Active Record to automatically add a suite of methods to make working with the association easier. Table 6-2 shows a summary of the methods that are added when you declare a `has_one` and `belongs_to` relationship between `User` and `Profile`, where `user` is a `User` instance.

Table 6-2. Methods Added by the `has_one` Association in the User/Profile Example

Method	Description
<code>user.profile</code>	Returns the associated (<code>Profile</code>) object; <code>nil</code> is returned if none is found
<code>user.profile=(profile)</code>	Assigns the associated (<code>Profile</code>) object, extracts the primary key, and sets it as the foreign key
<code>user.profile.nil?</code>	Returns true if there is no associated <code>Profile</code> object
<code>user.build_profile(attributes={})</code>	Returns a new <code>Profile</code> object that has been instantiated with attributes and linked to <code>user</code> through a foreign key but hasn't yet been saved
<code>user.create_profile(attributes={})</code>	Returns a new <code>Profile</code> object that has been instantiated with attributes and linked to <code>user</code> through a foreign key and that has already been saved

Although you're using the `User.has_one :profile` example here, the rules work for any object associated to another using `has_one`. Here are some examples, along with sample return values:

```
user.profile
#=> #<Profile id: 2, user_id: 1, ...>
user.profile.nil?
#=> false
user.build_profile(:bio => 'eats leaves')
#=> #<Profile id: nil, user_id: 1, ...>
user.create_profile(:bio => 'eats leaves')
#=> #<Profile id: 3, user_id: 1, ...>
```

The `has_one` declaration can also include an options hash to specialize its behavior if necessary. Table 6-3 lists the most common options. For a complete list of all options, consult the Rails API documentation (http://api.rubyonrails.org/classes/ActiveRecord/Associations/ClassMethods.html#method-i-has_one).

Table 6-3. Common has_one Options

Option	Description	Example
:class_name	Specifies the class name of the association. Used when the class name can't be inferred from the association name.	has_one :profile, :class_name => 'Account'
:conditions	Specifies the conditions that the associated object must meet in order to be included as a WHERE SQL fragment.	has_one :profile, :conditions => "active = 1"
:foreign_key	Specifies the foreign key used for the association in the event that it doesn't adhere to the convention of being the lowercase, singular name of the target class with _id appended.	has_one :profile, :foreign_key => 'account_id'
:order	Specifies the order in which the associated object is picked as an ORDER BY SQL fragment.	has_one :profile, :order => "created_at DESC"
:dependent	Specifies that the associated object should be removed when this object is. If set to :destroy, the associated object is deleted using the destroy method. If set to :delete, the associated object is deleted without calling its destroy method. If set to :nullify, the associated object's foreign key is set to NULL.	has_one :profile, :dependent => :destroy

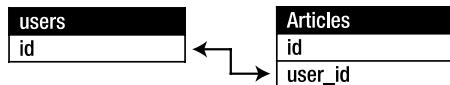
Creating One-to-Many Associations

One-to-many associations describe a pattern where a row in one table is related to one or more rows in another table. Examples are an Email that has many Recipients, or a Magazine that has many Subscriptions.

Up until now, your articles have been orphaned—they don't belong to anyone. You remedy that now by associating users with articles. In your system, each article belongs to a user, and a user may have many articles. Figure 6-3 illustrates this association.

has_many

User has_many :articles
Article belongs_to :user

**Figure 6-3.** The one-to-many relationship between users and articles

Associating User and Article Models

Just as you associated users and profiles, you want to have a similar relationship between users and articles. You need to add a foreign key user_id in the articles table that points to a record in the users table.

Fire up the migration generator:

```
$ rails generate migration add_user_id_to_articles user_id:integer
```

Open db/migrate/20130406221123_add_user_id_to_articles.rb, and you should see the code in Listing 6-5.

Listing 6-5. Migration to Add user_id Field, db/migrate/20130406221123_add_user_id_to_articles.rb

```
class AddUserIdToArticles < ActiveRecord::Migration
  def change
    add_column :articles, :user_id, :integer
  end
end
```

Pay attention to the user_id column in the change section. It's the foreign-key reference column you've heard so much about. Also note that its type is :integer. That's important, because it's referring to a numeric id.

Now, all you need to do is run the migration using the db:migrate Rake task. Run the migration with the following command:

```
$ rake db:migrate
```

```
--> AddUserIdToArticles: migrating =====
-- add_column(:articles, :user_id, :integer)
 -> 0.0012s
--> AddUserIdToArticles: migrated (0.0015s) =====
```

With the foreign key in place, Listings 6-6 and 6-7 show how you declare the one-to-many association in your Article and User models, respectively. Add these to the relevant models.

Listing 6-6. The Article Model, belongs_to Declaration in app/models/article.rb: <http://gist.github.com/323981>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user

  def long_title
    "#{title} - #{published_at}"
  end
end
```

Listing 6-7. The User Model, has_many Declaration in app/models/user.rb: <http://gist.github.com/323982>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles
end
```

That's all there is to it. This bit of code has endowed your Article and User models with a lot of functionality.

Note For has_one and has_many associations, adding a belongs_to on the other side of the association is always recommended. The rule of thumb is that the belongs_to declaration always goes in the class with the foreign key.

Creating a New Associated Object

Your associations are in place; so, let's get back into the code to put what you've learned to the test. Do this exercise on the console: either run `rails console` to start a new console session or type `reload!` if you still have a console window open from the previous section.

Let's test whether the association between users and articles is set up correctly. If it is, you should be able to ask the user object for its associated articles, and it should respond with a collection. Even though you haven't created any articles for this user yet, it should still work, returning an empty collection:

```
>> reload!
Reloading...
=> true
>> user = User.first
=> #<User id: 1, email: "user@example.com", password: "secret",
  created_at: "2013-04-02 15:10:07", updated_at: "2013-04-02 15:10:07">
>> user.articles
=> []
```

Great! The `has_many` association is working correctly, and User instances now has an `articles` method, which was created automatically by Active Record when it noticed the `has_many` declaration.

Let's give this user some articles. Enter the following commands:

```
>> user.articles << Article.first
=> [#<Article id: 6, ..., user_id: 1>]
>> user.articles.size
=> 1
>> user.articles
=> [#<Article id: 6, ..., user_id: 1>]
```

By using the append (`<<`) operator, you attach `Article.first` onto your `User` object. When you use `<<` with associations, it automatically saves the new association. Some things in Active Record don't happen until you say `save`, but this is one of the examples where that part is done automatically.

What did that do exactly? Let's look into the article and find out:

```
>> Article.first.user_id
=> 1
```

See how this article's `user_id` points to the user with an `id` of 1? This means you've successfully related the two objects. You can even ask an `Article` instance for its user:

```
>> Article.first.user
=> #<User id: 1, email: "user@example.com", password: "secret",
  created_at: "2013-04-02 15:10:07", updated_at: "2013-04-02 15:10:07">
```

Voilà! Your models can really start to express things now. The `has_many` and `belongs_to` declarations create more methods, as you did earlier with the `long_title` method. Let's look at what else these happy little helpers brought along to the party. Table 6-4 shows a summary of the methods that are added when you declare a `has_many` and `belongs_to` relationship between `User` and `Article` (`user` represents a `User` instance).

Table 6-4. Methods Added by the `has_many` Association in the User and Article Models

Method	Description
<code>user.articles</code>	Returns an array of all the associated articles. An empty array is returned if no articles are found.
<code>user.articles=(articles)</code>	Replaces the <code>articles</code> collection with the one supplied.
<code>user.articles << article</code>	Adds one or more articles to the collection and saves their foreign keys.
<code>user.articles.delete(articles)</code>	Removes one or more articles from the collection by setting their foreign keys to <code>NULL</code> .
<code>user.articles.empty?</code>	Returns true if there are no associated Article objects for this user.
<code>user.articles.size</code>	Returns the number of associated Article objects for this user.
<code>user.article_ids</code>	Returns an array of associated article ids.
<code>user.articles.clear</code>	Clears all associated objects from the association by setting their foreign keys to <code>NULL</code> .
<code>user.articles.find</code>	Performs a <code>find</code> that is automatically scoped off the association; that is, it finds only within items that belong to <code>user</code> .
<code>user.articles.build(attributes={})</code>	Returns a new Article object that has been instantiated with attributes and linked to <code>user</code> through a foreign key but hasn't yet been saved. Here's an example: <code>user.articles.build(:title => 'Ruby 1.9')</code> .
<code>user.articles.create(attributes={})</code>	Returns a new Article object that has been instantiated with attributes and linked to <code>user</code> through a foreign key and has already been saved. Here's an example: <code>user.articles.create(:title => 'Hoedown')</code> .

You're using the `User.has_many :articles` example here, but the rules work for any object associated with another using `has_many`. Here are some examples, along with sample return values:

```
>> user.articles
=> [#<Article id: 6, ...>]
>> user.articles << Article.new(:title => 'One-to-many associations',
:body => 'One-to-many associations describe a pattern ..')
=> [#<Article id: 6, ...>, #<Article id: 7, ...>]
>> user.article_ids
=> [6, 7]
>> user.articles.first
=> #<Article id: 6, ...>
>> user.articles.clear
=> []
>> user.articles.count
=> 0
>> Article.count
=> 2
>> user.articles.create :title => 'Associations',
```

```
:body => 'Active Record makes working with associations easy..'
=> #<Article id: 8, ...>
```

You can also pass in options to your association declaration to affect the way you work with those associations. Table 6-5 lists some of the most common options.

Table 6-5. Common has_many Options

Option	Description	Example
:class_name	Specifies the class name of the association. Used when the class name can't be inferred from the association name.	has_many :articles, :class_name => 'Post'
:conditions	Specifies the conditions that the associated objects must meet in order to be included as a WHERE SQL fragment.	has_many :articles, :conditions => "active = 1"
:foreign_key	Specifies the foreign key used for the association in the event that it doesn't adhere to convention of being the lowercase, singular name of the target class with _id appended.	has_many :articles, :foreign_key => 'post_id'
:order	Specifies the order in which the associated objects are returned as an ORDER BY SQL fragment.	has_many :articles, :order => "published_at DESC"
:dependent	Specifies that the associated objects should be removed when this object is. If set to :destroy, the associated objects are deleted using the destroy method. If set to :delete, the associated objects are deleted without calling their destroy method. If set to :nullify, the associated objects' foreign keys are set to NULL.	has_many :articles, :dependent => :destroy

There's much more to has_many associations than can possibly be covered here, so be sure to check out the Rails API documentation (http://api.rubyonrails.org/classes/ActiveRecord/Associations/ClassMethods.html#method-i-has_many) for the full scoop.

Applying Association Options

It's time to apply what you've learned to your domain model. Specifically, you use the :order option to apply a default order to the User.has_many :articles declaration, and you use the :dependent option to make sure when you delete a user, all their articles are deleted as well.

Specifying a Default Order

When you access a user's articles, you want to make sure they come back in the *order* in which they've been published. Specifically, you want the oldest to be at the bottom of the list and the newest to be at the top. You can do this by configuring the has_many association with a default order using the :order option. Add the :order option to the has_many :articles declaration, as shown in Listing 6-8.

Listing 6-8. The :order Option Added to has_many: <http://gist.github.com/324010>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles, -> { order('published_at DESC') }
end
```

You give the name of the field that you want to order by, and then you say either ASC (ascending) or DESC (descending) to indicate the order in which the results should be returned. Because time moves forward (to bigger numbers), you want to make sure you're going back in time, so you use the DESC keyword here.

Note ASC and DESC are SQL keywords. You're actually specifying an *SQL fragment* here, as discussed in the “Advanced Finding” section later in this chapter.

You can also specify a secondary order by adding a comma between arguments. Let's say you want to sort by the title of the article *after* you sort by the date. If two articles have been published on the same day, they are ordered first by the date and then by the lexical order of the title. Listing 6-9 shows the article title added to the :order option.

Listing 6-9. Adding the Title to the :order Option for has_many: <http://gist.github.com/324019>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') }
end
```

Notice that you use ASC for ordering on the title. This is because as letters go up in the alphabet, their value goes up. So, to sort alphabetically, use the ASC keyword.

Specifying Dependencies

Frequently, dependencies exist between models. For instance, in your blog application, if you delete users, you want to make sure they don't have articles in the system. Said another way, an Article is dependent on its User. You can let Active Record take care of this for you automatically by specifying the :dependent option to your association. Listing 6-10 shows all the options to has_many :articles, including the :dependent option.

Listing 6-10. The :dependent Option Added to has_many: <http://gist.github.com/324020>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :destroy
end
```

By passing in the symbol :destroy, you declare not only that articles are dependent, but also that when the owner is deleted, you want to call the destroy method on every related article. This ensures that any *_destroy callbacks on the article instances are called (callbacks are discussed later, in the “Making Callbacks” section). If you want to skip the callbacks, you can use the :delete option instead of :destroy, which deletes the records directly via SQL.

Let's say you want to set the foreign-key column (`user_id`) to `NULL` in the `articles` table, instead of completely destroying the article. Doing so essentially orphans the articles. You can do this by using the `:nullify` option instead of `:destroy`. If you don't use the `:dependent` option and you delete a user with associated articles, you break foreign-key references in your `articles` table. For this application, you want to keep the `:nullify` option, as per Listing 6-11.

Listing 6-11. The `:dependent` Option Set to `:nullify`: <http://gist.github.com/324023>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC')},
    :dependent => :nullify
end
```

Creating Many-to-Many Associations

Sometimes, the relationship between two models is many-to-many. This describes a pattern where two tables are connected to multiple rows on both sides. You use this in the blog application to add categories to articles. If you wanted to allow only one category to be selected for a given article, you could use `has_many`. But you want to be able to apply multiple categories.

Think about this for a minute: an article can have many categories, and a category can have many articles—where does the `belongs_to` go in this situation? Neither model belongs to the other in the traditional sense. In Active Record-speak, this kind of association is `has_and_belongs_to_many` (often referred to as `habtm` for short).

The `has_and_belongs_to_many` association works by relying on a join table that keeps a reference to the foreign keys involved in the relationship. The join table sits between the tables you want to join: `articles` and `categories`. Not surprisingly, then, the join table in this case is called `articles_categories`. Pay particular attention to the table name. It's formed from the names of each table in alphabetical order, separated by an underscore. In this case, the `a` in `articles` comes before the `c` in `categories`—hence, `articles_categories`. Figure 6-4 illustrates this relationship.

has_and_belongs_to_many

```
Category has_and_belongs_to_many :articles
Article has_and_belongs_to_many :categories
```



Figure 6-4. The many-to-many relationship between `articles` and `categories`

Let's start by adding the `Category` model. This is a simple matter of generating the model, consisting of just a `name` column. Run the following command inside your application root:

```
$ rails generate model Category name:string
```

Look at the generated migration in `db/migrate/20130407001138_create_categories.rb`; it's pretty familiar territory at this point. You need another migration to create the join table. Do that now by running the following command:

```
$ rails generate migration create_articles_categories
```

Remember that when you use `create_table` inside a migration, you don't need to specify the primary key, because it's created automatically. Well, in the case of a join table, you don't want a primary key. This is because the join table isn't a first-class entity in its own right. Creating tables without primary keys is the exception and not the rule, so you need to explicitly tell `create_table` that you don't want to create an `id`. Take a close look at the call to `create_table` in Listing 6-12. You pass in option `:id => false`, which prevents `create_table` from creating the primary key; also, you add the appropriate fields.

Listing 6-12. The db/migrate/20130407002156_create_articles_categories.rb: File

```
class CreateArticlesCategories < ActiveRecord::Migration
  def change
    create_table :articles_categories, :id => false do |t|
      t.references :article
      t.references :category
    end
  end
  def self.down
    drop_table :articles_categories
  end
end
```

You use the `references` method in the `create_table` block instead of using `integer`. It's just another notation that uses the association name as a parameter instead of a field name; so, `t.references :article` is the same as `t.integer :article_id`. Feel free to pick the syntax you prefer. Go ahead and run the migrations:

```
$ rake db:migrate
```

```
--> CreateCategories: migrating =====
-- create_table(:categories)
-> 0.0032s
--> CreateCategories: migrated (0.0034s) =====
--> CreateArticlesCategories: migrating =====
-- create_table(:articles_categories, {:id=>false})
-> 0.0011s
--> CreateArticlesCategories: migrated (0.0013s) =====
```

With the `Category` model and the join table in place, you're ready to let Active Record in on your association. Open the `Article` and `Category` models and add the `has_and_belongs_to_many` declarations to them, as shown in Listings 6-13 and 6-14, respectively.

Listing 6-13. Adding the `has_and_belongs_to_many` Declaration in the Article Model app/models/article.rb:
<http://gist.github.com/324055>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body
```

```

belongs_to :user
has_and_belongs_to_many :categories

def long_title
  "#{title} - #{published_at}"
end
end

```

Listing 6-14. Adding the has_and_belongs_to_many Declaration in Category Model app/models/category.rb:
<http://gist.github.com/324057>

```

class Category < ActiveRecord::Base
  has_and_belongs_to_many :articles
end

```

Seeding Data

As part of creating an application skeleton, Rails added a file called db/seeds.rb, which defines some data you always need in your database. The seeds file contains Ruby code, so you can use the classes and methods—including associations—available in your models, such as create and update. Open it and create one user and a few categories so that it looks like Listing 6-15.

Listing 6-15. The db/seeds.rb File: <http://gist.github.com/324072>

```

user = User.create :email => 'mary@example.com', :password => 'guessit'
Category.create [{:name => 'Programming'},
                 {:name => 'Event'},
                 {:name => 'Travel'},
                 {:name => 'Music'},
                 {:name => 'TV'}]

```

That should do nicely. You can load your seed data using the Rake task db:seed:

```
$ rake db:seed
```

If you need to add more default categories later, you can append them to the seeds file and reload it. If you want to rerun the seed data, the trick lies in the fact that the seeds file doesn't know whether the records already in the database have to be cleaned up; running rake db:seed again adds all records one more time, and you end up with duplicate user and categories. You should instead call rake db:setup, which re-creates the database and adds the seed data as you may expect.

Let's give this a test run. Get your console ready, reload!, and run the following commands:

```

>> article = Article.last
=> #<Article id: 8, title: "Associations", ...>
>> category = Category.find_by_name('Programming')
=> #<Category id: 1, name: "Programming", ..>
>> article.categories << category
=> [#<Category id: 1, name: "Programming", ..>]
>> article.categories.any?
=> true
>> article.categories.size
=> 1

```

Here, you automatically associate a category with an article using the `<<` operator. You can even do this from the category side of the association. Try the following:

```
>> category.articles.empty?
=> false
>> category.articles.size
=> 1
>> category.articles.first.title
=> "Associations"
```

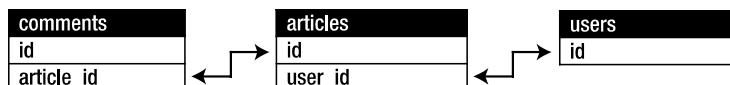
You just did the opposite of the previous test: `has_and_belongs_to_many` works in both directions, right? So, you found your category and asked it for its first article titled “Associations” because that’s what you associated in the other direction.

Using `has_and_belongs_to_many` is a very simple way to approach many-to-many associations. However, it has its limitations. Before you’re tempted to use it for more than associating categories with articles, note that it has no way of storing additional information on the join. What if you want to know *when* or *why* someone assigns a category to an article? This kind of data fits naturally in the *join table*. Rails includes another type of association called `has_many :through`, which allows you to create rich joins like this.

Creating Rich Many-to-Many Associations

Sometimes, when you’re modeling a many-to-many association, you need to put additional data on the join model. But because Active Record’s `has_and_belongs_to_many` uses a join table (for which there is no associated model), there’s no model on which to operate. For this type of situation, you can create *rich* many-to-many associations using `has_many :through`. This is really a combination of techniques that ends up performing a similar but more robust version of `has_and_belongs_to_many`.

The basic idea is that you build or use a full model to represent the join table. Think about the blog application: articles need to have comments, so you create a `Comment` model and associate it with `Article` in a one-to-many relationship using `has_many` and `belongs_to`. You also want to be able to retrieve all the comments added to users’ articles. You could say that users have many comments that belong to their articles, or users have many comments through articles. Figure 6-5 illustrates this relationship.



User has_many :replies,
`:through => :articles,`
`:source => :comments`

Figure 6-5. The rich many-to-many relationship between comments and users, through articles

Let’s generate the model and migration for the `Comment` model:

```
$ rails generate model comment article_id:integer name:string email:string body:text
```

```
invoke    active_record
create    db/migrate/20130408213941_create_comments.rb
create    app/models/comment.rb
invoke    test_unit
create    test/unit/comment_test.rb
create    test/fixtures/comments.yml
```

Migrate by issuing the `rake db:migrate` command:

```
$ rake db:migrate
```

```
--> CreateComments: migrating =====
-- create_table(:comments)
-> 0.0016s
--> CreateComments: migrated (0.0017s) =====
```

Update your models to reflect the one-to-many association between comments and articles. Listings 6-16 and 6-17 show the updated `Comment` and `Article` models, respectively.

Listing 6-16. The `Comment` Model in `app/models/comment.rb`: <http://gist.github.com/325793>

```
class Comment < ActiveRecord::Base
  belongs_to :article
end
```

Listing 6-17. The `Article` Model in `app/models/article.rb`: <http://gist.github.com/325796>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  def long_title
    "#{title} - #{published_at}"
  end
end
```

Nothing is new here—what you implement is very similar to the users and articles relationship you saw earlier, but instead of a user having many articles, an article has many comments.

Let's get back to the relationship between users and comments. You need to tell your user model that a user has many comments through its articles. Basically, you use the article model as a join table between users and comments. You achieve the linking using the `has_many :through` method. Listing 6-18 shows the updated User model.

Listing 6-18. The Updated User Model, has_many :through Declarations in app/models/user.rb:
<http://gist.github.com/325809>

```
class User < ActiveRecord::Base
  has_one :profile
  has_many :articles, -> {order('published_at DESC, title ASC')},
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments
end
```

Notice that you rework how you name associations. One aspect of the Rails philosophy is that you should always be questioning and refactoring your code to work with best practices. In this incarnation, comments that users receive on their articles are called replies.

As an added benefit, has_many :through allows you to easily have nice names for your associations. The :source option lets you define the source name of the association. In this case, the replies are the articles' comments, so you set the :source option accordingly.

Let's play with this on the console to see how it works—don't forget to reload!. You first find the first user, find the user's first article, and create a comment on it. Then, you see that comment directly from the user object:

```
>> user = User.first
=> #<User id: 1, email: "user@example.com", ...>
>> user.replies.empty?
=> true
>> article = user.articles.first
=> #<Article id: 8, title: "Associations", ..., user_id: 1>
>> article.comments.create(:name => 'Guest',
:email => 'guest@example.com', :body => 'Great article!')
=> #<Comment id: 1, article_id: 8, name: "Guest", ...>
>> user.replies
=> [#<Comment id: 1, article_id: 8, name: "Guest", ...>]
>> user.replies.size
=> 1
```

Advanced Finding

Chapter 5 covered use of the `find` class method in Active Record. This section expands on different find operations using the `where` method. Building advanced finder methods is one of the most important things you do with your models.

Using the *where* Method

The most basic condition style is the hash syntax. Active Record takes the Hash passed to the `where` method and turns the keys into column names and the values into parameters to match. The hash syntax is useful only if you're trying to find an exact match. Run the following command in a console window to try out the hash syntax:

```
>> Article.where(:title => 'Advanced Active Record')
=> [#<Article id: 6, title: "Advanced Active Record", ...>]
```

The hash syntax works well for straightforward `where` operations where you use only ANDs to join together the conditions (that is, all conditions must match). However, sometimes you need more flexibility than exact matches.

Using an SQL Fragment

To specify conditions, you can pass in an SQL fragment as a string that is sent directly to the query. You need to have a pretty decent knowledge of SQL to use this kind of syntax; but it provides a lot of flexibility, and you can create arbitrarily complex SQL fragments if you're an SQL ninja.

Try the same find operation as in the previous section, but use a pure SQL condition fragment:

```
>> Article.where("title = 'Advanced Active Record'")
=> [#<Article id: 6, title: "Advanced Active Record", ...]
```

Let's try something more complicated that only SQL is able to do:

```
>> Article.where("created_at > '23-04-2013' OR body NOT LIKE '%model%'")
=> [#<Article id: 7, title: "One-to-many associations", ...,
#<Article id: 8, title: "Associations", ...]
```

Instead of using the `=` sign, you use the greater-than (`>`) symbol to make sure the date occurs after April 23, 2013. This is followed by the SQL `OR` operator, which says “if this first part isn’t a match, then try the right-hand side and give it a second chance at matching.” Therefore, you check the right-hand side only if the left-hand side fails. If an item fails the `created_at` match, you check to see if the body is `NOT LIKE` code. You can think of `OR` as a more permissive joining operator. It only cares that one of the conditions is a match. `OR` has a sister named `AND`, which requires that both conditions are true:

```
>> Article.where("created_at > '23-04-2013' AND body NOT LIKE '%model%'")
=> []
```

You also use the SQL `LIKE` (modified using `NOT`, for negation) operator, which allows you to make partial matches. Normally, when using `=`, SQL requires that the string match *perfectly*. However, `LIKE` is more permissive and allows partial matches when used with the `%` wildcard. The `%` symbols are SQL wildcard characters that apply in `LIKE` clauses. A `%` at the beginning of a pattern says that the pattern must match at the end of the field (the beginning can be any sequence of characters); a `%` at the end means that the pattern must match at the beginning, where the end can be any sequence of characters. Using a `%` on both sides of the pattern means that it must match anywhere in the field. Using `%model%` means that the word `model` must occur somewhere (anywhere) in the body of the article. In the previous example, you don't want articles that have the word `model`; therefore, an article with the sentence “I don't have your match” is accepted as a match.

As you can see, this usage has all the flexibility of SQL, but it also has SQL's natural limitations. For instance, you may need to find information based on what the user passes into the application via the request parameters in your application (Chapter 7 covers request parameters). If you aren't careful, those data can be very dangerous to your application, because they are open to SQL injection attacks. In such an attack, a user submits malicious code that tricks your database server into doing far more than you intended. For more information about SQL injection, check out the Wikipedia article at http://en.wikipedia.org/wiki/SQL_injection. Fortunately, Rails gives you a way to avoid such threats by using the array condition syntax, which performs correctly quoted replacements.

Using an Array Condition Syntax

The array condition syntax gives you the ability to specify conditions on your database calls in a safer way than using SQL syntax. Also, you don't need to worry as much about SQL specifics like quoting and other concerns, because it does automatic conversions for you on the inputs you give it. This is how it protects against SQL injection—it ensures that the substituted values are safely quoted, thereby preventing malicious users from injecting arbitrary SQL into your queries.

The following example requires the use of a nice little Ruby method called `Time.now`. Basically, it returns a `Time` object that is set to the current time. Let's see if you can find all the articles that were published before today:

```
>> Article.where("published_at < ?", Time.now)
=> [#<Article id: 6, title: "Advanced Active Record", ...>]
```

It doesn't look like much was published before today. Instead of writing in the date, you put a `?` in the spot where you'd normally write the value you want to find. The `where` method takes the second element in the array, `Time.now`, and replaces it where the first `?` appears. Additionally, the array syntax automatically takes your time and converts it into something your database likes. You can invoke the `to_sql` method after the `where` method to inspect the issued SQL statement:

```
>> Article.where("published_at < ?", Time.now).to_sql
=> "SELECT      \"articles\".* FROM      \"articles\""
WHERE      (published_at < '2013-04-02 16:27:51.059277')"
```

You give it a `Time` object, and it turns the object into the format that pleases your database. If you had passed it a string, it wouldn't have converted. You can even pass some information from another model:

```
>> Article.where("created_at = ?", Article.last.created_at)
=> [#<Article id: 8, title: "Associations", ...>]
```

That condition returns all the articles created at the same moment as the last article. You can pass as many conditions as you want, as long as they occur in the same order as the question marks:

```
>> Article.where("created_at = ? OR body LIKE ?", Article.last.created_at, 'model')
=> [#<Article id: 8, title: "Associations", ...>]
```

MONITORING THE LOGS

You can see the SQL statements issued by your application in the file `log/development.log`. It's often useful to monitor what the server is doing. You may have already noticed that when you run `rails server`, it tells you about what is going on in your application. However, different web servers (depending on what you've installed) give different outputs, some more descriptive than others.

Fortunately, Rails prints all of its activities to a log file. If you look in your `log` directory, you see `log/development.log`. This is the file where all the activities of your application are output. If you're running in production mode, the log file is `log/production.log`.

This file is written to live by your server. Sometimes it's useful (especially on a live server) to monitor the events occurring on your server. If you're on a UNIX system, you can run the command `tail -f log/development.log` to get a live feed from your logs. If you're on a Windows system, you can find several applications that behave like `tail` with a quick Google search.

During debugging, it can be useful to output messages to the log to see what's going on with your application. Almost anywhere in your application, you can type this:

```
Rails.logger.debug "This will only show in development"
Rails.logger.warn "This will show in all environments"
```

Both of these messages print directly to the log file and can be extremely useful for figuring out what is happening with your server.

The main disadvantage with the array syntax is that it can become confusing to remember the order of the elements you're passing in for the conditions.

Instead of adding a list of things at the end of the array, you can pass in a hash and change the question marks to actual named replacements. This can help you keep the order of your arguments straight:

```
>> Article.where("title LIKE :search OR body LIKE :search",
{:search => '%association%'})
=> [#<Article id: 7, title: "One-to-many associations", ...>,
#<Article id: 8, title: "Associations", ...>]
```

As you can see, you can reuse the same term in multiple places in your condition. If you were using the regular array syntax, you'd have to pass the same value '%association%' twice. This is especially useful if you have many, many conditions.

Using Association Proxies

Association proxy is a fancy term for the ability to chain together multiple calls to Active Record. You've been using this technique throughout the book, but it hasn't received special mention. Here is a basic example of association proxies:

```
>> User.first.articles.all
=> [#<Article id: 8, title: "Associations", ...>]
```

This code returns all the articles of the first user. The `all` method (off `articles`) is automatically scoped to the user, which is to say it finds articles that belong to that user. If you recall, `articles` is a `has_many` relationship on the `User` model. The cool part is that this isn't two queries to the database. It does all this in one request to the database.

Scoped finders are also more secure. Imagine a multiple-user system where data owned by one user shouldn't be accessible by another user. Finding an associated object (say, an article) by its `id` doesn't restrict it to articles owned by a particular user. You could pass in the `article_id` and the `user_id` as conditions, but that's sloppy and prone to error. The correct way to do this is to scope all find operations off the user in question. For example, assuming you have a `User` object stored in the variable `current_user`; `current_user.articles.find(1)` ensures that the article with `id 1` is returned only if it belongs to the `current_user`.

Anyone who has done database work will realize that this incredibly simple syntax is far easier than the SQL queries that need to be created to achieve similar results. If you play around with these chains, you can check out the log to see the SQL that's generated—be happy that you didn't have to write it yourself!

This technique doesn't just apply to finding. You can use it to automatically assign ownership with `build` and `create` constructors by setting the appropriate foreign keys. Consider the following example, which creates a new article for the `current_user`. It automatically sets the article's `user_id` to that of the current user:

```
current_user.articles.create(:title => 'Private', :body => 'Body here..')
```

This is much better than the alternative, which is to go through the `Article` model directly and set the `user_id` as an attribute (`Article.create(:user_id => current_user.id)`). As a rule, whenever you need to restrict find operations to an owner, or if you're assigning ownership, you should use the power of the association proxy.

Other Finder Methods

Active Record ships with other finder methods that complement the `where` method and can be used on their own as well. Table 6-6 lists some of those methods with a brief description and a quick example.

Table 6-6. Some Active Record Finder Methods

Method	Description	Example
where(conditions)	Specifies the conditions in which the records are returned as a WHERE SQL fragment	Article.where("title = 'Advanced Active Record'")
order	Specifies the order in which the records are returned as an ORDER BY SQL fragment	Article.order("published_at DESC")
limit	Specifies the number of records to be returned as a LIMIT SQL fragment	Article.limit(1)
joins	Specifies associated tables to be joined in as a JOIN SQL fragment	Article.joins(:comments)
includes	Specifies associated tables to be joined and loaded as Active Record objects in a JOIN SQL fragment	Article.includes(:comments)

You've used the `where` method before. Let's take the rest for a spin:

```
>> Article.all
=> [#<Article id: 6, title: "Advanced Active Record", ...>,
#<Article id: 7, title: "One-to-many associations", ...>,
#<Article id: 8, title: "Associations", ...>]
>> Article.order("title ASC")
=> [#<Article id: 6, title: "Advanced Active Record", ...>,
#<Article id: 8, title: "Associations", ...>,
#<Article id: 7, title: "One-to-many associations", ...>]
>> Article.limit(1)
=> [#<Article id: 6, title: "Advanced Active Record", ...>]
>> Article.order("title DESC").limit(2)
=> [#<Article id: 7, title: "One-to-many associations", ...>,
#<Article id: 8, title: "Associations", ...>]
```

You first retrieve a list of articles with `all`; then, you retrieve all articles ordered alphabetically by their title using the `order` method. After that, you retrieve a single article using the `limit` method. Finally, you chain the `limit` method to `order` to retrieve a couple of articles after sorting them. All methods listed in Table 6-6 are chainable; when you chain finder methods to one another, Rails combines their specifics to form a single query to the database.

Default Scope

As you write applications, you may notice that you repeat certain conditions many times throughout your code. For the blog application, it would make sense to display categories in alphabetical order, as the user would expect. Rails provides a technique called *scope* to encapsulate commonly used find operations. Rails doesn't enforce a default order; it lets the database take care of sorting the results, which in most cases is done on the primary key `id`. Let's look at how your `Category` records are returned now:

```
>> Category.all
=> [#<Category id: 1, name: "Programming", ...>, #<Category id: 2, name: "Event", ...>,
#<Category id: 3, name: "Travel", ...>, #<Category id: 4, name: "Music", ..>,
#<Category id: 5, name: "TV", ...>]
```

As you can see, categories are returned according to their primary key `id`. Let's make sure categories are always listed alphabetically, regardless of the conditions you use for the query. The code in Listing 6-19 tells the `Category` class that you always want records to be ordered by the `name` field.

Listing 6-19. The `default_scope` Declaration in `app/models/category.rb`: <http://gist.github.com/325932>

```
class Category < ActiveRecord::Base
  has_and_belongs_to_many :articles

  default_scope lambda { order('categories.name') }
end
```

Notice that what you pass to the `default_scope` method is identical to the `order` method call in the previous section. As you may expect, you can pass any finder method to `default_scope`. Let's see the order in which your categories are retrieved now:

```
>> reload!
Reloading...
>> Category.all
=> [#<Category id: 2, name: "Event", ...>, #<Category id: 4, name: "Music", ...>,
#<Category id: 1, name: "Programming", ...>, #<Category id: 5, name: "TV", ...>,
#<Category id: 3, name: "Travel", ...>]
```

As you can see, your categories are sorted alphabetically by default.

Note When specifying the order for the default scope, you use the fully qualified name `categories.name` instead of just `name`. Whenever you use associations to retrieve categories, Rails runs an SQL command to join categories with other tables, which may also have a `name` field; using the table name ensure that your database always knows the field you really want to sort by.

Named Scope

The default scope is useful. But in most cases, the only code you want to have there is default ordering for your application, because adding a condition to `default_scope` would cause that condition be applied every time. For queries that you run often, you should create named scopes that make your code easier to read and maintain.

Let's create two named scopes: the first one lists all the articles with a `published_at` date and is named `:published`; the second scope lists all the articles without a `published_at` date and is named `:draft`. You create both scopes using the `scope` method, which takes the name of the scope as its first parameter and a finder method call as its second. Listing 6-20 shows the updated `Article` model.

Listing 6-20. Named Scopes Declarations in `app/models/article.rb`: <http://gist.github.com/325944>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments
```

```

scope :published, lambda { where("articles.published_at IS NOT NULL") }
scope :draft, lambda { where("articles.published_at IS NULL") }

def long_title
  "#{title} - #{published_at}"
end
end

```

As in a regular `where` method, you can use arrays as parameters. In fact, you can chain finder methods with other named scopes. You define the `recent` scope to give you articles recently published: first, you use the `published` named scope, then you chain to it a `where` call (Listing 6-21).

Listing 6-21. Recent Named Scope Declaration in `app/models/article.rb`: <http://gist.github.com/326095>

```

class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  scope :published, lambda { where("articles.published_at IS NOT NULL") }
  scope :draft, lambda { where("articles.published_at IS NULL") }
  scope :recent, lambda { published.where("articles.published_at > ?",
1.week.ago.to_date) }

  def long_title
    "#{title} - #{published_at}"
  end
end

```

Note `lambda` is a keyword in Ruby that defines a self-contained standalone method, which is executed only when you invoke it. You must use a `lambda` or some other object that responds to `call` when defining a scope.

To make scopes even more useful, you can define scopes that can receive parameters, instead of hard coding the values you want to query with. You need search functionality that allows the end user to look up articles by title; so, let's add another scope called `where_title` that accepts an argument and searches by it (Listing 6-22).

Listing 6-22. The `where_title` Named Scope Declaration in `app/models/article.rb`: <http://gist.github.com/326098>

```

class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

```

```

scope :published, lambda { where("articles.published_at IS NOT NULL") }
scope :draft, lambda { where("articles.published_at IS NULL") }
scope :recent, lambda { published.where("articles.published_at > ?",
1.week.ago.to_date) }
scope :where_title, lambda { |term| where("articles.title LIKE ?", "%#{term}%") }

def long_title
  "#{title} - #{published_at}"
end
end

```

Now that you've added those scopes, let's see them in action in a console session. When you look at the results of running the methods, you get an English-like syntax that makes the code easy to read and expand. Pay special attention to the line that uses Article.draft.where_title("one"), which shows how you chain scopes to get the exact data you want:

```

>> reload!
Reloading...
>> Article.published
=> [#<Article id: 6, title: "Advanced Active Record", ...>]
>> Article.draft
=> [#<Article id: 7, title: "One-to-many associations", ...>,
#<Article id: 8, title: "Associations", ...>]
>> Article.recent
=> [#<Article id: 6, title: "Advanced Active Record", ...>]
>> Article.draft.where_title("one")
=> [#<Article id: 7, title: "One-to-many associations", ...>]
>> Article.where_title("Active")
=> [#<Article id: 6, title: "Advanced Active Record", ...>]

```

Applying Validations

It's probably a safe bet that you don't want every field in your tables to be optional. Certain fields need to be required, terms of service agreements need to be accepted, and passwords need to be confirmed. That's just the way it is when you're building web applications, and Rails understands this. Consider this example of an Account model:

```

class Account < ActiveRecord::Base
  validates_presence_of :login
  validates_confirmation_of :password
  validates_acceptance_of :terms_of_service
end

```

Like associations, validations are sets of high-level macros that let you selectively apply common validation requirements to your model's attributes. In this section, you create a full set of validations for your blog application, and you see firsthand how easy it is to perform basic validations with Active Record. You start by applying some of the built-in validations, and then you build a couple custom validation methods.

Using Built-In Validations

Rails has myriad built-in validators, all of which are accessible through the `validates` method. Here you will learn about some of the options the `validates` method accepts as you apply them to your blog application. Check the API for details of all the Rails validators (<http://api.rubyonrails.org/classes/ActiveModel/Validations/ClassMethods.html>).

As a reference to get you started, you can pass two common options into any built-in validator. These are described in Table 6-7.

Table 6-7. Default Options for All Validators

Option	Description	Example
<code>:message</code>	Specifies the error message shown if validation fails.	<code>:message => "too long"</code>
<code>:on</code>	Specifies when this validation happens. The default is <code>:save</code> . Other options are <code>:create</code> and <code>:update</code> .	<code>:on => :create</code>

Validating That a Value Has Been Entered

You can use the `:presence` option to make sure a user has entered *something* into a field. This is very useful in many cases. You have those validations in the Article model for the title and body fields, as shown in Listing 6-23.

Listing 6-23. The Article Model, `validates_presence_of` Method in `app/models/article.rb`

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  scope :published, lambda { where("articles.published_at IS NOT NULL") }
  scope :draft, lambda { where("articles.published_at IS NULL") }
  scope :recent, lambda { published.where(["articles.published_at > ?",
    1.week.ago.to_date]) }
  scope :where_title, lambda { |term| where("articles.title LIKE ?", "%#{term}%") }

  def long_title
    "#{title} - #{published_at}"
  end
end
```

The default message is “can’t be blank.”

Validating That a Value Is Unique

Often, you want to ensure that a certain field is unique. The `:uniqueness` option validates whether the value of the specified attribute is unique across the system. You use this method in the User model to make sure each e-mail is unique, as shown in Listing 6-24.

Listing 6-24. The `validates_uniqueness_of` Method in `app/models/user.rb`:

```
class User < ActiveRecord::Base
  validates_uniqueness_of :email

  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments
end
```

When the record is created, a check is performed to ensure no record exists in the database with the given value for the specified attribute `email` (that maps to a column). When the record is updated, the same check is made, disregarding the record itself. The default error message is “`#{value}` has already been taken.”

The `:scope` option can also validate whether the value of the specified attributes is unique based on multiple parameters. For example, you can use it to ensure that a teacher is on the schedule only once per semester for a particular class:

```
class Schedule < ActiveRecord::Base
  validates_uniqueness_of :teacher_id, :scope => [:semester_id, :class_id]
end
```

Validating Length or Size

Sometimes you want to validate the length, or size, of a field entry. You can do this by using the `:length` option. You use this method in the User model to specify a valid number of characters for an e-mail address, as shown in Listing 6-25. The option for specifying a size range is `:within`.

Listing 6-25. The `validates_length_of` Method in `app/models/user.rb`:

```
class User < ActiveRecord::Base
  validates_uniqueness_of :email
  validates_length_of :email, :within => 5..50

  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments
end
```

If you want to ensure only the minimum or maximum, you can use the `:minimum` or `:maximum` option. Table 6-8 lists the most common `:length` validator’s options.

Table 6-8. Options for Validating :length

Option	Description
:minimum	Specifies the minimum size of the attribute
:maximum	Specifies the maximum size of the attribute
:is	Specifies the exact size of the attribute
:within	Specifies the valid range (as a Ruby Range object) of values acceptable for the attribute
:allow_nil	Specifies that the attribute may be nil; if so, the validation is skipped
:too_long	Specifies the error message to add if the attribute exceeds the maximum
:too_short	Specifies the error message to add if the attribute is below the minimum
:wrong_length	Specifies the error message to add if the attribute is of the wrong size
:message	Specifies the error message to add if :minimum, :maximum, or :is is violated

Validating the Format of an Attribute

The :format option checks whether a value is in the correct format. Using this method requires familiarity with regular expressions (regex) or being able to steal other people's regular expressions. The classic example (and the one you need) is e-mail. Update the validates method as shown in Listing 6-26.

Listing 6-26. Update validates :format Method in app/models/user.rb: <http://gist.github.com/326146>

```
class User < ActiveRecord::Base
  validates_uniqueness_of :email
  validates_length_of :email, :within => 5..50
  validates_format_of :email, :with => /^[^@][\w.-]+@[^\w.-]+[.][a-z]{2,4}$/i

  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments
end
```

Don't be put off by how complicated this looks. You pass in the :with option and a regex object to say what patterns you want to match.

TIP If you want to learn more about using regular expressions, you can find many tutorials and books on the subject. One good reference is Nathan Good's *Regular Expression Recipes* (Apress, 2004).

Validating Confirmation

Whenever a user changes an important piece of data (especially the password), you may want the user to confirm that entry by typing it again. This is the purpose of the validates_confirmation_of method. When you use this helper, you create a new virtual attribute called `#{field_name}_confirmation`. Add this to the User model for password confirmation, as shown in Listing 6-27.

Listing 6-27. The validates :confirmation Method in app/models/user.rb: <http://gist.github.com/326153>

```
class User < ActiveRecord::Base
  validates_uniqueness_of :email
  validates_length_of :email, :within => 5..50
  validates_format_of :email, :with => /^[^@][\w.-]+@[\\w.-]+[.][a-z]{2,4}$/i
  validates_confirmation_of :password

  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments
end
```

The password attribute is a column in the users table, but the password_confirmation attribute is virtual. It exists only as an in-memory variable for validating the password. This check is performed only if password_confirmation isn't nil and runs whenever the object is saved.

Other Validations

There is one other important validation method: validates_acceptance_of, which validates the acceptance of a Boolean field.

Building Custom Validation Methods

In the blog application, you'd like to make sure no one creates a comment for an article that hasn't been published yet. First, you need to create a method so you can ask an Article whether its published_at field is null by using the present? method, which returns true if a values exists and false otherwise. This method is useful outside validations, because you may want to indicate on the administration interface later whether an article has been published. Let's add that method now and call it published?. Add the code shown in Listing 6-28 to the Article model.

Listing 6-28. Adding the published? Method in app/models/article.rb: <http://gist.github.com/326170>

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  scope :published, lambda { where("articles.published_at IS NOT NULL") }
  scope :draft, lambda { where("articles.published_at IS NULL") }
  scope :recent, lambda { published.where("articles.published_at > ?",
    1.week.ago.to_date) }
  scope :where_title, lambda { |term| where("articles.title LIKE ?", "%#{term}%") }

  def long_title
    "#{title} - #{published_at}"
  end

  def published?
    published_at.present?
  end
end
```

This gets you a step closer to your goal. When building validations, Active Record gives you nice objects called `errors` to use. Whenever you want to add a validation error to the list of errors, you just type `errors.add(column_name, error_message)`. So, let's implement a method called `article_should_be_published` in the `Comment` class that uses this functionality, as shown in Listing 6-29.

Listing 6-29. Adding the `article_should_be_published` Method in `app/models/comment.rb`:

```
class Comment < ActiveRecord::Base
  belongs_to :article

  def article_should_be_published
    errors.add(:article_id, "is not published yet") if article
    && !article.published?
  end
end
```

This checks whether you should apply the error by evaluating the `if` statement. If that `if` statement is true, you want to add an error into the `errors` object. Note that before you test whether the article is published, you make sure `article` isn't `nil`. This is so your test doesn't throw an error. If `article` is `nil`, that should be handled by another validator: the `validates_presence_of` method.

How do you tell Active Record that this method should be run before a save? You use the `validate` class method and pass it a symbol with the name of the method. At the top of your `Comment` class, add the code shown in Listing 6-30. Note that we also expect comments to have values for `name`, `email`, and `body`; so, we add a presence validation call.

Listing 6-30. The `validate` Method in `app/models/comment.rb`:

```
class Comment < ActiveRecord::Base
  belongs_to :article

  validates_presence_of :name, :email, :body
  validate :article_should_be_published

  def article_should_be_published
    errors.add(:article_id, "is not published yet") if article && !article.published?
  end
end
```

This advises Active Record to pay attention to your new `article_should_be_published` method. In Chapter 11 you write tests to make sure this is working. But you can also go to the console—if you have it open already, don't forget to `reload!`—and try to create an invalid object to see if it reports errors for you. The easiest way to get to errors in an Active Record object is with `comment.errors.full_messages`, as shown here:

```
>> article = Article.draft.first
=> #<Article id: 7, title: "One-to-many associations", ...>
>> comment = article.comments.create :name => 'Dude',
:email => 'dude@example.com', :body => 'Great article!'
=> #<Comment id: nil, article_id: 7, name: "Dude", email: "dude@example.com",
body: "Great article!", created_at: nil, updated_at: nil>
>> comment.errors.full_messages
=> ["Article is not published yet"]
```

Making Callbacks

You often want to have certain things happen during the lifecycle of the model. Certain actions need to happen during certain events pertaining to a particular model. For instance, what if you want to send an e-mail to your administrator whenever someone cancels an account? Or perhaps you want to make sure to create a new model because some other model was also created. Sometimes, certain actions in the life of a model should execute associated actions.

To implement this, Active Record has *callbacks*. Six callbacks are commonly used in Active Record models:

- `before_create`
- `after_create`
- `before_save`
- `after_save`
- `before_destroy`
- `after_destroy`

As you can see, the names of the Rails callbacks describe their purpose. When you create a method with any of these names in your model, the method is called automatically by the model during the time the name suggests. For instance, if you make a `before_save` method, that method is called right before the model object is saved.

Any callback that starts with `before_` can stop the execution chain if it returns `false`. For instance, if you define `before_create`, you ensure that this model object will *never* be created:

```
def before_create
  false
end
```

This can be a gotcha later if you're doing something like an assignment of `false` to a variable. If you're ever confused why a model won't save, check your `before_filters`.

In the blog application, you'd like to make sure that when a user creates a comment, an e-mail is automatically sent to the article author. Although you don't send an e-mail here, this chapter goes over the steps required to put together code to eventually send the e-mail in Chapter 10. To set this up, you add an `after_create` method to the `Comment` class that will eventually have the code to send an e-mail. Add the method shown in Listing 6-31 to the `Comment` model.

Listing 6-31. Adding `after_create` Method in `app/models/comment.rb`

```
class Comment < ActiveRecord::Base
  belongs_to :article

  validates_presence_of :name, :email, :body
  validate :article_should_be_published

  def article_should_be_published
    errors.add(:article_id, "is not published yet") if article
    && !article.published?
  end

  def after_create
    puts "We will notify the author in Chapter 9"
  end
end
```

You use the code you want to be executed directly in the code of the `after_create` method. This is nice and simple, but you should probably use the pattern as you did for `validate` in Listing 6-30, where you pass in a symbol that references the method to run when the validation is performed. This helps keep the code readable and easier to augment in the future, because you can supply an arbitrary number of methods to run on a callback, separated by a comma. Name the method `email_article_author`, and tell Active Record to run it after a record is created, as shown in Listing 6-32.

Listing 6-32. The `email_article_author` Method Specified as an `after_create` Callback in `app/models/comment.rb`: <http://gist.github.com/326211>

```
class Comment < ActiveRecord::Base
  belongs_to :article

  validates_presence_of :name, :email, :body
  validate :article_should_be_published

  after_create :email_article_author

  def article_should_be_published
    errors.add(:article_id, "is not published yet") if article
    && !article.published?
  end

  def email_article_author
    puts "We will notify #{article.user.email} in Chapter 9"
  end
end
```

Active Record provides many more callbacks than are mentioned here, but those listed at the beginning of this section are the ones you'll find yourself using often. Some of the others are used in extremely rare cases (for instance, `after_initialize`, which is called after an object is initialized). These callbacks can help you with just about anything you need to do during the lifecycle of a model. They're part of smart models, which know how to deal with their own birth, life, and death.

Updating the User Model

You still need to do a little work on your `User` model. You can apply many of the techniques described in this chapter, such as custom methods to allow you to perform user authentication, and validation methods to make sure your data stay clean.

When you created the user migration (Listing 6-2), you added a field called `password`. This field stores a plain-text password, which, if you think about it, isn't very secure. It's always a good idea to encrypt any sensitive data so they can't be easily read by would-be intruders. You deal with the encryption in the `User` model itself, but the first thing you do is rename the field in the database from `password` to `hashed_password`. This is so you can create a custom accessor called `password` with which to set the password while maintaining a field to store the encrypted version in the database. The plain-text password is never saved.

To accomplish this, you create a migration. From the terminal, issue the following command to create the new migration:

```
$ rails generate migration rename_password_to_hashed_password
```

Next, fill in the migration as shown in Listing 6-33.

Listing 6-33. Migration to Rename password to hashed_password in db/migrate/20130409052120_rename_password_to_hashed_password.rb: <http://gist.github.com/326261>

```
class RenamePasswordToHashedPassword < ActiveRecord::Migration
  def change
    rename_column :users, :password, :hashed_password
  end
end
```

Run the migration using the `rake db:migrate` command, as follows:

```
$ rake db:migrate
```

```
--> RenamePasswordToHashedPassword: migrating =====
-- rename_column(:users, :password, :hashed_password)
-> 0.0558s
--> RenamePasswordToHashedPassword: migrated (0.0560s) =====
```

Next, update your User model so it looks like that in Listing 6-34. You program all the user-authentication methods you need for allowing users to log in. Let's look at the code first and then see in detail what you've done.

Listing 6-34. Current User Model in app/models/user.rb: <http://gist.github.com/326271>

```
require 'digest'
class User < ActiveRecord::Base
  attr_accessor :password
  validates_uniqueness_of :email
  validates_length_of :email, :within => 5..50
  validates_format_of :email, :with => /^[^@][\w.-]+@[^\w.-]+\.[a-z]{2,4}$/i

  validates_confirmation_of :password
  validates_length_of :password, :within => 4..20
  validates_presence_of :password, :if => :password_required?

  has_one :profile
  has_many :articles, -> { order('published_at DESC, title ASC') },
    :dependent => :nullify
  has_many :replies, :through => :articles, :source => :comments

  before_save :encrypt_new_password

  def self.authenticate(email, password)
    user = find_by_email(email)
    return user if user && user.authenticated?(password)
  end

  def authenticated?(password)
    self.hashed_password == encrypt(password)
  end
```

```

protected
  def encrypt_new_password
    return if password.blank?
    self.hashed_password = encrypt(password)
  end

  def password_required?
    hashed_password.blank? || password.present?
  end

  def encrypt(string)
    Digest::SHA1.hexdigest(string)
  end
end

```

Note The SHA1 hashing algorithm used in this example is weak and was only used for an example. For production web sites you should take a look at the bcrypt gem (<https://github.com/codahale/bcrypt-ruby>).

Whenever you store something sensitive like a password, you should encrypt it. To encrypt the password in your User model, you use a simple algorithm called a *hash* that creates a random-looking string from the provided input. This hashed output can't be turned back into the original string easily, so even if someone steals your database, they will have a prohibitively difficult time discovering your users' passwords. Ruby has a built-in library called Digest that includes many hashing algorithms.

Let's go through the additions to the User model:

- `require 'digest'`: You start by requiring the Digest library you use for encrypting the passwords. This loads the needed library and makes it available to work within your class.
- `attr_accessor :password`: This defines an accessor attribute, `password`, at the top of the class body. It tells Ruby to create reader and writer methods for `password`. Because the `password` column doesn't exist in your table anymore, a `password` method isn't created automatically by Active Record. Still, you need a way to set the password before it's encrypted, so you make your own attribute to use. This works like any model attribute, except that it isn't persisted to the database when the model is saved.
- `before_save :encrypt_new_password`: This `before_save` callback tells Active Record to run the `encrypt_new_password` method before it saves a record. That means it applies to all operations that trigger a save, including `create` and `update`.
- `encrypt_new_password`: This method should perform encryption only if the `password` attribute contains a value, because you don't want it to happen unless a user is changing his password. If the `password` attribute is blank, you return from the method, and the `hash_password` value is never set. If the `password` value isn't blank, you have some work to do. You set the `hashed_password` attribute to the encrypted version of the password by laundering it through the `encrypt` method.
- `encrypt`: This method is fairly simple. It uses Ruby's Digest library, which you included on the first line, to create an SHA1 digest of whatever you pass it. Because methods in Ruby always return the last thing evaluated, `encrypt` returns the encrypted string.

- `password_required?`: When you perform validations, you want to make sure you're validating the presence, length, and confirmation of the password only if validation is required. And it's required only if this is a new record (the `hashed_password` attribute is blank) or if the `password` accessor you created has been used to set a new password (`password.present?`). To make this easy, you create the `password_required?` predicate method, which returns true if a password is required or false if it's not. You then apply this method as an `:if` condition on all your password validators.
- `self.authenticate`: You can tell this is a class method because it's prefixed with `self` (it's defined on the class *itself*). That means you don't access it via an instance; you access it directly off the class, just as you would `find`, `new`, or `create` (`User.authenticate`, not `@user = User.new; @user.authenticate`). The `authenticate` method accepts an e-mail address and an unencrypted password. It uses a dynamic finder (`find_by_email`) to fetch the user with a matching e-mail address. If the user is found, the `user` variable contains a `User` object; if not, it's `nil`. Knowing this, you can return the value of `user` if, and only if, it isn't `nil` and the `authenticated?` method returns true for the given password (`user && user.authenticated?(password)`).
- `authenticated?`: This is a simple predicate method that checks to make sure the stored `hashed_password` matches the given password after it has been encrypted (via `encrypt`). If it matches, `true` is returned.

Let's play with these new methods from the console so you can get a better idea of how this comes together:

```
>> user = User.first
=> #<User id: 1, email: "user@example.com", ...>
>> user.password = 'secret'
=> "secret"
>> user.password_confirmation = 'secret'
=> "secret"
>> user.save
=> true
>> user.hashed_password
=> "e5e9fa1ba31ecd1ae84f75caaa474f3a663f05f4"
>> User.authenticate('user@example.com', 'secret')
=> #<User id: 1, email: "user@example.com", ...>
>> User.authenticate('user@example.com', 'secret2')
=> nil
>> second_user = User.last
=> #<User id: 2, email: "mary@example.com", ...>
>> second_user.update_attributes(:password => 'secret',
:password_confirmation => 'secret')
=> true
>> User.authenticate('mary@example.com', 'secret')
=> #<User id: 2, email: "mary@example.com", ...>
```

When you ask the `User` model to authenticate someone, you pass in the e-mail address and the plain-text password. The `authenticate` method hashes the given password and then compares it to the stored (hashed) password in the database. If the passwords match, the `User` object is returned, and authentication was successful. When you try to use an incorrect password, `nil` is returned. In Chapter 8 you write code in your controller to use these model methods and allow users to log in to the site. For now, you have a properly built and secure backend for the way users authenticate.

With the validation in the User model, the db/seeds.rb file also needs to be updated to make sure it follows the rules expected in the model. While we are at it, we also add some code to create a few articles. Update your db/seeds.rb file so that it looks like Listing 6-35.

Listing 6-35. Current Seeds File in db/seeds.rb: <http://gist.github.com/387374>

```
user = User.create :email => 'mary@example.com',
                  :password => 'guessit',
                  :password_confirmation => 'guessit'
Category.create [{:name => 'Programming'},
                 {:name => 'Event'},
                 {:name => 'Travel'},
                 {:name => 'Music'},
                 {:name => 'TV'}]
user.articles.create :title => 'Advanced Active Record',
                     :body => "Models need to relate to each other. In the real world, ..",
                     :published_at => Date.today
user.articles.create :title => 'One-to-many associations',
                     :body => "One-to-many associations describe a pattern ..",
                     :published_at => Date.today
user.articles.create :title => 'Associations',
                     :body => "Active Record makes working with associations easy..",
                     :published_at => Date.today
```

Reviewing the Updated Models

You've made a lot of changes to your models, so let's make sure we're on the same page before you move on. Look at the Article, Category, and Comment models in Listings 6-36, 6-37, and 6-38, respectively, and make sure yours match.

Listing 6-36. Current Article Model in app/models/article.rb

```
class Article < ActiveRecord::Base
  validates_presence_of :title
  validates_presence_of :body

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  scope :published, lambda { where("articles.published_at IS NOT NULL") }
  scope :draft, lambda { where("articles.published_at IS NULL") }
  scope :recent, lambda { published.where("articles.published_at > ?",
1.week.ago.to_date) }
  scope :where_title, lambda { |term| where("articles.title LIKE ?", "%#{term}%") }

  def long_title
    "#{title} - #{published_at}"
  end
```

```
def published?  
  published_at.present?  
end  
end
```

Listing 6-37. Current Category Model in app/models/category.rb

```
class Category < ActiveRecord::Base  
  has_and_belongs_to_many :articles  
  
  default_scope lambda { order('categories.name') }  
end
```

Listing 6-38. Current Comment Model in app/models/comment.rb

```
class Comment < ActiveRecord::Base  
  belongs_to :article  
  
  validates_presence_of :name, :email, :body  
  validate :article_should_be_published  
  
  def article_should_be_published  
    errors.add(:article_id, "is not published yet") if article  
    && !article.published?  
  end  
end
```

Summary

After reading this chapter, you should have a complete understanding of Active Record models. The chapter covered associations, conditions, validations, and callbacks at breakneck speed. Now the fun part starts. In the next chapter, you get to use all the groundwork established in this chapter to produce the web interface for the data structures you've created. This is when you get to reap the benefits of your hard work.



Action Pack: Working with the View and the Controller

When you type a URL into your browser's address bar and press Enter, a few things happen behind the scenes. First, the domain name is translated into a unique address by which the server that hosts the application can be identified. The request is then sent to that server, which begins a chain of events that culminates in a response. The response is usually, but not always, in the form of an HTML document, which is essentially a text document full of special codes that your browser understands and can render visually on your screen. At this point, the request cycle is complete, and the browser waits for further input from you. If you click a link somewhere on the page or type a new URL in the address bar, the cycle begins all over again: the request is sent, the server processes it, and the server sends back the response.

When you make a request to a Rails application, this request cycle is the responsibility of a component of Rails called Action Pack. The Action Pack library is an integral component of the Rails framework and one that you need to be familiar with if you intend to master Rails.

This chapter begins with an overview of Action Pack. Then, you get to work using it in your sample blog application.

Note If you need to get the code at the exact point where you finished Chapter 6, download the source code zip file from the book's page on <http://www.apress.com> and extract it on your computer.

Action Pack Components

You've been introduced to the MVC pattern, but if you need a refresher, here it is. The *model* is your application's world, most often represented by database objects like articles, comments, and subscribers. The *controller* is the grand orchestrator, dealing with requests and issuing responses. The *view* is the code that contains instructions for rendering visual output for a browser, like HTML.

Armed with this refresher, you may be able to guess what roles are played by Action Pack. This isn't a test, so here's the answer: Action Pack is the controller and the view. The controller performs the logic, and the view renders the template that is given back to the requesting browser. Not surprisingly, two of the modules that make up the Action Pack are named accordingly: Action Controller and Action View.

At this point, you may be wondering why the view and the controller are wrapped up in a single library, unlike models, which have a library of their own. The answer is subtle and succinct: controllers and views are very closely related. The sections that follow paint a more complete picture of both the role and the relationship of controllers and views, how they work, and how they work together to create and control the interface of a Rails application.

Action Controller

Controllers orchestrate your application's flow. Every time a user requests a page, submits a form, or clicks a link, that request is handled—in one way or another—by a controller. When you're programming your application, you spend a lot of time building controllers and giving them instructions on how to handle requests.

The concept of controllers can sometimes be difficult for newcomers to grasp. Even if you've built web applications before, say in ASP (Active Server Pages) or PHP (Hypertext Preprocessor), you may not be used to this form of separation, where the mechanics of flow are controlled by a separate entity and not embedded in the pages themselves.

Let's look at the example of the CD player in a car to illustrate the concept of controllers. The player is required to respond to certain events, such as the user pressing the play button, fast forwarding, or rewinding a track. When you push a button, you expect something to happen—you've made a request, and you wait for the subsequent response.

If your CD player were a Rails application, the instructions for what to do when a certain event takes place, such as the pressing the eject button, would be contained in a controller. If you were to sketch it on paper, it might look something like this:

- CD Player
- Play
- Stop
- Fast-forward
- Rewind
- Eject

These events, or actions, describe what the player should be capable of *doing*. Obviously, each of these actions would need to be programmed to do something with the disc inside the player. When someone presses eject, you would first call on the stop action (if the disc is playing) and then arrange for the player to spit out the disc. You would code all the instructions for dealing with an eject event into the controller—specifically, inside the eject action. The same would apply for play, fast-forward, and rewind.

It's worth noting that this type of logic has nothing to do with the CD itself, nor does it have anything to do with the music on the CD. If this were a Rails application, the CD would be the model. It can be used independently of the player. In fact, it can be used in all sorts of players, not just the one in your car.

The stereo in your car is probably capable of more than just playing CDs. Most stereos have a radio receiver built in as well. The radio would have its own set of events that would likewise need to be handled. These actions might include things like changing stations, setting presets, and switching between AM and FM. To keep things well organized, you would probably want to group these actions inside their own controller, separate from the CD controller. After all, the radio and the CD player do different things.

When you're dealing with a Rails application, it's not much different. You separate the things that you need your application to do with an object from the object itself. Even when you're not dealing directly with an object (adjusting the volume on your car stereo has little to do with either the CD in the player or the station on the radio), you still handle the event inside a controller.

Each controller in Rails is designed as a Ruby class. Without getting too technical, Listing 7-1 shows how the CD player example would look if it were a Ruby class.

Listing 7-1. CDPlayer Class

```
class CDPlayer
  def play
  end

  def stop
  end
```

```

def fast_forward
end

def rewind
end

def eject
end
end

```

Inside the `CDPlayer` class, you define a method for each action, or each thing you want your CD player to be able to do. So, if you were to send the message “play” to an instance of the `CDPlayer` class, it would know how to handle it (of course, because the `play` method is empty in this example, nothing would happen). On the other hand, if you sent the message “pause,” Ruby would raise an exception and tell you that the method wasn’t found. If you wanted `CDPlayer` objects to respond to that message, you would need to add a method called (you guessed it) `pause`.

All the methods in this class are public, which means they can be invoked by anyone. You don’t need to do anything special to a method to make it public. *Unless otherwise declared, all methods in a Ruby class are public by default.* If you were to mark an action as private, though, it could be used only internally by the class. For example, if the `stop` method were private, it would raise a `NoMethodError` if you called it from outside the `CDPlayer` class. However, the `eject` method is free to call on `stop`, because it does so internally. Although the usefulness of this feature will become apparent as you continue to learn about controllers, consider this: if your CD player needed to display the time remaining for a given track, it might need to perform a few calculations to figure that out. You might create a method for doing these internal calculations, but would you want that method to be accessible from the outside? Would you have a button called `Calculate` on your player?

It’s time for a working definition: *Action Controllers are Ruby classes containing one or more public methods known as actions.* Each action is responsible for responding to a request to perform some task. A typical controller is most often a collection of actions that relates to a specific area of concern. For example, consider the `blog` application you’ve been building in the previous chapters. The controller that manages articles has the class name `ArticlesController` and has action methods for listing, creating, reading, updating, and deleting articles.

The example of the CD player worked well to illustrate the basic concept of controllers, but it won’t take you much further when dealing with web applications. If you were really dealing with a CD player, you would press play, the disc would start playing, and that would be the end of it. But because Rails was specifically designed for building web applications, it makes a fair number of assumptions about what you want your actions to do when they’re finished firing. Chief among these is the rendering of a view.

Imagine that you’re reading a list of posts on someone’s blog. You click the title of a post, and you expect to be taken to a new screen that shows you just that post. You requested an action (`show`), and in response, you receive a new screen. This happens all the time in the world of web applications: when you click a link, you expect to go to a new page.

In Rails, it’s the general case that when actions have completed their work, they respond by rendering a view. The concept of actions rendering views is so common that Rails has internalized it as a convention: *unless otherwise stated, when an action is finished firing, it renders a view.* How does Rails know what view to render if you don’t tell it? It looks for a view whose name matches that of the requested action. This should give you some insight as to why Action Controller and Action View are bundled together in Action Pack. Because of the way controller actions relate to views, a few other mechanisms facilitate their communication, all of which are covered shortly.

Action View

The Action View library is the second half of Action Pack. Given that controllers are responsible for handling the request and issuing a response, views are responsible for rendering the output of a response in a way a browser (or any other user agent) can understand. Let’s say you request the `index` action from the `ArticlesController`. After performing the logic to retrieve a list of articles, the controller hands off to the view, which formats the list of articles to make them look pretty. The controller then collects the results of the render, and the HTML is sent back to the browser, thus completing the request cycle.

Although the controller and the view are separate entities, they need to communicate with each other. The primary mechanism by which they do this is through shared variables. These shared variables are called *instance variables* and are easy to spot in Ruby because they're prefixed with the @ symbol. Keep this in mind as you look at the view example in Listing 7-2, which uses an instance variable called @articles to produce an article listing.

Listing 7-2. An Example View

```
<html>
<body>
<ul>
  <% @articles.each do |article| %>
    <li><%= article.title %></li>
  <% end %>
</ul>
</body>
</html>
```

Even without knowing any Ruby, you should be able to guess what this code does: it iterates over the collection of articles stored in the variable @articles and prints the title of each between HTML list-item () tags. If @articles contained three articles whose titles were One, Two, and Three, respectively, the preceding code would be compiled to the following:

```
<html>
<body>
<ul>
  <li>One</li>
  <li>Two</li>
  <li>Three</li>
</ul>
</body>
</html>
```

You may wonder where the variable @articles came into being. If you guessed in the controller, you would be right. The controller sets up instance variables that the view can access. In this case, the controller created a variable called @articles, and the view was given automatic access to it. Notice that the view doesn't perform any logic to fetch the list of articles; it relies on the controller to have set up the variable and performs the display logic necessary to turn the collection into a browser-ready HTML list.

Embedded Ruby

The code you see mixed into the HTML markup is Ruby. Because templates that are capable of dealing only with static HTML wouldn't be very useful, Action View templates have the benefit of being able to use Embedded Ruby (ERb) to programmatically enhance them.

Using ERb, you can embed Ruby into your templates and give them the ability to deal with data from the controller to produce well-formed HTML representations. ERb is included in the Ruby standard library, and Rails makes extensive use of it. You trigger ERb by using embeddings such as <% %> and <%= %> in your template files to evaluate or print Ruby code, respectively. If you've ever worked with ASP, JSP (Java Server Page), or PHP, this style of embedding should be familiar to you.

In the example in the preceding section, the loop is constructed within *evaluation embedding* tags (<% %>), and the article's title is printed using *output embedding* tags (<%= %>). Pay close attention to the subtle difference between

the two embedding types: output embedding includes an equals sign; regular embedding doesn't. When you use output embedding, you're effectively saying *print the results of the Ruby code when it's evaluated*. Regular embedding doesn't print results; it evaluates whatever is between the tags and goes on its merry way. If you mistakenly omit the equals sign, no errors are raised, but nothing is printed either. You have a set of empty list tags.

Note Following the Model behavior, Rails is modular and can be used with other templating libraries. A popular alternative is the Haml (<http://haml-lang.com>) template language used by many Rails developers.

Helpers

The terms of the MVC are fairly strict in the way they advocate the separation of components. Controllers really shouldn't concern themselves with the generation of view code, and views shouldn't concern themselves with anything but the simplest of logic. Although it's possible to use ERb to execute arbitrary Ruby code inside a view, and although controllers are certainly capable of generating markup, it's generally considered in violation of the MVC pattern to do so. This is where helpers come in to play.

Action Pack's *helpers* do exactly what their name implies: they help views by providing a convenient location to encapsulate code that would otherwise clutter the view and violate the terms of the MVC. They offer a middle ground between controllers and views and help to keep your application organized and easy to maintain.

If you think about it, ERb tags really aren't the best place for performing complex logic, and templates can quickly become unwieldy when creating markup programmatically. For this reason, Action Pack includes a large suite of built-in helpers for generating all sorts of HTML fragments—from creating forms and formatting dates to making hyperlinks and image tags. And when the built-in helpers aren't enough, you can write your own. Each controller gets its own helper module that's mixed in automatically, ready to lend your templates a hand when they need it.

Routing

All the information pertaining to which controller and action to call on comes in the form of the request URL. Action Pack includes a specialized component called *routing*, which is responsible for dissecting the incoming URL and delegating control to the appropriate controller and action.

Every request that comes into your web application originates in the form of a URL. The routing system allows you to write the rules that govern how each URL is picked apart and handled.

A traditional URL contains the path to a file on the server, relative to the server's home directory. Here's an example:

<http://example.com/articles/show.asp?id=1037>

You can tell a lot from this URL. First, you know the server technology being used is Microsoft's ASP. Given that, you also know that this URL resolves to the `show.asp` script, which is inside the `/articles` directory. In this case, there is no URL rewriting going on; the mapping of the URL to the script that handles it is one to one.

The problem with this kind of mapping is that you have no control over the URL. The URL is coupled to the script. What if you want to invoke the `show.asp` script but want the URL to read `articles/detail.asp` instead of `show.asp`? Or better yet, what if you don't want to expose the underlying script implementation (ASP) at all and use just `articles/detail`? There's no way. The lack of flexibility in this kind of URL mapping is a problem. If you ever need to change the name of the script being invoked, you instantly break all the URL references. This can be a major pain if you need to update all your code, especially if your pages are indexed by search engines.

Action Pack's routing solves this problem by decoupling the URL from the underlying program implementation. In Rails, the URL is related to the specific resource being requested, and it can be formatted to correctly identify that resource without having to conform to the name of the script that does the handling. When

thought of in this way, URLs become part of the interface of an application, unrelated to the files that are ultimately invoked to process a request.

There are myriad reasons why a routing system is a good idea. Here are just a few of them:

- Decoupled URLs can convey meaning, becoming part of the interface.
- Clean, readable URLs are more user-friendly and easier to remember.
- URLs can be changed without affecting the underlying implementation.

Of course, like most things in Rails, the routing system is open to configuration; and one of the great benefits of routes is that because they're decoupled, they can be customized to create meaningful URLs without much effort. This chapter teaches you how to build and customize routes for your application, understand the default routes that Rails creates for you, create named routes, and use routes when creating links and redirects in your code.

RESTful Resources

Rails adapted RESTful design as a convention in Rails 1.2 onward. Representational State Transfer (REST) is a principle used mainly over the HTTP protocol to offer a better interface for client-server operations. This section first discusses the REST concept and then explains how Rails implemented it through RESTful controllers and resources.

The REST principle is based on working with information in the form of *resources*. Each piece of information is dealt with as a *resource*, each resource has a unique interaction point for every *action* that can be performed on it, and each interaction point (action) is normally represented using a URL and a request method.

For example, think of a blog, which is a collection of information resources. Every article is a resource, and every action you perform on it, such as read, edit, or delete, has its own interaction point, mainly identified by a URL and a request method.

HTTP protocol, which is the main web protocol you normally use in browsers, has several request methods. These are the primary ones used in RESTful design:

- POST: Normally used to *submit* forms and *new* resource data
- GET: Mainly used to request a page to *view* a resource or more
- PATCH/PUT: Used to modify specific resource
- DELETE: Used to delete a resource

Do those methods remind you of anything? If you're thinking of CRUD, then you're right. Taking the main database operations create, read, update, and delete (CRUD) in REST design and tying them up with HTTP's main methods gives you what's called a *RESTful* web service.

RESTful web services are commonly used in APIs (referred to as REST APIs) by associating every CRUD method with its corresponding HTTP method:

- POST/Create: Creates of a resource
- GET/Read: Requests a specific resource or group of resources
- PATCH/PUT/Update: Edits attributes of a resource
- DELETE/Delete: Deletes a resource

Rails implemented RESTful design for controllers by introducing the concept of resources. Every model in your application is dealt with via a controller as a resources set, and that RESTful controller has certain methods that handle your regular operations on that model. We'll examine that in depth after you understand the Action Pack request cycle.

Action Pack Request Cycle

The entire request-to-response process is called the Action Pack *request cycle*. The request cycle consists of the following steps:

1. Rails receives a request from the outside world (usually a browser).
2. Routing picks apart the request to determine the controller and action to invoke.
3. A new controller object is instantiated, and an action method is called.
4. The controller interacts with the model (usually performing a CRUD operation).
5. A response is sent back to the browser, in the form of either a render or a redirect.

Figure 7-1 illustrates the process.

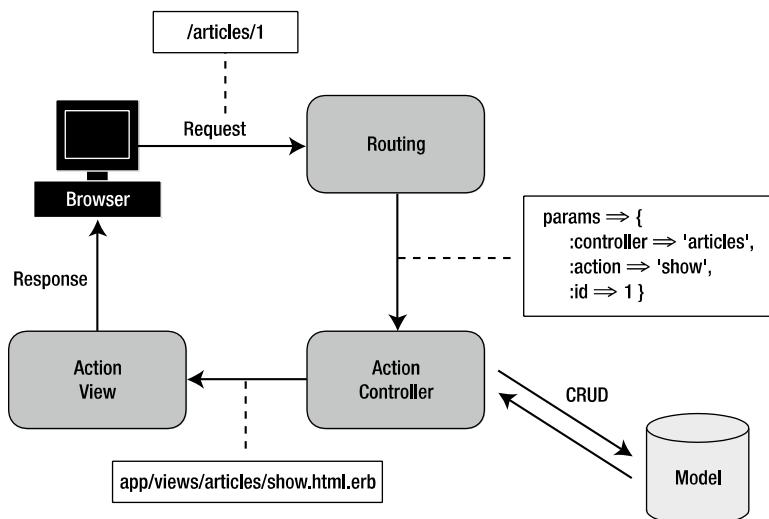


Figure 7-1. The Action Pack request cycle

Not long ago (and still today), developers used to construct *server pages*. Such a page had a bunch of code at the top of an otherwise static page, just above the opening HTML tag. The markup was littered with different sorts of code: it wasn't unusual to see the database being accessed, forms being processed, sessions being set, and all manner of logic being performed in line. The web server was responsible for controlling the application—one page redirecting to another, running the code, and then dumping the results to the screen.

We won't get into the multitude of reasons why this is a bad idea, except to say that it presents the problem of coupling. In this scenario, the business logic and the view are mashed together, making the code more difficult to maintain and debug. ASP and PHP pages are notable offenders, and if you're coming from either of these camps, the concept of separating concerns may be foreign at first. Here's a way to think about it that may help. Imagine taking the code and logic from the top of each page and sticking it in one place, leaving only the HTML behind. Then, instead of using the web server to invoke each page as you would with a static site, have the web server call on a single dispatcher, which finds the code you want to execute and calls it. The code it invokes—the file that contains the processing logic extracted from the server page—is called the controller. Instead of logic being divided among pages, it's divided into actions.

The single biggest advantage of this pattern is that the processing logic is decoupled from the view and safely contained in one place. As you can see, it's a lot easier to work this way. The interplay between actions is considerably easier to visualize and understand when it isn't spread out over a host of locations. Your server pages become lightweight views, left to handle only the simplest of instructions, if any.

A Controller Walk-Through

Instead of boring you with more theory about controllers, views, and MVC, let's dig in and start writing some real-world code. You'll continue building your blog application, examining the finer points and the places where convention wins out over configuration. Along the way, this section touches on some of the most essential controller and view concepts. By the end of this walk-through, you should have a complete grasp of how the Rails request cycle works and a working example to refer to and expand on in the subsequent chapters. The purpose of this walk-through isn't to examine each and every aspect of Action Pack in detail, but rather to give you a practical overview of how the components—routes, controllers, helpers, views, layouts, and partials—work together to control your application and construct its interface.

Setting Up Routes

Links and URLs are important in web applications. They serve as the entry point to the application and contain all the information required to route an incoming request to the controller and action that will handle it. Before you get into the meat of understanding controllers and their actions, you need to spend a few minutes learning how to get from request to response. It all starts with routing.

Routing Basics

In Rails, all the rules for mapping URLs to controllers are a matter of configuration. You find the `routes.rb` file in the `config` directory. If you open that file in your editor now and examine it, you'll see lots of commented code with samples of routes you could possibly define. Look at the last commented line to understand how routes can be defined:

```
get ':controller(/:action(/:id(.:format)))'
```

Routes work based on pattern matching and can include variables to set directly within the pattern. Here, the pattern is a three-segment string partitioned by slashes (/) and containing variables to set, where the variables are prefixed by a colon (:). The first segment sets the `:controller` variable, the second the `:action` variable, and the third the `:id` and `:format` variables. These variables are used to determine the controller and action to invoke and the parameters to send along. The pattern is matched from left to right, and the variables are optional. If the `:id` variable is empty, only the controller and action are set. If the `:action` variable is empty, only the `:controller` variable is set, and so on.

Routing priority is based on the order in which routes exist in `routes.rb`, so that the first route defined has the highest priority. If an incoming URL matches the first route defined, the request is sent along, and no other routes are examined.

Here's an example that matches a specific pattern and sets the controller and action in response:

```
get '/teams/home' => 'teams#index'
```

This route matches a URL like <http://example.com/teams/home> and routes the request to the `index` action on the `teams` controller. The names of the controller and action are separated by the `#` symbol. You can also set arbitrary parameters when using the route. For example, let's say you want to set a parameter called `query` that you can access and use in your controller:

```
get '/teams/search/:query' => 'teams#search'
```

This route matches a URL like <http://example.com/teams/search/toronto>, routing the request to the `teams` controller and the `search` action. The third segment in the URL is assigned to the `:query` parameter because you specify `:query` as an in-line variable.

Routes can be complex, and it's possible to apply conditions and other forms of logic to them. For the most part, though, you can get a lot of mileage from the general cases outlined here. The Rails API documentation (<http://guides.rubyonrails.org/routing.html>) contains details on using the more complex routing features.

Named Routes

One of the coolest things about routing in Rails is a feature known as *named routes*. You can assign a name to a given route to make referring to it in code easier. You still define the route the same way as a regular route, but you need a new hash pair, where the key is `:as` and the value is the name of the route.

For example, let's take the search route defined in the previous section and turn it into a named route:

```
get '/teams/search/:query' => 'teams#search', :as => 'search'
```

With this definition in place, Rails creates helper methods that allow you to reference this particular route using its name: `search_url` and `search_path`. The `*_url` variant returns a full URL including the protocol and hostname (<http://example.com/teams/search>), whereas the `*_path` variant returns just the path (`/teams/search`).

Later in this chapter, we'll cover redirection methods and hyperlink-generation helpers. For now, note that you can use them with named routes:

```
link_to "Search", search_path
```

outputs

```
<a href="/teams/search">Search</a>
```

Named routes are shorter, DRYer, and impervious to changes made at the routing level. So if you change the controller name from `teams` to `cities`, you don't need to update links that use the named route; for the unnamed version, you do.

RESTful Routes and Resources

Earlier, we said that RESTful design information is dealt with in the form of resources. Rails makes it easy for you to do that: for every action in your controller, you have an associated named route to call.

Resources are configured in the `routes.rb` file using the `resources` method. If you look at the `routes` file in your blog application, you see `resources :articles` at the top: it was added when you generated the `articles` scaffold in Chapter 3. The `resources :articles` method defines the following named routes for the `articles` controller:

```
article_path => /articles/:id
articles_path => /articles
edit_article_path => /articles/:id/edit
new_article_path => /articles/new
```

The `resources` method generated four named routes for you; but when you open the `ArticlesController`, you have seven actions (Table 7-1). How can you access the remaining actions? Remember that when you learned about REST earlier, you saw that every operation is identified by both a URL *and* a request method. Using different request methods with the generated named routes, Rails routes them to the appropriate controller actions.

Table 7-1. Articles Named Routes

Request Method	Named Routes	Parameters	Controller Action
GET	articles_path		index
POST	articles_path	Record hash	create
GET	article_path	ID	show
PUT	article_path	ID and Record hash	Update
PATCH	article_path	ID and Record hash	Update
DELETE	article_path	ID	destroy
GET	edit_article_path	ID	edit
GET	new_article_path		new

Note You can list all the available routes in your application by calling the Rake routes command from the terminal. You can also view a list of routes by going to <http://localhost:3000/rails/info> in your browser while you have your Rails server running in development mode.

By following the REST convention, instead of defining a named route for every action, you use the resources method in your routes file. To give some examples, if you want to access the index action in your articles controller, you go to /articles in your browser; the default request method when you type a URL in your browser is GET. What if you want to create a new article? You can do that by submitting a form to /articles with the default request method for forms, POST. To get a specific article, type /articles/:id, where :id is your article id. It's that simple.

Configuring Routes for the Blog Application

Let's configure the routes to be used in your blog application. You haven't built all the controllers and actions yet (you do that next), but that shouldn't stop you from getting the routes in place.

You can handle an empty request for the root of your application's domain using the root method. In the blog application, you want the root URL (<http://localhost:3000>) to connect to the list of articles. To accomplish this, you add a root declaration to your routes file and make it the first route. Make sure your config/routes.rb file looks like Listing 7-3 (note that all comments have been deleted).

Listing 7-3. The config/routes.rb File: <http://gist.github.com/330822>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles
end
```

Now that you have some routes defined, let's move back to the articles controller and try to understand its actions and templates.

Revisiting the Scaffold Generator

You generated a scaffold for your articles in Chapter 3, and this scaffold generated a RESTful controller for the Article model in addition to all the required templates. The generator also added the resources declaration to your route.rb file. Listing 7-4 shows the ArticlesController that your scaffold generated.

Listing 7-4. The ArticlesController app/controllers/articles_controller.rb

```
class ArticlesController < ApplicationController
  before_action :set_article, only: [:show, :edit, :update, :destroy]

  # GET /articles
  # GET /articles.json
  def index
    @articles = Article.all
  end

  # GET /articles/1
  # GET /articles/1.json
  def show
  end

  # GET /articles/new
  # GET /articles/new.json
  def new
    @article = Article.new
  end

  # GET /articles/1/edit
  def edit
  end

  # POST /articles
  # POST /articles.json
  def create
    @article = Article.new(article_params)

    respond_to do |format|
      if @article.save
        format.html { redirect_to(@article,
notice: 'Article was successfully created.') }
        format.json { render action: 'show',
status: :created, location: @article }
      else
        format.html { render action: "new" }
        format.json { render json: @article.errors,
status: :unprocessable_entity }
      end
    end
  end
end
```

```

# PATCH/PUT /articles/1
# PATCH/PUT /articles/1.json
def update
  respond_to do |format|
    if @article.update_attributes(article_params)
      format.html { redirect_to(@article,
        notice: 'Article was successfully updated.' ) }
      format.json { head :no_content }
    else
      format.html { render action: "edit" }
      format.json { render json: @article.errors,
        status: :unprocessable_entity }
    end
  end
end

# DELETE /articles/1
# DELETE /articles/1.json
def destroy
  @article.destroy

  respond_to do |format|
    format.html { redirect_to articles_url }
    format.json { head :no_content }
  end
end

private
# Use callbacks to share common setup or constraints between actions.
def set_article
  @article = Article.find(params[:id])
end

# Never trust parameters from the scary internet, only allow the white list through.
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at)
end

```

This may look like a lot of code to swallow, but in reality it's simple. The scaffold generator creates the articles controller with the default seven actions discussed earlier for RESTful controllers: index, show, new, edit, create, update, and destroy.

Before your action renders a view, you arrange for it to set an instance variable that the view can use. To refresh your memory, an instance variable is a special kind of Ruby variable that is unique to a given instance of a class, serving as a way for an object to maintain its state. Because views are, in essence, extensions of the controller object, they can access its instance variables directly (although not without some behind-the-scenes Ruby magic that Rails takes care of for you). For all intents and purposes, however, you can consider instance variables to be shared between controllers and views.

You can store any Ruby object in an instance variable, including strings, integers, models, hashes, and arrays. If you reexamine each action in the articles controller, notice that it always starts by setting an instance variable to be called later in that action's view. Let's take the index method as an example (Listing 7-5).

Listing 7-5. The Index Action in app/controllers/articles_controller.rb

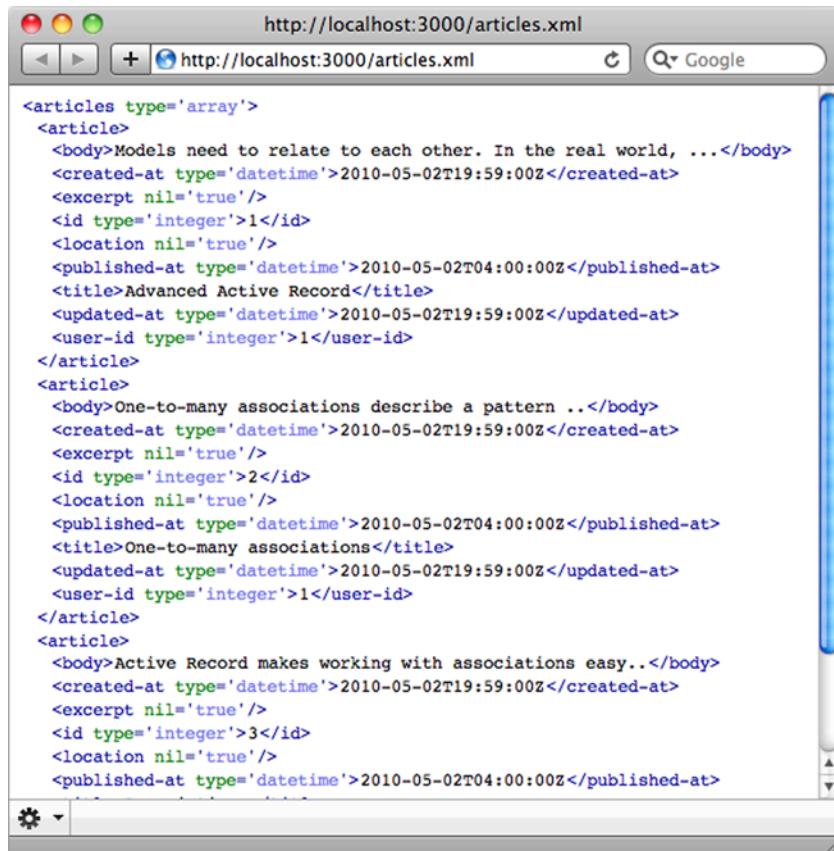
```
# GET /articles
# GET /articles.json
def index
  @articles = Article.all
end
```

You define and set an instance variable named `@articles`, which holds the array of all your articles.

Let's step back a bit. When you call the `index` method by typing the URL (<http://localhost:3000/articles>) into your browser—don't forget to start your local server using the `rails server` command—the request goes first to your routes file, where it's forwarded to the controller. Then, the controller responds to this request by setting an instance variable and rendering *something* back to the browser.

What the controller renders is based on what has been requested. Normally, it's an HTML page request, but it can also be an XML or an Ajax request. It's the responsibility of the `respond_to` method to define how to respond to each of those requests. In the `index` action, you accept two formats: HTML, where Rails renders the index template using the path `(/articles)`; and JSON (JavaScript Object Notation), where Rails renders the articles in JSON format using the path `(/articles.json)`. In this case the `respond_to` method is implicit, which means that since we didn't need to change any options, it will just use the defaults. You will see the `respond_to` method actually used when we look at later actions.

Try that in the browser. Visit <http://localhost:3000/articles> to see the list of articles you know and saw earlier, and visit <http://localhost:3000/articles.json> to see the result shown in Figure 7-2.



The screenshot shows a Mac OS X desktop with a web browser window open. The address bar displays `http://localhost:3000/articles.xml`. The page content is an XML document representing a collection of articles. The XML structure includes a root element `<articles type='array'>`, followed by three `<article>` elements. Each article contains fields such as `<body>`, `<created-at>`, `<excerpt nil='true'>`, `<id type='integer'>1</id>`, `<location nil='true'>`, `<published-at>`, `<title>`, `<updated-at>`, and `<user-id type='integer'>1</user-id>`. The XML is color-coded with green for tags and blue for attributes and values.

```

<articles type='array'>
  <article>
    <body>Models need to relate to each other. In the real world, ...</body>
    <created-at type='datetime'>2010-05-02T19:59:00Z</created-at>
    <excerpt nil='true'>
    <id type='integer'>1</id>
    <location nil='true'>
    <published-at type='datetime'>2010-05-02T04:00:00Z</published-at>
    <title>Advanced Active Record</title>
    <updated-at type='datetime'>2010-05-02T19:59:00Z</updated-at>
    <user-id type='integer'>1</user-id>
  </article>
  <article>
    <body>One-to-many associations describe a pattern ..</body>
    <created-at type='datetime'>2010-05-02T19:59:00Z</created-at>
    <excerpt nil='true'>
    <id type='integer'>2</id>
    <location nil='true'>
    <published-at type='datetime'>2010-05-02T04:00:00Z</published-at>
    <title>One-to-many associations</title>
    <updated-at type='datetime'>2010-05-02T19:59:00Z</updated-at>
    <user-id type='integer'>1</user-id>
  </article>
  <article>
    <body>Active Record makes working with associations easy..</body>
    <created-at type='datetime'>2010-05-02T19:59:00Z</created-at>
    <excerpt nil='true'>
    <id type='integer'>3</id>
    <location nil='true'>
    <published-at type='datetime'>2010-05-02T04:00:00Z</published-at>
  
```

Figure 7-2. Output of <http://localhost:3000/articles.json>

GET AN API FOR FREE

Using RESTful controllers in Rails gives you the ability to have an API for your application. An API is a set of functions that enables other applications to talk to your application. On the Web, this is normally done using JSON, and REST is one of the main architectures used for that.

With Rails and its RESTful controllers, defining your API is a seamless process; basically, you just need to tell your controller to respond to JSON requests, and you have an API. What's neat in Rails is that the scaffold generator adds the JSON part by default to all your controller actions, providing you with an API for free. Rails also supports XML, but JSON is the default.

Rendering Responses

When an action has completed, it attempts to render a template of the same name. That's the case with the `index` action just discussed: it renders the `index.html.erb` template by default. The same applies to `edit`, `new`, and `show` actions. But sometimes you want to *render* something else.

If you look at the `create` and `update` actions, notice that if the `@article.save` succeeds, you redirect to the saved `@article` show page with a friendly message. However, if the save fails, you want to render the `new` or the `edit` template. If you didn't explicitly render those templates, the actions would fall through to their default behavior and attempt to render their default `create` and `update` templates, which doesn't exist.

The `render` method takes several options for its first argument: `:text`, `:nothing`, `:inline`, `:template`, `:action`, `:xml`, `:json`, `:js`, and `:update`.

Note The `:update` and `:js` responses are fairly specialized. You use them when you're rendering Ajax or JavaScript responses, as you'll learn in Chapter 9

Redirecting

It may not sound like it, but a redirection is a response. Redirects don't happen on the server side. Instead, a response is sent to your browser that tells it to perform a redirection to another URL. The specifics of issuing a redirect aren't something you need to worry about, though, because Rails provides a specialized method to take care of the internals. That method is called `redirect_to`, and it's one you'll find yourself using a lot, so it's a good idea to get familiar with it.

The `redirect_to` method usually takes a URL as a parameter, which in most cases is represented by one of your routes. Let's say that you want to redirect the user to the articles' index page, and the path you use is `articles_path`—a route added by `resources :articles` in `config/routes.rb`; so, you execute `redirect_to(articles_path)`. If you look at the `destroy` action, the user is redirected to `articles_url` after an article is deleted.

As you can see from the `create` and `update` actions, `redirect_to` can also take an object as a parameter, in which case it redirects to a path that represents that object. This means Rails uses a convention to translate objects to their `show` action named route. In this case, `redirect_to(@article)` is a shortcut equivalent to `redirect_to(article_path(:id => @article))`.

WHAT MAKES A CLASS AN ACTION CONTROLLER?

If you're the curious sort (and, of course, you are), you may wonder how `ArticlesController`, a seemingly normal Ruby class, becomes a full-fledged Action Controller. Well, if you look closely, you'll notice that `ArticlesController` inherits from another class: `ApplicationController`. To get a better picture of what's going on, let's take a peek at the `ApplicationController` class in `app/controllers/application_controller.rb`:

```
class ApplicationController < ActionController::Base
  # Prevent CSRF attacks by raising an exception.
  # For APIs, you may want to use :null_session instead.
  protect_from_forgery with: :exception
end
```

The mystery is quickly solved. The simple controller becomes an Action Controller by subclassing the `ApplicationController` class, itself a subclass of `ActionController::Base`. This is an example of inheritance and is common in object-oriented programming. When one class subclasses another, it inherits all the behavior and methods of the parent. In the case of the articles controller, it inherits all the capabilities of the application controller. Likewise, `ApplicationController` inherits all the capabilities of its parent, `ApplicationController::Base`. The `ApplicationController::Base` class effectively endows your articles controller with its special abilities.

The application controller is the base from which all the controllers you make inherit. Because it's the parent of all controllers in your application, it's a great place to put methods that you want accessible in every controller.

By looking at the articles controller, you now understand the basic conventions and common concepts of how a RESTful controller normally behaves. You have seven default actions, and in every one of them you do the following:

- Set an instance variable to be used later in the rendered action or template
- Handle the response using the `respond_to` method to either do a `render` or `redirect_to` another path, depending on the behavior you want to achieve

Understanding Templates

The next step is to look at the actions' templates. Look in the `app/views/articles` directory, and you see seven templates:

- `form.html.erb`
- `index.html.erb`
- `index.json.jbuilder`
- `show.html.erb`
- `show.json.jbuilder`
- `new.html.erb`
- `edit.html.erb`

The basic convention of Action Pack is as follows: *templates are organized by controller name, and a template with the same name as the action being invoked is rendered automatically*. You don't need to wire up anything. Merely by requesting an action from a controller, Rails renders the corresponding template in that controller's directory inside `app/views/` that has the same name.

Let's try an example. Make sure your local web server is running (`rails server`), and open <http://localhost:3000/articles/> in your browser. You see the articles index page shown in Figure 7-3.

Title	Location	Excerpt	Body	Published at
Advanced Active Record		Models need to relate to each other. In the real world, ...		2010-05-02 04:00:00 UTC Show Edit Destroy
One-to-many associations		One-to-many associations describe a pattern ..		2010-05-02 04:00:00 UTC Show Edit Destroy
Associations		Active Record makes working with associations easy..		2010-05-02 04:00:00 UTC Show Edit Destroy

[New Article](#)

Figure 7-3. Output of <http://localhost:3000/articles/>

The articles listing is actually rendered from `app/views/articles/index.html.erb`, which follows the convention discussed earlier. It's the articles controller, so it goes to the `articles` directory in `app/views`. After determining which controller to invoke, Rails proceeds to instantiate it and calls its `index` method. Its default response after running the `index` action is to perform a render. Rails looks for a template named `index.html.erb` in the `app/views/articles` directory and loads it. The same applies to the `show` action: the `show.html.erb` template is rendered.

At this point, the request cycle is complete. If you refresh your browser, the cycle begins anew, and the same result is rendered. Notice how all the internals are taken care of for you. All you need to do is create an appropriately named route, controller, action, and view, stick them in the right place, and request the URL in your browser. Rails takes care of making sure everything is knit together properly.

Before you go any further, use your browser's View Source command to see the HTML that was produced. If you know anything about HTML (and chances are you do), you'll quickly realize that some additional HTML code has been rendered around the code in `index.html.erb`; it came from a *layout*. Most web pages have headers, footers, sidebars, and other page elements that, when styled, make the page look pretty. Rails has a built-in facility for dealing with page layouts.

Working with Layouts

Rails uses layouts to interpolate the output of an individual template into a larger whole—a reversal of the common pattern of including a shared header and footer on every page (which, if you've done any work in languages like PHP and ASP, is all too familiar). The scaffold generator you ran in Chapter 3 created a layout file and placed it in `app/views/layouts/application.html.erb`. The `application.html.erb` layout is applied to all controllers. However, if you like your layout to apply to a specific controller, you can create a layout file named after the controller you want. For example, a layout that applies only to the `articles` controller should be created in `app/views/layouts/articles.html.erb`. That's the way it works in Rails. Just as an action tries to render itself using a view that matches its name, a controller attempts to use a layout that matches its name.

Note Layouts always default to the most specific declaration. If your controller inherits from ApplicationController and doesn't specify a layout directly, Rails will look for a layout named after your controller first. If that layout isn't found, it will look for a layout declaration on ApplicationController, and if that isn't found, it will look for a layout named application. In other words, layout declaration follows normal class inheritance.

Open the file app/views/layouts/application.html.erb in your editor. You should see something like the file shown in Listing 7-6.

Listing 7-6. The app/views/layouts/application.html.erb File

```
<!DOCTYPE html>
<html>
<head>
  <title>Blog</title>
  <%= stylesheet_link_tag "application", media: "all", "data-turbolinks-track" => true %>
  <%= javascript_include_tag "application", "data-turbolinks-track" => true %>
  <%= csrf_meta_tag %>
</head>
<body>

<%= yield %>
</body>
</html>
```

At rendering time, the layout yields the results of the template fragment's execution in place. See the `<%= yield %>` bit that's highlighted in bold? That's the important part. Wherever you put the `yield` keyword is where your content goes.

One more thing to note: Rails is all about convention over configuration. Here, the convention is that a layout with the name `application.html.erb` is automatically applied to all templates unless an alternate is specified. This means that if you change the name of the layout as it stands, it won't be automatically applied. If you want to apply a different layout to a given controller, you can either have a layout named after the controller or specify it in the controller using the class method `layout`:

```
class ExampleController < ApplicationController
  layout 'my_layout' # Will use a layout in app/views/layouts/my_layout.html.erb
end
```

COMMON LAYOUT CONVENTIONS

A few conventions apply to working with layouts:

- A layout named `application.html.erb` is applied automatically unless a more specific candidate exists or is explicitly specified in the controller.
- A layout that matches the name of a controller is automatically applied if present. Controller-specific layouts take precedence over the application-level layout.
- You can use the `layout` directive at the class level in any controller (that is, not inside an action) to set the layout for the entire controller: `layout 'my_layout'`.

- You can include a layout for a specific action with an explicit call to `render` inside the action:
`render :layout => 'my_layout'.`
- Sometimes, you want to render an action without a layout. In that case, you can pass `false` in place of the layout name: `render :layout => false`.

In practice, you usually use `application.html.erb` and rarely take advantage of the controller-specific layout functionality. On the occasions when you need to use a different layout for a particular controller, use the `layout` directive.

Looking at the Article Form

Let's look at the new template in action. The new action has a single purpose: to initialize and display the form for creating a new article. The actual creation of a new `Article` object is the responsibility of the `Article` model (remember the discussions of the model in Chapters 5 and 6), but it's orchestrated by the controller. Moreover, it needs data (like a title and body), which it must procure from somewhere. The `edit` action isn't any different, except that it finds and displays a form of an existing article object rather than a new one.

You can extract this information from HTML form elements placed in the view and handled in the controller. Open `new.html.erb` and `edit.html.erb`, which look like Listings 7-7 and 7-8, respectively.

Listing 7-7. Content of `app/views/articles/new.html.erb`

```
<h1>New article</h1>

<%= render 'form' %>

<%= link_to 'Back', articles_path %>
```

Listing 7-8. Content of `app/views/articles/edit.html.erb`

```
<h1>Editing article</h1>

<%= render 'form' %>

<%= link_to 'Show', @article %> |
<%= link_to 'Back', articles_path %>
```

Notice the similarity between the templates, especially the `render 'form'` part highlighted in bold. The `render` method renders a partial named `form` in this context. The upcoming section "Staying DRY with Partials" discusses partials in more depth; for now, let's focus on the content of the template in `app/views/articles/_form.html.erb` (Listing 7-9).

Listing 7-9. Content of `app/views/articles/_form.html.erb`

```
<%= form_for(@article) do |f| %>
  <% if @article.errors.any? %>
    <div id="error_explanation">
      <h2><%= pluralize(@article.errors.count, "error") %>
      prohibited this article from being saved:</h2>
      <ul>
        <% @article.errors.full_messages.each do |msg| %>
          <li><%= msg %></li>
```

```

<% end %>
</ul>
</div>
<% end %>

<div class="field">
  <%= f.label :title %><br />
  <%= f.text_field :title %>
</div>
<div class="field">
  <%= f.label :location %><br />
  <%= f.text_field :location %>
</div>
<div class="field">
  <%= f.label :excerpt %><br />
  <%= f.text_field :excerpt %>
</div>
<div class="field">
  <%= f.label :body %><br />
  <%= f.text_area :body %>
</div>
<div class="field">
  <%= f.label :published_at %><br />
  <%= f.datetime_select :published_at %>
</div>
<div class="actions">
  <%= f.submit %>
</div>
<% end %>

```

You use form helpers for each of your fields. Visit the article's new page at <http://localhost:3000/articles/new> in your browser and you'll see that the helpers function to produce a nicely formatted HTML form. Use your browser's View Source command to look at the HTML that was generated. Here's part of the generated HTML:

```

<h1>New article</h1>

<form accept-charset="UTF-8" action="/articles" class="new_article"
id="new_article" method="post">
  <div style="margin:0;padding:0;display:inline">
    <input name="_snowman" type="hidden" value="#9731;" />
    <input name="authenticity_token" type="hidden"
      value="Dhq8NmXwgsAPMVHJPfoRWu8UJN1XurF8ngdzaksC48=" />
  </div>

  <div class="field">
    <label for="article_title">Title</label><br />
    <input id="article_title" name="article[title]" size="30" type="text" />
  </div>
  <div class="field">
    <label for="article_location">Location</label><br />
    <input id="article_location" name="article[location]" size="30" type="text" />
  </div>

```

```

<div class="field">
  <label for="article_excerpt">Excerpt</label><br />
  <input id="article_excerpt" name="article[excerpt]" size="30" type="text" />
</div>
<div class="field">
  <label for="article_body">Body</label><br />
  <textarea cols="40" id="article_body" name="article[body]" rows="20"></textarea>
</div>
...
<div class="actions">
  <input id="article_submit" name="commit" type="submit" value="Create Article" />
</div>
</form>

```

Note the way in which Rails formats the name attribute of each form element: *model[attribute]*. This helps when it comes to parsing the parameters from the form, as you'll see shortly. If you manually create your form elements (which you need to do sometimes), you can use this naming convention to make sure your form values are easy to parse in the controller. Most of the time, though, you use form helpers when working with forms, especially when you're dealing with Active Record objects. Let's spend some time discussing form helpers.

Using Form Helpers

One of the best things about working with templates in Rails is the presence of helpers. Rails comes with a bunch of helper methods that take the tedium out of generating the bits of HTML that your views need. Let's face it, nothing is more of a drag to build than HTML forms. Fortunately, Rails understands the plight of the web developer all too well and provides a suite of easy ways to build forms.

Two basic varieties of form helpers are available:

- **FormHelper**: Active Record-aware tag helpers for creating forms that hook into models.
- **FormTagHelper**: Helpers that output tags. They aren't integrated with Active Record.
The names of these helpers are suffixed with `_tag`.

The `FormHelper` type is aware of Active Record objects assigned to the template; the `FormTagHelper` (note the `Tag` type isn't). The advantage of the Active Record-aware, `FormHelper`, helpers is that they know how to populate themselves with data and can automatically be highlighted in the event of validation errors from the model. But not every form element you make corresponds directly to a model attribute. That's where the `FormTagHelper` group comes in handy. These have no special relationship with Active Record; they just output form tags.

In your article's form template (Listing 7-9), you use six helpers: `form_for`, `label`, `text_field`, `text_area`, `datetime_select`, and `submit`.

The `form_for` helper is of the `FormHelper` variety. It creates an HTML form tag for the passed object (`@article`, in this case) and places everything in the `do..end` block inside the resulting form. It also produces and sets a *form local variable* to the form block. The form local variable, in this case called `f`, is aware of the `@article` object and uses its attributes' names and values when calling the other form helpers: `label`, `text_field`, `text_area`, `datetime_select`, and `submit`.

By default, forms use the HTTP POST method. If you want to use a different method, you need to specify it manually using the `:method` option (for example, `:method => :get`). If you recall, POST is the request method you used for the `create` action in your RESTful designed controller. HTTP verbs.

The HTTP protocol defines several request methods, the most popular of which are GET and POST. Both are methods for requesting a web page; the difference is in how the request is sent. GET is the simpler of the two.

It includes all the information about the request as part of the URL. POST sends information invisibly, which is to say, as part of the request header and not part of the URL. So, you can't type a POST request into your browser's location bar. Every time you request a web page via the location bar in your browser, you're using GET. When you submit a form, say to register on a web site, the form is usually submitted via a POST.

How do you know when to use each? The best way to think of this is to consider GET a read method. It should never do anything destructive, such as modifying a database record. POST, on the other hand, can be thought of as a write method. When you need to create data, use POST. PATCH is used when you need to update a record partially, for instance, only changing your e-mail address. PUT is used to update a record completely. There has been a lot of controversy over these verbs on the Internet, but they are effectively used interchangeably in Rails. The DELETE verb is used to destroy a record.

A small note: Most browsers only support the GET and POST verbs. Rails gets around this by using an actual POST request but inserting hidden form fields specifying which actual verb to use. Rails automatically removes this field and converts the request into the specified type. Once the request has reached the controller, it will appear as the intended verb.

Remember that you should never put a state-changing action behind a GET request. For more information, see <http://www.w3.org/2001/tag/doc/whenToUseGet.html>.

The label helper is a FormHelper method that outputs an HTML label tag for the provided attribute. Here's an example of the output for title:

```
<label for="article_title">Title</label>
```

The text_field helper is of the FormHelper variety, meaning that it corresponds to Active Record objects. It creates an HTML input tag whose type is set to "text" and assigns it a name and an ID that match the given object and method (title in this case). Here's what the rendered output looks like:

```
<input id="article_title" name="article[title]" size="30" type="text" />
```

The text_area helper is also of the FormHelper variety. It's similar to text_field, except it returns a text area instead of a text input. Here's what the HTML output looks like for the body field:

```
<textarea cols="40" id="article_body" name="article[body]" rows="20"></textarea>
```

The datetime_select helper is a FormHelper that outputs a set of HTML select tags to input a date and time value.

The submit helper is a FormHelper that creates an input element whose type is set to "submit". It accepts the name of the submit button as its first argument. If you don't provide a name to the submit method, it generates a name based on the @article object. For example, in the New Article form, the generated name is Create Article, whereas in the Edit Article form, the name is Update Article. Here's the HTML output from the example:

```
<input id="article_submit" name="commit" type="submit" value="Create Article" />
```

All these helpers (and, to be sure, most helpers in Rails) accept a hash of options as their last argument to customize the resulting HTML. For example, to give your title field a class of large, you type f.text_field :title, :class => 'large', which adds the class attribute to the output:

```
<input class="large" id="article_title" name="article[title]"  
size="30" type="text" />
```

You can pass arbitrary options in this way, all of which end up as attributes on the resulting tag. For example, to apply an in-line style attribute, you can use `:style => 'background: #fab444'`. Here's a list of some of the most common FormHelper helpers:

- `text_field`
- `hidden_field`
- `password_field`
- `file_field`
- `text_area`
- `check_box`
- `radio_button`

All these methods can be suffixed with `_tag` to create standard HTML tags (with no Active Record integration).

For a full list of FormHelper and FormTagHelper methods, consult the Rails API, where you can find a complete reference along with usage examples:

- <http://api.rubyonrails.org/classes/ActionView/Helpers/FormHelper.html>
- <http://api.rubyonrails.org/classes/ActionView/Helpers/FormTagHelper.html>

Now, back to your form. Let's see what happens when you submit it. (Make sure your server is still running.) Click the Create Article button, and you see the screen shown in Figure 7-4.

The screenshot shows a web browser window titled "Blog". The address bar displays "<http://localhost:3000/articles>". The main content area shows a form titled "New article". A red box highlights an error message: "2 errors prohibited this article from being saved:" followed by a list: "Title can't be blank" and "Body can't be blank". Below the error message are four input fields: "Title" (empty), "Location" (empty), "Excerpt" (empty), and "Body" (empty). The "Body" field has a red border around its entire area. On the right side of the form, there is a vertical scroll bar.

Figure 7-4. New article form with validation errors

What happened? Well, as the message says, Rails couldn't create an article for you. Of course it couldn't—you set validation rules in your Article model to prevent the creation of a new article object with an empty title or body field. But let's look at the output from the server running in the command prompt and see what happened:

```
Started POST "/articles" for 192.168.1.105 at 2013-07-06 02:46:45 -0500 Processing by
ArticlesController#create as HTML

Parameters: {"utf8"=>"✓",
"authenticity_token"=>"fZ2oh3Tt6wx3mjxXB0H3kIZie/Jzap3NDDdCUZsjxk=", "article"=>{"title"=>"", "location"=>"", "excerpt"=>"", "body"=>"", "published_at(1i)"=>"2013", "published_at(2i)"=>"7", "published_at(3i)"=>"6", "published_at(4i)"=>"07", "published_at(5i)"=>"46"}, "commit"=>"Create Article"}

(0.1ms) begin transaction
(0.1ms) rollback transaction
Rendered articles/_form.html.erb (6.7ms)
Rendered articles/new.html.erb within layouts/application.html.erb (7.2ms)
Completed 200 OK in 14ms (Views: 9.9ms | ActiveRecord: 0.1ms)
```

See the section titled `Parameters` above? You may recognize this as a Ruby hash. This hash contains all the form values you submitted. Notice that there's an entry for the button name (`commit`), called Create Article, and for `authenticity_token`, which is used for security in Rails to prevent anonymous form posts. The `article` portion of the hash looks like this:

```
"article"=>{"title"=>"", "location"=>"", "excerpt"=>"", "body"=>"", "published_at(1i)"=>"2010", "published_at(2i)"=>"5", "published_at(3i)"=>"2", "published_at(4i)"=>"22", "published_at(5i)"=>"35"
```

If you're thinking that this looks a lot like the options hashes you passed to article objects when you were working with Active Record on the console, you're right. Rails automatically turns form elements into a convenient hash that you can pass into your models to create and update their attributes. In the sections that follow, you'll put this feature to use in the next action, `create`. First, let's take a deeper look at `params`.

Processing Request Parameters

Request parameters—whether they originate from requests of the GET or POST variety—are accessible via the `params` hash. To be specific, `params` is a method that returns a Hash object so you can access it using hash semantics. Hashes in Ruby are similar to arrays but are indexed by arbitrary keys—unlike arrays, which are indexed by number. (If you need a quick review of the Hash object, flip to Chapter 4 for a Ruby primer.)

The value of any request variable can be retrieved by its symbolized key. So, if there's a variable called `id` in the request parameters, you can access it with `params[:id]`. Just to drive this concept home, let's look at a sample URL and display the `params` hash that it populates. Point your browser to <http://localhost:3000/articles?title=rails&body=great> and check the server output. You should see something similar to this:

```
Parameters: {"title"=>"rails", "body"=>"great"}
```

Revisiting the Controller

With an understanding of `params` under your belt, let's go back to your controller. The `create` action is the target of the form submission. The method code shown in Listing 7-10 is from the `articles` controller, just under the new method.

Listing 7-10. The Create Action in `app/controllers/articles_controller.rb`

```
# POST /articles
# POST /articles.json
def create
  @article = Article.new(article_params)

  respond_to do |format|
    if @article.save
      format.html { redirect_to @article, notice: 'Article was successfully created.' }
      format.json { render action: 'show', status: :created, location: @article }
    else
      format.html { render action: 'new' }
      format.json { render json: @article.errors, status: :unprocessable_entity }
    end
  end
end
```

Let's walk through this. First, you initialize a new `Article` object with whatever attributes come in via the `article_params` method. This is something that is new to Rails 4.0. You can imagine that taking raw input from the user and putting it directly into your model without filtering it could be dangerous. Imagine that you were letting users sign up using the `User#create` action. If you had an attribute on the user model called `admin` that determined whether or not the user had admin access to the system, a user could just add that parameter in themselves and make themselves an admin. You can see how vital it is that we filter the parameters now! Let's take a look at the `article_params` method:

```
# Never trust parameters from the scary internet, only allow the white list through.
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at)
end
```

The syntax for this is simple. We are telling rails that we require the `article` param, and permit `title`, `location`, `excerpt`, `body`, and `published_at`. Any other params will be filtered out before they get to the model. If you try to just pass `params[:article]` to the `new` or `create` method, an error will be returned.

After we pass the filtered params to the new method, we attempt to save the model. If the save is successful, you use a facility that Rails provides called the `flash` to set a message—by passing the `:notice` option to `redirect_to`—before redirecting to the `show` action on the same `articles` controller. The flash is a special kind of storage mechanism provided by Rails for convenience. It encapsulates the pattern of wanting to set a message on one action and have that message persist to the next, only to disappear after that action is rendered. This is useful for providing user feedback, as you do here to say "Article was successfully created." If you look at the `show` article file in `app/views/articles/show.html.erb`, you have access to the `notice` variable, allowing the message to be displayed:

```
<p class="notice"><%= notice %></p>
```

The flash message you set is available to the controller and action you redirect to (the `show` action on the `articles` controller). There are two special flash cases, `notice` and `alert`, which you can use just as you did in the previous example by passing them as arguments to `redirect_to`.

Note When you pass notice: "Article was successfully created" to redirect_to, it's identical to calling flash[:notice] = "Article was successfully created" in a separate line. Also, when you retrieve, in any view template, the message using notice, you could as well use flash[:notice]. So, you can use any named key when calling flash because it's implemented as a Ruby hash. You store values in it based on a key. The key can be anything you like: you can use any symbol, such as flash[:warning] =, in your controller and later retrieve it in your views using the same flash[:warning] call.

If the save fails, you render the new action again so that any errors can be corrected.

Displaying Error Messages in Templates

Let's try submitting the form empty one more time to explore it again. Sure enough, the form doesn't save. Notice that you're still on the same screen and that the form elements are highlighted in red, as shown in Figure 7-4.

If you look at the HTML source, you see that the input and label tags are surrounded by div elements with the class name `field_with_errors`:

```
<div class="field_with_errors">
  <label for="article_title">Title</label>
</div><br />
<div class="field_with_errors">
  <input id="article_title" name="article[title]" size="30" type="text" value="" />
</div>
```

Rails does this automatically for any fields that fail validation. You can use these classes to style invalid elements.

Note The style rules that turn the invalid fields red are generated by the scaffold generator and are in `app/assets/stylesheets/scaffold.css.scss`.

The formatted list of errors that appears at the top of the page is rendered using the code snippet below, which is a part of `app/views/articles/_form.html.erb`:

```
<% if @article.errors.any? %>
<div id="error_explanation">
  <h2><%= pluralize(@article.errors.count, "error") %> prohibited this article from being
  saved:</h2>
  <ul>
    <% @article.errors.full_messages.each do |msg| %>
      <li><%= msg %></li>
    <% end %>
  </ul>
</div>
<% end %>
```

Now that you understand this, let's submit the form with valid data. If all goes according to plan, the new article should be created, and you're redirected to that article's show action, where you see the friendly notice message you set. Notice that if you refresh the page using your browser's Refresh button, the notice message disappears.

The *edit* and *update* Actions

The edit and update actions look almost identical to the new and create actions. The main difference is that instead of instantiating a new Article object, you fetch an existing one. This happens with a callback called `before_action`. This is similar to the Active Record callbacks we looked at in Chapter 7. In this case, the `set_article` method is called before the show, edit, update, and destroy actions are run. The `set_article` method loads the article using the `id` passed in params. This allows your code to stay DRY by keeping you from typing that line multiple times throughout the controller. It works exactly as if the code from the `set_article` method were typed at the very beginning of your action.

Looking at our action again. We used Active Record's `update` method to update all the Article attributes with those from the `article_params` method. If the update fails, `update` returns false and your `if` statement takes the `else` path (Listing 7-11).

Listing 7-11. The Update Action in `app/controllers/articles_controller.rb`

```
# PATCH/PUT /articles/1
# PATCH/PUT /articles/1.json
def update
  respond_to do |format|
    if @article.update(article_params)
      format.html { redirect_to @article, notice: 'Article was successfully updated.' }
      format.json { head :no_content }
    else
      format.html { render action: 'edit' }
      format.json { render json: @article.errors, status: :unprocessable_entity }
    end
  end
end
```

Revisiting the Views

Let's get back to the views. If you look at the new and edit templates, you can't help but notice they render almost the same HTML: only the header and navigation are slightly different. Remember from the RESTful discussion that the HTTP request methods for create and update should be POST and PUT, respectively. Rails once more takes care of that for you. You're rendering the same `app/view/articles/_form.html.erb` partial, but Rails knows the request method to use based on the `@article` variable passed to the `form_for` helper.

Try editing one of the articles. The URL should be something like `http://localhost:3000/articles/1/edit`; it looks similar to the new form, but with the record information already populated (Figure 7-5).

The screenshot shows a web browser window titled "Blog". The address bar displays "http://localhost:3000/articles/1/edit". The main content area is titled "Editing article". It contains several form fields: "Title" (with value "Advanced Active Record"), "Location" (empty), "Excerpt" (empty), "Body" (containing "Models need to relate to each other. In the real world, ..."), and a "Published at" date-time selector set to "2010 May 2 04:00". Below the form is an "Update Article" button and links to "Show" and "Back".

Figure 7-5. Editing an existing article

Thanks to the `form_for` helper, the form fields are populated with their respective `@article` attributes. If you try to submit this form and look at the output from the server running on the command prompt, you'll see the following:

```
Started PATCH "/articles/1" for 192.168.1.105 at 2013-07-06 03:11:48 -0500
Processing by ArticlesController#update as HTML
Parameters: {"utf8"=>"✓",
"authenticity_token"=>"fZoh3Tt6wx3mjXXBOH3kIZie/Jzap3NDDdCUZsjxk=", "article"=>{"title"=>"Advanced Active Record", "location"=>"", "excerpt"=>"", "body"=>"Models need to relate to each other. In the real world, ..."}, "published_at(1i)"=>"2013", "published_at(2i)"=>"7", "published_at(3i)"=>"6", "published_at(4i)"=>"00", "published_at(5i)"=>"00"}, "commit"=>"Update Article", "id"=>"1"}
Article Load (0.1ms)  SELECT "articles".* FROM "articles" WHERE "articles"."id" = ? LIMIT 1
[["id", "1"]]
(0.4ms)  begin transaction
```

```
SQL (1.9ms) UPDATE "articles" SET "location" = ?, "excerpt" = ?, "updated_at" = ? WHERE
"articles"."id" = 1 [[{"location": "", "excerpt": ""}, {"updated_at": "Sat, 06 Jul 2013
08:11:48 UTC +00:00"}]

(4.1ms) commit transaction

Redirected to http://192.168.1.104:3000/articles/1

Completed 302 Found in 52ms ( ActiveRecord: 6.5ms)
```

Notice the bold line: the update action of the articles controller was called as expected. Rails recognizes that the `@article` variable passed to `form_for` isn't a new record; therefore, it calls the update action for you. This is yet another example of convention over configuration in Rails.

Staying DRY with Partials

A typical web application is rife with view code and often suffers from a lot of needless duplication. The HTML forms for adding and modifying articles are good examples of forms that are very similar. Wouldn't it be nice if there were a way to reuse the common elements from one form in more than one place? That's where partial templates come in.

Partial templates, usually referred to as *partials*, are similar to regular templates, but they have a more refined set of capabilities. Partials are used quite often in a typical Rails application, because they help cut down on duplication and keep the code well organized. They follow the naming convention of being prefixed with an underscore, thus distinguishing them from standard templates (which are meant to be rendered on their own).

Rather than creating two separate forms, Rails keeps your code DRY by using a single partial and including it from both the new and edit templates. Let's look at the code from `new.html.erb` and `edit.html.erb`, shown in Listings 7-12 and 7-13, respectively.

Listing 7-12. The `app/views/articles/new.html.erb` File

```
<h1>New article</h1>

<%= render 'form' %>

<%= link_to 'Back', articles_path %>
```

Listing 7-13. The `app/views/users/edit.html.erb` File

```
<h1>Editing article</h1>

<%= render 'form' %>

<%= link_to 'Show', @article %> |
<%= link_to 'Back', articles_path %>
```

Let's take a closer look at the `render` method. When referencing the partial in the `render` method, you don't include the leading underscore:

```
<%= render 'form' %>
```

A single string argument is passed to the render method. The render method also accepts a second argument in the form of a hash. (Have you noticed that Rails is a big fan of the options hash?) The string argument is the partial's name. Upon seeing this, the render method searches the current directory for a file named `_form.html.erb`. Notice that you don't need to include the leading underscore or the file extension when specifying the partial's name; Rails knows to look for a file in the same directory as the calling template with a leading underscore.

Let's take a brief detour to discuss a few things about partials. One of the things that makes partials unique is a special convenience: automatic local variable assignment by way of convention over configuration. That's was a mouthful. The next section explains.

Local Variable Assignment in Partials

The render method accepts a hash of local variables as part of the options hash. This is an example of what a render partial with local variables looks like:

```
<%= render 'header', :title => 'My Blog' %>
```

Any number of local variables can be assigned this way, and any object can be set as the value. In the preceding example, the partial has access to the local variable `title`.

Rendering an Object Partial

Following the same convention of local variable assignment in partials, Rails makes it easier to render a partial that represents a specific object. For example, suppose you have the following render call in your code:

```
<%= render @article %>
```

Rails looks for a partial in `app/views/articles/_article.html.erb` and automatically assigns a local variable called `article`. It's a shortcut for:

```
<%= render 'articles/article', :article => @article %>
```

Rendering a Collection of Partials

Another common pattern of rendering partials renders a collection of objects. Rails has a convention for rendering collections where you pass the collection as the first argument of the render method; Rails automatically loops across this collection and renders the partial of every object inside that array accordingly. Here's an example:

```
<%= render @articles %>
```

This behaves exactly like the previous call, but it performs more magic under the hood. For example, if the `@articles` array contains different Active Record objects, such as two articles and two comments, the render call renders the right partial template for each of those objects. It renders `/app/views/comments/_comment.html.erb` for the comment objects and `/app/views/articles/_article.html.erb` for the article objects. It translates as follows:

```
<% @articles.each do |object| %>
  <%= render object %>
<% end %>
```

Summary

This chapter covered a lot of ground. It began with a general introduction to the components that compose Action Pack, the Rails library responsible for the controller and the view. Then, it launched into a controller walk-through, where you visited your scaffold-generated controller. In doing so, you learned about routes, what happens when you generate a scaffold, how actions relate to views, and how to work with layouts. You were introduced to Rails' form helpers, and you learned how easily forms integrate with Active Record objects. The chapter also introduced partials, and you learned how to keep your templates DRY and easy to maintain.

This chapter gave you your first taste of Rails outside the model. You now have a complete understanding of how Rails divides its concerns and a firsthand look at MVC in action. You started by modeling your domain in Chapters 5 and 6, and now you've completed the first iteration of building a web application around your domain.

You should be proud of yourself. At this stage, you know a lot about Rails. The next chapter builds on this knowledge, starting with more advanced topics like building a controller from scratch, sessions, and state, and sprucing up the application with some CSS.



Advanced Action Pack

Now that you have a very good understanding of how the components of Action Pack work, it's time to dig a little deeper. You start by generating the user controller from scratch, writing its actions, and creating its templates. Then you'll add some functionality to the blog application: you allow users to leave comments when reading an article and make sure only logged-in users have access to adding and editing content. Finally, you give your application some styling so it looks better and more like a real application.

Note If you need to get the code at the exact point where you finished Chapter 7, download the source code zip file from the book's page on <http://www.apress.com> and extract it on your computer.

Generating a Controller

It's time to create your first controller from scratch. If you haven't noticed already, Rails ships with generators for most common tasks, and controllers are no exception. The syntax for the controller generator is as follows:

```
$ rails generate controller ControllerName [actions] [options]
```

As a minimum, the controller generator takes the name of the controller as an argument, which you can specify using either CamelCase (sometimes called MixedCase) or snake_case. The generator also takes an optional list of actions to generate. For every action you specify, you'll get an empty method stub in the controller and a template in `app/views/#controller_name`. To see a list of all available options, you can run the generator `rails generate controller` without arguments.

Tip The help output for the controller generator contains sample usage and options that you're sure to find interesting. All of the generators (and most UNIX commands, for that matter) respond to the `--help` argument (or variations thereof), so you're encouraged to try it whenever you're issuing a system command.

Generate the Users controller using the following command:

```
$ rails generate controller Users
```

```
create  app/controllers/users_controller.rb
invoke  erb
create  app/views/users
invoke  test_unit
create  test/controllers/users_controller_test.rb
invoke  helper
create  app/helpers/users_helper.rb
invoke  test_unit
create  test/helpers/users_helper_test.rb
invoke  assets
invoke  coffee
create  app/assets/javascripts/users.js.coffee
invoke  scss
create  app/assets/stylesheets/users.css.scss
```

Take the time to read the output of the generator so you get a sense of all the files that were just created. Notice where the templates are located: in the `app/views` directory, inside a subdirectory named after the controller. In this case, because your controller is called `users`, your templates go in `app/views/users`. Open the newly minted controller file in `app/controllers/users_controller.rb` and take a look (Listing 8-1).

Listing 8-1. Users Controller in `app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
end
```

Tip It's a convention in Rails that controller names should always be plural.

As you can see, all the generator gives you is an empty stub. If you want your `users` controller to do anything useful, you'll need to add a few actions and give it something to do. Let's add the actions you need to the controller now. Edit `users_controller.rb` so that it looks like the code in Listing 8-2.

Listing 8-2. Updated `app/controllers/users_controller.rb`: <http://gist.github.com/337147>

```
class UsersController < ApplicationController
  before_action :set_user, only: [:show, :edit, :update, :destroy]

  def new
    @user = User.new
  end

  def create
    @user = User.new(user_params)
    if @user.save
      redirect_to articles_path, notice: 'User successfully added.'
    else
      render action: :new
    end
  end
end
```

```

def edit
end

def update
  if @user.update(user_params)
    redirect_to articles_path, notice: 'Updated user information successfully.'
  else
    render action: 'edit'
  end
end

private
def set_user
  @user = User.find(params[:id])
end

def user_params
  params
end
end

```

You add four actions: new, create, edit, and update. The actions you add look very similar to the ones you saw in the articles controller in Chapter 7. The main difference is that you aren't using the respond_to block; therefore, Rails directly renders the default erb templates. Let's create those templates: Listings 8-3 and 8-4 show the new and edit templates, respectively.

Listing 8-3. New User Template in app/views/users/new.html.erb: <http://gist.github.com/337207>

```

<h1>New user</h1>

<%= render 'form' %>

<%= link_to 'Back', articles_path %>

```

Listing 8-4. Edit User Template in app/views/users/edit.html.erb: <http://gist.github.com/337206>

```

<h1>Editing user</h1>

<%= render 'form' %>

<%= link_to 'Back', articles_path %>

```

In both the new and edit templates, you render a form partial, which is expected to be in app/views/users/_form.html.erb. Create the form partial and make sure it looks like the code in Listing 8-5.

Listing 8-5. User Form Partial in app/views/users/_form.html.erb: <http://gist.github.com/337209>

```
<%= form_for(@user) do |f| %>
  <% if @user.errors.any? %>
    <div id="error_explanation">
      <h2><%= pluralize(@user.errors.count, "error") %>
prohibited this user from being saved:</h2>
      <ul>
        <% @user.errors.full_messages.each do |msg| %>
          <li><%= msg %></li>
        <% end %>
      </ul>
    </div>
  <% end %>

<div class="field">
  <%= f.label :email %><br />
  <%= f.text_field :email %>
</div>
<div class="field">
  <%= f.label :password %><br />
  <%= f.password_field :password %>
</div>
<div class="field">
  <%= f.label :password_confirmation %><br />
  <%= f.password_field :password_confirmation %>
</div>
<div class="actions">
  <%= f.submit %>
</div>
<% end %>
```

You use the same form helpers discussed in Chapter 7: `text_field` for text input and `password_field` for password inputs. Before you go to the browser to try what you've created, you need to add users as a resource in your routes file. Edit `config/routes.rb` so it looks like the code in Listing 8-6.

Listing 8-6. Adding Users to `routes.rb` in `config/routes.rb`: <http://gist.github.com/337211>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles
  resources :users
end
```

To see it all in action, try adding a new user by visiting <http://localhost:3000/users/new>. The form should look like Figure 8-1.

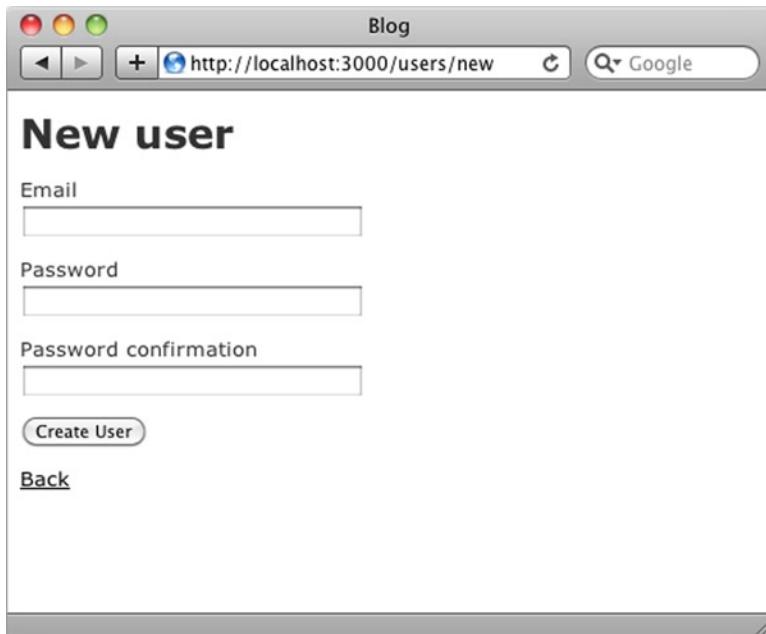


Figure 8-1. Adding a new user

When you try to actually create a user, you should receive an error message `ActiveModel::ForbiddenAttributesError`. It helpfully highlights the line in the code where the error occurred. Line 9 is where the params actually are added to the user. What happened? If you recall back in Chapter 7, when the scaffold generated the articles controller for us, it was very specific about what params should and shouldn't be sent. In our controller, we're just passing params into it. We need to specify which parameters are acceptable so nefarious users can't hack our system. Modify the `user_params` method to look like this:

```
 UsersController in app/controllers/users_controller.rb

def user_params
  params.require(:user).permit(:email, :password, :password_confirmation)
end
```

Run the server again, and retry the user creation. Now you can create a new user, and you can also edit that user if you have the user's ID. In fact, right now anyone can create and edit users; but shortly, you'll change the edit and update actions' implementation to make sure only users can edit their own profile.

Nested Resources

You added support for comments earlier, but only at the model level. You didn't implement a controller or view for the Comment model, and that's what you'll do now.

Comments are special because they aren't regular resources that you can implement in a regular RESTful controller, like articles or users are. Comments are more dependent on an article; they never exist on their own because they're conceptually meaningless if they're not tied to an article.

Instead of defining comments as standalone resources, as you did for articles, you define them as nested resources of articles. Go to the routes file and update the `:resources :article` call to look like the code in Listing 8-7.

Listing 8-7. Adding Comments to routes.rb in config/routes.rb: <http://gist.github.com/338747>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles do
    resources :comments
  end
  resources :users
end
```

To define a nested resource, you use the `resources` method passed inside a block to the parent resource. Notice how `resources :comments` is passed as a block to the `resources :articles` call; therefore, comments become a nested resource of articles. The named routes for nested resources are different from standalone ones; they're built on top of a singular article named route, requiring an article ID every time they're called. Table 8-1 lists the generated named routes for comments.

Table 8-1. Comments' Named Routes

Request Method	Nested Named Routes	Parameters	Controller Action
GET	<code>article_comments_path</code>	Article ID	<code>index</code>
POST	<code>article_comments_path</code>	Record hash	<code>create</code>
GET	<code>article_comment_path</code>	ID, article ID	<code>Show</code>
PUT/PATCH	<code>article_comment_path</code>	ID, article ID, and record hash	<code>update</code>
DELETE	<code>article_comment_path</code>	ID, article ID	<code>destroy</code>
GET	<code>edit_article_comment_path</code>	ID, article ID	<code>edit</code>
GET	<code>new_article_comment_path</code>	Article ID	<code>new</code>

Every time you call comment-named routes, you must provide an article ID. Let's generate the comments controller and see how you take care of that:

```
$ rails generate controller Comments
```

```
create app/controllers/comments_controller.rb
invoke erb
create app/views/comments
invoke test_unit
create test/controllers/comments_controller_test.rb
invoke helper
create app/helpers/comments_helper.rb
invoke test_unit
create test/helpers/comments_helper_test.rb
invoke assets
invoke coffee
create app/assets/javascripts/comments.js.coffee
invoke scss
create app/assets/stylesheets/comments.css.scss
```

Of the default seven actions for which Rails generates named routes, you need only two for comments: create and destroy. You don't need index, new, or show actions because comments are listed, shown, and added from the article's show page. You don't want to support editing or updating a comment, so you don't need edit or update either. Listing 8-8 shows how the comments controller looks with only those two actions.

Listing 8-8. Comments Controller in app/controllers/comments_controller.rb: <http://gist.github.com/338756>

```
class CommentsController < ApplicationController
  before_action :load_article

  def create
    @comment = @article.comments.new(comment_params)
    if @comment.save
      redirect_to @article, notice: 'Thanks for your comment'
    else
      redirect_to @article, alert: 'Unable to add comment'
    end
  end

  def destroy
    @comment = @article.comments.find(params[:id])
    @comment.destroy
    redirect_to @article, notice: 'Comment Deleted'
  end

  private
  def load_article
    @article = Article.find(params[:article_id])
  end

  def comment_params
    params.require(:comment).permit(:name, :email, :body)
  end
end
```

Notice the `before_action` call at the beginning of the controller; it runs the method `load_article` before all the actions in your comments controller. That's all you'll need to know for now. We'll explain more about controller callbacks shortly.

The `load_article` method does a simple task: it finds the article from the passed `article_id` and assigns it to the `@article` instance variable. Remember that you always have the `article_id` in your parameters because it's always included in your nested named routes. With `load_article` in `before_action`, you'll always have `@article` loaded and accessible in your comments controller's actions and templates.

Also notice how you find and assign `@comment`: you do so using `@article.comments`. This way, you'll make sure you're dealing only with `@article` comments and you don't create or delete comments from another article.

Now let's update the views and create some templates. As mentioned earlier, you *list*, *show*, and add *new* comments from the article's show page; so let's update the article show page, make it a little nicer, and then add new code to display comments. Listing 8-9 shows how `app/views/articles/show.html.erb` looks after the update.

Listing 8-9. Updated Article Show Template in app/views/articles/show.html.erb: <http://gist.github.com/338770>

```
<%= render @article %>

<h3>Comments</h3>
<div id="comments">
  <%= render @article.comments %>
</div>

<%= render :file => 'comments/new' %>
```

That's a lot of cleaning. First, you extract the displaying attributes into a partial named app/views/articles/_article.html.erb, which you call using render @article. One of the benefits of creating a partial is that you can use it in other pages, such as the articles' index page, which you'll implement shortly.

Notice that the flash notice is removed from the article show template. To make sure the flash messages show in any view template, you move it to the application layout in app/views/layouts/application.html.erb (Listing 8-10).

Listing 8-10. Updated Application Layout Template in app/views/layouts/application.html.erb:

<http://gist.github.com/388446>

```
<!DOCTYPE html>
<html>
<head>
  <title>Blog</title>
  <%= stylesheet_link_tag :all %>
  <%= javascript_include_tag :defaults %>
  <%= csrf_meta_tag %>
</head>
<body>

<%= content_tag :p, notice, :class => 'notice' if notice.present? %>
<%= content_tag :p, alert, :class => 'alert' if alert.present? %>

<%= yield %>

</body>
</html>
```

Then you list comments using the collection render on @article.comments. To refresh your memory, this loops through the article comments, rendering the app/views/comments/_comment.html.erb partial for every comment.

Finally, you render the app/views/comments/new.html.erb template as a file; you could render it as a partial too, but it's more like a template than a partial. In some projects, you may need to render a regular template file—corresponding to an action—which can also be used like a partial in other parts of the application. In this case, to render the app/views/comments/new.html.erb file in the context of another template, you use render :file.

None of the files mentioned have been created yet. Let's do that now. Create app/views/articles/_article.html.erb, app/views/comments/_comment.html.erb, and app/views/comments/new.html.erb, as shown in Listings 8-11, 8-12, and 8-13, respectively.

Listing 8-11. Article Partial in app/views/articles/_article.html.erb: <http://gist.github.com/338785>

```
<%= div_for article do %>
  <h3>
    <%= link_to article.title, article %>
    <span class='actions'>
```

```

<%= link_to "Edit", edit_article_path(article) %>
<%= link_to "Delete", article, confirm: "Are you sure?", method: :delete %>
</span>
</h3>
<%= article.body %>
<% end %>

```

Listing 8-12. Comment Partial in app/views/comments/_comment.html.erb: <http://gist.github.com/338789>

```

<%= div_for comment do %>
  <h3>
    <%= comment.name %> &lt;%= comment.email %>&gt; said:
    <span class='actions'>
      <%= link_to 'Delete', article_comment_path(:article_id => @article,
:id => comment), confirm: 'Are you sure?', method: :delete %>
    </span>
  </h3>
  <%= comment.body %>
<% end %>

```

Listing 8-13. New Comment Template in app/views/comments/new.html.erb: <http://gist.github.com/338791>

```

<%= form_for([@article, @article.comments.new]) do |f| %>
  <div class="field">
    <%= f.label :name %><br />
    <%= f.text_field :name %>
  </div>
  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :body %><br />
    <%= f.text_area :body %>
  </div>
  <div class="actions">
    <%= f.submit 'Add' %>
  </div>
<% end %>

```

The article and comment partials are pretty straightforward; aside from the markup, you display the attributes and link to actions. You also wrap the content inside a `div_for` helper, which renders a `div` element with a unique ID for the object passed as a parameter.

The new comment form calls `form_for`: instead of passing a single object as you did for `article` when you called `form_for(@article)`. You pass an array of both the `article` object `@article` and a new comment object by calling `@article.comments.new`. This is a short way to tell Rails that you're dealing with a nested resource. Calling `form_for([@article, @article.comments.new])` is identical to

```
form_for(:comment, @article.comments.new, url: [@article, @article.comments.new])
```

which is also identical to

```
form_for(:comment, @article.comments.new, url: article_comments_path(article_id: @article))
```

Passing the array of both the parent object and the nested child object translates into the corresponding nested named route. You can replace the named route path of the delete link in `app/views/comments/_comment.html.erb` with `[@article, comment]`. Change the comment partial to look like the code in Listing 8-14.

Listing 8-14. Delete Link Changed in `app/views/comments/_comment.html.erb`: <http://gist.github.com/338801>

```
<%= div_for comment do %>
<h3>
  <%= comment.name %> &lt;%= comment.email %>&gt; said:
  <span class='actions'>
    <%= link_to 'Delete', [@article, comment], :confirm => 'Are you sure?', :method => :delete %>
  </span>
</h3>
<%= comment.body %>
<% end %>
```

Now that you've created the missing templates and added the required code to the controller, let's go to the browser and see how it looks in the article show page. Run your server, go to your browser, and click your way through to an article; you should see something very similar to Figure 8-2.

The screenshot shows a web browser window titled "Blog". The address bar displays the URL `http://localhost:3000/articles/1`. The main content area shows an article with the title "Advanced Active Record Edit Delete". Below the title, there is a note: "Models need to relate to each other. In the real world, ..". A section titled "Comments" contains three input fields: "Name" (with an empty input box), "Email" (with an empty input box), and "Body" (a large text area). At the bottom of the form is a "Add" button.

Figure 8-2. Article show page with new comment form

Try adding a few comments and see how the form behaves. Congratulations! You just added comment support to your blog application using nested resources. Now that you have comments and users in the system, let's add some authorization logic to make sure only logged-in users can create and update articles.

Sessions and the Login/Logout Logic

The whole point of having users in your blog application is to allow them to create their *own* articles. But you also need to be able to recognize them when they create an article. Web applications normally do that by using *sessions*. Let's talk a little more about that before you implement it in your application.

Lying in State

HTTP is *stateless*. The web server has no idea that it has talked to your browser before; each request is like a blind date. Given this tidbit of information, you may wonder how you can stay logged in to a given site. How can the application remember you're logged in if HTTP is stateless? The answer is that you fake state.

You've no doubt heard of browser cookies. In order to simulate state atop HTTP, Rails uses cookies. When the first request comes in, Rails sets a cookie on the client's browser. The browser remembers the cookie locally and sends it along with each subsequent request. The result is that Rails is able to match the cookie that comes along in the request with session data stored on the server.

Using the Session

Secure in the knowledge that Rails will take care of all the low-level details of sessions for you, using the `session` object couldn't be easier. The `session` is implemented as a hash, just like `flash`. We should come clean here—`flash` is a session in disguise (you can think of it as a specialized session due to its autoexpiring properties). Not surprisingly then, the `flash` and `session` interfaces are identical. You store values in the `session` according to a key:

```
session[:account_id] = @account.id
session[:account_id] # => 1

session['message'] = "Hello world!"
session['message'] # => "Hello world!"
```

Session as a Resource

Now that you understand sessions, you can go back to your main task: allowing users to log in and log out. You create a session when the user logs in and clear (destroy) it when they're done. Of course, you do that in a RESTful way, by treating the session as a resource:

Start by generating a sessions controller:

```
$ rails generate controller Sessions
```

```
create  app/controllers/sessions_controller.rb
invoke  erb
create  app/views/sessions
invoke  test_unit
create  test/controllers/sessions_controller_test.rb
invoke  helper
```

```
create app/helpers/sessions_helper.rb
invoke test_unit
create test/helpers/sessions_helper_test.rb
invoke assets
invoke coffee
create app/assets/javascripts/sessions.js.coffee
invoke scss
create app/assets/stylesheets/sessions.css.scss
```

Now define this as a resource in your routes file in config/routes.rb, as shown in Listing 8-15.

Listing 8-15. Adding session to routes.rb in config/routes.rb: <http://gist.github.com/338904>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles do
    resources :comments
  end
  resources :users
  resource :session
end
```

Notice that you define session as a resource and not resources, because you never deal with a set of sessions at once. You never list sessions in an index or anything like that—you just need to create or destroy a *single* session at a time.

Let's step back and try to explain the difference between resource and resources definitions. The main benefit you get from defining resources in your routes file is the named routes that are generated for you. In case of a single resource definition, you get different named routes: none of them are pluralized, all are singular, and there's no index action. Rails maps six actions instead of the seven in a resources definition. Table 8-2 provides a quick comparison between resources named routes and resource named routes.

Table 8-2. Named Routes: resources vs. resource

Request Method	resources Named Routes	resource Named Routes	Controller Action
GET	articles_path	Not available	index
POST	articles_path	session_path	create
GET	article_path	session_path	show
PATCH/PUT	article_path	session_path	update
DELETE	article_path	session_path	destroy
GET	edit_article_path	edit_session_path	edit
GET	new_article_path	new_session_path	new

Note Although a singular name is used for the resource, the controller name is still taken from the plural name, so sessions_controller is the controller for the session resource in this case.

To avoid confusion, let's map this in your mind; to log in, you need to create a session; to log out, you clear that session. You use `new_session_path` as your login path, and the new template is your login page. POSTing the form in the new session page to `session_path` creates the session. Finally, submitting a DELETE request to `session_path` clears that session, performing a log out. Now, let's map it in the routes file, as shown in Listing 8-16.

Listing 8-16. Adding session to routes.rb in config/routes.rb: <http://gist.github.com/338912>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles do
    resources :comments
  end
  resources :users
  resource :session, :only => [:new, :create, :destroy]
  get '/login' => "sessions#new", :as => "login"
  get '/logout' => "sessions#destroy", :as => "logout"
end
```

You basically define two named routes, `login_path` and `logout_path`, which are more meaningful than `new_session_path` and `session_path` when referring to those actions.

Logging In a User

As you did for Active Record resources, in the `create` action, you first check the validity of the resource—in this case through authentication—and you save the state if all is good. If the validity check fails, you return the user to the login page with an error message. In this controller, you never save a record to the database—you save a session object. Listing 8-17 shows the `create` action.

Listing 8-17. The create Method in app/controllers/sessions_controller.rb: <http://gist.github.com/338919>

```
class SessionsController < ApplicationController
  def create
    if user = User.authenticate(params[:email], params[:password])
      session[:user_id] = user.id
      redirect_to root_path, :notice => "Logged in successfully"
    else
      flash.now[:alert] = "Invalid login/password combination"
      render :action => 'new'
    end
  end
end
```

First, you use the `authenticate` class method from the `User` model to attempt a log in (see Listing 6-37 in Chapter 6). Remember that `authenticate` returns a `User` object if the authentication succeeds; otherwise, it returns `nil`. Therefore, you can perform your conditional and your assignment in one shot using `if user = User.authenticate(params[:email], params[:password])`. If the assignment takes place, you want to store a reference to this user so you can keep the user logged in—a perfect job for the session:

```
session[:user_id] = user.id
```

Notice that you don't need to store the entire User object in session. You store just a reference to the user's ID. Why not store the entire User object? Well, think about this for a minute: what if the user is stored in session and later changes their login? The old login would remain in the session and would therefore be *stale*. This can cause problems if the underlying User model changes. Your entire object could become stale, potentially causing a NoMethodError when accessing attributes that didn't exist on the model at the time it was placed in session. The best bet is to just store the id.

With a reference to the logged-in user safely stored in session, you can redirect to the root path, corresponding to the articles controller.

If the assignment doesn't take place and the User.authenticate method returns nil, you know the provided login and password are invalid, and you return to the login page with an alert message using flash.now. RESTfully speaking, the login page is where you enter the *new* session information, so it's basically the new action.

Note flash.now differs from the regular flash call by setting a flash message that is only available to the current action. If you recall, regular flash makes messages available after a redirect.

But wait: you don't have a new action yet. Don't you need to define it first? The truth is, you don't need to initialize anything there—all you need is its template. By having the template, Rails automatically renders that template when it doesn't find the action definition. Let's create the new template, as shown in Listing 8-18.

Listing 8-18. The new Session Template in app/views/sessions/new.html.erb: <http://gist.github.com/338925>

```
<h1>Login</h1>

<%= form_tag session_path do %>
  <div class="field">
    <%= label_tag :email %><br />
    <%= text_field_tag :email %>
  </div>
  <div class="field">
    <%= label_tag :password %><br />
    <%= password_field_tag :password %>
  </div>
  <div class="actions">
    <%= submit_tag 'Login' %>
  </div>
<% end %>
```

Notice that you use form_tag instead of form_for, which you used earlier with Active Record objects; that's because session isn't an Active Record object. You also submit to session_path because it's a resource, not resources, as explained earlier. Again, because you aren't dealing with an Active Record object, you use label_tag, text_field_tag, and password_field_tag helpers.

Logging Out a User

The user is logged in when a session is created, so in order to log out the user, you need to clear that session. You do so in the destroy action. The destroy action is fairly straightforward. You clear the session by using the reset_session method that comes with Rails, which does exactly as it says: it resets the session by clearing all the values in it. After you clear the session, you redirect back to the login_path, which is your login screen.

Another way to do this is to specifically clear the `user_id` key from the session hash, but it's safer for the logout in particular to clear all the session values. Listing 8-19 shows how the sessions controller looks after you add the `destroy` method.

Listing 8-19. Updated Sessions Controller in `app/controllers/sessions_controller.rb`:
<http://gist.github.com/338944>

```
class SessionsController < ApplicationController
  def create
    if user = User.authenticate(params[:email], params[:password])
      session[:user_id] = user.id
      redirect_to root_path, :notice => "Logged in successfully"
    else
      flash.now[:alert] = "Invalid login/password combination"
      render :action => 'new'
    end
  end

  def destroy
    reset_session
    redirect_to root_path, :notice => "You successfully logged out"
  end
end
```

Go ahead and try it. Create a user by going to <http://localhost:3000/users/new>. Then log in by visiting the login path at <http://localhost:3000/login> (Figure 8-3). Finally, if you want to log out, go to <http://localhost:3000/logout>.

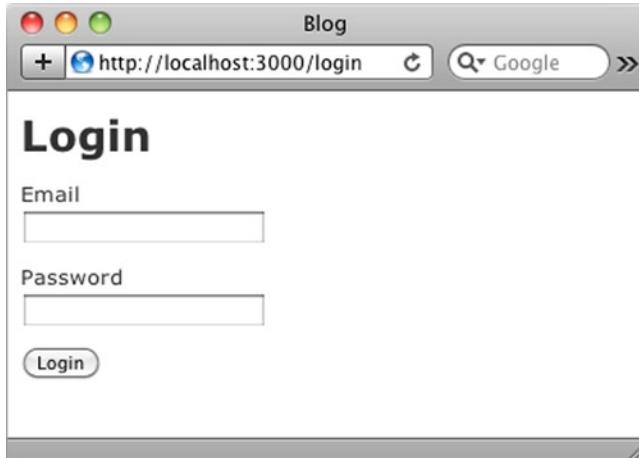


Figure 8-3. Login page

Don't worry about remembering all the URLs. You can link to them when you update your application layout.

Improving Controllers and Templates

Chapter 7 and earlier parts of this chapter covered generating controllers, creating templates and layouts, and DRYing up with partials. Let's take this a step forward: first, you update article views, and then you add callbacks to some of your controllers, making sure some actions require authorization.

Cleaning Up the Articles Index Page

The current articles' index page uses a table markup to list articles. If you've ever visited a blog, you know you've never seen one like that; so, let's change the table markup and loop to a friendlier markup that uses the article's partial in `app/views/articles/_article.html.erb`. Listing 8-20 shows the updated articles index.

Listing 8-20. Updated Articles Index in `app/views/articles/index.html.erb`: <http://gist.github.com/338960>

```
<h1>Listing articles</h1>

<div id="articles">
  <%= render @articles %>
</div>

<br />

<%= link_to 'New article', new_article_path %>
```

Caution Be careful with reusing partials. In some cases, you may prefer to keep separate files. You reuse the article partial here just to simplify things.

Visit your root path at <http://localhost:3000>. If all goes right, you should see something similar to Figure 8-4. That looks like a real blog!

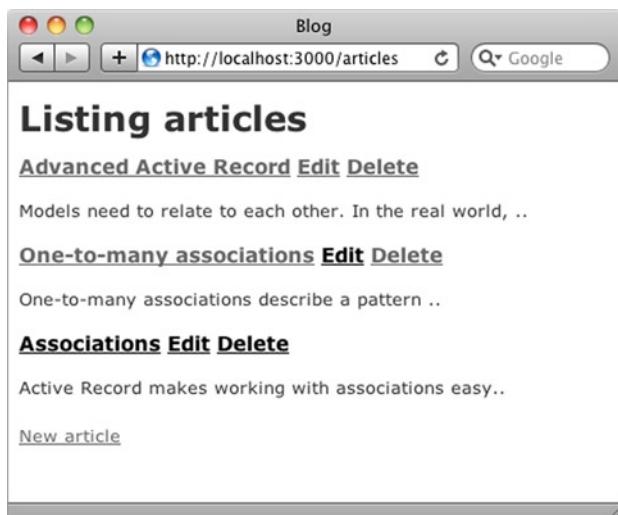


Figure 8-4. Blog-like home page

Adding Categories to the Article Form

In Chapter 6, you added categories to the Article model, but neither your controller nor your templates know about this yet. Let's remedy that now, starting with the article form. Add the code shown in bold in Listing 8-21 to the form partial in `app/views/articles/_form.html.erb`.

Listing 8-21. Modified `app/views/articles/_form.html.erb`: <http://gist.github.com/341499>

```
<%= form_for(@article) do |f| %>
  <% if @article.errors.any? %>
    <div id="error_explanation">
      <h2><%= pluralize(@article.errors.count, "error") %>
      prohibited this article from being saved:</h2>
      <ul>
        <% @article.errors.full_messages.each do |msg| %>
          <li><%= msg %></li>
        <% end %>
      </ul>
    </div>
  <% end %>
  <div class="field">
    <%= f.label :title %><br />
    <%= f.text_field :title %>
  </div>
  <div class="field">
    <%= f.label :location %><br />
    <%= f.text_field :location %>
  </div>
  <div class="field">
    <%= f.label "Categories" %><br />
    <% Category.all.each do |category| %>
      <%= check_box_tag 'article[category_ids][]', category.id,
@article.category_ids.include?(category.id), id: dom_id(category) %>
        <%= label_tag dom_id(category), category.name, class: "check_box_label" %>
    <% end %>
  </div>
  <div class="field">
    <%= f.label :excerpt %><br />
    <%= f.text_field :excerpt %>
  </div>
  <div class="field">
    <%= f.label :body %><br />
    <%= f.text_area :body %>
  </div>
  <div class="field">
    <%= f.label :published_at %><br />
    <%= f.datetime_select :published_at %>
  </div>
  <div class="actions">
    <%= f.submit %>
  </div>
<% end %>
```

To offer articles the chance to be part of one or more categories, you show all the categories as check boxes. But how do you associate those check boxes with the article?

Remember that Chapter 6 talked about the methods that each association adds to your model when you use them. In the case of the Article model, the `has_and_belong_to_many :categories` association adds the `category_ids` method, which returns an array of the associated category IDs; it also adds the `category_ids=(category_ids)` method, which replaces the current associated categories with the ones supplied.

Knowing that, look back at the new code added to the form: you loop through all the categories and draw a check box for each one. The `check_box_tag` call deserves a closer look:

```
check_box_tag 'article[category_ids][]', category.id,
@article.category_ids.include?(category.id), :id => dom_id(category)
```

Notice how you use a form tag helper `check_box_tag` instead of an Active Record form helper `check_box`. You do that so you can customize the name of the generated check-box form elements.

The `check_box_tag` method is an Action View form helper that generates a check box. The `check_box_tag` method takes four arguments: the first defines the name of the check box, the second specifies the value of the check box, and the third defines whether the check box is checked. Finally, an optional fourth argument can be passed; it's a hash of HTML options. In this case, you pass `id: dom_id(category)` as the last parameter to give the check box an ID you refer to later.

The name you use is `article[category_ids][]`. The `article[category_ids]` part associates this field with the `params[:article]` hash of attributes; the `[]` at the end of the name makes sure all the check-box values are passed as arrays. For example, if categories 5 and 2 are selected, `params[:article][:category_ids]` is `[5, 2]` when the form is submitted.

The only thing left to do is go back to the `articles` controller and make sure Rails knows that you want to allow categories to be saved to the article. Otherwise it would just discard this information and never save it to the article.

Code listing app/controllers/articles.rb

```
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at,
:category_ids => [])
end
```

When Rails generates the HTML markup from a form tag helper, it autogenerateds the element's ID from the element's name. For example, the HTML ID of `check_box_tag('article[category_ids]')` is `article_category_ids`. If you loop over all the categories and depend on Rails to generate the IDs of the check boxes, all the rendered check boxes have the same HTML ID. Not only are it standards incompliant to have multiple elements with the same ID, but it also makes your life harder if you want to refer to any of those elements later. You solve this problem by setting the ID of every check box by passing the ID to `check_box_tag` as the last argument.

The `dom_id(category)` method is a helper method that generates an HTML-friendly ID for your object. For example, if a category's id is 3, the `dom_id` of this category is `category_3`.

That's about it. Now that you have category integration for articles, try adding a new article; you should see a form similar to that in Figure 8-5.

The screenshot shows a web browser window titled "Blog" with the URL "http://localhost:3000/articles/new". The page displays a "New article" form. The "Title" field contains "Attending Rails Workshop". The "Location" field contains "Miami". Under "Categories", "Event" and "Programming" are checked, while "Music", "TV", and "Travel" are unchecked. The "Excerpt" field contains "Rails workshop in Miami". The "Body" field contains the text "I'm attending a Ruby on Rails workshop in Miami this summer.". Below the form, a "Published at" date selector is set to "2010 May 3 : 19 : 26". A "Create Article" button is visible, along with a "Back" link.

Figure 8-5. Updated article form with category check boxes

Fill in the mandatory fields, select a couple of categories, and submit the form. Check the parameters output in your rails server window. You should see something similar to the following output, depending on the values you entered—pay attention to the category array:

```
Parameters: {"authenticity_token"=>"FqpxvxGqPz1Q/biDC+muT2KGyFHp10CIbm+4KjBi1jI=",
"article"=>{"title"=>"Attending Rails Workshop", "location"=>"Miami",
"category_ids"=>["2", "1", "3"], "excerpt"=>"Rails workshop in Miami",
"body"=>"I'm attending a Ruby on Rails workshop in Miami this summer.",
"published_at(1i)"=>"2010", "published_at(2i)"=>"5", "published_at(3i)"=>"3",
"published_at(4i)"=>"19", "published_at(5i)"=>"26"}, "commit"=>"Create Article"}
```

If you try to edit the article you just created, you'll see that your categories are selected, and you can modify them like any other article attribute. The `category_ids=` method that the `has_and_belong_to_many` association added for you does all the magic behind the scenes.

Using Controller Callbacks

Callbacks provide a way for you to perform operations either before or after an action is invoked. There's even an around callback that can wrap the execution of an action. Of the three, the before action is the most commonly used, so this section focuses on it.

All the code you place in `before_action` is run before the action in question is called. Pretty simple, really. But there's a catch: if `before_action` returns `false`, the action isn't executed. We often use this to protect certain actions that require a log in. If you have an events controller and you want the new and create actions to remain open (anyone can access them), but you want to restrict all other actions to logged-in users, you can do so using callbacks:

```
class EventsController < ApplicationController
  before_action :authenticate, :except => [:new, :create]
end
```

This causes the `authenticate` method to be run before every action except those listed. Assume the `authenticate` method is defined in the `application_controller` controller and is therefore available to every other controller in the system. If the `authenticate` method returns `false`, the requested action isn't executed, thereby protecting it from unauthorized visitors.

You can also use the `:only` modifier to specify that the filter is to run for *only* the given actions. You can write the preceding example more concisely as follows:

```
before_action :authenticate, :only => :destroy
```

Without the `:only` or `:except` modifier, the filter runs for all actions.

Controller inheritance hierarchies share filters downward, but subclasses can also add or skip filters without affecting the superclass. Let's say you apply a global filter to the `application_controller`, but you have a particular controller that you want to be exempt from filtration. You can use `skip_before_action`, like this:

```
class ApplicationController < ActionController::Base
  before_action :authenticate_with_token
end

class PublicController < ApplicationController
  # We don't want to check for a token on this controller
  skip_before_action :authenticate_with_token
end
```

Callbacks are a fairly involved topic, and we've only scratched the surface here. Still, you've seen the most common usage pattern: protecting actions. For more information about filters, including usage examples, check out the Rails guide at http://guides.rubyonrails.org/action_controller_overview.html.

Requiring Authentication with Filters

In your blog application, you want to protect blog creation and modification, restricting access to registered users. To do this, you use callbacks that call specific methods and check for the `user_id` session you set on user log in. Recall that any methods you add to the `application_controller` are available to all other controllers (because it's the superclass of all controllers).

Open the `application_controller` in `app/controllers/application_controller.rb` and add the protected methods that enforce your authentication requirement, as shown in Listing 8-22.

Listing 8-22. Modified app/controllers/application_controller.rb: <http://gist.github.com/341622>

```
class ApplicationController < ActionController::Base
  protect_from_forgery

  protected
    # Returns the currently logged in user or nil if there isn't one
    def current_user
      return unless session[:user_id]
      @current_user ||= User.find_by_id(session[:user_id])
    end

    # Make current_user available in templates as a helper
    helper_method :current_user

    # Filter method to enforce a login requirement
    # Apply as a before_action on any controller you want to protect
    def authenticate
      logged_in? || access_denied
    end

    # Predicate method to test for a logged in user
    def logged_in?
      current_user.is_a? User
    end

    # Make logged_in? available in templates as a helper
    helper_method :logged_in?

    def access_denied
      redirect_to login_path, notice: "Please log in to continue"
    and return false
    end
  end
```

The `current_user` method acts like an *accessor* for the currently logged-in user. Because it returns a `User` object, you can call instance methods of `User` on it, such as `current_user.email`. The `authenticate` method is your callback method (the one you call from individual controllers). It checks whether there is a currently logged-in user via `logged_in?` (which, in turn, checks that there is actually a `User` returned by `current_user`) and calls `access_denied` if there isn't; `access_denied` redirects to the `login_path` in the sessions controller with a notice message in the flash.

You want two of these methods available in templates as well: `logged_in?` and `current_user`. Having `logged_in?` available allows you to make dynamic decisions about whether a user is logged in. You can use this to show or hide administrative controls (such as adding or editing a given article). Having `current_user` around also proves useful in templates, allowing you to access information about users, such as their e-mail addresses. Rails provides a handy way to extend the visibility of methods to templates by declaring them as helpers. You can use `helper_method` followed by a symbolic reference to the method in question, as you would in `helper_method :logged_in?`. You can also pass an array of method references to `helper_method` if you want to declare them all at once.

Let's apply the filter to the `articles` controller now. You also apply a filter to the `users` controller to restrict who can edit user profiles.

Applying Filters to Controllers

You apply filters using a declarative syntax. In this case, you'll want to check that a user is authenticated *before* you process a protected action, so you use `before _filter`. Add the filter to the `articles` controller, just inside the class body, as shown in Listing 8-23.

Listing 8-23. Before Filter Added in `app/controllers/articles_controller.rb`: <http://gist.github.com/341628>

```
class ArticlesController < ApplicationController
  before_action :authenticate, except: [:index, :show]

  ...
end
```

Notice how you're able to selectively apply the filter to specific actions. Here, you want every action to be protected *except* `index` and `show`. The `:except` modifier accepts either a single value or an array. You'll use an array here. If you want to protect only a few actions, you can use the `:only` modifier, which, as you would expect, behaves the opposite of `:except`.

You'll also want to use a filter in the `users` controller. Right now, anyone can edit a user as long as they know the user's ID. This would be risky in the real world. Ideally, you want the `edit` and `update` actions to respond only to the currently logged-in user, allowing that user to edit their profile. To do this, instead of retrieving `User.find(params[:id])`, you retrieve `current_user` and apply a filter to protect the `edit` and `update` actions. Listing 8-24 shows the latest version of the `users` controller.

Listing 8-24. Before Filter Added in `app/controllers/users_controller.rb`: <http://gist.github.com/341632>

```
class UsersController < ApplicationController
  before_action :authenticate, only: [:edit, :update]
  before_action :set_user, only: [:show, :edit, :update, :destroy]

  def new
    @user = User.new
  end

  def create
    @user = User.new(user_params)
    if @user.save
      redirect_to articles_path, notice: 'User successfully added.'
    else
      render action: :new
    end
  end

  def edit
    @user = current_user
  end

  def update
    @user = current_user
    if @user.update(user_params)
      redirect_to articles_path, notice: 'Updated user information successfully.'
    end
  end
```

```

else
  render action: 'edit'
end
end

private
def set_user
  @user = current_user
end

def user_params
  params.require(:user).permit(:email, :password, :password_confirmation)
end
end

```

Try it. If you attempt to add, edit, or delete an article, you're asked to log in (Figure 8-6).

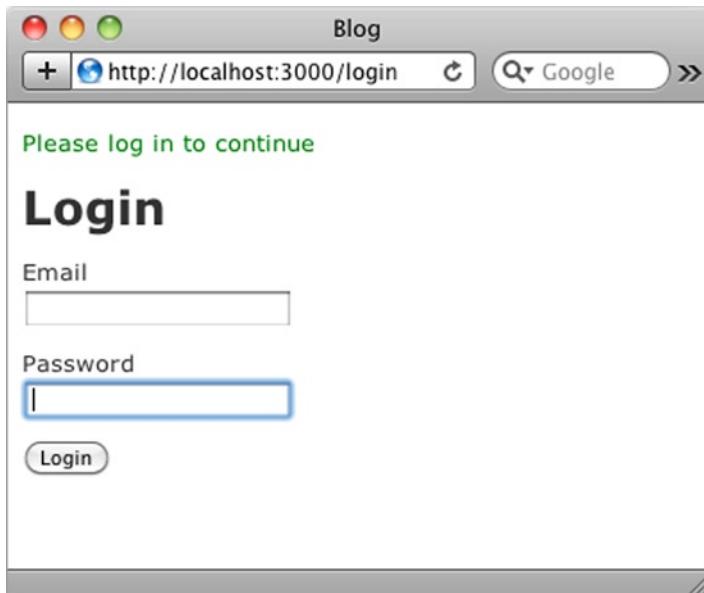


Figure 8-6. Authentication required

You don't want to keep deleting a comment as a public task; therefore, authorization code is required in the comments controller. First, you add a `before_filter` to authorize users before calling the `destroy` action. Next, in the `destroy` action, you find the article, making sure it belongs to the current user by using `current_user.articles.find`. Then, you find the comment on that article; and finally, you destroy it. Listing 8-25 shows the updated code, in bold, for the comments controller.

Listing 8-25. Authorization Before Deleting a Comment in `app/controllers/comments_controller.rb`:
<http://gist.github.com/388584>

```

class CommentsController < ApplicationController
  before_filter :load_article, :except => :destroy
  before_filter :authenticate, :only => :destroy

```

```

def create
  @comment = @article.comments.new(params[:comment])
  if @comment.save
    redirect_to @article, :notice => 'Thanks for your comment'
  else
    redirect_to @article, :alert => 'Unable to add comment'
  end
end

def destroy
  @article = current_user.articles.find(params[:article_id])
  @comment = @article.comments.find(params[:id])
  @comment.destroy
  redirect_to @article, notice: 'Comment deleted'
end

private
  def load_article
    @article = Article.find(params[:article_id])
  end

  def comment_params
    params.require(:comment).permit(:name, :email, :body)
  end
end

```

Adding Finishing Touches

You're almost finished with your work in this chapter. Only a few tasks remain. You need to spruce up your templates a bit and make them a little cleaner. You also need to make it possible for article owners to edit and delete their articles. Finally, you want to update the layout and apply some CSS styles to make things look pretty. Ready? Let's get started!

Using Action View Helpers

One of the ways you can clean up your templates is with helpers. Rails ships with a bevy of formatting helpers to assist in displaying numbers, dates, tags, and text in your templates. Here's a quick summary:

- *Number helpers*: The NumberHelper module provides methods for converting numbers into formatted strings. Methods are provided for phone numbers, currency, percentages, precision, positional notation, and file size. See <http://api.rubyonrails.org/classes/ActionView/Helpers/NumberHelper.html> for more information.
- *Text helpers*: The TextHelper module provides a set of methods for filtering, formatting, and transforming strings that can reduce the amount of in-line Ruby code in your views. See <http://api.rubyonrails.org/classes/ActionView/Helpers/TextHelper.html> for more information.
- *URL helpers*: Rails provides a set of URL helpers that makes constructing links that depend on the controller and action (or other parameters) ridiculously easy. For more information, see <http://api.rubyonrails.org/classes/ActionView/Helpers/UrlHelper.html> and <http://api.rubyonrails.org/classes/ActionController/Base.html>.

A very handy URL helper is `link_to`, which you've used several times already. It creates a hyperlink tag of the given name using a URL constructed according to the options hash given. It's possible to pass a string instead of an options hash to get a link tag that points to any URL. Additionally, if `nil` is passed as a name, the link itself becomes the name. Here's the fine print:

```
link_to(name, options={}, html_options={})
```

This generates an HTML anchor tag using the following parameters:

- The first argument is the link's name.
- The second argument is the URL to link to, given as a string, a named route, or a hash of options used to generate the URL. It can also be an object, in which case Rails replaces it with its show action named route.
- The third argument is a hash of HTML options for the resulting tag.

In Ruby, if the last argument to a method is a hash, the curly braces are optional. Most `link_to` helpers therefore look like this:

```
link_to 'New', new_article_path, id: 'new_article_link'
```

If you use all three arguments and pass in options for HTML (like a `class` or `id` attribute), you need to disambiguate them. Consider the following example, which uses two hashes: one for the URL generation and another for the HTML options:

```
link_to 'New', {controller: 'articles', action: 'new'}, class: 'large'
```

Notice that you need to use the curly braces for at least the first hash to inform Ruby that there are three arguments. Using braces on the last hash of options is still optional, and you can just as easily include them:

```
link_to 'New', {controller: 'articles', action: 'new'}, {class: 'large'}
```

Escaping HTML in Templates

You should always escape any HTML before displaying it in your views to prevent malicious users from injecting arbitrary HTML into your pages (which is how cross-site scripting attacks are often carried out). The rule of thumb is that whenever you have data that are provided by the user, you can't trust them blindly. You need to escape it. This includes model attributes as well as parameters. Fortunately, Rails escapes all rendered strings for you.

Try adding a new article with some HTML markup in the body, saving, and visiting the show page. If you enter an anchor HTML tag, for example, you see something like the screen shown in Figure 8-7. As you can see, Rails escapes the HTML entered in the body field.



Figure 8-7. Escaped HTML in the article page

If you check the source code, you'll see that the characters you entered have been escaped:

```
&lt;a href="quot;#quot;&gt;No link for you&lt;/a&gt;
```

Sometimes, you may want to display the strings entered by users without escaping them. To do that, Rails provides a method on strings named `html_safe` that skips the HTML escaping process. To display the article's body in its raw format, which you'll do shortly, you can call `article.body.html_safe` instead of `article.body` in the article partial in `app/views/articles/_article.html.erb`.

Formatting the Body Field

Let's improve the display of the body field. One of the aforementioned text helpers is `simple_format`. This helper converts text to HTML using simple formatting rules. Two or more consecutive newlines are considered a paragraph and wrapped in `<p>` tags. One newline is considered a line break, and a `
` tag is appended. Listing 8-26 shows the additions.

Listing 8-26. Formatting Helpers Added in app/views/articles/_article.html.erb:
<http://gist.github.com/341839>

```
<%= div_for article do %>
  <h3>
    <%= link_to article.title, article %>
    <span class='actions'>
      <%= link_to "Edit", edit_article_path(article) %>
      <%= link_to "Delete", article, :confirm => "Are you sure?", :method => :delete %>
    </span>
  </h3>
  <%= simple_format article.body %>
<% end %>
```

Adding Edit Controls

You've applied authentication filters, but you still don't have a way to prevent users from editing or deleting articles that belong to other users. To do this, you add a method to the Article model that can tell you whether the article in question is owned by the user you pass in. Open the Article model and add the `owned_by?` method, as highlighted in bold in Listing 8-27.

Listing 8-27. Updated app/models/article.rb:
<http://gist.github.com/388527>

```
class Article < ActiveRecord::Base
  validates :title, :presence => true
  validates :body, :presence => true

  belongs_to :user
  has_and_belongs_to_many :categories
  has_many :comments

  scope :published, where("articles.published_at IS NOT NULL")
  scope :draft, where("articles.published_at IS NULL")
  scope :recent, lambda { published.where("articles.published_at > ?", 1.week.ago.to_date) }
  scope :where_title, lambda { |term| where("articles.title LIKE ?", "%#{term}%") }

  def long_title
    "#{title} - #{published_at}"
  end

  def published?
    published_at.present?
  end

  def owned_by?(owner)
    return false unless owner.is_a?(User)
    user == owner
  end
end
```

Now, let's use this method in the article and comment partials in `app/views/articles/_article.html.erb` and `app/views/comments/_comment.html.erb`, respectively, by adding links to edit or delete *only* if the article is owned by the currently logged-in user, as shown in Listings 8-28 and 8-29.

Listing 8-28. Edit Controls for Article in `app/views/articles/_article.html.erb`: <http://gist.github.com/341845>

```
<%= div_for article do %>
<h3>
  <%= link_to article.title, article %>
  <% if article.owned_by? current_user %>
    <span class='actions'>
      <%= link_to "Edit", edit_article_path(article) %>
      <%= link_to "Delete", article, confirm: "Are you sure?", method: :delete %>
    </span>
  <% end %>
</h3>
<%= simple_format article.body %>
<% end %>
```

Listing 8-29. Edit Controls for Comment in `app/views/comments/_comment.html.erb`: <http://gist.github.com/388593>

```
<%= div_for comment do %>
<h3>
  <%= comment.name %> &lt; <%= comment.email %>&gt; said:
  <% if @article.owned_by? current_user %>
    <span class='actions'>
      <%= link_to 'Delete', [@article, comment], confirm: 'Are you sure?', method: :delete %>
    </span>
  <% end %>
</h3>
<%= comment.body %>
<% end %>
```

Note When you try this in your browser, you may not see the edit and delete links for any of the articles because their `user_id` field is `nil`. This is great console practice for you. Start your console with `rails console`, find your own user section using `user = User.find_by_email('email@example.com')`, and update all articles in the system using `Article.update_all(["user_id = ?", user.id])`.

Making Sure Articles Have Owners

You need to make sure that when you add an article, a user is assigned. To do that, you update the `create` method in the `articles` controller to use the association between `User` and `Article`. When creating the `@article` variable, instead of using `Article.new`, you use `current_user.articles.new`: it instantiates an `article` object with the `user_id` field set to the ID of `current_user`. That's exactly what you need.

Applying the same logic, you can change the `edit`, `update`, and `destroy` actions to retrieve only articles belonging to the logged-in user. In code parlance, you'll use `current_user.articles.find` wherever you were using `Article.find`. Listing 8-30 shows the changes to make in `app/controllers/articles_controller.rb`.

Listing 8-30. Updated app/controllers/articles_controller.rb: <http://gist.github.com/341925>

```
class ArticlesController < ApplicationController
  before_action :authenticate, except: [:index, :show]
  before_action :set_article, only: [:show]

  # GET /articles
  # GET /articles.json
  def index
    @articles = Article.all
  end

  # GET /articles/1
  # GET /articles/1.json
  def show
  end

  # GET /articles/new
  def new
    @article = Article.new
  end

  # GET /articles/1/edit
  def edit
    @article = current_user.articles.find(params[:id])
  end

  # POST /articles
  # POST /articles.json
  def create
    @article = current_user.articles.new(article_params)

    respond_to do |format|
      if @article.save
        format.html { redirect_to @article, notice: 'Article was successfully created.' }
        format.json { render action: 'show', status: :created, location: @article }
      else
        format.html { render action: 'new' }
        format.json { render json: @article.errors, status: :unprocessable_entity }
      end
    end
  end

  # PATCH/PUT /articles/1
  # PATCH/PUT /articles/1.json
  def update
    @article = current_user.articles.find(params[:id])
    respond_to do |format|
      if @article.update(article_params)
        format.html { redirect_to @article, notice: 'Article was successfully updated.' }
        format.json { head :no_content }
      end
    end
  end
end
```

```

else
  format.html { render action: 'edit' }
  format.json { render json: @article.errors, status: :unprocessable_entity }
end
end
end

# DELETE /articles/1
# DELETE /articles/1.json
def destroy
  @article = current_user.articles.find(params[:id])
  @article.destroy
  respond_to do |format|
    format.html { redirect_to articles_url }
    format.json { head :no_content }
  end
end

private
# Use callbacks to share common setup or constraints between actions.
def set_article
  @article = Article.find(params[:id])
end

# Never trust parameters from the scary internet, only allow the white list through.
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at,
:category_ids => [])
end
end

```

Adding Custom Helpers

Your blog application is looking pretty good, but let's make it a bit more user friendly. One thing you can do is add a helpful cancel link beside each submit button on the forms, so users can back out of editing. You could do this by adding a `link_to` helper beside each button, but you'd need to do this for every form. Because you probably want to repeat this pattern throughout the application, this could end up being a lot of duplication. Why not create a custom helper to do this for you? Listing 8-31 shows the method `submit_or_cancel` added to the `application_helper`.

Listing 8-31. The `submit_or_cancel` Method in `app/helpers/application_helper.rb`:
<http://gist.github.com/341856>

```

module ApplicationHelper
  # Creates a submit button with the given name with a cancel link
  # Accepts two arguments: Form object and the cancel link name
  def submit_or_cancel(form, name='Cancel')
    form.submit + " or " +
      link_to(name, 'javascript:history.go(-1);', :class => 'cancel')
  end
end

```

Now, let's use this helper on your forms. Open both the user and the article form partials in `app/views/users/_form.html.erb` and `app/views/articles/_form.html.erb`, and update them so they look like Listings 8-32 and 8-33, respectively.

Listing 8-32. Updated `app/views/users/_form.html.erb`: <http://gist.github.com/341857>

```
<%= form_for(@user) do |f| %>
  <% if @user.errors.any? %>
    <div id="error_explanation">
      <h2><%= pluralize(@user.errors.count, "error") %>
prohibited this user from being saved:</h2>
      <ul>
        <% @user.errors.full_messages.each do |msg| %>
          <li><%= msg %></li>
        <% end %>
      </ul>
    </div>
  <% end %>

  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :password %><br />
    <%= f.password_field :password %>
  </div>
  <div class="field">
    <%= f.label :password_confirmation %><br />
    <%= f.password_field :password_confirmation %>
  </div>
  <div class="actions">
    <%= submit_or_cancel(f) %>
  </div>
<% end %>
```

Listing 8-33. Updated `app/views/articles/_form.html.erb`: <http://gist.github.com/341858>

```
<%= form_for(@article) do |f| %>
  <% if @article.errors.any? %>
    <div id="error_explanation">
      <h2><%= pluralize(@article.errors.count, "error") %>
prohibited this article from being saved:</h2>
      <ul>
        <% @article.errors.full_messages.each do |msg| %>
          <li><%= msg %></li>
        <% end %>
      </ul>
    </div>
  <% end %>
```

```

<div class="field">
  <%= f.label :title %><br />
  <%= f.text_field :title %>
</div>
<div class="field">
  <%= f.label :location %><br />
  <%= f.text_field :location %>
</div>
<div class="field">
  <%= f.label "Categories" %><br />
  <% for category in Category.all %>
    <%= check_box_tag 'article[category_ids][]', category.id,
@article.category_ids.include?(category.id), id: dom_id(category) %>
      <%= label_tag dom_id(category), category.name, class: "check_box_label" %>
    <% end %>
  </div>
<div class="field">
  <%= f.label :excerpt %><br />
  <%= f.text_field :excerpt %>
</div>
<div class="field">
  <%= f.label :body %><br />
  <%= f.text_area :body %>
</div>
<div class="field">
  <%= f.label :published_at %><br />
  <%= f.datetime_select :published_at %>
</div>
<div class="actions">
  <%= submit_or_cancel(f) %>
</div>
<% end %>

```

As in the earlier examples, every time you copy and paste view code in more than one template, it means that you very likely can extract it into a helper method.

Giving It Some Style

Your blog application could use a little varnish. Let's update the layout and apply a style sheet.

Updating the Layout

Let's update the main layout and add some style hooks that you can target via CSS. You also add some pieces to allow the user to log in, log out, edit their password, and add a new article. The final result looks like the code in Listing 8-34, with changes in bold.

Listing 8-34. Updated app/views/layouts/application.html.erb: <http://gist.github.com/341867>

```
<!DOCTYPE html>
<html>
<head>
  <title>Blog</title>
  <%= stylesheet_link_tag :all %>
  <%= javascript_include_tag :defaults %>
  <%= csrf_meta_tag %>
</head>
<body>
  <div id="header">
    <h1><%= link_to "Blog", root_path %></h1>
    <div id="user_bar">
      <% if logged_in? %>
        <%= link_to "New Article", new_article_path %> | 
        <%= link_to "Edit Password", edit_user_path(current_user) %> | 
        <%= link_to "Logout", logout_path %>
      <% else %>
        <%= link_to "Login", login_path %>
      <% end %>
    </div>
  </div>
  <div id="main">
    <%= content_tag(:p, notice, class: 'notice') if notice.present? %>
    <%= content_tag(:p, alert, class: 'alert') if alert.present? %>
    <%= yield %>
  </div>
  <div id="footer">
    A simple blog built for the book
    <a href="http://beginningrails.com">Beginning Rails 4</a>
  </div>
</body>
</html>
```

You add a link here to add a new article in the application layout; therefore, you no longer need that link on the articles' index page. Update the app/views/articles/index.html.erb file to remove the new article link. It should look like the code in Listing 8-35.

Listing 8-35. Remove New Article Link from app/views/articles/index.html.erb: <http://gist.github.com/341893>

```
<h1>Listing articles</h1>

<div id="articles">
  <%= render @articles %>
</div>
```

Applying a Style Sheet

We've prepared a simple CSS that you can apply to make the application look pretty. Listing 8-36 shows the application.css file, which you should create in public/stylesheets/application.css. You're no longer using the public/stylesheets/scaffold.css file; remove it to avoid any styling conflicts.

Listing 8-36. The app/assets/stylesheets/application.css File: <http://gist.github.com/341881>

```
/*
 * This is a manifest file that'll be compiled into application.css, which will include all the
files
 * listed below.
 *
 * Any CSS and SCSS file within this directory, lib/assets/stylesheets, vendor/assets/stylesheets,
* or vendor/assets/stylesheets of plugins, if any, can be referenced here using a relative path.
*
* You're free to add application-wide styles to this file and they'll appear at the top of the
* compiled file, but it's generally better to create a new file per style scope.
*
*= require_self
*= require_tree.
*/
* {
margin: 0 auto;
}

body {
background-color: #fff;
color: #333;
}

body, p, ol, ul, td {
font-family: verdana, arial, helvetica, sans-serif;
font-size: 13px;
line-height: 18px;
}

pre {
background-color: #eee;
padding: 10px;
font-size: 11px;
}

p {
padding: 5px;
}

a {
color:#D95E16;
padding:0 2px;
text-decoration:none;
}

a:hover {
background-color:#FF813C;
color:#FFFFFF;
}
```

```
.notice { color: green; }
.alert { color: red; }

#header, #main, #footer {
  width: 800px;
}

#header {
  font-family: "Myriad Web Pro", Helvetica, Arial, sans-serif;
  letter-spacing: 1px;
  border-bottom: 5px solid #333333;
  color: #333333;
  padding: 15px 0;
  height: 35px;
}

#header #user_bar {
  float: right;
  font-size: 10px;
}

#footer {
  border-top: 5px solid #C1C1C1;
  margin-top: 10px;
  clear: both;
  padding: 10px 0;
  text-align: center;
  font-size: 11px;
}

#header h1 {
  padding-top: 14px;
  float: left;
  font-size: 30px;
}

#header h1 a{
  color: black;
}

#header h1 a:hover {
  background-color: white;
  color: black;
  border-bottom: 4px solid #ccc;
}

#header p {
  float: right;
}
```

```
#main h1 {  
  font-size: 16px;  
  padding: 10px 0;  
  border-bottom: 1px solid #bbb;  
  margin-bottom: 10px;  
}  
  
#main table{  
  margin: 0;  
}  
  
#main form{  
  text-align: left;  
}  
  
#main form br{  
  display: none;  
  float: left;  
}  
  
#main form label {  
  width: 150px;  
  display: block;  
  text-align: right;  
  padding-right: 10px;  
  float: left;  
  line-height: 21px;  
  vertical-align: center;  
  background-color: #F0F0F0;  
  border: 2px solid #ccc;  
  margin-right: 10px;  
}  
  
#main form label.check_box_label {  
  width: auto;  
  display: inline;  
  text-align: right;  
  padding-right: 10px;  
  line-height: 21px;  
  vertical-align: center;  
  background-color: #FFF;  
  border: none;  
}  
  
#main form .field, #main form .actions {  
  padding-top: 10px;  
  clear: both;  
}
```

```
#main form input[type=text], #main form input[type=password], #main form textarea {  
  float: left;  
  font-size: 14px;  
  width: 250px;  
  padding: 2px;  
  border: 2px solid #ccc;  
}  
  
#main form input[type=checkbox] {  
  margin: 4px;  
  float: left;  
}  
  
#main form textarea {  
  height: 150px;  
}  
  
#main form input[type=submit] {  
  margin-left: 175px;  
  float:left;  
  margin-right: 10px;  
  margin-bottom: 10px;  
}  
  
#main h3 {  
  padding-top: 10px;  
  height: 20px;  
}  
  
#main h3 .actions{  
  display:none;  
  font-weight: normal;  
  font-size: 10px;  
}  
  
#main h3:hover .actions{  
  display: inline;  
}  
  
.field_with_errors {  
  display:table;  
  float:left;  
  margin:0;  
  width:100px;  
  margin-right: 10px;  
}  
  
#main form .field_with_errors label{  
  border: 2px solid red;  
  margin-right: 0px;  
}
```

```
#main form .field_with_errors input, #main form .field_with_errors textarea{
  width: 250px;
  border: 2px solid red;
}

#error_explanation {
  width: 413px;
  border: 2px solid red;
  padding: 7px;
  padding-bottom: 12px;
  margin-bottom: 20px;
  background-color: #f0f0f0;
  margin: 0;
}

#error_explanation h2 {
  text-align: left;
  font-weight: bold;
  padding: 5px 5px 5px 15px;
  font-size: 12px;
  margin: -7px;
  background-color: #c00;
  color: #fff;
}

#error_explanation p {
  color: #333;
  margin-bottom: 0;
  padding: 5px;
  margin: 0;
}

#error_explanation ul li {
  font-size: 12px;
  list-style: square;
}
```

Yikes! That's a lot of CSS! Don't worry, though. Like all the other listings in the book, you can get the code from the Gist URL in the listing caption. The code is also available on the book's web site (<http://www.apress.com>) so you can download it and copy it into your project. We certainly don't expect you to type it all in.

With the CSS in place, your application is starting to look nice. If you've done everything correctly, it should look a lot like Figure 8-8.

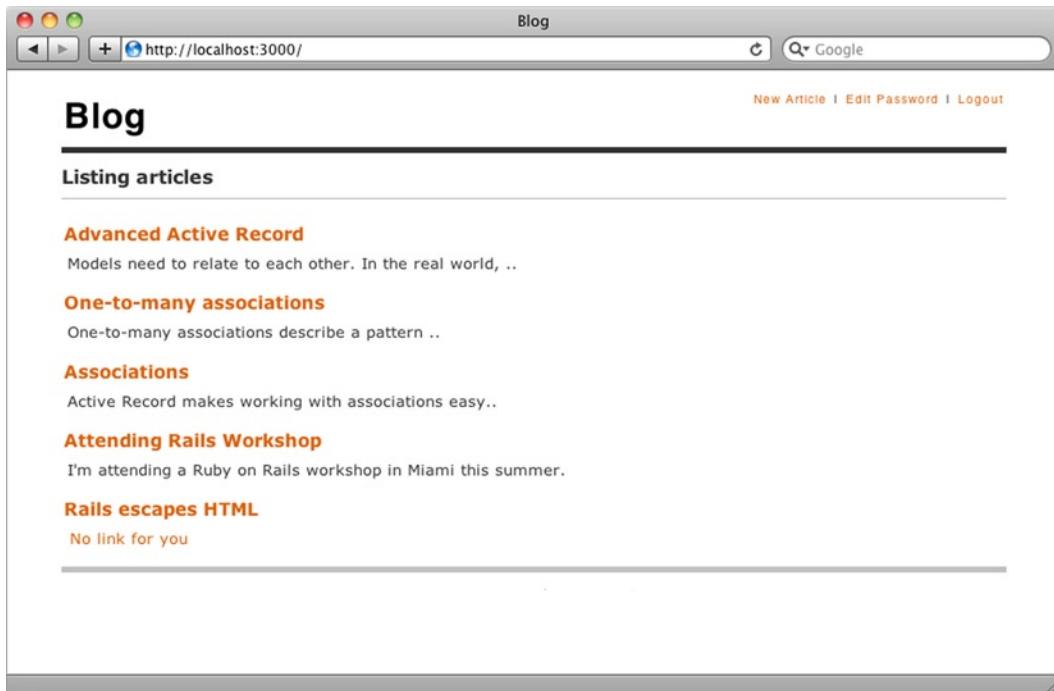


Figure 8-8. Final layout with CSS

Summary

This chapter discussed a fair number of advanced topics, including nested resources, sessions and state, and how to protect actions using filters. But we didn't stop there. You also learned how to sanitize HTML to protect your application from defacement by malicious users, and how to create your own helpers to improve your interface. You even took the time to make your application look pretty, sprucing up the layout and adding some CSS.

The next chapters build on this knowledge, and we'll start by going over how Rails handles JavaScript and CSS.



JavaScript and CSS

JavaScript and CSS (Cascading Style Sheets) are no longer relegated to supporting roles on the Web. They have become first-class citizens, and it would seem pretty strange to build a web application without them. It's now common to have hundreds or thousands of lines of JavaScript, and just as many lines of style sheets. Rails applications have always supported serving JavaScript and CSS, but it became clear that a better solution was needed. In Rails 3.1, the asset pipeline was introduced. This made JavaScript and CSS a first-class citizen of the Rails application and allowed for proper organization and handling of your scripts and styles. It handles combining and minifying files to support quick download times and provides support for languages that compile down to JavaScript (CoffeeScript) and CSS (SASS).

Note If you need to get the code at the exact point where you finished Chapter 8, download the source code zip file from the book's page on www.apress.com and extract it onto your computer.

Asset Concatenation and Compression

Applications that have a heavy JavaScript component can have hundreds of individual .js files. If a browser has to download all of these files, it causes a lot of overhead just starting and stopping the transfer of files. The asset pipeline concatenates both your JavaScript and style sheets into files so that a browser only has to download one or two files instead of hundreds. It also compresses, or minifies, the files. This removes things like comments, whitespace, and long variable names from the final output. The final product is functionally equivalent, but usually much smaller. Both of these features combine to make web applications load much faster and are transparent to the user.

Secondary Languages

Browsers have very strong support for both JavaScript and CSS, but if you want to use another language on the front end, you'd be out of luck. The browser would at best ignore it and at worse throw errors all over the screen. The asset pipeline allows you to use other languages that compile down to code that browsers understand. The two languages that the asset pipeline supports out of the box are CoffeeScript, which compiles to JavaScript, and SASS, which compiles to CSS.

Detailed description of CoffeeScript and SASS is out of the scope of this book, but you should know what they are if you encounter them. For more information on CoffeeScript, visit <http://coffeescript.org/>, and for more information on SASS, visit <http://sass-lang.com/>.

Asset Pipeline Locations

The asset pipeline allows you to place files in three different locations, which have different purposes but are functionally equivalent (Table 9-1).

Table 9-1. Locations for Files in the Asset Pipeline

File Location	Description
app/assets	This is for assets that are owned by the application. You can include images, style sheets, and JavaScript.
lib/assets	This location is for assets that are shared across applications but are owned by you. These assets don't really fit into the scope of this specific application but are used by it.
vendor/assets	This location is for assets that are from an outside vendor, like JavaScript or CSS frameworks.

In general, the asset pipeline stays out of the way, but it can empower you to do impressive things with your assets. For more information on the asset pipeline, visit http://guides.rubyonrails.org/asset_pipeline.html.

Turbolinks

Rails 4 now includes the Turbolinks gem. This gem will speed up your application by using Ajax to request pages instead of the more traditional page requests. It tracks files that are commonly shared across requests, like JavaScript and style sheets, and only reloads the information that changes. It attaches itself to links on your page instead of making those requests the traditional way. It makes an Ajax request and replaces the body tag of your document. Turbolinks also keeps track of the URL and manages the back and forward buttons. It's designed to be transparent to both users and developers.

Turbolinks is turned on by default in Rails 4. It is included in the default application.js file. If you needed to remove Turbolinks for some reason, you could do so by removing it from app/assets/javascripts/application.js, as shown in Listing 9-1.

Listing 9-1. Remove This Line from app/assets/javascripts/application.js to Turn Turbolinks Off

```
//= require turbolinks
```

By default Turbolinks attaches itself to every link on the page, but you can disable it for specific links by attaching a data-no-turbolink attribute to the link, as shown in Listing 9-2. This causes the link to behave in a traditional fashion.

Listing 9-2. Rails link_to Helper with a No-Turbolink Attribute Attached

```
link_to "Some Link", "/some-location", 'data-no-turbolink' => true
```

Note Some JavaScript libraries aren't compatible with Turbolinks. Listing these is out of the scope of this book, but you can find more information at <https://github.com/rails/turbolinks/blob/master/README.md>. If you continue to have problems, you can always disable Turbolinks.

Let's Build Something!

We've talked about the features of Rails that support JavaScript and CSS, but let's actually put JavaScript to work. We've added our style sheets in Chapter 8, but this chapter will focus on making our application use Ajax to load and submit forms.

Ajax and Rails

Ajax is a combination of technologies centered around the `XMLHttpRequest` object, a JavaScript API originally developed by Microsoft but now supported in all modern browsers. Of course, you could interface with the `XMLHttpRequest` API directly, but it wouldn't be fun. A far better idea is to use one of several libraries that abstracts the low-level details and makes cross-browser support possible.

Rails makes Ajax easier for web developers to use. Toward that end, it implements a set of conventions that enable you to implement even the most advanced techniques with relative ease.

Most of the Ajax features you implement in Rails applications are coded using JavaScript; so, familiarity with JavaScript code always helps and is pretty important for today's web developers.

jQuery and DOM

jQuery is a JavaScript library that lets you build JavaScript in a cross-browser way. It provides functionality for everything from traversing the Document Object Model (DOM) to making web requests, animation, and everything in between. Although there are multiple JavaScript libraries available on the Web, jQuery is easily the most widely used.

jQuery provides several utility functions that make working with JavaScript better. Perhaps the most important of these is the `$()` function, which is a selector function to access one or more elements in a web page. Truth be told, `$()` is far more than a simple selector wrapper. The real magic stems from the fact that any element accessed using it is automatically extended by jQuery.

The biggest part of the jQuery framework is its Document Object Model (DOM) extensions. These DOM extensions allow you to write things like `$('.comment').removeClass('active').hide();`, which gets all the elements with the CSS class name `comment`, removes the `active` class name, and hides all elements from the view. The elements with the class name `comment` wouldn't have these methods natively. Because they were fetched using `$()`, jQuery makes those functions available.

Tip Wikipedia defines DOM as follows: “The Document Object Model (DOM) is a cross-platform and language-independent convention for representing and interacting with objects in HTML, XHTML, and XML documents” (http://en.wikipedia.org/wiki/Document_Object_Model).

The `$()` function finds elements in a web page by matching them to a CSS selector. Table 9-2 lists some examples of the most commonly used CSS selectors. For a complete list, see www.w3.org/TR/CSS2/selector.html.

Table 9-2. CSS Selectors

Function	Description
<code>\$('#article_123')</code>	Returns the element matching the given ID <code>article_123</code>
<code>\$('.comment')</code>	Returns a list of elements with the class name <code>comment</code>
<code>\$('.div.article')</code>	Returns a list of <code>div</code> elements with the class name <code>article</code>

Moving to Practice

Now that you know what Ajax is, how it works, and the reasons behind using a JavaScript framework such as jQuery, you can apply some of this knowledge to enhance the usability of your application. Mainly, you use Ajax in your pages when you think a snappier interaction is possible and recommended. Let's begin Ajaxifying the blog application in the article page.

Not All Users Comment

If you look at the article page, you quickly notice that every time users read a post, they're presented with a form for entering comments. Although reader participation is paramount, most users are only interested in reading the content. You can modify the article page to not load the comment form automatically; instead, it will load the form only after a user clicks the new comment link.

Loading a Template via Ajax

One of the rules of good interface design is to make things snappy. That is to say, the interface should be responsive and quick to load. A good way to achieve this is to load elements (like forms or content areas) onto the page whenever the user requests them. Modify the article's show template, as shown in Listing 9-3.

Listing 9-3. The Article Partial in app/views/articles/show.html.erb: <http://gist.github.com/353475>

```
<%= render @article %>
<h3>Comments</h3>
<div id="comments">
  <%= render @article.comments %>
</div>
<%= link_to "new comment", new_article_comment_path(@article, :format => :js),
:remote => true, :id => 'new_comment_link' %>
```

The template hasn't changed a lot: you no longer directly render the comment form, and you add a link called new comment. The new link still uses the well-known `link_to` helper to generate a link; however, you pass in the `:remote => true` option, which tells Rails that you want the request triggered by this link to hit the server using Ajax.

There are a couple of things to note in the use of `link_to` in Listing 9-3. First, you send the request to a URL that already exists; the `new_article_comment_path` route identifies a path to a new comment. However, you pass in the `:format => :js` argument, which tells Rails you're requesting a JavaScript representation of your page, fitting your intention. Second, you use the `:id => 'new_comment_link'` option to give the rendered HTML element an ID that you can refer to later.

On the server side, you don't need to make any changes to the comments controller. As currently implemented, you don't explicitly implement a new action; the default behavior in this case is to render the new template in `app/views/comments/new.html.erb`. That isn't really what you want—this template shouldn't be the result of this JavaScript call. Instead, you want a separate JavaScript template to be used as response for this action.

Responding to Requests with `:format => :js`

Your controller reads the format parameter passed in from the browser. To make sure you send a response that includes JavaScript code, you must create a template with the `.js.erb` template extension. Create the `app/views/comments/new.js.erb` template as per Listing 9-4. The following text explains all the lines in the template to make sure you know what's happening.

Listing 9-4. The .js.erb New Comment Template in app/views/comments/new.js.erb:
<http://gist.github.com/353484>

```
$("=<%= escape_javascript render(file: 'comments/new.html.erb') %>").insertAfter('#comments');
$('#new_comment_link').hide();
```

The first line renders the existing app/views/comments/new.html.erb template into a string variable that you dynamically insert into the HTML page, right after the comments div—that's achieved by using jQuery's `insertAfter` method, which, as its name implies, inserts a block of HTML after an existing HTML element. Table 9-3 lists similar jQuery methods that you can use in place of `insertAfter`.

Table 9-3. jQuery Methods for Inserting HTML into a Page

Method	Description
<code>insertAfter(target)</code>	Inserts the preceding element after the target element passed as a parameter
<code>insertBefore(target)</code>	Inserts the preceding element before the target element passed as a parameter
<code>append(content)</code>	Appends the preceding element with the passed content element
<code>prepend(content)</code>	Prepends the preceding element with the passed content element

Going back to Listing 9-4, the last line hides the `new_comment_element`, which contains the link to add a new comment. Because you already have the comment form in your page, it makes little sense to keep that link around. Hiding an element is achieved by calling jQuery's `hide` method on the element you want to hide.

Note In a similar fashion, you can call jQuery's `show` method to display a hidden element. As you can see, it's very important that you know your way around jQuery to be able to build complex user interfaces.

Let's see what you built in practice. Open your browser to any existing article, such as <http://localhost:3000/articles/2>, and notice that the comments form is no longer there (Figure 9-1).



Figure 9-1. The article page without the comment form

As soon as you click the new comment link, the comment form pops into place, and you can add comments (Figure 9-2). You achieved your goal of keeping the user interface cleaner, while allowing users to quickly access functionality without having to move to a new page. That's a good start.

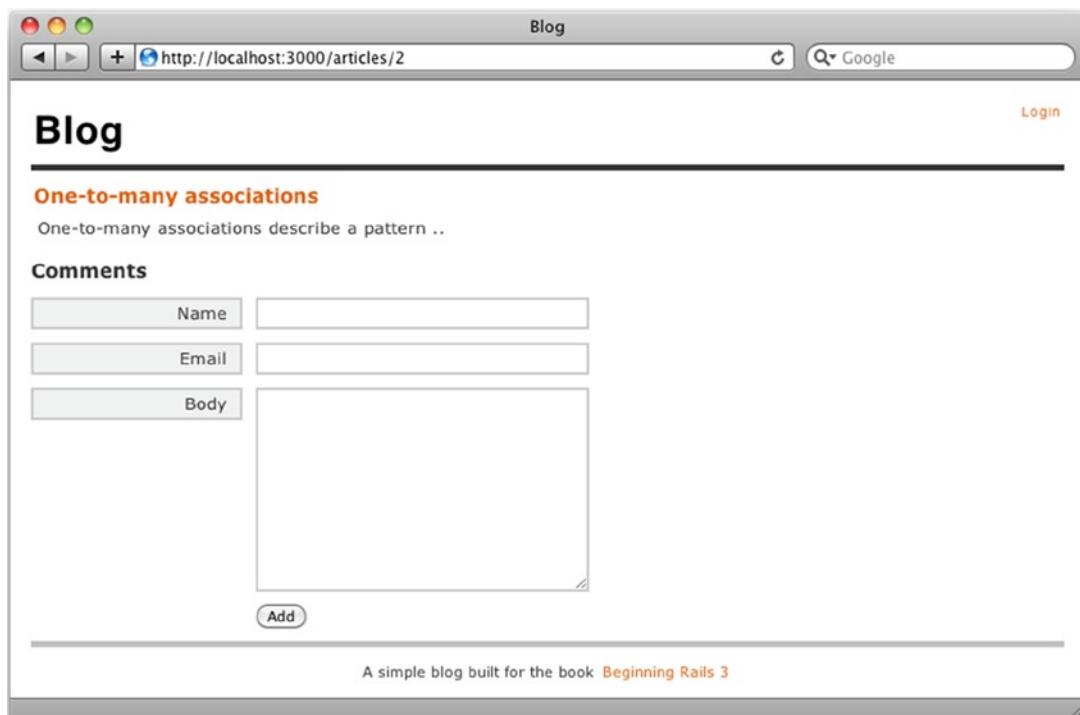


Figure 9-2. The article page with the comment form and without the new comment link

Making a Grand Entrance

In the previous section, you added an element to the screen via Ajax—the comment form. It's a pretty big form—it's a very obvious inclusion on the page and your users won't miss it; however, sometimes you may want to add just an extra link or highlight some text on a page. jQuery is so awesome that it provides a set of methods for adding animation to a web page.

You use a simple one-line command to make sure users understand that something is being added to the page. Be sure the template at `app/views/comments/new.js.erb` looks like the code in Listing 9-5.

Listing 9-5. The Updated New Comment Template in `app/views/comments/new.js.erb`:

<http://gist.github.com/353497>

```
$("<%= escape_javascript render(file: 'comments/new.html.erb') %>").insertAfter('#comments');

$('#new_comment').hide().slideDown();
$('#new_comment_link').hide();
```

The only change adds the `$('#new_comment').hide().slideDown();` line, which makes the newly added `new_comment` element appear in the page with a sliding motion.

Open your browser at any article page and look at the shiny effect that is being applied.

Note You very likely want to learn more about all the available effects in jQuery. To do so, head over to <http://api.jquery.com/category/effects/>.

Using Ajax for Forms

Another user-interaction improvement is to not refresh the page after a user adds a new record. In quite a few applications, users may be required to enter a considerable amount of data in forms; so, this technique is important to grasp.

In the same way that you made a link submit data via Ajax, you can add the `remote: true` option to the `form_for` helper you're using; this tells Rails you want the form data to be submitted via Ajax (Listing 9-6). By sticking with the `form_for` helper, you don't have to change the way you read parameters in controllers. The `remote: true` option is one you'll definitely use frequently.

Listing 9-6. The Updated Comment Form in `app/views/comments/new.html.erb`: <http://gist.github.com/353505>

```
<%= form_for([\@article, @article.comments.new], remote: true) do |f| %>
  <div class="field">
    <%= f.label :name %><br />
    <%= f.text_field :name %>
  </div>
  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :body %><br />
    <%= f.text_area :body %>
  </div>
```

```
<div class="actions">
  <%= f.submit 'Add' %>
</div>
<% end %>
```

Although the changes in the view are minimal, you have to make a few more changes in your controller layer. You want to respond to JavaScript and HTML requests in different ways. Change the `create` method in your comments controller to look like the code in Listing 9-7.

Listing 9-7. The Updated Comments Controller in `app/controllers/comments_controller.rb`:
<http://gist.github.com/353507>

```
class CommentsController < ApplicationController
  before_action :load_article
  before_action :authenticate, only: :destroy
  def create
    @comment = @article.comments.new(comment_params)
    if @comment.save
      respond_to do |format|
        format.html { redirect_to @article, notice: 'Thanks for your comment' }
        format.js
      end
    else
      respond_to do |format|
        format.html { redirect_to @article, alert: 'Unable to add comment' }
        format.js { render 'fail_create.js.erb' }
      end
    end
  end
  def destroy
    @article = current_user.articles.find(params[:article_id])
    @comment = @article.comments.find(params[:id])
    @comment.destroy
    redirect_to @article, notice: 'Comment Deleted'
  end
  private
  def load_article
    @article = Article.find(params[:article_id])
  end
  def comment_params
    params.require(:comment).permit(:name, :email, :body)
  end
end
```

The main method in this code is the `respond_to` helper. By using `respond_to`, you can have some code in the `format.html` block that's called when you receive a regular request and some code in the `format.js` block that's called when a JavaScript request is received. Hang on! There is no code in `format.js`! When no code is added to a `format` block, Rails looks for a template named after the view, just like regular views, which means it looks for `create.js.erb`. When a submitted comment fails validation, you also want to warn the user by displaying error messages; for that, you use `format.js { render 'fail_create.js.erb' }` to explicitly render a specific template.

The new `apps/views/comments/create.js.erb` and `app/views/comments/fail_create.js.erb` templates are shown in Listings 9-8 and 9-9, respectively.

Listing 9-8. The Template in app/views/comments/create.js.erb: <http://gist.github.com/353510>

```
$('#comments').append("<%= escape_javascript(render(@comment)) %>");
$('#new_comment')[0].reset();
```

Listing 9-9. The Template in app/views/comments/fail_create.js.erb: <http://gist.github.com/354866>

```
alert("<%= @comment.errors.full_messages.to_sentence.html_safe %>");
```

In the `create.js.erb` template, you run a couple of simple JavaScript commands. First, you render the template for a new comment—using `render(@comment)`—and append that HTML to the end of the `comments` div, achieved by using jQuery's `append` method. The `$('#new_comment')[0].reset()` line is a simple call to reset all the elements of the new comment form, which is blank and ready to accept another comment from your user.

In the `fail_create.js.erb` template, you use the `alert` JavaScript function to display a dialog box with the validation error message, as shown in Figure 9-3.

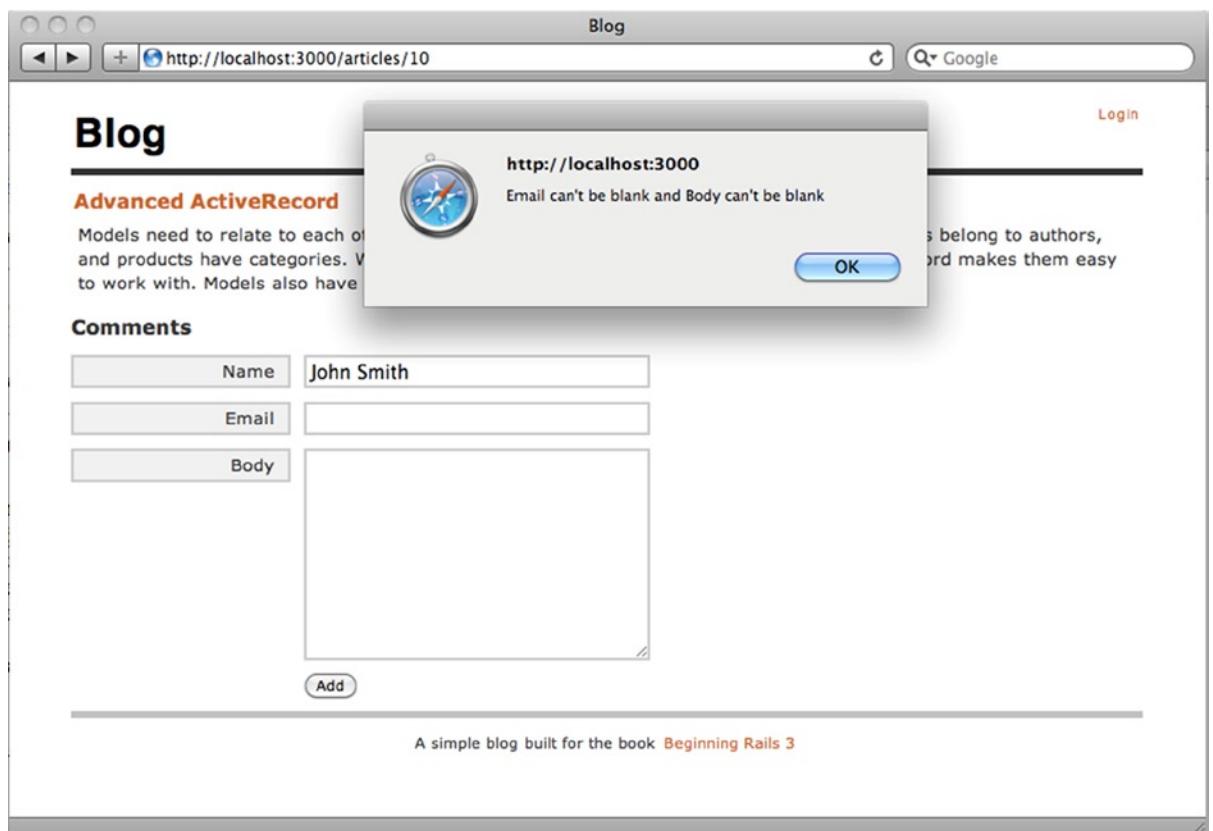


Figure 9-3. Displaying an error message

Give it a try: point your browser to an existing article, for example <http://localhost:3000/articles/2>, and enter a few—or lots of—comments. As you can see, you can interact with the page in a much more efficient way: there's no need to wait until a full page-reload happens.

Deleting Records with Ajax

To complete the “making things snappy” section, you may want to delete some of the comments that are added by users. You can combine the techniques you’ve learned in this chapter to let users delete comments without delay.

You already have a link to delete comments in the comment template at `app/views/comments/_comment.html.erb`. To use Ajax with that link, you again need to add the `:remote => true` option to the method call (Listing 9-10).

Listing 9-10. The Template in `app/views/comments/_comment.html.erb`: <http://gist.github.com/354807>

```
<%= div_for comment do %>
  <h3>
    <%= comment.name %> &lt;<%= comment.email %>&gt; said:
    <% if @article.owned_by? current_user %>
      <span class='actions'>
        <%= link_to 'Delete', [@article, comment], confirm: 'Are you sure?', method: :delete, remote: true %>
      </span>
    <% end %>
  </h3>
  <%= comment.body %>
<% end %>
```

The changes in the controller are also minimal. Use the `respond_to` and `format` block to make sure you support both regular and JavaScript requests, as shown in Listing 9-11.

Listing 9-11. The Comments Controller in `app/controllers/comments_controller.rb`: <http://gist.github.com/354809>

```
class CommentsController < ApplicationController
  before_action :load_article
  before_action :authenticate, only: :destroy
  def create
    @comment = @article.comments.new(comment_params)
    if @comment.save
      respond_to do |format|
        format.html { redirect_to @article, notice: 'Thanks for your comment' }
        format.js
      end
    else
      respond_to do |format|
        format.html { redirect_to @article, alert: 'Unable to add comment' }
        format.js { render 'fail_create.js.erb' }
      end
    end
  end
  def destroy
    @article = current_user.articles.find(params[:article_id])
    @comment = @article.comments.find(params[:id])
    @comment.destroy
    respond_to do |format|
      format.html { redirect_to @article, notice: 'Comment Deleted' }
    end
  end
end
```

```

      format.js
    end
  end
private
def load_article
  @article = Article.find(params[:article_id])
end
def comment_params
  params.require(:comment).permit(:name, :email, :body)
end
end

```

You wire up the delete link in the comment partial to send an Ajax request to the controller. The controller responds to those Ajax requests with the default action, which is to render the `app/views/comments/destroy.js.erb` file (Listing 9-12).

Listing 9-12. The `app/views/comments/destroy.js.erb` File: <http://gist.github.com/354811>

```
$("#<%= dom_id(@comment) %>").remove()
```

Open your browser to an article page—make sure you are logged in as the article owner—with some comments you want to delete—or add lots of spam-like comments. See how quickly you can get rid of comments now? It's a lot better than waiting for page reloads.

Summary

To be sure, Ajax is a large topic. Entire books and conferences are devoted to this subset of technology, so it goes without saying that this chapter only scratches the surface. Still, in short order, you've learned the basics of implementing Ajax in Rails applications, and you know where to go when you need to dig deeper.

You learned how to make remote Ajax calls using the `remote: true` option for links and forms. You also used a simple visual effect to show new elements on the page, thanks to the jQuery library.

Finally, you learned about using JavaScript templates—which have the `.js.erb` extension—to produce responses to Ajax requests using JavaScript code.

At this stage, you have a solid grasp of the Action Pack side of web development with Rails. Next, you'll learn how to conquer another common component of web application development: sending e-mail.



Sending and Receiving E-Mail

It's a rare web application that doesn't need to send e-mail from time to time. For example, you may want to send messages to welcome users who sign up on your web site, relay passwords, or confirm orders placed with an online store. Rails ships with a library called Action Mailer, which provides developers with an easy-to-use yet powerful tool to handle e-mail.

This chapter explains how Action Mailer works and how to use it in your applications. You first learn how to configure it, and then you'll see a few examples of how to send e-mail in various formats. In addition to sending e-mail, Action Mailer can also *receive* e-mail, an advanced topic that the chapter touches on briefly.

Note If you need to get the code at the exact point where you finished Chapter 9, download the source code zip file from the book's page on <http://www.apress.com> and extract it on your computer.

Setting Up Action Mailer

Like Active Record and Action Pack, Action Mailer is one of the components that makes up the Rails framework. It works much like the other components of Rails: mailers are implemented to behave like controllers, and mailer templates are implemented as views. Because it's integrated into the framework, it's easy to set up and use, and it requires very little configuration to get going.

When you send e-mail using an e-mail client such as Outlook or a web-based e-mail application like Gmail or Yahoo Mail, your messages are sent via a mail server. Unlike a web server, Rails doesn't provide a built-in mail server. You need to tell Action Mailer where your e-mail server is located and how to connect to it. This sounds a bit complicated, but it's really quite easy. Depending on the kind of computer you're using, you may have a mail server built in (this is true of most UNIX systems). If not, you can use the same server that you use to process your regular e-mail. If this is the case, you can find your server information in your e-mail client settings, as provided by your Internet service provider (ISP), or in the settings section of your web-based e-mail application, like Gmail.

Configuring Mail Server Settings

Before you can send e-mail from your Rails application, you need to tell Action Mailer how to communicate with your mail server. Action Mailer can be configured to send e-mail using either `sendmail` or a Simple Mail Transfer Protocol (SMTP) server. SMTP is the core Internet protocol for relaying e-mail messages between servers. If you're on Linux, OS X, or any other UNIX-based system, you're in luck: you can use `sendmail`, and as long as it's in the standard location (`/usr/bin/sendmail`), you don't need to configure anything. If you're on Windows or if you want to use SMTP, you have some work to do.

Action Mailer options are set at the class level on `ActionMailer::Base`. The best place to set these options is in your environment files, located in the `config` directory of your application. You can also add your configuration in an initializer file in `config/initializers`; doing so ensures that your settings apply for all environments. In most cases, though, you have different settings for the development and production environments; so it may be wiser to add settings in any of the environment-specific configuration files (`config/environments/*.rb`), because this takes precedence over the global configuration.

This section describes how to set up Action Mailer to use SMTP, because it works on all systems and is the default delivery method. To do this, you supply the SMTP settings via the `smtp_settings` option. The `smtp_settings` method expects a hash of options, most of which are shown in Table 10-1.

Table 10-1. Server Connection Settings

Setting	Description
Address	The address of your mail server. The default is <code>localhost</code> .
port	The port number of your mail server. The default is port 25.
domain	If your e-mail server responds to different domain names, you may need to specify your domain name here.
authentication	If your mail server requires authentication, you need to specify the authentication type here. This can be one of <code>:plain</code> , <code>:login</code> , or <code>:cram_md5</code> .
user_name	The username you use to authenticate when you connect to the mail server, if your server requires authentication.
password	The password you use to authenticate when you connect to the mail server, if your server requires authentication.

Listing 10-1 shows a typical configuration for a server that requires authentication, in this case Gmail. You can use this sample configuration as a starting point to configure your connection. Change each of the settings (authentication, username, password, and address) to connect to your own SMTP server. If you’re using `sendmail` as the delivery method, add `ActionMailer::Base.delivery_method = :sendmail`; then, everything should “just work.”

Listing 10-1. Sample Action Mailer Configuration Using SMTP, in `config/environments/development.rb`:
<https://gist.github.com/adamgamble/6190788>

```
Blog::Application.configure do
  # Settings specified here will take precedence over those in config/environment.rb

  # In the development environment your application's code is reloaded on
  # every request. This slows down response time but is perfect for development
  # since you don't have to restart the webserver when you make code changes.
  config.cache_classes = false

  # DO not eager load code on boot.
  config.eager_load = false

  # Show full error reports and disable caching
  config.consider_all_requests_local      = true
  config.action_controller.perform_caching = false
```

```

# Don't care if the mailer can't send
config.action_mailer.raise_delivery_errors = false

# Gmail SMTP server setup
ActionMailer::Base.smtp_settings = {
  :address => "smtp.gmail.com",
  :enable_starttls_auto => true,
  :port => 587,
  :authentication => :plain,
  :user_name => "beginningrails@gmail.com",
  :password => 'pleasechange'
}

# Raise an error on page load if there are pending migrations
config.active_record.migration_error = :page_load

# Debug mode disables concatenation and preprocessing of assets.
# This option may cause significant delays in view rendering with a large
# number of complex assets
config.assets.debug = true
end

```

Make sure you modify the options to your own credentials. Restart your server if it's running, and your application is ready to send e-mail.

Note If you need to use any advanced Action Mailer settings, the Rails API has a good chunk of information at <http://api.rubyonrails.org/classes/ActionMailer/Base.html>.

Configuring Application Settings

In addition to the mail server settings, Action Mailer has a set of configuration parameters you can tweak to make the library behave in specific ways according to the application or the environment. You'll stick with the defaults here, so you don't need to set up any special application settings. For reference, Table 10-2 lists the most common configuration options. Just like the server settings, these can be specified in an initializer file or in the environment-specific configuration files (`config/environments/*.rb`).

Table 10-2. Common Action Mailer Application Settings

Option	Description
<code>template_root</code>	Indicates the base folder from which template references will be made. The default is <code>app/views</code> .
<code>raise_delivery_errors</code>	Allows you to indicate whether you want errors to be raised when an error occurs while trying to deliver e-mail.
<code>perform_deliveries</code>	Indicates whether messages should really be delivered to the mail server.

(continued)

Table 10-2. (continued)

Option	Description
deliveries	Keeps an array of all delivered e-mail when the delivery method is set to :test. This is useful when you're in testing mode.
default_charset	Specifies the default character set to be used when sending messages. The default is UTF-8.
default_content_type	Specifies the default content type that will be used for outbound mail. The default is text/plain.

Note When you create a new Rails application, the configuration files automatically use sensible defaults for each of the development, test, and production environments. Take a quick look in config/environments to see how Action Mailer behaves in development, production, and test mode to make sure you understand your application's behavior.

Sending E-Mail

Now that you have Action Mailer configured, it's time to see it in action. This section explores all the possibilities in the Action Mailer world, starting with basic text-only e-mail and then adding extra e-mail options such as attachments.

To demonstrate Action Mailer, let's enhance the blog application by allowing users to send e-mail to their friends, so they can share information about a specific article. This is a common feature in today's web applications, affectionately referred to as "send to friend."

By now, you know that Rails provides helpful generators to get started writing your own code. You saw generators in action when you created models and controllers in previous chapters. The mailer generator works just like the other generators.

Enter the following command to generate the Notifier class with one method named email_friend:

```
$ rails generate mailer Notifier email_friend
```

create	app/mailers/notifier.rb
invoke	erb
create	app/views/notifier
create	app/views/notifier/email_friend.text.erb
invoke	test_unit
create	test/functional/notifier_test.rb

The generator creates a mailer class named Notifier, containing the email_friend method you specified on the command line. Notice that the generated Action Mailer class is created in the app/mailers directory; Notifier and any Action Mailer class are subclasses of the ActionMailer::Base class. The generator also creates a template file in the views directory (app/views/notifier/email_friend.text.erb) that corresponds to the email_friend method (action) you use to set up the mailer message.

Just like controllers, Action Mailer classes contain methods that, when triggered, execute some code and render a related view of the same name, unless otherwise specified.

Listing 10-2 shows the Notifier class located in app/mailers/notifier.rb. The email_friend method has some code, which will be the starting point for most of the methods you write using Action Mailer.

Listing 10-2. Notifier Class in app/mailers/notifier.rb

```
class Notifier < ActionMailer::Base
  default :from => "from@example.com"

  # Subject can be set in your I18n file at config/locales/en.yml
  # with the following lookup:
  #
  #   en.actionmailer.notifier.email_friend.subject
  #
  def email_friend
    @greeting = "Hi"

    mail :to => "to@example.org"
  end
end
```

Action Mailer classes have a class-wide configuration hash that you can modify using the `default` method. Notice the first line in the `Notifier` class definition: it sets the sender address (using `:from`) for all e-mails sent from this class to from@example.com.

In the `email_friend` method body, the first line defines an instance variable named `@greeting`; just like in controllers, those variables are available in your views. Also in the `email_friend` method body, the `mail` method is called with a parameter of `:to => "to@example.org"`, specifying the e-mail address that will receive this message. The `mail` method accepts an options hash that specifies the various headers of the message. Table 10-3 lists the available methods you use to configure an individual message.

Table 10-3. Mailer Instance Variables

Method name	Description	Example
<code>subject</code>	The subject of the e-mail message to be sent.	<code>:subject => "Action Mailer is powerful"</code>
<code>To</code>	A string or array of e-mail addresses to which the message will be sent.	<code>:to => "friend@example.com"</code>
<code>from</code>	A string specifying the sender of the e-mail message.	<code>:from => "sender@example.com"</code>
<code>reply_to</code>	A string specifying the reply-to e-mail address.	<code>:reply => "sender@example.com"</code>
<code>date</code>	The date header. The default is the current date.	<code>:date => Time.now</code>
<code>cc</code>	A string or array of e-mail addresses to carbon copy with the message.	<code>:cc => "admin@example.com"</code>
<code>bcc</code>	A string or array of e-mail addresses to blind carbon copy with the message.	<code>:bcc => ["support@example.com", "sales@example.com"]</code>

The mailer generator creates a template named after the action in the `Notifier` class in the `/app/views/notifier` folder: `email_friend.text.erb`. This is the template that generates the body of the e-mail sent when using the `email_friend` method. This template works similarly to the templates used for regular views in Action Pack. It's an ERb file, which can contain text and markup mixed with some Ruby code. It also has the same one-to-one relationship between action and view exhibited by Action Pack—each action in your mailer class expects one template in the `app/views` directory.

Handling Basic E-Mail

Let's start enhancing the blog application by adding "notify a friend" functionality to the article page. The first iteration is a very basic example that sends a text e-mail message containing a brief message.

The first piece of the puzzle is to make a change to the routes file, to include a route for the action that will be called after the user submits the form. You add a member route to articles using the `member` method to give a `notify_friend_article` route. Make sure your `config/routes.rb` file looks like the code in Listing 10-3.

Listing 10-3. Added a `notify_friend` Action to `config/routes.rb`: <http://gist.github.com/355014>

```
Blog::Application.routes.draw do
  root :to => "articles#index"
  resources :articles do
    member do
      post :notify_friend
    end
    resources :comments
  end
  resources :users
  resource :session
  match '/login' => "sessions#new", :as => "login"
  match '/logout' => "sessions#destroy", :as => "logout"
end
```

Note Using the `member` method inside your `resources` block helps define a route that requires the `id` of the resource. Custom member routes are similar to the default member routes, such as `edit_article_path` and `article_path`. Following the same convention, you can define collection routes using the `collection` method. Custom collection routes are similar to the default collection routes, such as `articles_path`, which don't require an `id`.

You want to show users a link that slides down a form where they can enter the e-mail address of the friend to whom they want to send a message. Let's update the article's show view to include the new link directly after rendering the article's partial. Add the code shown in Listing 10-4 in `app/views/articles/show.html.erb`.

Listing 10-4. "Notify a Friend" Functionality Added to `app/views/articles/show.html.erb`:

<http://gist.github.com/355029>

```
<%= render @article %>

<%= link_to 'Email a friend', '#',
:onclick => "$( '#notify_friend_form' ).slideDown()" %>
<%= render 'notify_friend' %>

<h3>Comments</h3>
<div id="comments">
  <%= render @article.comments %>
</div>

<%= link_to "new comment", new_article_comment_path(@article, :format => :js),
:remote => true, :id => 'new_comment_link' %>
```

Notice that you also render a partial named `notify_friend`. Create the partial in `app/views/articles/_notify_friend.html.erb` so it looks like the code in Listing 10-5.

Listing 10-5. “Notify a Friend” Partial in `app/views/articles/_notify_friend.html.erb`:
<http://gist.github.com/355030>

```
<%= form_tag(notify_friend_article_path(@article), :id => 'notify_friend_form',
:style => 'display:none') do %>
  <div class="field">
    <%= label_tag :name, 'Your name' %><br />
    <%= text_field_tag :name %>
  </div>
  <div class="field">
    <%= label_tag :email, "Your friend's email" %><br />
    <%= text_field_tag :email %>
  </div>
  <div class="actions">
    <%= submit_tag 'Send' %> or
    <%= link_to 'Cancel', '#', :onclick => "$( '#notify_friend_form' ).slideUp()" %>
  </div>
<% end %>
```

Now, when you go to any article page, you’ll see a link to e-mail a friend. Because you don’t want to show the form all the time, you made the form hidden. If users are interested in recommending the article by sending an e-mail to a friend, they can click the link, and the form will be revealed through the help of some clever JavaScript. The end result is shown in Figures 10-1 and 10-2.



Figure 10-1. Article page without “notify a friend” form

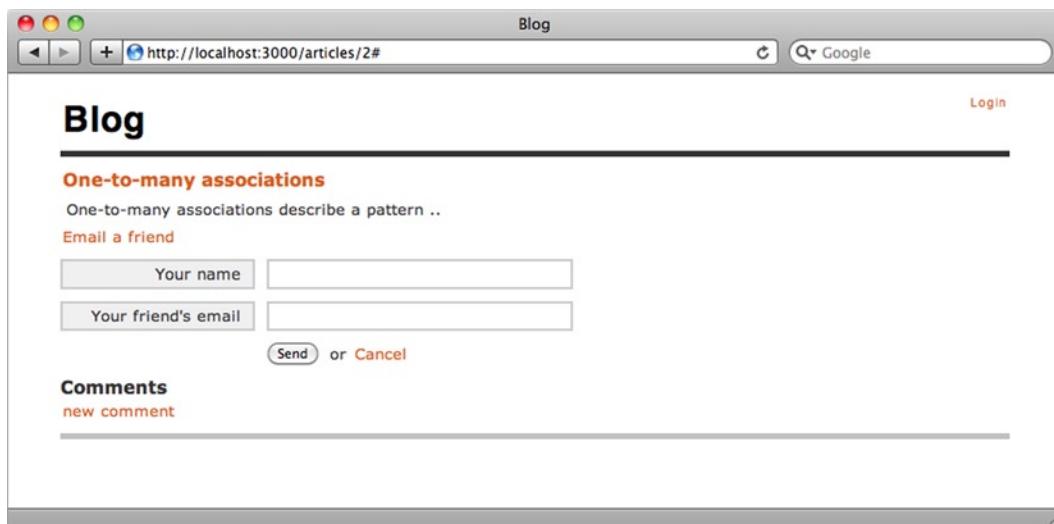


Figure 10-2. Article page with visible “notify a friend” form

The interface is ready to go, but the `articles` controller doesn’t know how to handle this request yet. Your form is configured to submit to an action called `notify_friend`, but that action doesn’t exist. Update the `articles` controller and add the `notify_friend` method, as shown in Listing 10-6.

Listing 10-6. The `notify_friend` Action Added to `app/controllers/articles_controller.rb`:
<https://gist.github.com/adamgamble/19c6a5082bc625aecac4>

```
def notify_friend
  @article = Article.find(params[:id])
  Notifier.email_friend(@article, params[:name], params[:email]).deliver
  redirect_to @article, :notice => "Successfully sent a message to your friend"
end
```

The action you just added is short and concise, but there’s something that deserves a closer look. To perform the delivery, you call a class method on the `Notifier` class called `email_friend`; by calling `deliver` after `email_friend`, Action Mailer executes the `email_friend` method in the `Notifier` class and generates an e-mail message as per that method. After the e-mail message is created, it’s passed to the `deliver` method, which performs the actual delivery.

This is important. Every time you create a mail action, you call it directly from the mailer class (in this case, `Notifier`), and then you call `deliver`. If you had a mailer class called `BlogMailer` and a mail action called `invitation`, you’d call it using `BlogMailer.invitation.deliver`.

Before you try this, the `email_friend` method still needs a bit of work. You need to augment it so that it sets a meaningful title and uses the `name` and `email` parameters you collect from the form and pass into the method. Listing 10-7 shows these changes.

Listing 10-7. Updated Notifier Mailer in `app/mailers/notifier.rb`:
<http://gist.github.com/390971>

```
class Notifier < ActionMailer::Base
  default :from => "from@example.com"

  def email_friend(article, sender_name, receiver_email)
    @article = article
    @sender_name = sender_name
```

```

mail :to => receiver_email, :subject => "Interesting Article"
end
end

```

Notice how you add three arguments: article, sender_name, and receiver_email. You start by defining two instance variables @article and @sender_name to make them available in your view template. Then, you modify the mail method call to send the e-mail to the receiver_email address with the subject “Interesting Article.” Next, you want your e-mail message to have some sort of formatting and include the URL of your application. Change the template file located in app/views/notifier/email_friend.text.erb and make it look like the code in Listing 10-8.

Listing 10-8. Notifier Template in app/views/notifier/email_friend.text.erb: <http://gist.github.com/355053>

Your friend <%= @sender_name %> thinks you may like the following article:

```
<%= @article.title %>: <%= article_url(@article, :host => "localhost:3000") %>
```

Now you can give this a try in your browser. Fill out the e-mail form using your own e-mail address, so you can see what the e-mail message looks like when it is received. If all goes according to plan, you should receive a message that looks something like Figure 10-3.



Figure 10-3. Message delivered to user's inbox

This is a plain text message, the default content type. The next section shows how to send e-mail messages that use rich, HTML formatting.

Sending HTML E-Mail

So far, your e-mail message is pretty plain. To make it more interesting and informative, you can add a link to the specific article being recommended. You can also make it more visually appealing for users with rich e-mail clients (like Gmail) by adding some HTML formatting.

Making another analogy to controllers, whenever you want to render an e-mail template with HTML content, you need to change the template extension from .text.erb to .html.erb; Rails takes care of setting the appropriate e-mail headers.

In this case, if you keep the `.text.erb` file and add a template with a `.html.erb` extension, Rails sends both templates, representing the same e-mail, known as a *multipart* message. The e-mail client that receives a multipart message recognizes that there are alternatives for the same message and chooses the most appropriate part based on the environment or the user's preferences.

Let's put together an HTML template for the message in `app/views/notifier/email_friend.html.erb`. The file should look like the code in Listing 10-9.

Listing 10-9. HTML `email_friend` Template in `app/views/notifier/email_friend.html.erb`:

<http://gist.github.com/355076>

```
<html>
<body>
<p>
  One of your friends, <%= @sender_name %>,
  thinks you like an article we have written.
</p>

<p>
  Come check all the information about <strong><%= @article.title %></strong> at
  <%= article_url @article, :host => 'localhost:3000' %>
</p>
</body>
</html>
```

Use the application to send yourself another e-mail. It looks pretty good, as shown in Figure 10-4. If your users don't have a rich e-mail client and can't read HTML mail, they get the plain text version.

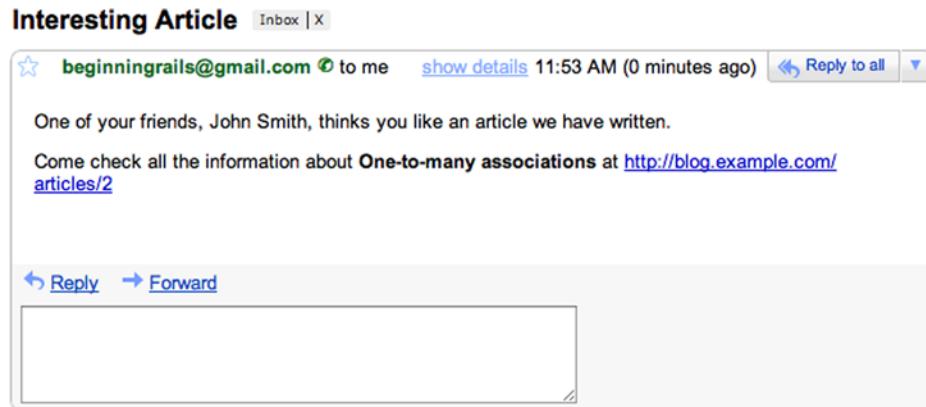


Figure 10-4. HTML message delivered to a user's inbox

Note If you think maintaining both text and HTML versions of an e-mail message is a lot of work, it may be safer to stick with the HTML message.

Adding Attachments

In some cases, you may want to add attachments to an e-mail message. Action Mailer makes this a straightforward task by providing the attachment helper. You tell attachment which file you want to attach to the e-mail, and it does its magic.

Let's walk through an example of attaching a file to an e-mail message. Assume that you want to send a photo related to one of the articles in the blog application every time a user sends an e-mail about that article to a friend. Just for demonstration, use the default Rails logo, which should be in the public/images directory (a vestige of the "Welcome to Rails" index.html file you deleted in Chapter 7). To attach this image file to the e-mail you created in the previous section, add a call to the attachments method in the email_friend method in the Notifier class, as shown in Listing 10-10.

Listing 10-10. Adding an Attachment to the Mailer in app/models/notifier.rb: <http://gist.github.com/355081>

```
class Notifier < ActionMailer::Base
  default :from => "from@example.com"

  def email_friend(article, sender_name, receiver_email)
    @article = article
    @sender_name = sender_name

    attachments["rails.png"] = File.read(Rails.root.join("public/images/rails.png"))
    mail :to => receiver_email, :subject => "Interesting Article"
  end
end
```

When you call the attachments method, you first identify the name of the file the e-mail receiver sees; in this case, the receiver sees a file named rails.png because you call attachments["rails.png"]. Next, you tell Rails where the file you want to attach is located: you load the Rails logo from the project using the Rails.root method. The resulting message looks like Figure 10-5.



Figure 10-5. Message with an attachment delivered to a user's inbox

Tip You can specify the body of an attachment by using the `File.read` method if the file you're sending exists on disk. Alternatively, you can generate the file on the fly if it's a dynamic file, like a personalized PDF or Word document.

Letting Authors Know About Comments

Just to make sure you've grasped how to send e-mail from your Rails applications, this section quickly goes over the complete flow to add another mailer action. If you recall, in Chapter 6 you added an observer to be invoked every time a new comment is added.

In total, you change three files. First, you add a new action to the `Notifier` mailer class; second, you add a new template with the contents of the e-mail to send; third, you change the observer to invoke the mailer when new comments are added. Listings 10-11 to 10-13, respectively, show the code for these actions.

Listing 10-11. Adding the `comment_added` Method to `app/models/notifier.rb`: <http://gist.github.com/358011>

```
class Notifier < ActionMailer::Base
  default :from => "from@example.com"

  def email_friend(article, sender_name, receiver_email)
    @article = article
    @sender_name = sender_name

    attachments["rails.png"] = File.read(Rails.root.join("public/images/rails.png"))
    mail :to => receiver_email, :subject => "Interesting Article"
  end

  def comment_added(comment)
    @article = comment.article
    mail :to => @article.user.email,
    :subject => "New comment for '#{@article.title}'"
  end
end
```

Listing 10-12. The `comment_added` Mailer Template in `app/views/notifier/comment_added.html.erb`: <http://gist.github.com/358023>

```
<html>
<body>
  <p>
    Someone added a comment to one of your articles <i><%= @article.title %></i>.
    Go read the comment: <%= link_to @article.title,
    article_url(@article, :host => 'localhost:3000') %>.
  </p>
</body>
</html>
```

Listing 10-13. Updates to app/model/comment.rb: <https://gist.github.com/adamgamble/6190838>

```
class Comment < ActiveRecord::Base
  belongs_to :article

  validates_presence_of :name, :email, :body
  validate :article_should_be_published

  after_create :email_article_author

  after_create :send_comment_email

  def article_should_be_published
    errors.add(:article_id, "is not published yet") if article && !article.published?
  end

  def email_article_author
    puts "We will notify #{article.user.email} the author in Chapter 9"
  end

  def send_comment_email
    Notifier.comment_added(self).deliver
  end
end
```

Now that those changes have been made, create an article with your user account and add some comments. You should receive one e-mail message per comment. If you want to, you could add the comment text to the e-mail; that way, you wouldn't need to go to the article page to read the comment. You could easily implement that by changing the mailer view.

Receiving E-Mail

So far, you've seen that Action Mailer has extensive support for sending all types of e-mail messages. But what if your application needs to receive e-mail? You can handle incoming e-mail in a Rails application a few different ways. This section explains how to use a Rails process and how to read e-mail from your mail server. The approach you choose depends a lot on your operating system and e-mail server.

Using a Rails Process

In an Action Mailer class, you can write a `receive` method that receives a `Mail::Message` object as a parameter, which corresponds to an incoming e-mail message your code can process. Inside the `receive` method, it's easy to extract details about the incoming e-mail, such as header, subject, body text, and attachments.

For example, the blog application can have a special e-mail address (such as `new@blog.example.com`) that can be monitored to create a new article whenever an e-mail arrives. This way, users can send an e-mail to `new@blog.example.com` and create a new article without needing to open their browsers. The implementation of this feature looks something like the code in Listing 10-14.

Listing 10-14. Example Mailer Class with receive Method

```
class ExampleMailer < ActionMailer::Base
  def receive(email)
    article = Article.new
    article.title = email.subject
    article.body = email.body
    article.save
  end
end
```

This code is pretty simple and takes care of receiving e-mail; however, this is just the first part of the solution. The remaining part is tricky and may demand some research and system administration skills. You need to tell your mail server that it should redirect messages sent to a specific address to a special process. In this case, the process is the Rails runner script, which executes the Ruby code passed as a parameter as if it were running from within a Rails application. You can see this technique in action by saving an e-mail message to any location on your disk and invoking the receive method using the following command (POSIX only):

```
rails runner ExampleMailer.receive(STDIN.read) < email.txt
```

This chapter doesn't go into the implementation details of configuring a script to route incoming e-mail to your Rails process because it's impossible to cover all set ups. There is a project called Griddler that makes this relatively easy to do (<https://github.com/thoughtbot/griddler>).

Reading E-Mail Using POP or IMAP

If you don't have control over the mail server being used and can't write a server-side script, you can still read e-mail from your mail server as your regular e-mail client does. To do this, you can create a separate Ruby script that fetches e-mail and run it as a background process that polls for new messages.

The code in Listing 10-15 connects to a mail server through POP3 and checks a specific mailbox to see if any new e-mail has arrived. If so, the script reads that message and passes it to the ExampleMailer.receive method for processing. This example uses a POP3 (Post Office Protocol) server, but it could just as easily use an IMAP (Internet Message Access Protocol) server. The only difference would be that you would use the Net::IMAP class instead of the Net::POP3 class to connect to the mail server. Both classes are part of the Ruby Standard library.

Listing 10-15. Sample Ruby Script for Reading E-Mail Messages

```
Net::POP3.start("mail.example.com", nil, "username", "password") do |pop|
  if pop.mails.empty?
    logger.info "NO MAIL"
  else
    pop.mails.each do |email|
      begin
        logger.info "receiving mail..."
        ExampleMailer.receive(email.pop)
        email.delete
      rescue Exception => e
        logger.error "Error receiving email: #{Time.now.to_s} - #{e.message}"
      end
    end
  end
end
```

This script starts by trying to connect to the POP3 server with the credentials indicated on the first line. As soon as the connection is established, it checks to see if there are any new e-mail messages by using the `pop.mails.empty?` method. If there are new e-mail messages, it iterates through each of them, calling `ExampleMailer.receive(email.pop)`. After a message is processed, it's deleted from the server to avoid reprocessing the same message the next time the script is called.

Summary

This chapter explained how to send e-mail from your web applications using Action Mailer. You configured Action Mailer to talk to your mail server and learned the most common configuration parameters you can use to fine tune how Action Mailer works with your application.

You learned that Action Mailer allows you to send e-mail messages based on templates and how to use implicit parts for text and HTML messages, as well as how to use the `attachment` helper to add attachments to your messages.

You also learned the basics on receiving e-mail using Action Mailer. This chapter only scratched the surface, this being a rather advanced technique. Still, you have a good starting point should your application ever need to perform this task, and you know where to look when you need more information.

This chapter brings us to the end of your tour of the main Rails libraries: Active Record, Action Pack, and Action Mailer. The next chapter covers one of the most important techniques to improve the quality of your code: testing.



Testing Your Application

Smart developers test their code. Take a minute to read that sentence again and let it sink in: smart developers test their code. The fact is, testing is one of the most important things you can do to improve the quality of your code, reduce the cost of change, and keep your software bug-free. Rails (and the Ruby community at large) takes testing seriously. Not surprisingly, Rails goes out of its way to make testing hassle-free.

The basic idea of testing is simple: you write code that exercises your program and tests your assumptions. Instead of just opening a browser and adding a new user manually to check whether the process works, you write a test that automates the process—something repeatable. With a test in place, every time you modify the code that adds a new user, you can run the test to see if your change worked—and, more important, whether your seemingly innocuous change broke something else.

If you stop and think about it, you're already testing your software. The problem is that you're doing it manually, often in an ad hoc fashion. You may make a change to the way users log in, and then you try it in your browser. You make a change to the sign-up procedure, and then you take it for a spin. As your application grows in size, it becomes more and more difficult to manually test like this, and eventually you miss something important. Even if you're not testing, you can be sure that your users are. After all, they're the ones using the application in the wild, and they'll find bugs you never knew existed. The best solution is to replace this sort of visual, ad hoc inspection with automatic checking.

Testing becomes increasingly important when you're refactoring existing code. *Refactoring* is the process of improving the design of code without changing its behavior. The best way to refactor is with a test in place acting as a safety net. Because refactoring shouldn't result in an observable change in behavior, it shouldn't break your tests either. It's easy, therefore, to see why many programmers won't refactor without tests.

Given the importance placed on testing, it may seem odd that this book leaves a discussion of this until Chapter 11. Ideally, you should be writing tests as you go, never getting too far ahead without testing what you've written. But we decided that explaining how to test would be overwhelming while you were still learning the basics of Ruby and the Rails framework. Now that you have a good deal of knowledge under your belt, it's time to tackle testing.

Note If you need to get the code at the exact point where you finished Chapter 10, download the zip file from www.apress.com and extract it onto your computer.

How Rails Handles Testing

Because Rails is an integrated environment, it can make assumptions about the best ways to structure and organize your tests. Rails provides:

- Test directories for controller, model, mailer, helper, and integration tests
- Fixtures for easily working with database data
- An environment explicitly created for testing

The default Rails skeleton generated by the `rails` command creates a directory just for testing. If you open it, you'll see subdirectories for each of the aforementioned test types:

```
test
|-- models           <-- model tests
|-- controllers     <-- controller tests |-- fixtures      <-- test data
|-- helpers          <-- helper test
|-- mailers          <-- mailer tests
|-- integration      <--
integration tests
```

In addition to the other directories, Rails creates the `integration` directory when you generate your first `integration` test. There is also a `fixtures` directory. What are these for?

Fixtures are textual representations of table data written in YAML—a data-serialization format. Fixtures are loaded into the database before your tests run; you use them to populate your database with data to test against. Look at the `users` fixtures file in `test/fixtures/users.yml`, as shown in Listing 11-1.

Listing 11-1. Users Fixture in `test/fixtures/users.yml`

```
# Read about fixtures at http://guides.rubyonrails.org/testing.html#the-low-down-on-fixtures

one:
  email: MyString
  password: MyString

two:
  email: MyString
  password: MyString
```

Rails generates the `users` fixtures file for you when you generate the user model. As you can see, the file has two fixtures: `one` and `two`. Each fixture has both attributes `email` and `password` set to `MyString`; but, as you recall, you renamed the `password` column `hashed_password` back in Chapter 6. Let's update the `users` fixtures file to reflect the new column name and use meaningful data. Listing 11-2 shows the updated fixture.

Listing 11-2. Updated Users Fixtures in `test/fixtures/users.yml`: <http://gist.github.com/358067>

```
eugene:
  email: 'eugene@example.com'
  hashed_password: 'e5e9fa1ba31ecd1ae84f75caaa474f3a663f05f4' # => secret

lauren:
  email: 'lauren@example.com'
  hashed_password: 'e5e9fa1ba31ecd1ae84f75caaa474f3a663f05f4' # => secret
```

Remember that every time you generated a model or a controller while building the blog application, Rails automatically generated test files for you. This is another example of its opinionated nature—Rails thinks you should test, so it goes out of its way to remind you.

You may also recall that Rails created three SQLite databases for the blog application: one for development (which is all you've been using thus far), one for production, and one for testing. Not surprisingly, Rails uses the testing database just for testing.

Rails drops and re-creates this test database on every run of the test suite. Make sure you don't list your development or production database in its place, or all your data will be gone.

Unit Testing Your Rails Application

You know that Rails generated some tests automatically. Let's open one of them now and take a look. Let's start with the Article test, located in `test/model/article_test.rb`, as shown in Listing 11-3.

Listing 11-3. Generated Article Unit Test in `test/model/article_test.rb`

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  # test "the truth" do
  #   assert true
  # end
end
```

Although there's not much to it (all it does is show you how to test that `true` is, in fact, `true`), this test gives you a template from which to build your real tests. It has the following elements:

- The test class is a subclass of Rails's enhanced version of Ruby's built-in testing framework, `Test::Unit`. It's called `ActiveSupport::TestCase`.
- Tests are implemented as blocks using the `test` method, with the first parameter as the description of that test—"the `truth`" in this case.
- Within a test case, *assertions* are used to test expectations. The "Testing with Assertions" section later in this chapter explains how these work.

If you peek inside the `test/model` directory, you'll see a similar test case for every model you've generated so far: Article, Comment, Category, User, and Profile. Each looks almost exactly the same as the Article test. Let's run the unit tests now using the `rake test:models` command from your command prompt, and see what happens:

```
$ rake test:models
Run options: --seed 45237
```

```
# Running tests:

Finished tests in 0.020474s, 0.0000 tests/s, 0.0000 assertions/s.

0 tests, 0 assertions, 0 failures, 0 errors, 0 skips
```

In this case, there are no tests yet (the ones generated are commented out). If there were tests, and the test passed you would see a `.` (dot) character. When the test case produces an error, you would see an `E`. If any assertion fails to return `true`, you would see an `F`. Finally, when the test suite is finished, it prints a summary.

Testing the Article Model

Let's test the Article model. If you recall from Chapter 5, one of the first things you did with your Article model was basic CRUD operations. Well, testing that you can create, read, update, and delete articles is a great place to start. Here's a quick summary of the specific things you test in this chapter:

- Creating a new article
- Finding an article
- Updating an article
- Destroying an article

Before you begin, you'll need to create a few fixtures (remember that a fixture is a textual representation of test data).

Creating Fixtures

You create fixtures for articles. Open the `test/fixtures/articles.yml` file and replace its content with the code as shown in Listing 11-4.

Listing 11-4. Articles Fixtures in `test/fixtures/articles.yml`: <http://gist.github.com/358400>

```
welcome_to_rails:
  user: eugene
  title: "Welcome to Rails"
  body: "Rails is such a nice web framework written in Ruby"
  published_at: <%= 3.days.ago %>
```

That's all you need to do. The data in the fixtures are inserted automatically into your test database before your tests run. With fixtures in place, you're ready to start creating test cases.

Tip Fixtures are parsed by ERb before they're loaded, so you can use ERb in them just as you can in view templates. This is useful for creating dynamic dates, as you do in `published_at: <%= 3.days.ago %>`.

The following sections present the test cases one at a time, beginning with *create*.

Adding a Create Test

Open the `test/models/article_test.rb` file and create the first test case by deleting the test "the truth" method and replacing it with a test called test "should create article". Your file should look like Listing 11-5.

Listing 11-5. The Create Article Test in `test/models/article_test.rb`: <http://gist.github.com/358401>

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  test "should create article" do
    article = Article.new
```

```

article.user = users(:eugene)
article.title = "Test article"
article.body = "Test body"

assert article.save
end
end

```

The test “should create article” case is standard article-creation fare. You create a new article in the same way you’d create one from the console. The only real difference is on the last line of the test case:

```
assert article.save
```

Note Fixtures can be accessed in your test cases by name. Use `fixture(:name)`, where `fixture` is the plural name of the model and `:name` is the symbolized name of the fixture you’re after. This returns an Active Record object on which you can call methods. Here, you get at the `eugene` user fixture using `users(:eugene)`.

Before you go any further, let’s take a deeper look at assertions as they pertain to `Test::Unit` and `ActiveSupport::TestCase`.

Testing with Assertions

Assertions are statements of expected outcome. As the README for `Test::Unit` states, assertions are like saying “I assert that `x` should be equal to `y`. If the assertion turns out to be correct, the assertion passes. If the assertion turns out to be false, the assertion fails, and `Test::Unit` reports a failure.

`Test::Unit` ships with a bevy of built-in assertions, and Rails adds a bunch of its own. You see the Rails-added assertions as you look at each test case; but first, here’s the standard set of `Test::Unit` assertions for reference:

```

assert(boolean, message=nil)
assert_block(message="assert_block failed.") do ... end
assert_equal(expected, actual, message=nil)
assert_in_delta(expected_float, actual_float, delta, message="")
assert_instance_of(klass, object, message="")
assert_kind_of(klass, object, message="")
assert_match(pattern, string, message="")
assert_nil(object, message="")
assert_no_match(regexp, string, message="")
assert_not_equal(expected, actual, message="")
assert_not_nil(object, message="")
assert_not_same(expected, actual, message="")
assert_nothing_raised(*args) do ... end
assert_nothing_thrown(message "") do ... end
assert_operator(object1, operator, object2, message="")
assert_raise(expected_exception_klass, message "") do ... end
assert_respond_to(object, method, message="")
assert_same(expected, actual, message="")
assert_send(send_array, message="")
assert_throws(expected_symbol, message "") do ... end

```

The `assert` method is perhaps the most basic of the lot. It asserts that the return value of its first argument is true. And you know that `article.save` returns true if the article saves and returns false otherwise. So, by asserting `article.save`, you successfully test that the article was saved. Pretty easy, isn't it?

Tip Geoffrey Grosenbach (a.k.a. Topfunk) has a useful cheat sheet that summarizes all available assertions. You can download it from <http://nubyonrails.com/articles/ruby-rails-test-rails-cheat-sheet>.

Let's run the test. This time, run only the models tests:

```
$ rake test:models
```

```
Started
.
Finished in 0.168879 seconds.

1 tests, 1 assertions, 0 failures, 0 errors, 0 skips
```

Just as the output from the test says, you ran one test (`test "should create article"`), which included one assertion (`assert article.save`), and everything passed. Life is good!

Adding a Find Test

Next on the list is testing that you can successfully find an article. You use the data in the fixture you created to help you. Add the method shown in Listing 11-6 after test `"should create article"`.

Listing 11-6. Test Case for Finding an Article in `test/models/article_test.rb`: <http://gist.github.com/358402>

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  test "should create article" do
    article = Article.new

    article.user  = users(:eugene)
    article.title = "Test article"
    article.body  = "Test body"

    assert article.save
  end

  test "should find article" do
    article_id = articles(:welcome_to_rails).id
    assert_nothing_raised { Article.find(article_id) }
  end
end
```

Here, you test that you can find an article of the given `id`. First, you grab the `id` attribute from the fixture, then you test that you can use `Article.find` to retrieve it. You use the assertion `assert_nothing_raised` because you know that `find` raises an exception if the record can't be found. If no exception is raised, you know that finding works. Again, run the test and see what happens:

```
$ rake test:models
Run options: --seed 36921
# Running tests:
...
Finished tests in 0.068310s, 29.2781 tests/s, 14.6391 assertions/s.

2 tests, 1 assertions, 0 failures, 0 errors, 0 skips
```

Sure enough, finding works! So far, so good.

Adding an Update Test

Your next move is to test updating. Add the test "should update article" case, as shown in Listing 11-7.

Listing 11-7. Test Case for Updating an Article in `test/models/article_test.rb`: <http://gist.github.com/358403>

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  test "should create article" do
    article = Article.new

    article.user  = users(:eugene)
    article.title = "Test article"
    article.body   = "Test body"

    assert article.save
  end

  test "should find article" do
    article_id = articles(:welcome_to_rails).id
    assert_nothing_raised { Article.find(article_id) }
  end

  test "should update article" do
    article = articles(:welcome_to_rails)
    assert article.update_attributes(:title => 'New title')
  end
end
```

First, you find the “Welcome to Rails” article from your fixture, and then you assert that changing the title via `update_attributes` returns true. Once again, run the test and see what happens:

```
$ rake test:models
```

```
Run options: --seed 2804
```

```
# Running tests:
```

```
...
```

```
Finished tests in 0.073883s, 40.6050 tests/s, 27.0700 assertions/s.
```

```
3 tests, 2 assertions, 0 failures, 0 errors, 0 skips
```

Adding a Destroy Test

Only one more test to go: `destroy`. You find an article, destroy it, and assert that Active Record raises an exception when you try to find it again. Listing 11-8 shows the test.

Listing 11-8. Test Case for Destroying an Article in `test/models/article_test.rb`: <http://gist.github.com/358404>

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  test "should create article" do
    article = Article.new

    article.user = users(:eugene)
    article.title = "Test article"
    article.body = "Test body"

    assert article.save
  end

  test "should find article" do
    article_id = articles(:welcome_to_rails).id
    assert_nothing_raised { Article.find(article_id) }
  end

  test "should update article" do
    article = articles(:welcome_to_rails)
    assert article.update_attributes(:title => 'New title')
  end

  test "should destroy article" do
    article = articles(:welcome_to_rails)
    article.destroy
    assert_raise(ActiveRecord::RecordNotFound) { Article.find(article.id) }
  end
end
```

The `assert_raise` assertion takes as an argument the class of the exception you expect to be raised for whatever you do inside the given block. Because you've deleted the article, you expect Active Record to respond with a `RecordNotFound` exception when you try to find the article you just deleted by `id`. Run the test and see what happens:

```
$ rake test:models
```

```
Run options: --seed 64949
```

```
# Running tests:
```

```
....
```

```
Finished tests in 0.078197s, 51.1526 tests/s, 38.3644 assertions/s.
```

```
4 tests, 3 assertions, 0 failures, 0 errors, 0 skips
```

And there you have it. You've successfully tested article CRUD.

Testing Validations

You have a few validations on your `Article` model, specifically for the presence of a title and body. Because you want to make sure these are working as expected, you need to test them. Add the method shown in Listing 11-9 to test that you can't create invalid articles.

Listing 11-9. Test Case for Validations in `test/models/article_test.rb`: <http://gist.github.com/358405>

```
require 'test_helper'

class ArticleTest < ActiveSupport::TestCase
  test "should create article" do
    article = Article.new

    article.user  = users(:eugene)
    article.title = "Test article"
    article.body   = "Test body"

    assert article.save
  end

  test "should find article" do
    article_id = articles(:welcome_to_rails).id
    assert_nothing_raised { Article.find(article_id) }
  end

  test "should update article" do
    article = articles(:welcome_to_rails)
    assert article.update_attributes(:title => 'New title')
  end
end
```

```

test "should destroy article" do
  article = articles(:welcome_to_rails)
  article.destroy
  assert_raise(ActiveRecord::RecordNotFound) { Article.find(article.id) }
end

test "should not create an article without title nor body" do
  article = Article.new
  assert !article.valid?
  assert article.errors[:title].any?
  assert article.errors[:body].any?
  assert_equal ["can't be blank"], article.errors[:title]
  assert_equal ["can't be blank"], article.errors[:body]
  assert !article.save
end

```

This is pretty straightforward, although you may have to read it a few times before it clicks. First, you instantiate a new Article object in the local variable `article`. Without having given it any attributes, you expect it to be invalid. So, you assert that it's not valid using `assert !article.valid?` (notice the `!`, which negates truth). Next, you access the `errors` hash to explicitly check for the attributes you expect to be invalid:

```

assert article.errors[:title].any?
assert article.errors[:body].any?

```

You also want to check that the validation responses are what you expect. To do this, you use the `assert_equal` assertion. Here's its basic syntax:

```
assert_equal(expected, actual)
```

To check the error messages, you again access the `errors` hash, but this time you ask for the specific messages associated with the given attribute:

```

assert_equal ["can't be blank"], article.errors[:title]
assert_equal ["can't be blank"], article.errors[:body]

```

Finally, you assert that `article.save` returns false using `!article.save`. Run the test one more time:

```
$ rake test:models
```

```
Run options: --seed 49599
```

```
# Running tests:
```

```
.....
```

```
Finished tests in 0.108590s, 46.0446 tests/s, 82.8802 assertions/s.
```

```
5 tests, 9 assertions, 0 failures, 0 errors, 0 skips
```

Feels good, doesn't it? Life isn't all roses, though, and requirements change. What if one day you decide to make a change to the Article model and remove the validation requirements for the `title` attribute? If that were to happen, your test would fail. If you want to try it, open the Article model in `app/models/article.rb` and remove `validates :title, :presence => true`, and then run the tests again.

When your requirements change, you often need to update your tests. We recommend updating the tests *first* (which should make them fail) and then updating your code (which makes them pass). This is also known as test-driven development (TDD).

Functional Testing Your Controllers

Tests to check your controllers are called *functional* tests. When you tested your models, you didn't test them in the context of the web application—there were no web requests and responses nor were there any URLs to contend with. This focused approach lets you hone in on the specific functionality of the model and test it in isolation. Alas, Rails is great for building web applications; and although unit-testing models is important, it's equally important to test the full request/response cycle.

Testing the Articles Controller

Functional tests aren't that much different from unit tests. The main difference is that Rails sets up request and response objects for you; those objects act just like the live requests and responses you get when running the application via a web server. If you open the articles controller test in `test/functional/articles_controller_test.rb` and examine the first few lines, as shown in Listing 11-10, you can see how this is done.

Listing 11-10. Setup of a Functional Test in `test/controllers/articles_controller_test.rb`

```
require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  # ...
end
```

Just as in the unit test, the first thing you do is require the `test_helper`. The `test_helper.rb` file sets up some common environment variables and generally endows `Test::Unit` with specific methods that make testing Rails applications easier.

Note You can think of `test_helper` as being akin to `application_helper`. Any methods you define here are available to all your tests.

Notice that `ArticlesControllerTest` is actually a subclass of `ActionController::TestCase`, which performs some magic for you behind the scenes. It prepares three instance variables for you to use in your tests: the first is `@controller` as an instance variable of `ArticlesController`, after which it instantiates both `@request` and `@response` variables, which are instances of `ActionController::TestRequest` and `ActionController::TestResponse`, respectively. As you can no doubt tell by their names, these objects are made specifically for testing, and they're designed to simulate the Action Controller environment as closely as possible.

Most of the time, you don't need to worry about all this. Still, it's important to know what's going on. Because the test you're looking at was created by the scaffold generator, it has quite a bit more code than you would get from the standard controller generator. There's a problem with this code, though: not all the test cases will pass—at least, not without some modification. Warts and all, this gives you a good start and serves well as a template.

As you look over the articles controller functional test file, notice that each test case tests a specific request for an action on the controller. There's a test for every action: index, show, new, create, edit, update, and destroy. Let's walk through each test case, making adjustments as you go.

Creating a Test Helper

Before you test your actions, a little foresight tells you that in order to create an article, your application expects a logged-in user. So you need to simulate a logged-in user for your tests. This is a perfect job for a test helper. You can create a helper method called `login_as` that accepts the name of the user for login. This method sets up the session object, just as your controller expects. You can use this method for any test case that requires a login.

To begin, open the `test_helper` file in your editor and add the `login_as` method, as shown in Listing 11-11. You can find the test helper file in `test/test_helper.rb`; the method you're adding is highlighted in bold.

Listing 11-11. The `login_as` Test Helper in `test/test_helper.rb`: <https://gist.github.com/adamgamble/6198360>

```
ENV["RAILS_ENV"] = "test"
require File.expand_path(File.dirname(__FILE__) + '/../config/environment')
require 'rails/test_help'

class ActiveSupport::TestCase
  ActiveRecord::Migration.check_pending!

  # Setup all fixtures in test/fixtures/*.yml for all tests in alphabetical order.
  #
  # Note: You'll currently still have to declare fixtures explicitly in integration tests
  # -- they do not yet inherit this setting
  fixtures :all

  # Add more helper methods to be used by all tests here...
  def login_as(user)
    @request.session[:user_id] = users(user).id
  end
end
```

The `login_as` method is simple. All it does is manually set `user_id` in the `@request.session` object (just like your `login` action does) to the `id` of the given user, as obtained from the fixture. If you give it the name of one of your users fixtures, say, `:eugene`, it sets `session[:user_id]` to `users(:eugene).id`.

Now that you've created a way to simulate a logged-in user, you're ready to proceed with your tests, beginning with the `index` action.

Testing the Index Action

The updated setup method and test "should get index" are shown in Listing 11-12. Make sure yours looks like this before you proceed.

Listing 11-12. Updated Test Case for the Index Action in test/controllers/articles_controller_test.rb:
<http://gist.github.com/358409>

```
require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end

  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    get :new
    assert_response :success
  end

  test "should create article" do
    assert_difference('Article.count') do
      post :create, :article => @article.attributes
    end
    assert_redirected_to article_path(assigns(:article))
  end

  test "should show article" do
    get :show, :id => @article.to_param
    assert_response :success
  end

  test "should get edit" do
    get :edit, :id => @article.to_param
    assert_response :success
  end

  test "should update article" do
    put :update, :id => @article.to_param, :article => @article.attributes
    assert_redirected_to article_path(assigns(:article))
  end

  test "should destroy article" do
    assert_difference('Article.count', -1) do
      delete :destroy, :id => @article.to_param
    end

    assert_redirected_to articles_path
  end
end
```

The setup method is executed before every test case. In this case, the setup method assigns the `:welcome_to_rails` record from the fixtures to an instance variable `@article`; the `@article` variable is available to all test cases in the `ArticlesControllerTest` class.

Functional tests define methods that correspond to HTTP verbs (GET, POST, PUT, and DELETE), which you use to make requests. The first line of the test "should get index" method makes a GET request for the index action using `get :index`. Here's the full syntax you use for these requests:

```
http_method(action, parameters, session, flash)
```

In the case of test "should get index", you have no parameters to submit along with the request, so your call is simple. It makes a GET request to the `index` action just as if you had done so with a browser. After the request has been made, you need to assert your expectations:

```
assert_response :success
```

The `assert_response` assertion is a custom assertion defined by Rails (that is, it's not part of the standard `Test::Unit` library), and it does exactly what its name implies: it asserts that there was a successful response to the request.

Every time you make an HTTP request, the server responds with a status code. When the response is successful, the server returns a status code of 200. When an error occurs, it returns 500. And when the browser can't find the resource being requested, it returns 404. In your assertion, you use the shortcut `:success`, which is the same as 200. You could have used `assert_response(200)`, but it's easier to remember words like *success* or *error* than HTTP status codes, which is why we avoid using the latter whenever possible. Table 11-1 lists the shortcuts available when using `assert_response`.

Table 11-1. Status-Code Shortcuts Known to `assert_response`

Symbol	Meaning
<code>:success</code>	Status code was 200
<code>:redirect</code>	Status code was in the 300–399 range
<code>:missing</code>	Status code was 404
<code>:error</code>	Status code was in the 500–599 range

Tip You can pass an explicit status-code number to `assert_response`, such as `assert_response(501)` or its symbolic equivalent `assert_response(:not_implemented)`. See <http://iana.org/assignments/http-status-codes> for the full list of codes and default messages you can use.

You also want to assert that the proper template was rendered in response to the request, for which you use another of Rails's custom assertions: `assert_template`. Here, you expect to see the `index` template (from `app/views/articles/index.html.erb`) to be rendered, so test this expectation:

```
assert_template 'index'
```

You need to do one more thing: assert that the correct instance variables were assigned. If you look at the `articles` controller, you see that you set an instance variable called `@articles` that contains the `articles` collection. Rails gives you the ability to test whether this assignment was successful by way of the `assigns` method:

```
assert_not_nil assigns(:articles)
```

This asserts that `@articles` was, in fact, assigned (by virtue of the fact that it shouldn't be `nil`). You can use this technique to test for the existence of any instance variable set within your controllers. Useful, isn't it? Using `assigns` gives you access to the instance variable, so you can do with it as you please.

Testing the Show Action

Listing 11-13 shows the test "should show article" case.

Listing 11-13. Updated Test Case for the show Action in `test/functional/articles_controller_test.rb`:
<http://gist.github.com/358410>

```
require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end

  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    get :new
    assert_response :success
  end

  test "should create article" do
    assert_difference('Article.count') do
      post :create, :article => @article.attributes
    end

    assert_redirected_to article_path(assigns(:article))
  end

  test "should show article" do
    get :show, :id => @article.to_param

    assert_response :success
    assert_template 'show'
```

```

assert_not_nil assigns(:article)
assert assigns(:article).valid?
end

test "should get edit" do
  get :edit, :id => @article.to_param
  assert_response :success
end

test "should update article" do
  put :update, :id => @article.to_param, :article => @article.attributes
  assert_redirected_to article_path(assigns(:article))
end

test "should destroy article" do
  assert_difference('Article.count', -1) do
    delete :destroy, :id => @article.to_param
  end

  assert_redirected_to articles_path
end
end

```

The test "should show article" case is almost the same as test "should get index", but with one notable difference: you need to identify the record you want to show. If you were requesting this in a browser, the URL would look like /articles/1. Therefore, you need to pass in the :id parameter with a value of 1. If you look closely at the test request, you can see how this is done:

```
get :show, :id => @article.to_param
```

`@article.to_param` by default gets the `id` attribute of the `:welcome_to_rails` article you're getting from the fixtures that are assigned to the `@article` variable. You should always use the `to_param` method instead of `id` because it's the method Rails uses to generate URLs internally. You can pass arbitrary parameters in this fashion. You see more of this when you test the `create` and `update` actions, both of which require a set of article parameters.

One more thing to notice here: you can treat the results of `assigns(:article)` as you would any `Article` object and call methods on it:

```
assert assigns(:article).valid?
```

So, not only can you assert that there is an instance variable named `@article`, but you can also assert that it contains a valid `Article` object.

Testing the New Action

Listing 11-14 shows the test "should get new" case.

Listing 11-14. Updated Test Case for the new Action in test/functional/articles_controller_test.rb:
<http://gist.github.com/358413>

```
require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end

  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    login_as(:eugene)
    get :new
    assert_response :success
  end

  test "should create article" do
    assert_difference('Article.count') do
      post :create, :article => @article.attributes
    end

    assert_redirected_to article_path(assigns(:article))
  end

  test "should show article" do
    get :show, :id => @article.to_param

    assert_response :success
    assert_template 'show'

    assert_not_nil assigns(:article)
    assert assigns(:article).valid?
  end

  test "should get edit" do
    get :edit, :id => @article.to_param
    assert_response :success
  end
```

```

test "should update article" do
  put :update, :id => @article.to_param, :article => @article.attributes
  assert_redirected_to article_path(assigns(:article))
end

test "should destroy article" do
  assert_difference('Article.count', -1) do
    delete :destroy, :id => @article.to_param
  end

  assert_redirected_to articles_path
end
end

```

The test "should get new" case is pretty simple. First, you use the `login_as` helper method you created earlier to log in as a user; then, you get the new action; and finally, you assert the success of the response.

Testing the Create Action

Listing 11-15 shows the test "should create article" case. Notice how you use the `login_as` helper, because this action expects a logged-in user.

Listing 11-15. Test Case for the create Action in `test/functional/articles_controller_test.rb`:

<http://gist.github.com/358414>

```

require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end

  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    login_as(:eugene)
    get :new
    assert_response :success
  end

  test "should create article" do
    login_as(:eugene)
    assert_difference('Article.count') do
      post :create, :article => { :title => 'Post title',
                                  :body  => 'Lorem ipsum..'}
    end
  end

```

```

assert_response :redirect
assert_redirected_to article_path(assigns(:article))
end

test "should show article" do
  get :show, :id => @article.to_param

  assert_response :success
  assert_template 'show'

  assert_not_nil assigns(:article)
  assert assigns(:article).valid?
end

test "should get edit" do
  get :edit, :id => @article.to_param
  assert_response :success
end

test "should update article" do
  put :update, :id => @article.to_param, :article => @article.attributes
  assert_redirected_to article_path(assigns(:article))
end

test "should destroy article" do
  assert_difference('Article.count', -1) do
    delete :destroy, :id => @article.to_param
  end

  assert_redirected_to articles_path
end
end

```

To test the `create` action, you need to submit some form parameters to create a valid article. Fortunately, this is easy. All you need to do is pass a hash of parameters that contains a valid set of article attributes, just as you would using an HTML form (remember that HTML form parameters are converted into a hash object by Rails). Here's how it's done:

```
post :create, :article => { :title => 'Post title',
                            :body => 'Lorem ipsum..' }
```

Unlike the other test cases you've looked at so far, the test `"should create article"` case uses a request method other than GET. To create a new article, you need to use the POST method. You formulate a POST request that includes a `params` hash with a valid article.

Notice how you request the post to create an article inside the `assert_difference` method block? The `assert_difference` method does a simple job: it takes a parameter and compares it with itself after running the block content. It expects the difference to be 1 by default. For this case, it expects `Article.count` to return the same count plus 1 after running the post request to create an article, which is the case in a successful create.

Next, you come across another of Rails's additions to `Test::Unit`, `assert_redirected_to`. As you can gather from its name, it lets you assert that a redirect to the expected location took place in response to the request:

```
assert_redirected_to(options, message)
```

You're really rolling now. The test "should get edit" case is straightforward, and test "should update article" is similar to test "should create article". Listing 11-16 shows the updated methods.

Listing 11-16. Test Case for the `create` Action in `test/functional/articles_controller_test.rb`:
<http://gist.github.com/358416>

```
require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end
  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    login_as(:eugene)
    get :new
    assert_response :success
  end

  test "should create article" do
    login_as(:eugene)
    assert_difference('Article.count') do
      post :create, :article => { :title => 'Post title',
                                  :body   => 'Lorem ipsum..' }
    end

    assert_response :redirect
    assert_redirected_to article_path(assigns(:article))
  end

  test "should show article" do
    get :show, :id => @article.to_param

    assert_response :success
    assert_template 'show'

    assert_not_nil assigns(:article)
    assert assigns(:article).valid?
  end
end
```

```

test "should get edit" do
  login_as(:eugene)
  get :edit, :id => @article.to_param
  assert_response :success
end

test "should update article" do
  login_as(:eugene)
  put :update, :id => @article.to_param, :article => { :title => 'New Title' }
  assert_redirected_to article_path(assigns(:article))
end

test "should destroy article" do
  assert_difference('Article.count', -1) do
    delete :destroy, :id => @article.to_param
  end

  assert_redirected_to articles_path
end
end

```

Testing the Destroy Action

Listing 11-17 shows the test "should destroy article" case. Again, because this action expects a logged-in user, you use the `login_as` helper to log in as `:eugene`.

Listing 11-17. Test Case for the destroy Action in `test/functional/articles_controller_test.rb`:
<http://gist.github.com/358417>

```

require 'test_helper'

class ArticlesControllerTest < ActionController::TestCase
  setup do
    @article = articles(:welcome_to_rails)
  end

  test "should get index" do
    get :index
    assert_response :success
    assert_template 'index'
    assert_not_nil assigns(:articles)
  end

  test "should get new" do
    login_as(:eugene)
    get :new
    assert_response :success
  end

```

```

test "should create article" do
  login_as(:eugene)
  assert_difference('Article.count') do
    post :create, :article => { :title => 'Post title',
                                :body   => 'Lorem ipsum..' }
  end

  assert_response :redirect
  assert_redirected_to article_path(assigns(:article))
end

test "should show article" do
  get :show, :id => @article.to_param

  assert_response :success
  assert_template 'show'

  assert_not_nil assigns(:article)
  assert assigns(:article).valid?
end

test "should get edit" do
  login_as(:eugene)
  get :edit, :id => @article.to_param
  assert_response :success
end

test "should update article" do
  login_as(:eugene)
  put :update, :id => @article.to_param, :article => { :title => 'New Title' }
  assert_redirected_to article_path(assigns(:article))
end

test "should destroy article" do
  login_as(:eugene)
  assert_nothing_raised { Article.find(@article.to_param) }

  assert_difference('Article.count', -1) do
    delete :destroy, :id => @article.to_param
  end
  assert_response :redirect
  assert_redirected_to articles_path

  assert_raise(ActiveRecord::RecordNotFound) { Article.find(@article.to_param) }
end

```

First, you test that you can find the article in question, knowing full well that `find` will raise an exception if the article doesn't exist. Then, you formulate a `DELETE` request to the `destroy` action, passing in the `id` of the article to destroy. You assert that the response is a redirect to the `index` action; and finally, you ensure that the article has been deleted by asserting that Active Record raises a `RecordNotFound` exception.

Now that you've fixed the articles controller test suite, let's get Ruby to execute all methods and verify whether you've really fixed everything:

```
$ ruby -Itest test/controllers/articles_controller_test.rb
```

```
Loaded suite test/controllers/articles_controller_test
Started
.....
Finished in 0.432005 seconds.

7 tests, 20 assertions, 0 failures, 0 errors, 0 skips
```

One Missing Test

Before you can run the complete functional test suite, you still have to fix one test. When you generated the Notifier mailer, Rails added a mailer test file in the `test/mailers` directory, with tests that are now failing because of changes you made in that class. Listing 11-18 shows an updated version of the test.

Listing 11-18. Updated `test/mailers/notifier_test.rb` File: <http://gist.github.com/358418>

```
require 'test_helper'

class NotifierTest < ActionMailer::TestCase
  test "email_friend" do
    article = articles(:welcome_to_rails)
    message = Notifier.email_friend(article, 'John Smith', 'dude@example.com')

    assert_equal "Interesting Article", message.subject
    assert_equal ["dude@example.com"], message.to
    assert_equal ["from@example.com"], message.from
  end
end
```

In this code, you use the `message` variable to represent the e-mail message generated by a call to `Notifier.email_friend(article, 'John Smith', 'dude@example.com')`. You later use the `message` variable to verify that some properties of the e-mail are as expected, using `assert_equal`: in this case, subject, sender, and receiver e-mail addresses.

Running the Full Test Suite

Now that your functional testing tour is complete, run the `rake test:functionals` command, which runs your entire suite of functional tests:

```
$ rake test:functionals
```

```
Loaded suite
Started
.....
Finished in 0.665524 seconds.

8 tests, 23 assertions, 0 failures, 0 errors, 0 skips
```

Not bad—23 assertions all in less than one second! You've got to admit, this is a lot more efficient than manually clicking through your application to test it. Moreover, because this uses the test database, you don't risk polluting your production database with bogus data while you test. Whenever you make a change to your articles controller, you can run this test to see if you've broken any of your expectations.

Now that you have unit and functional tests for articles in place, let's run the entire test suite, which runs both your unit and functional tests. To do this, use the built-in Rake task, `test`:

```
$ rake test
```

```
Loaded suite
Started
.........
Finished in 0.230277 seconds.

13 tests, 32 assertions, 0 failures, 0 errors, 0 skips
```

Note Some of the output is omitted here to cut down on the clutter. Your actual output will look slightly different, but the summary (tests, assertions, failures, and errors) should be the same.

It may interest you to know that the default Rake task is to run all tests. That means running `rake` with no arguments is the same as running `rake test`.

Integration Testing

Rails defines one more type of test, and it's the highest level of the bunch. Integration tests go a little further than their functional equivalents. Unlike functional tests, which test a specific controller, integration tests can span multiple controllers and actions with full session support. They're the closest you can get to simulating actual interaction with a web application. They test that the individual pieces of your application integrate well together.

Integration Testing the Blog Application

Let's get started by generating the integration test. Given that Rails ships with a generator for just about everything, it shouldn't surprise you that it includes one for generating integration tests. It works like the others you've already used:

```
$ rails generate test_unit:integration UserStories
```

```
create  test/integration/user_stories_test.rb
```

Open the newly generated file and take a peek, as shown in Listing 11-19.

Listing 11-19. User Stories Test in test/integration/user_stories_test.rb

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest
  # test "the truth" do
  #   assert true
  # end
end
```

At this stage, it looks a lot like the other test files you've seen so far. Notice, however, that it's a subclass of `ActionController::IntegrationTest`. That's about the only difference, but not for long.

Test cases are added to integration tests in exactly the same way unit and functional test cases are added: using the `test` method with a description string and a block where you write your assertions. Integration test cases tend to look deceptively like functional tests, but they have a few subtle differences, which we point out as you add them.

`Listing 11-20` shows a test case that goes through the process of logging in a user.

Listing 11-20. Login Integration Test in test/integration/user_stories_test.rb: <http://gist.github.com/358420>

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest

  test "should login user and redirect" do
    get login_path

    assert_response :success
    assert_template 'new'

    post session_path, :email => 'eugene@example.com', :password => 'secret'

    assert_response :redirect
    assert_redirected_to root_path

    follow_redirect!

    assert_response :success
    assert_template 'index'
    assert session[:user_id]
  end
end
```

When your response is a redirect, you're able to follow it, even when it redirects to another controller. The `follow_redirect!` method does exactly what you might think: it lets you follow a single redirect response, as long as the last response was, in fact, a redirect. If the last response wasn't a redirect, an exception is raised.

Let's add a test case for the logout action as well, as shown in Listing 11-21.

Listing 11-21. Logout Integration Test in test/integration/user_stories_test.rb: <http://gist.github.com/358421>

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest

  test "should login user and redirect" do
    get login_path

    assert_response :success
    assert_template 'new'

    post session_path, :email => 'eugene@example.com', :password => 'secret'
    assert_response :redirect
    assert_redirected_to root_path

    follow_redirect!

    assert_response :success
    assert_template 'index'
    assert session[:user_id]
  end

  test "should logout user and redirect" do
    get logout_path

    assert_response :redirect
    assert_redirected_to root_path
    assert_nil session[:user]

    follow_redirect!

    assert_template 'index'
  end
end
```

Again, you can follow the redirect and test that the correct template was rendered.

Let's get a little fancier by testing that you can log in, create a new article, and log out, all in a single test. Combine the login and logout tests you've already written and sandwich an article-creation test in the middle, as shown in Listing 11-22.

Listing 11-22. Article Creation Integration Test in test/integration/user_stories_test.rb:

<http://gist.github.com/358422>

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest

  test "should login create article and logout" do
    # Login
    get login_path
```

```
assert_response :success
assert_template 'new'

post session_path, :email => 'eugene@example.com', :password => 'secret'

assert_response :redirect
assert_redirected_to root_path

follow_redirect!

assert_response :success
assert_template 'index'
assert session[:user_id]

# Create New Article
get new_article_path

assert_response :success
assert_template 'new'

post articles_path, :article => { :title => 'Integration Tests',
:body => 'Lorem Ipsum..'}

assert assigns(:article).valid?
assert_response :redirect
assert_redirected_to article_path(assigns(:article))

follow_redirect!

assert_response :success
assert_template 'show'

# Logout
get logout_path

assert_response :redirect
assert_redirected_to root_path
assert_nil session[:user]

follow_redirect!

assert_template 'index'
end
end
```

Here's what you get when you run the test:

```
$ ruby -Itest test/integration/user_stories_test.rb
```

```
Loaded suite test/integration/user_stories_test
Started
.
Finished in 0.490546 seconds.

1 tests, 21 assertions, 0 failures, 0 errors, 0 skips
```

Great—and you've just tested the whole stack from dispatcher to database. Not too shabby, is it? If you're thinking that what you've just done looks a lot like you're telling a story (Eugene logs in, Eugene creates article, Eugene logs out), you're right.

Story-Based Testing

Integration tests are great for creating story-based scenarios using a domain-specific language (DSL). They even go so far as to allow you to test multiple users interacting! What are story-based tests? Well, suppose you could do something like this:

```
test "creating an article" do
  eugene = registered_user
  eugene.logs_in 'eugene', 'secret'
  eugene.create_article
  eugene.logs_out
end
```

Here, you're telling an easy-to-understand story that requires no programming knowledge to follow. Eugene logs in and proceeds to create a new article. When he is finished, he logs out. Behind the scenes, you can test every request, response, and redirect, following Eugene's path through the entire process, just as you did in the test "should login create article and logout" case.

Integration tests provide a method called `open_session` that you can use to simulate a distinct user interacting with the application as if from a web browser. Although this lets you simulate multiple connections with ease, you use it to help create an object on which to define your custom story-based methods, like `logs_in` and `creates_article`.

Let's reshape test "should login create article and logout" into methods you can add straight onto a new `session` object. Listing 11-23 shows the updated `user_stories_test.rb` test.

Listing 11-23. Updated User Stories Integration Test in `test/integration/user_stories_test.rb`:
<http://gist.github.com/358423>

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest

  test "creating an article" do
    eugene = registered_user
    eugene.logs_in 'eugene@example.com', 'secret'
    eugene.create_article :title => 'Integration tests', :body => 'Lorem Ipsum...'
    eugene.logs_out
  end
end
```

```
private
  def registered_user
    open_session do |user|
      def user.logs_in(email, password)
        get login_path

        assert_response :success
        assert_template 'new'

        post session_path, :email => email, :password => password

        assert_response :redirect
        assert_redirected_to root_path

        follow_redirect!

        assert_response :success
        assert_template 'index'
        assert session[:user_id]
      end

      def user.logs_out
        get logout_path

        assert_response :redirect
        assert_redirected_to root_path
        assert_nil session[:user]

        follow_redirect!

        assert_template 'index'
      end

      def user.creates_article(article_hash)
        get new_article_path

        assert_response :success
        assert_template 'new'

        post articles_path, :article => article_hash

        assert assigns(:article).valid?
        assert_response :redirect
        assert_redirected_to article_path(assigns(:article))

        follow_redirect!

        assert_response :success
        assert_template 'show'
      end
    end
  end
end
```

See how you create a private method called `registered_user` and create a new integration session inside it? The `open_session` method yields a `session` object onto which you attach *singleton* methods (methods that exist only on a particular instance):

```
def registered_user
  open_session do |user|
    def user.logs_in(email, password)
      #...
    end

    def user.logs_out
      #...
    end

    def user.creates_article(article_hash)
      #...
    end
  end
end
```

The return value of `registered_user`, then, is a fresh integration session object that responds to the methods you've created. This means you can create as many user sessions as you want and simulate multiple connections to the application. Listing 11-24 updates the test "creating an article" method and renames it test "multiple users creating an article".

Listing 11-24. Updated Article-Creation Story Testing Multiple Users in `test/integration/user_stories_test.rb`:
<http://gist.github.com/358424>

```
require 'test_helper'

class UserStoriesTest < ActionDispatch::IntegrationTest
  fixtures :all

  test "multiple users creating an article" do
    eugene = registered_user
    lauren = registered_user

    eugene.logs_in 'eugene@example.com', 'secret'
    lauren.logs_in 'lauren@example.com', 'secret'

    eugene.creates_article :title => 'Integration Tests', :body => 'Lorem Ipsum...'
    lauren.creates_article :title => 'Open Session', :body => 'Lorem Ipsum...'

    eugene.logs_out
    lauren.logs_out
  end

  private
  def registered_user
    open_session do |user|
      def user.logs_in(email, password)
        get login_path
```

```
assert_response :success
assert_template 'new'
post session_path, :email => email, :password => password

assert_response :redirect
assert_redirected_to root_path

follow_redirect!

assert_response :success
assert_template 'index'
assert session[:user_id]
end

def user.logs_out
  get logout_path

  assert_response :redirect
  assert_redirected_to root_path
  assert_nil session[:user]

  follow_redirect!

  assert_template 'index'
end

def user.creates_article(article_hash)
  get new_article_path

  assert_response :success
  assert_template 'new'

  post articles_path, :article => article_hash

  assert assigns(:article).valid?
  assert_response :redirect
  assert_redirected_to article_path(assigns(:article))

  follow_redirect!

  assert_response :success
  assert_template 'show'
end
end
end
```

This is only the tip of the iceberg. The sky is the limit for how creative you can get with this style of testing.

Running the Full Test Suite

Now that you have a respectable number of tests for your application, let's use Rake to run the entire suite. Because the default Rake task is `test`, you can save yourself five keystrokes (a space counts as a keystroke!) and run your tests with just `rake`:

```
$ rake
```

```
Loaded suite
Started
.....
```

```
Finished in 0.237562 seconds.
```

```
14 tests, 32 assertions, 0 failures, 0 errors, 0 skips
```

Take a moment to bask in the glory of a successful test run.

Summary

This chapter introduced you to the Rails philosophy behind testing and stressed its importance as part of the development cycle. You've now been on a complete tour of the baked-in facilities Rails provides for testing. You learned about testing your models with unit tests, testing your controllers with functional tests, and testing the whole Rails stack with integration tests.

Testing is an important part of the development cycle. Despite the fact that we left it until near the end of this book, it's not something you should treat as an afterthought. Now that you know how to write a Rails application and how to test it, you can combine the steps: write some code, and then test it. As you get into the code/test rhythm (or better yet, test/code), you'll find that you can write better, more reliable software. And you may sleep a little better at night, too.

I'd be remiss if I didn't mention some of the other test frameworks that are available from the Ruby community. In fact, test frameworks become almost like a religion to developers. DHH famously loves the test unit, while many others disagree. Rspec (<http://rspec.info/>) is a popular choice, as is minitest (<https://github.com/seattlerb/minitest>). There is also cucumber (<http://cukes.info/>), which is built off of Rspec, but uses a language called Gherkin to let you write tests. No matter which framework you decide to use, make sure you test early and often. Not only does it ensure your application does what you expect, but it is also frequently used as a source of documentation by developers moving into your project for the first time.

The next chapter will look at preparing your applications for a global audience through the use of Rails's built-in internationalization and localization support.



Internationalization

Internationalization in Rails used to be a complex task until Rails version 2.2 came out with internationalization and localization support built in. Since then, launching an application in another language or even multiple languages has become a relatively simple task. According to Wikipedia, “Internationalization is the process of designing a software application so that it can be adapted to various languages and regions without engineering changes. Localization is the process of adapting software for a specific region or language by adding locale-specific components and translating text” (http://en.wikipedia.org/wiki/Internationalization_and_localization).

This chapter explains internationalization and localization support in Rails. You first set up internationalization in the blog application with English as the main language; then you localize it to another language; and, finally, you support both languages and allow users to pick the language they want.

Internationalization and *localization* are long words, so developers use short names for them. The short name for internationalization is i18n, which is the first and the last letters of the word with the count of how many characters are in between. Following the same logic, localization’s short name is l10n.

Note If you need to get the code at the exact point where you finished Chapter 11, download the source code zip file from the book’s page on <http://www.apress.com> and extract it on your computer.

Internationalization Logic in Rails

The i18n and l10n support in Rails is based on a single module that takes care of all the translation and locale changes for you; this module is called I18n. The I18n module’s main method is `translate`, which translates content by looking for a translation text in a locale file, normally located in `config/locales`.

Locales are like languages but are more specific to regions. For example, `en` represents English in general, whereas `en-us` represents US English and `en-uk` represents UK English. In Rails, those differences are reflected in the translation files, mainly for localization options like time, date formats, and currency.

If you look in the `config/locales` directory, you’ll see a file called `en.yml`; it’s a YAML file that defines the English translations for your application. Open the file and you’ll see something similar to the code in Listing 12-1.

Listing 12-1. The Default English Locale File in `config/locales/en.yml`

```
# Files in the config/locales directory are used for internationalization
# and are automatically loaded by Rails. If you want to use locales other
# than English, add the necessary files in this directory.
#
```

```
# To use the locales, use 'I18n.t':
#
#     I18n.t 'hello'
#
# In views, this is aliased to just 't':
#
#     <%= t('hello') %>
#
# To use a different locale, set it with 'I18n.locale':
#
#     I18n.locale = :es
#
# This would use the information in config/locales/es.yml.
#
# To learn more, please read the Rails Internationalization guide
# available at http://guides.rubyonrails.org/i18n.html.
```

en:
 hello: "Hello world"

A translation file is a YAML file. It starts with the locale symbol, which is also the translation file name; in this case, it's en. Then, the file lists the translations in a key-value pair style: the en.yml example defines the translation of hello as "Hello world".

Now, let's see the translation in action by trying the I18n module in the console. Launch it with rails console:

```
I18n.translate "hello"
=> "Hello world"
I18n.t "hello"
=> "Hello world"
```

You pass the *key* to the translate method, and it returns the corresponding *value* from the English locale file. The I18n module has the t method as an alias for the translate method, which you used in the previous example.

The i18n works in Rails by having a set locale. At any time, you can determine the current locale by calling the I18n.locale method. When you don't set the locale yourself, it's set to a default locale, normally en. You can access the default locale by calling I18n.default_locale. Let's check the current locale and the default locale in your application:

```
I18n.locale
=> :en
I18n.default_locale
=> :en
```

To change the locale or the default locale, you can use the I18n.locale= and I18n.default_locale= methods. Change the locale to Brazilian Portuguese, for which the locale symbol is pt-br, and try the translate method again:

```
>> I18n.locale = 'pt-br'
=> "pt-br"
>> I18n.translate "hello"
=> "translation missing: pt-br, hello"
```

You don't have a translation for `hello` in Brazilian Portuguese—the `translate` method returns a string saying `translation missing: pt-br, hello`. To define the translation for `hello` in Brazilian Portuguese, you create a new translation file named after the locale symbol `pt-br.yml` in `config/locales`, as shown in Listing 12-2.

Listing 12-2. The Brazilian Portuguese Locale File in `config/locales/pt-br.yml`: <http://gist.github.com/360760>

```
pt-br:
  hello: "Ola mundo"
```

Rails doesn't reload locale files automatically, unlike other files. So, exit the console, restart it to make sure Rails loads the new translation file, and try again:

```
>> I18n.locale
=> :en
>> I18n.locale = 'pt-br'
=> "pt-br"
>> I18n.t "hello"
=> "Ola mundo"
```

That's how simple it is. All you need are the translation files, each with several translations in key-value pairs. You access those translations by passing the corresponding key to the `I18n.translate` method or its alias `I18n.t`.

Rails manages all its internals using the `I18n` module. For example, all the Active Record validation messages you saw in Chapters 5 and 6 are called by using the `translate` method and referring to a translation key. If you change the locale, Rails has no translation for those error messages. Check it out:

```
>> I18n.locale = 'pt-br'
=> "pt-br"
>> article = Article.new
=> #<Article id: nil, title: nil, body: nil, published_at: nil, created_at: nil,
updated_at: nil, excerpt: nil, location: nil, user_id: nil>
>> article.save
=> false
>> article.errors.full_messages
=> ["Title translation missing: pt-br, activerecord, errors, models, article,
attributes, title, blank", "Body translation missing: pt-br, activerecord,
errors, models, article, attributes, body, blank"]
```

Active Record tried to get the translations for the error messages, but it couldn't find them in the Brazilian Portuguese translation file. Let's add them by updating `pt-br.yml` so it looks like the code in Listing 12-3.

Listing 12-3. Updated Brazilian Portuguese Locale File in `config/locales/pt-br.yml`: <http://gist.github.com/360764>

```
pt-br:
  hello: "Ola mundo"

  activerecord:
    errors:
      models:
        article:
          attributes:
            title:
```

```

blank: "não pode ficar em branco"
body:
blank: "não pode ficar em branco"

```

Notice how you nest the keys. The “translation missing” message you saw earlier in the console included a list of names: pt-br, activerecord, errors, models, article, attributes, body, and blank. Those names represent the path inside the pt-br translation file. I18n calls the blank key, for example, by using dots to connect it and its parents; the translate call is I18n.translate('activerecord.errors.models.article.attributes.title.blank'). Now that you've added the translation, try it from the console again (don't forget to restart your console):

```

>> I18n.locale = 'pt-br'
=> "pt-br"
>> article = Article.new
=> #<Article id: nil, title: nil, body: nil, published_at: nil, created_at: nil,
updated_at: nil, excerpt: nil, location: nil, user_id: nil>
>> article.save
=> false
>> I18n.translate('activerecord.errors.models.article.attributes.title.blank')
=> "não pode ficar em branco"
>> article.errors.full_messages
=> ["Title não pode ficar em branco", "Body não pode ficar em branco"]

```

Congratulations! You just translated the validates :title, :presence => true error message for both the title and body attributes of your Article model to Brazilian Portuguese. Now that you understand how I18n works, let's set it up in the blog application.

Setting Up i18n in the Blog Application

Rails at its core uses i18n. You'll also need to use it. To do so, you must make sure that all hard-coded text and strings are replaced with an I18n.translate method call.

This may sound like a lot of work, but it's fairly simple in this case because your application is still small. We encourage you to use the I18n.translate method in your project as early as possible; that way, you avoid doing the text replacement later—it gets more difficult as your project grows.

Let's begin with your models. The only one that uses strings is the Comment model, which includes a custom validation with an error message. Replace this error message with an I18n.t call to a key, and add that key to your en.yml translation file. Listing 12-4 shows how the comment model should look after you edit it.

Listing 12-4. Updated Comment Model in app/models/comment.rb: <http://gist.github.com/360774>

```

class Comment < ActiveRecord::Base
  belongs_to :article

  validates :name, :email, :body, :presence => true
  validate :article_should_be_published

  def article_should_be_published
    errors.add(:article_id, I18n.t('comments.errors.not_published_yet'))
  end
end

```

Notice how you use the dots notation in the `comments.errors.not_published_yet` key. It's good practice to keep the locale file organized; doing so helps you find the translation more easily when your file gets bigger. Because the error message is added to the comment object, you add it under `comments`; and because it's an error message, you drill a step deeper and place it under `errors`.

Don't forget to add the translation to your `en.yml` translation file. It should look like Listing 12-5 after you clean it up and update it with the new translation.

Listing 12-5. Updated English Locale File in `config/locales/en.yml`: <http://gist.github.com/360776>

```
en:
  comments:
    errors:
      not_published_yet: is not published yet
```

Now, let's move to controllers. If you check all the controllers, you see that you only need to translate the flash messages. Listing 12-6 shows how `ArticlesController` should look after you update it.

Listing 12-6. Updated `ArticlesController` in `app/controllers/articles_controller.rb`: <https://gist.github.com/adamgamble/6190882>

```
class ArticlesController < ApplicationController
  before_action :authenticate, except: [:index, :show]
  before_action :set_article, only: [:show]

  # GET /articles
  # GET /articles.json
  def index
    @articles = Article.all
  end

  # GET /articles/1
  # GET /articles/1.json
  def show
  end

  # GET /articles/new
  def new
    @article = Article.new
  end

  # GET /articles/1/edit
  def edit
    @article = current_user.articles.find(params[:id])
  end

  # POST /articles
  # POST /articles.json
  def create
    @article = current_user.articles.new(article_params)
```

```
respond_to do |format|
  if @article.save
    format.html { redirect_to @article, notice: t('articles.create_success') }
    format.json { render action: 'show', status: :created, location: @article }
  else
    format.html { render action: 'new' }
    format.json { render json: @article.errors, status: :unprocessable_entity }
  end
end

# PATCH/PUT /articles/1
# PATCH/PUT /articles/1.json
def update
  @article = current_user.articles.find(params[:id])
  respond_to do |format|
    if @article.update(article_params)
      format.html { redirect_to @article, notice: t('articles.update_success') }
      format.json { head :no_content }
    else
      format.html { render action: 'edit' }
      format.json { render json: @article.errors, status: :unprocessable_entity }
    end
  end
end

# DELETE /articles/1
# DELETE /articles/1.json
def destroy
  @article = current_user.articles.find(params[:id])
  @article.destroy
  respond_to do |format|
    format.html { redirect_to articles_url }
    format.json { head :no_content }
  end
end

def notify_friend
  @article = Article.find(params[:id])
  Notifier.email_friend(@article, params[:name], params[:email]).deliver
  redirect_to @article, :notice => t('articles.notify_friend_success')
end

private
  # Use callbacks to share common setup or constraints between actions.
  def set_article
    @article = Article.find(params[:id])
  end
```

```
# Never trust parameters from the scary internet, only allow the white list through.
def article_params
  params.require(:article).permit(:title, :location, :excerpt, :body, :published_at,
:category_ids => [])
end
```

There are two things to notice here. First, you call the `t` method without the `I18n` module, unlike what you did in the console and the model; that's because the `I18n` module is integrated with Action Pack to keep things cleaner in the controllers, helpers, and views. Second, you also nest the messages under `articles`—again, to keep things cleaner. You do the same with the other controllers, also nesting them under their corresponding name: `users` controller translations go under `users`, the `comments` controller goes under `comments`, the application controller goes under `application`, and the sessions controller goes under `session`.

Updating the rest of the controllers is fairly simple. You can do it yourself, or you can download the updated files from <http://www.apress.com>.

Finally, let's look at the views. They're very similar, so you look at a single view here: Listing 12-7 shows the `article` partial after using translations, with changes in bold. As with controllers, you can apply the rest of the changes yourself using the Gist URLs <http://gist.github.com/360814> and <http://gist.github.com/360819>.

Listing 12-7. Updated article Partial in `app/views/articles/_article.html.erb`:

<http://gist.github.com/360814>

```
<%= div_for article do %>
  <h3>
    <%= link_to article.title, article %>
    <% if article.owned_by? current_user %>
      <span class='actions'>
        <%= link_to t('general.edit'), edit_article_path(article) %>
        <%= link_to t('general.delete'), article,
:confirm => t('general.are_you_sure'), :method => :delete %>
      </span>
    <% end %>
  </h3>
  <%= simple_format raw(article.body) %>
<% end %>
```

After updating your code, you're one step away from completing the `i18n` setup. You still need to add the translations to the default locale file `en.yml`. Listing 12-8 shows the updated `config/locales/en.yml` translation file.

Listing 12-8. English Locale File After Implementing `i18n` Support in `config/locales/en.yml`:

<http://gist.github.com/360824>

```
en:
  general:
    are_you_sure: Are you sure?
    back: Back
    cancel: Cancel
    create: Create
    delete: Delete
    edit: Edit
    editing: Editing
    footer: A simple blog built for the book
```

```
email_a_friend: Email a friend
search: Search
send_email: Send email
show: Show
title: Blog
update: Update
your_name: Your name
your_friend_email: Your friend's email
or: or
application:
access_denied: Please log in to continue
articles:
editing_article: Editing Article
listing_articles: Listing Articles
new_article: New Article
article: article
create_success: Article was successfully created.
update_success: Article was successfully updated.
articles: articles
notify_friend_success: Successfully sent a message to your friend
users:
new_user: New user
edit_password: Edit Password
editing_user: Editing user
create_success: User successfully added.
update_success: Updated user information successfully.
sessions:
email: Email
password: Password
login: Login
logout: Logout
successful_login: Logged in successfully
invalid_login: Invalid login/password combination
logout_success: You successfully logged out
comments:
name: Name
email: Email
body: Body
comments: Comments
new_comment: New comment
create_success: Thanks for your comment
create_failure: Unable to add comment
destroy_success: Comment deleted
add: Add
errors:
not_published_yet: is not published yet
```

That's about it—the blog application has i18n support now.

You must restart your server for Rails to load the translation file. After you do that, browse the site: you won't see any differences yet. Although i18n support is in place, you're still using English as your locale. To see i18n in action, let's change the locale and try Brazilian Portuguese.

Localizing the Blog Application to Brazilian Portuguese

Localizing an i18n-ready Rails application is amazingly simple. All you have to do is add a new translation file and configure your Rails application to use that locale as the default locale, and you're good to go.

Brazilian Portuguese is the locale to which you localize the blog application in this section. The locale symbol for Brazilian Portuguese is `pt-br`; so, first you change the `config/locales/pt-br.yml` file, using the same keys as your English translation file, but with Brazilian Portuguese text instead of English.

The separation between the translation files and your application code is very helpful; it gives you the ability to send the translation file to a translator, for example. When you have it back, you plug it into your application and you're all set. Listing 12-9 shows the newly created Brazilian Portuguese translation file.

Listing 12-9. Brazilian Portuguese Locale File in `config/locales/pt-br.yml`: <http://gist.github.com/391492>

```
pt-br:
  general:
    are_you_sure: Tem certeza?
    back: Volta
    cancel: Cancelar
    create: Criar
    delete: Apagar
    edit: Editar
    editing: Editando
    footer: Um blog simples desenvolvido para o livro
    email_a_friend: Avisar um amigo
    search: Pesquisar
    send_email: Mandar email
    show: Mostrar
    title: Blog
    update: Atualizar
    your_name: Seu nome
    your_friend_email: O email do seu amigo
    or: ou
  application:
    access_denied: "Por favor, efetue o login para continuar"
  articles:
    editing_article: Editando Artigo
    listing_articles: Listando Artigos
    new_article: Novo Artigo
    article: artigo
    create_success: Artigo foi criado com sucesso.
    update_success: Artigo foi atualizado com sucesso.
    articles: artigos
    notify_friend_success: Seu amigo foi avisado a respeito desse artigo
  users:
    new_user: Novo Usuario
    edit_password: Editar senha
    editing_user: Editando usuario
    create_success: Usuario editado com sucesso.
    update_success: Usuario atualizado com sucesso.
```

```

sessions:
  email: Email
  password: Senha
  login: Logar
  logout: Desconectar
  successful_login: Logado com sucesso
  invalid_login: Senha ou Email invalidos
  logout_success: Voce desconectou do sistem com sucesso
comments:
  name: Nome
  email: Email
  body: Conteudo
  comments: Comentarios
  new_comment: Novo Comentario
  create_success: Obrigado pelo comentario
  create_failure: Nao foi possivel adicionar o comentario
  destroy_success: Comentario deletado
  add: Adicionar
  errors:
    not_published_yet: ainda nao foi publicado

activerecord:
  errors:
    models:
      article:
        attributes:
          title:
            blank: "não pode ficar em branco"
          body:
            blank: "não pode ficar em branco"

date:
  formats:
    default: "%d/%m/%Y"
    short: "%d de %B"
    long: "%d de %B de %Y"

  day_names: [Domingo, Segunda, Terça, Quarta, Quinta, Sexta, Sábado]
  abbr_day_names: [Dom, Seg, Ter, Qua, Qui, Sex, Sáb]
  month_names: [~, Janeiro, Fevereiro, Março, Abril, Maio, Junho, Julho,
  Agosto, Setembro, Outubro, Novembro, Dezembro]
  abbr_month_names: [~, Jan, Fev, Mar, Abr, Mai, Jun, Jul, Ago, Set, Out,
  Nov, Dez]
  order: [:day, :month, :year]

```

You still have to tell Rails to use pt-br as the default locale. You do that by adding a configuration to your config/application.rb file. Listing 12-10 shows the updated config/application.rb file with the added lines in bold.

Listing 12-10. Setting the Default Locale to Brazilian Portuguese in config/application.rb:
<https://gist.github.com/adamgamble/6190889>

```
require File.expand_path('../boot', __FILE__)

require 'rails/all'

# Require the gems listed in Gemfile, including any gems
# you've limited to :test, :development, or :production.
Bundler.require(:default, Rails.env)

module Blog
  class Application < Rails::Application
    # Settings in config/environments/* take precedence over those specified here.
    # Application configuration should go into files in config/initializers
    # -- all .rb files in that directory are automatically loaded.

    # Set Time.zone default to the specified zone and make Active Record auto-convert to this zone.
    # Run "rake -D time" for a list of tasks for finding time zone names. Default is UTC.
    # config.time_zone = 'Central Time (US & Canada)'

    # The default locale is :en and all translations from config/locales/*.rb,yml are auto loaded.
    # config.i18n.load_path += Dir[Rails.root.join('my', 'locales', '*.{rb,yml}').to_s]
    # config.i18n.default_locale = :de
    config.i18n.default_locale = 'pt-br'
  end
end
```

Restart the server, and check out your Brazilian Portuguese blog application. You've localized the application in two simple steps: adding the translation file and setting up the locale. Figure 12-1 shows the blog application with its Brazilian Portuguese face.

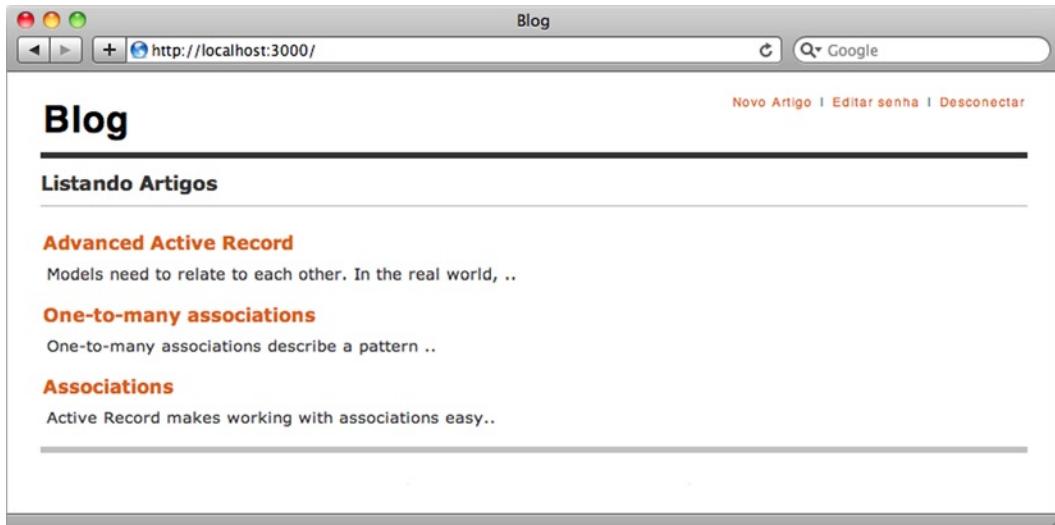


Figure 12-1. Brazilian Portuguese localized interface

Bilingual Blog

You know by now that all it takes to change the locale is to set the `I18n.locale` configuration to the locale of choice. How about giving users the power to do that themselves? To do so, you implement a controller filter that sets the locale depending on user input and provides the user with a language selector from which to choose the locale.

Let's create a helper in the application helper called `language_selector` that shows the available locales for the user to choose from. Listing 12-11 shows `application_helper.rb` with the new helper method in bold.

Listing 12-11. `language_selector` Helper Method in `app/helpers/application_helper.rb`:
<http://gist.github.com/360849>

```
module ApplicationHelper
  # Creates a submit button with the given name with a cancel link
  # Accepts two arguments: Form object and the cancel link name
  def submit_or_cancel(form, name=t('general.cancel'))
    form.submit + " #{t('general.or')}" + link_to(name, 'javascript:history.go(-1);', :class =>
'cancel')
  end

  def language_selector
    if I18n.locale == :en
      link_to "Pt", url_for(:locale => 'pt-br')
    else
      link_to "En", url_for(:locale => 'en')
    end
  end
end
```

In the language selector method, you show a link to the language that isn't currently selected. You do that by linking to the URL you're at, with an extra `:locale` parameter using the `url_for` helper with the `:overwrite_params` option. The user should always be able to change the language; so you use this function in the application layout. Listing 12-12 shows the updated application layout with the new helper call in bold.

Listing 12-12. Calling `language_selector` in `app/views/layouts/application.html.erb`:
<http://gist.github.com/360851>

```
<!DOCTYPE html>
<html>
<head>
  <title>Blog</title>
  <%= stylesheet_link_tag :all %>
  <%= javascript_include_tag 'jquery-1.4.2.min', 'rails', 'application' %>
  <%= csrf_meta_tag %>
</head>
<body>
  <div id="header">
    <h1><%= link_to t('general.title'), root_path %></h1>
    <%= language_selector %>
    <div id="user_bar">
      <% if logged_in? %>
        <%= link_to t('articles.new_article'), new_article_path %> |
        <%= link_to t('users.edit_password'), edit_user_path(current_user) %> |
        <%= link_to t('sessions.logout'), logout_path %>
    </div>
  </div>
</body>
```

```

<% else %>
  <%= link_to t('sessions.login'), login_path %>
<% end %>
</div>
</div>
<div id="main">
  <%= content_tag(:p, notice, :class => 'notice') if notice.present? %>
  <%= content_tag(:p, alert, :class => 'alert') if alert.present? %>
  <%= yield %>
</div>
<div id="footer">
  <%= t('general.footer') %> <a href="http://beginningrails.com">
Beginning Rails 4</a>
</div>
</body>
</html>

```

Finally, let's create a before action in the application controller that sets the locale to the passed parameter. Listing 12-13 shows the updated application controller with the new additions in bold.

Listing 12-13. Before Filter to Set the Locale in app/controllers/application_controller.rb:
<http://gist.github.com/360856>

```

class ApplicationController < ActionController::Base
  protect_from_forgery with: :exception
  before_action :set_locale

  protected
    # Set the locale from parameters
    def set_locale
      I18n.locale = params[:locale] unless params[:locale].blank?
    end

    # Returns the currently logged in user or nil if there isn't one
    def current_user
      return unless session[:user_id]
      @current_user ||= User.find_by_id(session[:user_id])
    end

    # Make current_user available in templates as a helper
    helper_method :current_user

    # Filter method to enforce a login requirement
    # Apply as a before_filter on any controller you want to protect
    def authenticate
      logged_in? || access_denied
    end

    # Predicate method to test for a logged in user
    def logged_in?
      current_user.is_a? User
    end

```

```
# Make logged_in? available in templates as a helper
helper_method :logged_in?

def access_denied
  redirect_to login_path,
:notice => t('application.access_denied') and return false
end
```

A before action in an application controller runs before any request reaches any controller action. You call the `set_locale` method, which checks whether the locale parameter is provided and assigns it to `I18n.locale`.

Change the application locale configuration back to English—so users can select their language of choice—by removing the `config.i18n.default_locale` line from `config/application.rb` (Listing 12-14).

Listing 12-14. Remove the Locale Configuration in `config/application.rb`: <https://gist.github.com/adamgamble/6190907>

```
require File.expand_path('../boot', __FILE__)

require 'rails/all'

# Require the gems listed in Gemfile, including any gems
# you've limited to :test, :development, or :production.
Bundler.require(:default, Rails.env)

module Blog
  class Application < Rails::Application
    # Settings in config/environments/* take precedence over those specified here.
    # Application configuration should go into files in config/initializers
    # -- all .rb files in that directory are automatically loaded.

    # Set Time.zone default to the specified zone and make Active Record auto-convert to this zone.
    # Run "rake -D time" for a list of tasks for finding time zone names. Default is UTC.
    # config.time_zone = 'Central Time (US & Canada)'

    # The default locale is :en and all translations from config/locales/*.rb,yml are auto loaded.
    # config.i18n.load_path += Dir[Rails.root.join('my', 'locales', '*.{rb,yml}').to_s]
    # config.i18n.default_locale = :de
  end
end
```

Restart your server and try the application, as shown in Figure 12-2.

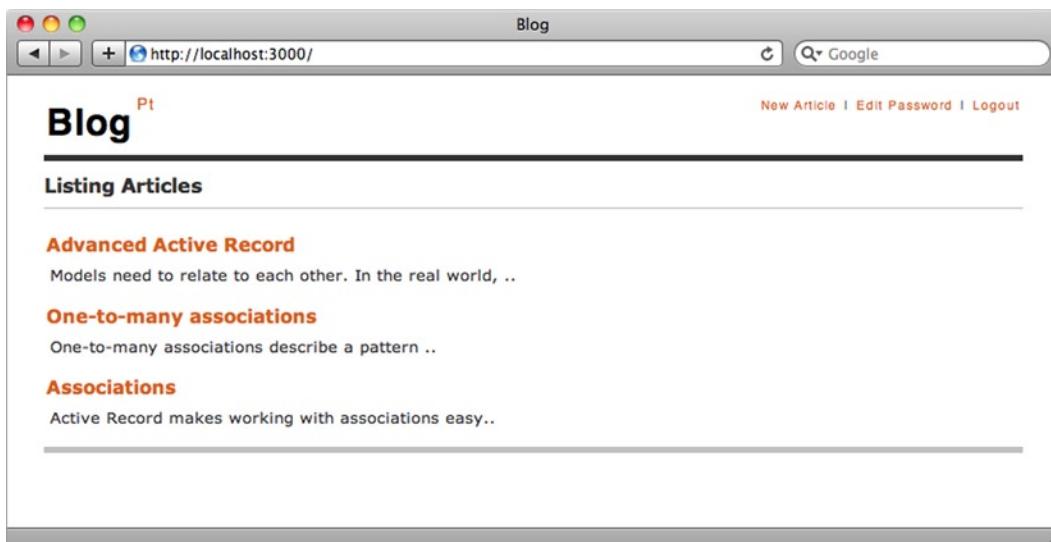


Figure 12-2. Language selector in the English interface

The application loads in English because it's the default locale. Click the Pt link, and see how everything switches to Brazilian Portuguese, as shown in Figure 12-3.

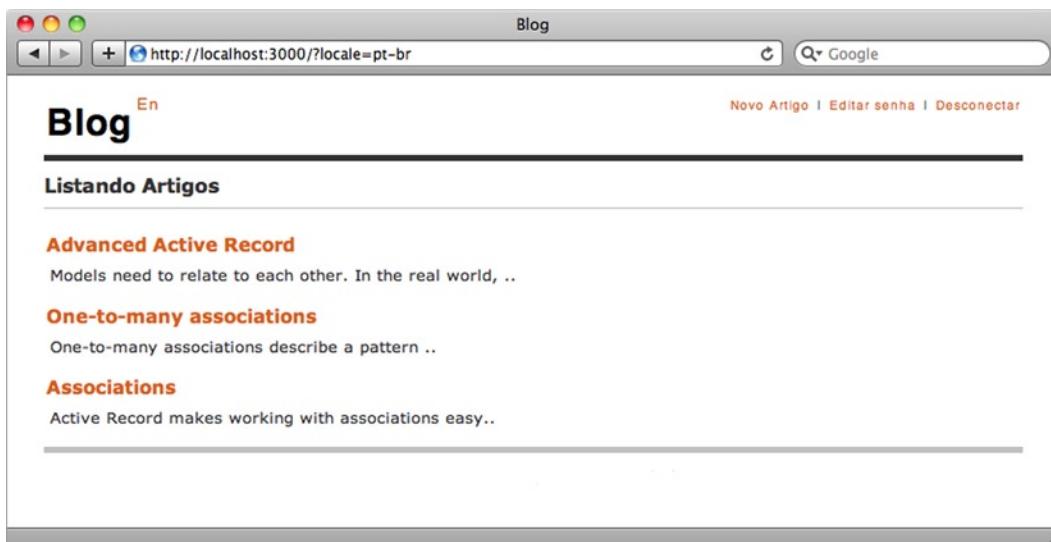


Figure 12-3. Brazilian Portuguese interface with the language selector link

Congratulations! Not only do you have a bilingual blog application, but you also know how easy it is to add more languages.

Summary

After reading this chapter, you have what it takes to launch a multilingual Rails application. You understand the logic behind the I18n module, how the `translate` method works, and how to create a translation file.

You worked on preparing your application for i18n support; you extracted the hard-coded text and strings into translation keys, and you placed them in a locale translation file. Then, you localized the application to another language. You wrapped up your effort by making the application bilingual so users can choose the language they want. The next, and final, chapter in the book will cover how to deploy your Rails apps, and let the world enjoy the fruits of your labors.



Deploying Your Rails Applications

If you're ready to turn the world on its head by unleashing your million-dollar web application to the public, then let this chapter be your starting point.

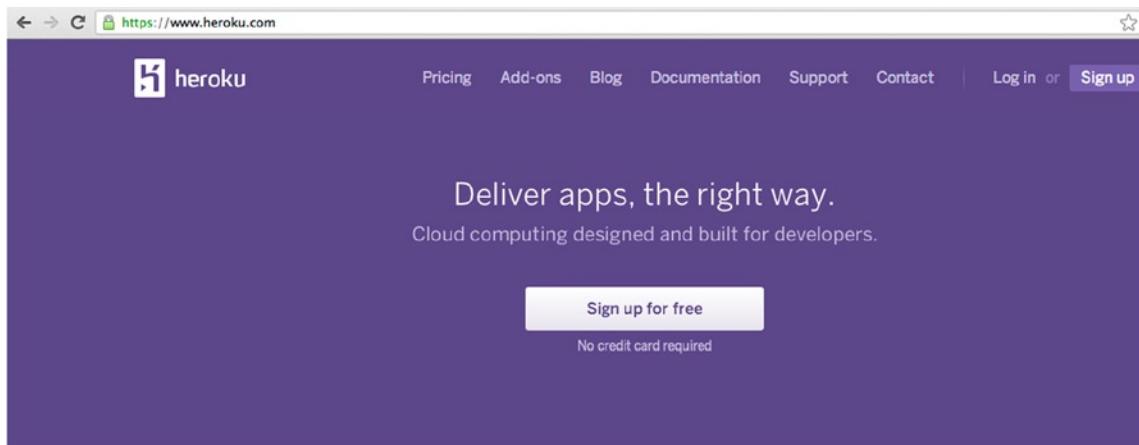
The various web-application development platforms in use today have very different deployment methods. If you're using PHP, deployment is usually as simple as dropping the right files into a directory on the remote server (usually via FTP) and then visiting them with your web browser. For PHP applications, there are thousands of hosts, and deployment is simple enough for your parents to accomplish. However, as you all know, with simplicity comes a lack of options. And as anyone who has developed in a language like PHP knows, simplicity can lead to some pretty complicated situations with your development. Luckily, Rails deployment is nothing like PHP deployment.

In the opposite camp, Java/Struts people have some deployment schemes that would make a PhD candidate panic and run away in a cloud of network diagrams. Java deployments can be terribly complex, with lots of little details to worry about and huge amounts of memory required. Fortunately, Rails is much simpler than that.

With Rails there are a lot of different ways to deploy your application. This chapter will focus on the one that will get your app out the door as simply and quickly as possible. Instead of configuring servers and building deploy scripts, you'll have your app running and live for the entire world to use in no time. You'll use a PAAS (platform as a service) called Heroku, which handles your servers and databases so you don't have to think about it. You can concentrate on building great applications. Although for a large production deployment Heroku can get pricey, for a developer with a small app it's completely free. Let's dive in.

Set Up an Account with Heroku

The first step in this process is to set up an account with Heroku, don't worry, it doesn't cost anything and you won't need a credit card. Point your browser to www.heroku.com and you should see something like Figure 13-1.



Build apps, not infrastructure.

Heroku provides you with all the tools you need to iterate quickly, and adopt the right technologies for your project. Build modern, maintainable apps and instantly extend them with functionality from hundreds of cloud services providers without worrying about infrastructure.

 "Heroku lets us focus on our app and deliver more value to customers."
— Aaron Peckham, *Urban Dictionary*



Figure 13-1. Setting up a Heroku account

Click the sign up link and enter an e-mail address. Heroku will then send you an e-mail with a confirmation link. Once you click that link, you'll be asked for a password and password confirmation. You're now the proud owner of a shiny new Heroku account. This will let you deploy Rails (and other) apps to your heart's content.

Heroku has a piece of software that facilitates interacting with your Heroku apps on your computer. It's called the Heroku Toolbelt (<https://toolbelt.heroku.com/>). Go to that URL and follow the instructions for installing the Heroku Toolbelt. For Windows and OS X, it will be a package you download; for Linux there is an install script.

Preparing Your Git Repository

Now that you have a Heroku account, you need to make a couple of small changes to your app so you can deploy it. Heroku's method of deployment is Git, a tool most developers are already using. If you are unfamiliar with Git, you can check out Appendix C to get up to speed. Let's set up a Git repository for your application now. Go to the terminal and make sure you're in the directory where your project is stored. Once there, type the following command:

```
$ git init
```

```
Initialized empty Git repository in /Users/adamgamble/rails/beginning_rails_4/.git/
```

You told Git that you want this directory to be a repository, meaning that Git will now keep track of the files you want it to. Let's tell it to keep track of all the files in this directory and make an initial commit:

```
$ git add .
$ git commit -m 'Initial Commit'
```

```
[master (root-commit) 7485c9c] Initial Commit
 121 files changed, 1980 insertions(+)
 create mode 100644 .gitignore
...
```

Your output might be slightly different, but it should add all files to your Git repository. This means that Git is now keeping track of the files and will notice when you make changes. You can then either decide to commit those changes or get rid of them. Once you have committed changes, Git has built-in support for pushing those changes to a remote server. Likewise, it can pull the changes from a remote server to your local repository. This is why so many developers use a source control system like Git; it makes it so easy to collaborate. It also happens to be the way you deploy your application to Heroku.

So let's tell Heroku we are ready to create an app. The first step is to create the app on Heroku. You can do this on their web control panel, or you can do it from the command line. We prefer to use the command line interface. You need to authenticate the Heroku Toolbelt you previously installed, but that's a simple task. Once that is done, you can move straight into creating your application on Heroku:

```
$ heroku login
```

```
Enter your Heroku credentials.
Email: adamgamble@gmail.com
Password (typing will be hidden):
Authentication successful.
```

```
$ heroku create
```

```
Creating floating-shore-4610... done, stack is cedar
http://floating-shore-4610.herokuapp.com/ | git@heroku.com:floating-shore-4610.git
Git remote heroku added
```

Notice that it called the app `floating-shore-4610`. If you don't specify a name, Heroku will choose a random name for you. If you would like to specify a name, type it after the Heroku create line, like so:

```
$ heroku create beginning-rails-4
```

```
Creating beginning-rails-4... done, stack is cedar
http://beginning-rails-4.herokuapp.com/ | git@heroku.com:beginning-rails-4.git
```

Since Heroku names must be unique, you obviously won't be able to use `beginning-rails-4`, but you can be creative and choose your own, like maybe `chunky-bacon`.

In your application's folder there is a file named Gemfile. This file stores a list of all the "gems" your project uses. Gems are little pieces of code that are easy to pull into your project to add features. Rails itself is a gem, and when you started this project a whole host of gems were pulled in. You need to edit one line and add two more to support deploying to Heroku. Open up your Gemfile and change the line that looks like this:

```
gem 'sqlite3'
```

to this:

```
gem 'sqlite3', group: [:development, :test]
```

Then below it, add these two lines:

```
gem 'rails_12factor', group: :production
gem 'pg', group: :production
```

Now save the Gemfile. Whenever you edit the Gemfile, you need to tell bundler (a gem manager) to make the changes you just requested in the Gemfile. To do that type:

```
bundle --without production
```

This tells bundler to install the gems requested but skip the ones that were in the production group. Bundler makes its changes and modifies a file called `Gemfile.lock`. You need to let Git know that you want to commit the changes to this file:

```
$ git add Gemfile
$ git add Gemfile.lock
$ git commit -m "Added gems to support heroku deployment"
```

```
[master a9afc76] added gems to support heroku deployment
2 files changed, 11 insertions(+), 1 deletion(-)
```

Now you're ready to deploy your app! It's as easy as one command:

```
$ git push heroku master
```

```
Counting objects: 138, done.
Delta compression using up to 8 threads.
Compressing objects: 100% (127/127), done.
```

```

Writing objects: 100% (138/138), 29.28 KiB, done.
Total 138 (delta 14), reused 0 (delta 0)

-----> Ruby/Rails app detected
-----> Using Ruby version: ruby-2.0.0
-----> Installing dependencies using Bundler version 1.3.2
      Running: bundle install --without development:test --path vendor/bundle
--binstubs vendor/bundle/bin --deployment
      Fetching gem metadata from https://rubygems.org/.....
      Fetching gem metadata from https://rubygems.org/..
      Installing rake (10.1.0)
      ...
      ...

-----> Compiled slug size: 33.9MB
-----> Launching... done, v5
http://floating-shore-4610.herokuapp.com deployed to Heroku

```

```
To git@heroku.com:floating-shore-4610.git
 * [new branch]      master -> master
```

There is a lot of output from the command, but at the end you will see:

<http://floating-shore-4610.herokuapp.com> deployed to Heroku

This means that the deployment worked. Now after running migrations on your app, you'll be able to see the app running on a real server!

```
$ heroku run rake db:migrate
```

```

Running 'rake db:migrate' attached to terminal... up, run.2873
==  CreateArticles: migrating =====
-- create_table(:articles)
  -> 0.0258s
==  CreateArticles: migrated (0.0358s) =====

```

Now point your web browser to the URL you picked (<the-name-of-your-app.herokuapp.com>, in my case its <http://floating-shore-4610.herokuapp.com/>) and you will see your app!

Notice that none of your articles or users made it into the deploy. That's because all of those data are stored locally in your database, while Heroku has its own copy on the database. You can create users manually in the console and then go from there filling your app with all the production data you need. To use the console on Heroku, it's as simple as typing:

```
$ heroku run console
```

You will then be presented with a console that looks just like the console you run locally.

That's it!

That's all there is to deploying your app with Heroku! Anyone with a web browser now has access to your application. This deployment is suitable for most small apps and even larger applications if you decide. Heroku allows you to purchase extra "dynos" or servers to scale your application to support heavier loads. You can do this by visiting Heroku's web console.

Deploying to Heroku is only one of the many different ways you can deploy your application. One of the more popular, but more difficult, solutions is called Capistrano (<https://github.com/capistrano/capistrano>). Capistrano deployment is out of the scope of this book, but you should be aware that it exists.

Whole books could be written on the topic of server set up and application deployment, but we tried to give you a quick way to let the masses use your great application with as little pain as possible.

Summary

This chapter showed you how to deploy your application quickly and easily using the Heroku platform as a service. It explained that you could develop web applications and not have to worry about your infrastructure. Although, for a full-scale enterprise deployment, it's probably best to call in the experts and build a solution from the ground up. Now that you know how to get your application in the hands of the public, you can go about the business of being the next billionaire startup, there isn't anything stopping you!



Databases 101

Let's begin with some simple definitions. A *database* is a piece of software that governs the storage, retrieval, deletion, and integrity of data. Databases are organized into *tables*. Tables have *columns* (or if you prefer, *fields*), and data are stored in *rows*. If you're familiar with spreadsheets, then the idea is fairly similar. Of course, databases blow spreadsheets out of the water in terms of power and performance.

Structured Query Language (SQL) is the standard way of communicating with databases. Using SQL, you can view column information, fetch a particular row or a set of rows, and search for rows containing certain criteria. You also use SQL to create, drop, and modify tables, as well as insert, update, and destroy the information stored in those tables. SQL is a fairly large topic, so a complete treatment is beyond the scope of this book. That said, you need to know the basics, so consider this a crash course.

Note The output in this appendix assumes you've followed the code in the book up to Chapter 5. If you read this appendix at a different point and implement the code, you may get different output.

Examining a Database Table

Here's an example definition for a table called `articles`. Note that the examples use SQLite, as they do throughout this book. If you're following along using a different piece of database software, the response you see may be slightly different or might not work at all. To start the SQLite utility tool, run the `rails dbconsole` command from the book's project folder on your computer:

```
sqlite> .tables
```

```
articles      schema_migrations
```

You use the SQLite `.tables` command to peek at the tables that exist in the database.

As you can see, the database has two tables: `articles` and `schema_migrations`. You don't get a lot of information about the tables from the `.tables` command, but that can be achieved by querying an internal SQLite table called `sqlite_master`:

```
sqlite> select sql from sqlite_master where name = 'articles';
```

```
CREATE TABLE "articles" ("id" INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
"title" varchar(255), "body" text, "published_at" datetime, "created_at" datetime,
"updated_at" datetime, "excerpt" varchar(255), "location" varchar(255))
```

The result of this command is an SQL statement that describes all the fields in the `articles` table. Each field has a *type*, which defines the kind of data it can store. The `id` field has a type of *integer*; `title` has a type of *varchar*, and `body` is a *text* field. Although it may sound strange, a type of *varchar* means the field has a variable number of characters up to a defined maximum. In this case, the maximum is 255 characters, which is the typical limit for *varchar* fields. (If you need to store more than 255 characters, use the *text* field type, like the `body` field.)

The `id` column is the one to pay attention to here. It's the *primary key*—a unique identifier for a particular row. Because this key is essential, it absolutely needs to be not null, and it must be unique; you let the database manage its value by automatically incrementing its number each time a new row is created. Notice how this is specified in the `articles` table column description: `PRIMARY KEY AUTOINCREMENT NOT NULL`. These are special commands that tell SQLite how to handle this particular field.

Let's look at some data from the `articles` table:

```
SELECT * FROM articles;
```

```
7|RailsConf|RailsConf is the official gathering for Rails developers..|2010-02-27..
8|RubyConf 2010|The annual RubyConf will take place in..|2010-05-19..
```

Here, you're using the SQL `SELECT` command to view this table's data. As you can see, you have two records in the table. You probably have different records in your database; the main point here is understanding the commands to see the data, not the data itself.

Working with Tables

The most common use of databases (not only in the context of Rails) is to implement CRUD functionality: create, read, update, and delete. Corresponding to the CRUD components are the most commonly used SQL commands: `INSERT`, `SELECT`, `UPDATE`, and `DELETE`, as shown in Table A-1.

Table A-1. Common SQL Commands

Operation	SQL Command
Create	<code>INSERT</code>
Read	<code>SELECT</code>
Update	<code>UPDATE</code>
Delete	<code>DELETE</code>

The following sections use the `articles` table presented in the previous section to show some examples of how these commands work. Remember that it's not necessary to have a complete understanding of SQL to work with Rails. The whole point of Active Record is to alleviate the tedium of needing to construct complex SQL statements to view and otherwise manipulate your data.

Selecting Data

The `SELECT` statement is a powerful and useful SQL command. Using `SELECT`, you can query (or request information from) the database and mine it for information. You can give `SELECT` any number of fields, a set of conditions to be applied to the data to be returned, a limit on the number of rows it returns, and instructions on how to order its results.

Earlier, you used the `SELECT` statement to see the data in the `articles` table:

```
SELECT * FROM articles;
```

The asterisk (*) character is a wildcard that means *every column*. This statement says, “Show me the values in every column for every row.” This is the easiest way to look at the contents of a table. But you don’t often need to see every single row; and for tables with a lot of data, you could end up with a really large list. So, sometimes it isn’t very efficient to select everything. Fortunately, you can also select specific columns by name. For example, to select only the `title` column, do this:

```
sqlite> SELECT title FROM articles;
```

```
RailsConf  
RubyConf 2010
```

Instead of returning all fields, this command returns only the one requested: `title`. To return both the `title` and the `body` fields, add `body` to the list of columns to select:

```
sqlite> SELECT title, body FROM articles;
```

```
RailsConf|RailsConf is the official gathering for Rails developers..  
RubyConf 2013|The annual RubyConf will take place in..
```

In both cases, the command returns all rows. If there were 100 rows in the table, they would all be returned.

But what if you need to find a particular row? This is where *conditions* come in to play. To supply conditions to a `SELECT` statement, you use the `WHERE` clause:

```
SELECT fields FROM table WHERE some_field = some_value;
```

Let’s apply this to the `articles` table by finding a row by its primary key:

```
sqlite> SELECT * FROM articles WHERE id = 7;
```

```
7|RailsConf|RailsConf is the official gathering for Rails developers..|2010-02-27  
00:00:00.000000|2013-04-02 22:45:03.777308|2013-04-02 22:45:03.777308||
```

This query returns only the row whose primary key, `id`, matches the condition. You can use this technique on any field—`id`, `title`, or `body`—or all of them combined. Conditions can be chained together using `AND` and further modified using `OR`. For example, the following query returns only records whose titles and authors match the specified criteria:

```
SELECT * FROM articles WHERE title = 'Beginning Rails' AND id = 2;
```

Inserting Data

To insert a row into a table, you use the `INSERT` command. `INSERT` requires a table name, a list of fields, and a list of values to insert into those fields. Here’s a basic `INSERT` statement for the `articles` table:

```
sqlite> INSERT INTO articles (title, body) VALUES ('Intro to SQL',  
'This is an introduction to Structured Query Language');
```

SQLite doesn't give any indication that something happened, which means your command was accepted and didn't generate any errors. To see what was inserted, you again use the SELECT command:

```
sqlite> SELECT * FROM articles;
```

```
7|RailsConf|RailsConf is the official gathering for Rails developers..|2013-04-27..
8|RubyConf 2013|The annual RubyConf will take place in..|2013-04-19..
9|Intro to SQL|This is an introduction to Structured Query Language|||||
```

You now have three rows in your table. Notice that in the INSERT statement, you don't specify the id field. That's because, as you recall, it's handled automatically by the database. If you were to insert a value, you wouldn't have a reliable way to guarantee its uniqueness and could cause an error if you attempted to insert a duplicate value. The database automatically inserts an id value into the field that's greater than the biggest existing id.

Updating Data

If you want to change the values in a row, you use the UPDATE statement. UPDATE is similar to INSERT, except that like SELECT, it can be modified (or constrained) by *conditions*. If you want to change the title for the "Intro to SQL" article, you can do so like this:

```
sqlite> UPDATE articles SET title = 'Introduction to SQL' WHERE id = 9;
```

Again, SQLite is silent, which means the command has been accepted. The fact that you use the primary key to find and update the row is significant. Although you can match any value in any column, the only surefire way to ensure you're updating the row you want is to use the primary key. You can confirm that the value was updated with another query:

```
sqlite> SELECT title FROM articles WHERE id = 9;
```

```
Introduction to SQL
```

Sure enough, the title field has been updated.

Deleting Data

Of course, not all information in a database will stay there forever. Sometimes you need to delete records, such as when a product goes out of stock or a user cancels their account. That's the purpose of the DELETE statement. It works a lot like the UPDATE statement, in that it accepts conditions and deletes the rows for any records that match the conditions. If you want to delete the article with the id of 9, the DELETE statement is as follows:

```
sqlite> DELETE FROM articles WHERE id = 9;
```

SQLite receives the command and deletes the record identified by the id you specified. And, of course, if you subsequently search for the record, you find that it no longer exists:

```
sqlite> SELECT * FROM articles WHERE id = 9;
```

Caution When you use either the UPDATE or DELETE command, you're making changes to existing data; so be careful to use a WHERE clause to limit the records you're updating or deleting. A good practice is to always run a SELECT command first to make sure your query returns the records you're expecting; then, later run the UPDATE or DELETE command.

Understanding Relationships

It's good practice to avoid duplication in your database by creating distinct tables to store certain kinds of information. You relate two tables to each other using an *association*. This makes more sense when you see it in action, so let's look at the articles table again. We added a column named author by running the ALTER TABLE articles ADD COLUMN author varchar(255); SQL command so the table contains more data:

```
sqlite> SELECT * FROM articles;
```

id	title	author
1	ActiveRecord Basics	Jeffrey Hardy
2	Advanced ActiveRecord	Cloves Carneiro Jr.
3	Setting up Subversion	Cloves Carneiro Jr.
4	Databases 101	Jeffrey Hardy

There's quite a bit of duplication in the author field. This can potentially create problems. Although you could search for all articles by a particular author using a standard SELECT query, what would happen if someone's name were misspelled? Any articles by the misspelled author wouldn't show up in the query. And if there were such a typo, you'd need to update a lot of records in order to fix it. Moreover, searching on a text field like author is both unreliable and rather slow when compared with searching using an integer type.

You can improve this design significantly by putting authors in their own table and referencing each author's unique id (primary key) in the articles table instead of the name. Let's do that now. Create a new table called authors, and change the author field in the articles table so it can store an integer instead of text. The new authors table looks like this:

```
sqlite> CREATE TABLE "authors" ("id" INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
"name" varchar(255));
sqlite> ALTER TABLE articles ADD COLUMN author_id integer;
```

Note how instead of a text field called author, you now have a numeric field called author_id that *references* the author's primary key from the authors table. This field holds a *foreign key*, which is a reference to the primary key of the table it relates to: in this case, the author who wrote the article. For each author in the articles table, you use INSERT to create a record in the authors table. You can then update the value of the author_id field using an UPDATE

statement for the dataset you’re working with. If you now look at the data from both tables, you’ll see that you’ve eliminated the duplication:

```
sqlite> SELECT id, author_id, title FROM articles;
```

```
1|1|ActiveRecord Basics  
2|2|Advanced ActiveRecord  
3|2|Setting up Git  
4|1|Databases 101
```

```
sqlite> SELECT * FROM authors;
```

```
1|Cloves Carneiro Jr.  
2|Jeffrey Hardy
```

You can now use this relationship in your SELECT queries by joining the two tables together using their association. In this association, the author_id in the articles table is equal to the id column in the authors table. Adding the JOIN directive requires only a slight change to the SQL:

```
sqlite> SELECT articles.id, title, name FROM articles  
JOIN authors ON articles.author_id = authors.id;
```

```
1|ActiveRecord Basics|Cloves Carneiro Jr.  
2|Advanced ActiveRecord|Jeffrey Hardy  
3|Setting up Git|Jeffrey Hardy  
4|Databases 101|Cloves Carneiro Jr.
```

Now you get the author names returned with the query, which effectively spans two tables. This is the crux of relational databases. Updating an author’s name is now easy because there is only one instance of a given author. Updating that author affects all of that author’s associated articles:

```
sqlite> UPDATE authors SET name = 'Packagethief' WHERE id = 2;
```

This changes the name of the author with the id of 2 to Packagethief. When you run the JOIN query again, you’ll see that all instances of the author’s name have been updated:

```
sqlite> SELECT articles.id, title, name FROM articles  
JOIN authors ON articles.author_id = authors.id;
```

```
1|ActiveRecord Basics|Cloves Carneiro Jr.  
2|Advanced ActiveRecord|Packagethief  
3|Setting up Git|Packagethief  
4|Databases 101|Cloves Carneiro Jr.
```

SQL and Active Record

This brings your database crash course to a close. This was by no means a complete reference, nor was it intended to be. Its purpose was to illustrate the basics of how databases work and to introduce you to their native language: SQL. Now that you have a taste, you can safely enter the world of Active Record, where most of this tedious work is handled for you.

Why did we bother showing you this if Active Record takes care of most of it for you? Because it's important to know what Active Record is doing behind the scenes. Although you can effectively use Active Record like a black box, you'll eventually need to debug your programs and figure out why something isn't working the way you expect. Having a basic understanding of SQL helps. Moreover, every bit of SQL that Active Record generates is logged by Rails. You can find the logs in the log/ directory of your application. Now, when you see these SQL commands in the logs, you'll have a good idea what they mean.



The Rails Community

Rails development is driven by a vibrant and passionate community of open source developers. The Rails community encourages its members to participate actively in Rails development. You can start by asking questions and discussing new features. As your knowledge increases, you can help others by writing about your own experiences in a personal blog, answering questions on the mailing list, contributing to the wiki, and fixing bugs and writing patches to make Rails even better. Whatever your intention, rest assured that participating in the community will help you get the most out of Rails.

Beginning Rails 4 Channels

As a companion to this book, we're opening a few communication channels with you. Feel free to contact us and other readers using the medium you feel most comfortable with:

- *Web site:* A resource to update you about changes in Rails 4 and later versions that may affect some of the code in the book and to direct you to more information in the Ruby and Rails worlds. You can check our web site at <http://beginningrails.com>.
- *Mailing list:* A mailing list for those interested in exchanging ideas or asking questions to the authors or other Rails developers. You can discuss changes in the Rails framework; or, if you find a bug in the framework, you can discuss proposed solutions. You can subscribe to this list at <http://groups.google.com/group/beginning-rails>.
- *Twitter:* The Twitter account is used to notify users about changes in Rails 4 and allows direct conversation between the authors and you. You can follow us at <http://twitter.com/beginningrails>.

Rails Mailing Lists

You can subscribe to several Rails-related mailing lists:

- *Talk mailing list:* A high-volume list where users can seek help, announce open source or commercial Rails projects, and discuss any miscellaneous matters about the Rails framework. You can subscribe to this list at <http://groups.google.com/group/rubyonrails-talk>.
- *Core mailing list:* A low-volume list for those interested in Rails development. This list is for developers interested in discussing changes in the Rails framework itself. You can expect to find technical threads about Active Record, Action Mailer, Action Pack, and Rails performance. You can subscribe to this list at <http://groups.google.com/group/rubyonrails-core>.

- *Security mailing list:* A list for those who want to keep abreast of Rails security concerns. You can subscribe to this read-only mailing list at <http://groups.google.com/group/rubyonrails-security>.

Rails IRC Channel

If you want to interact with other Rails developers live, you can try the Rails IRC channel. Open your favorite IRC client and connect to the Freenode IRC network at <irc.freenode.net>. Enter the #rubyonrails channel, and you'll find hundreds of Rails developers at any time of the day (or night) willing to help you and chat about their favorite web framework. If you want to be seen as a caring community participant, make sure you follow some basic etiquette on how to ask appropriate questions; look at www.slash7.com/pages/vampires for some guidelines.

■ **Note** Internet Relay Chat (IRC) is a type of real-time Internet chat, where users talk about their interests in topic-specific areas called *channels*. All you need to connect to IRC is IRC client software. The most commonly used IRC clients are the shareware mIRC (<http://mirc.com>) for Windows and the open source Colloquy (<http://colloquy.info>) for Mac.

Rails Blogs and Podcasts

The number of blogs dedicated to Rails information is rapidly growing and most of the new Rails features are covered in blogs or podcasts even before they're released to the public. You can subscribe to the blogs of your choice to keep up with news in the Rails world.

The following are some of the more rewarding Rails-related blogs you can visit, including the official Rails podcast:

- <http://weblog.rubyonrails.org>: The official Rails blog. You'll find information about upcoming releases, new functionality in Rails, and news that's considered important (such as documentation updates and Rails adoption worldwide).
- <http://rubylearn.com>: A Ruby community site where people post interesting and new things about Ruby or Rails.
- <http://railscasts.com>: A great web site by Ryan Bates with a series of informative screencasts that teach a wide range of Rails tricks, covering all aspects of the Rails framework.
- <http://rubylearn.enylabs.com>: The Ruby5 podcast. In five minutes or less—or a bit more—you can listen to Ruby and Rails news.

Rails Guides

The Rails community has started a documentation effort called Rails Guides, which you can find at <http://guides.rubyonrails.org>. It's a great effort to document various parts of the frameworks, from the basic beginner-oriented documentation to more advanced material.

Rails APIs

It's close to impossible to remember the names, methods, and possible parameters of all the functions and classes in Ruby and Rails. To help you with your coding tasks, we recommend that you keep the Ruby and Rails application programming interface (API) documentation open, or at least that you put them in your favorites. The API documentation contains all the information about specific functions you're trying to use, including the function source code.

You can find the Rails API documentation at <http://api.rubyonrails.org>. The Ruby API is at www.ruby-doc.org/core. For more user-friendly and searchable API documentation, go to <http://apidock.com/rails>.

Rails Source and Issue Tracking

You can find the Rails source code at <http://github.com/rails/rails>. It's powered by the GitHub application, a hosting service for projects using the Git revision-control system. GitHub allows you to download the Rails source code using a web interface. You can subscribe to the Git change log using RSS to be notified about changes to the Rails source code.

You can also participate in the development of Rails by submitting bug reports and patches to the GitHub account at <http://www.github.com/rails/rails> or by looking at the existing tickets and trying to fix them.



Git

Developers normally work in teams. You write plenty of code; sometimes you test some and decide to delete it, and other times you decide to stick to it. Managing this can be a painful process, which is why you can use Source Control Management (SCM) software: to help you focus on what you do best—writing beautiful code. Git is rapidly becoming the preferred SCM of developers everywhere.

What Is Source Control Management?

SCM software helps you keep track of code changes and gives you the ability to easily collaborate on that code with your teammates. The two main features of any SCM are:

- *Versioning*: When you’re using SCM for your project, files and directories in the project are tracked. Every time you make changes to your files, you can save those changes as a new version. Your project then has several versions—one for every change set—giving you the ability to browse those changes and revert to any one at any time.
- *File merging*: Let’s say you worked on a file and your colleague, John, worked on that same file, and you both committed (submitted) your files to the SCM system. Both files are merged by SCM of your involvement; if SCM can’t handle the merge for any reason, it lets you know and gives you some useful information about how to manually merge conflicting changes yourself.

How Does It Work?

Generally, when you add your code base to an SCM system, a *repository* is created, which is the store of all the versions of your code base. Then, you can take a copy of that repository and work on it; this is normally called your *working copy*. You can add files, change or delete some, and then *commit* those changes and send them back to the repository as a new revision. If your colleague John is working with you on the same code base, he can check out or pull those changes from the repository to update his working copy, letting the SCM take care of any necessary file merging (Figure C-1).

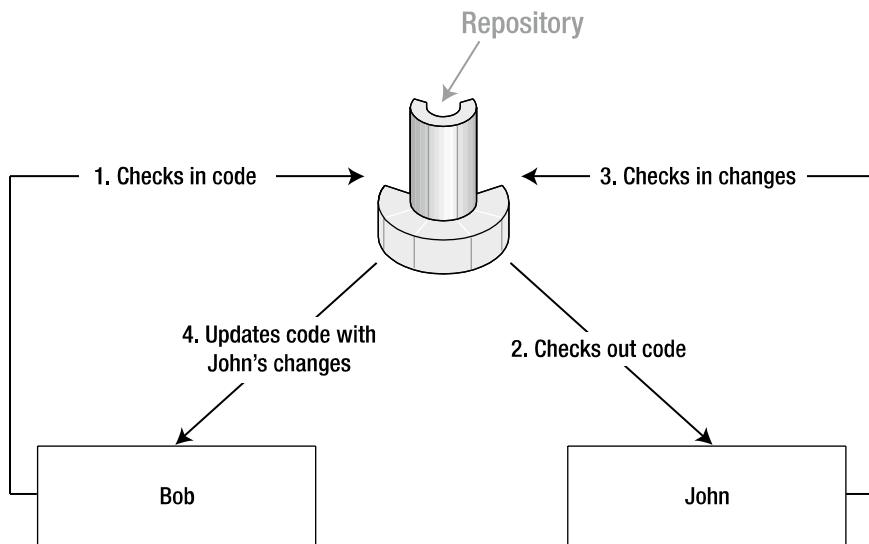


Figure C-1. SCM workflow

Git

The Git SCM was developed by Linus Torvalds for managing the Linux Kernel source code. It's also been used for several million open source projects, including Rails.

Git is different from other SCMs because it's a *distributed* source control system. This means that instead of having a single repository on your server that all your teammates use to check out working copies (*client-server* or *centralized SCM*), each team member has their own repository along with a working copy, and you all *push* a copy of that repository to a remote repository.

This approach has some great benefits, such as the ability to work and commit your code, even if you're offline, and being able to operate on your repository more quickly.

Now that you have a good understanding of what an SCM is and how it works, let's install Git and try it.

Installing Git

Installing Git is relatively easy. Thanks to open source contributions, several Git installation packages are available to facilitate a quick installation for most platforms.

Installing on Windows

If you're on Windows, you can install Git on your system using Git on Windows, also known as msysGit, which is an open source project available at <https://code.google.com/p/msysgit/downloads/list?q=full+installer+official+git>.

Download the Git installation executable from the download section of the project web site. The current version at the time of this writing is Git 1.8.1.2, and it's available to download at <https://msysgit.googlecode.com/files/Git-1.8.1.2-preview20130201.exe>.

Start the installation following the default options. You end up with Git installed on your system. The package also adds some interesting tools: Git GUI, which gives you a nice graphical interface to use Git; Git Bash, which you use to execute Git commands; and three context menu items integrated with Windows Explorer—Git Init Here, Git Bash, and Git GUI—to start the tools from the folder you’re calling them from.

Use the Git Bash tool to execute the commands mentioned in this appendix.

Installing on Mac OS X

To install Git on Mac OS X Leopard, you can either compile it from source using Macports or use Git OS X Installer, which you can find at <http://code.google.com/p/git-osx-installer/>.

Download the Git OS X Installer from the project’s web site. The current version at the time of this writing is Git 1.8.2.1, available to download from <https://git-osx-installer.googlecode.com/files/git-1.8.2.1-intel-universal-snow-leopard.dmg>.

The file you download is a disk image; it has a package installer named `git-1.8.2.1-intel-leopard.pkg`. Run the package installer, selecting the default options. When the installation is complete, Git is available from the terminal.

Installing on Linux

Most Linux distributions ship with a package manager. The most common one is the Debian package manager `apt`, and Git is part of its library.

To install Git using `apt`, run the following `apt-get` command from the terminal:

```
sudo apt-get install git
```

Accept if the package manager asks your permission to use additional disk space for this installation. When the installation is complete, Git is ready to use.

Setting Global Parameters

Every commit you make in your repository has flags for the user who did the commit; those flags are the user’s name and e-mail address. Now that you have Git installed on your system, it’s important to set a global username and e-mail address for Git to use for any new repository you work on.

To set global parameters, you use the `git config` command with the `--global` option, followed by the parameters you want to set. Listing C-1 shows the command to set up both the `user.name` and `user.email` parameters.

Listing C-1. Setting the Global Git Username and E-Mail

```
git config --global user.name "dude"
git config --global user.email my.email@address.com
```

These parameters can be set on a repository level as well; you can do that by using the same commands but without the `--global` option in your repository’s working directory.

Initializing a Repository

The first thing you want to do is to start a new repository for your application. Begin by creating a test application:

```
$ rails testapp
```

```
create
create README
create .gitignore
create Rakefile
.
.
.
create tmp/cache
create tmp/pids
create vendor/plugins
create vendor/plugins/.gitkeep
```

Now, initialize a local repository for that application by calling `git init` in the application directory:

```
$ cd testapp
$ git init
```

```
Initialized empty Git repository in /tmp/testapp/.git/
```

The `git init` command initializes an empty local repository for the application, but it doesn't add any files to the repository. To determine which files you can add to the repository, you call the `git status` command:

```
$ git status
```

```
# On branch master
#
# Initial commit
#
# Untracked files:
#   (use "git add <file>..." to include in what will be committed)
#
#       .gitignore
#       Gemfile
#       README
#       Rakefile
#       app/
#       config.ru
#       config/
#       db/
#       doc/
#       lib/
```

```
#    public/
#    script/
#    test/
#    vendor/
nothing added to commit but untracked files present (use "git add" to track)
```

As you can see, all the folders and files of the Rails application are in the untracked files list, which means they're still untracked. To start tracking those files, you need to *add* them to the track list; as this indicates, you can do this by using the `git add` command.

Ignoring Files

Before you add those files, let's think a little: Do you want all your files to be tracked? Are there any files you don't want to track? Normally, those would be configuration files that contain passwords, such as `database.yml`, the `tmp` folder, log files, and SQLite databases. If you add those files, your teammates will have this information, and it may even conflict with their files.

To skip those files in any `git add all` and `git status` commands, and to tell Git to never bother you about them again, you must configure Git to ignore them. You do that by declaring those files in a hidden configuration file called `.gitignore`, which is normally stored at the root of your working copy (in this case, at the root of the `testapp` directory). The `.gitignore` file is a regular text file; it is generated by Rails in all new projects. Edit it using your text editor of choice so it looks like the code in Listing C-2.

Listing C-2. The `.gitignore` File Content in `testapp/.gitignore`: <http://gist.github.com/287051>

```
.bundle
config/database.yml
log/*.log
db/*.sqlite3
tmp/**/*
```

As you can see, the files and folders listed in the `.gitignore` file weren't listed in the `git status` command you issued earlier.

Adding and Committing

You can add the untracked files to your repository by using the `git add` command and passing a dot to it, which refers to the current directory and all its content. Be cautious when using `git add .` to make sure you aren't adding files that don't belong in the repository.

```
$ git add .
```

Try the `git status` command again:

```
$ git status

---

# On branch master
#
# Initial commit
#
# Changes to be committed:
#   (use "git rm --cached <file>..." to unstage)
#
#       new file:   .gitignore
#       new file:   Gemfile
#       new file:   README
#       new file:   Rakefile
#
#       .
#       .
#       .
#       new file:   script/rails
#       new file:   test/performance/browsing_test.rb
#       new file:   test/test_helper.rb
#       new file:   vendor/plugins/.gitkeep
#

---


```

The `git status` command still shows all the files, because the `git add` command just added those files to be committed, but they aren't committed yet.

In order to commit the changes you added to the commit list, you have to call the `git commit` command. Use the `-m` argument to include a message describing the purpose of and the changes in this commit:

```
$ git commit -m "Empty Rails application"

---


```

```
[master (root-commit) 046116c] Empty Rails application
 41 files changed, 8434 insertions(+), 0 deletions(-)
  create mode 100644 .gitignore
  create mode 100644 Gemfile
  create mode 100644 README
  .
  .
  .
  create mode 100755 script/rails
  create mode 100644 test/performance/browsing_test.rb
  create mode 100644 test/test_helper.rb
  create mode 100644 vendor/plugins/.gitkeep

---


```

Congratulations—you've completed your first commit to your local repository! If you check the `git status` command now, you'll see that there are no changes to be added or committed:

```
$ git status
```

```
# On branch master
nothing to commit (working directory clean)
```

Starting from this point, if you edit any file in your working copy, rename it, move it, or delete it, it will show in your `git status` output, allowing you to pick which of those changes you would like to add to your commit list and then commit them to your repository.

Branching and Merging

Let's say you decide to work on a new feature for your project that you know will take you some time to finish. Meanwhile, you want to be able to keep working on the main project without changes from the new feature breaking your application. To implement this new feature in a separate copy of your project, you need to create a *branch*. A branch is a duplicate of your project that you can work on in parallel with the master copy of the same project.

When you called the `git init` command earlier, Git initialized a new repository for your application with a default branch called `master`. To create a new branch in the repository, use the `git checkout -b` command followed by the name of the new branch you want to create:

```
$ git checkout -b articles
```

This command creates a new branch named `articles` as a duplicate of the current branch `master`. To see a list of the branches in your project, you use the `git branch` command but without any parameters:

```
$ git branch
```

```
articles
* master
```

You have two branches, `articles` and `master`. The asterisk in front of `master` indicates that it's the *current* branch you're working on. To switch branches, use the `git checkout` command followed by the name of the branch you want to switch to:

```
$ git checkout articles
```

```
Switched to branch 'articles'
```

The `articles` branch is the current branch now. You can confirm this by listing the branches again:

```
$ git branch
```

```
* articles
  master
```

Now, let's implement a new feature—an articles scaffold:

```
$ rails generate scaffold Article title:string body:text
```

```
invoke  active_record
create  db/migrate/20130420235045_create_articles.rb
create  app/models/article.rb
invoke  test_unit
create  test/unit/article_test.rb
create  test/fixtures/articles.yml
route   resources :articles
invoke  scaffold_controller
create  app/controllers/articles_controller.rb
invoke  erb
create  app/views/articles
create  app/views/articles/index.html.erb
create  app/views/articles/edit.html.erb
create  app/views/articles/show.html.erb
create  app/views/articles/new.html.erb
create  app/views/articles/_form.html.erb
invoke  test_unit
create  test/functional/articles_controller_test.rb
invoke  helper
create  app/helpers/articles_helper.rb
invoke  test_unit
create  test/unit/helpers/articles_helper_test.rb
invoke  stylesheets
create  public/stylesheets/scaffold.css
```

You're done with the new feature changes. It's time to add the changes and commit them to the articles branch:

```
$ git add .
$ git commit -m "Adding Article scaffold"
```

```
[articles 33a538e] Adding Article scaffold
 15 files changed, 308 insertions(+), 0 deletions(-)
create mode 100644 app/controllers/articles_controller.rb
create mode 100644 app/helpers/articles_helper.rb
create mode 100644 app/models/article.rb
create mode 100644 app/views/articles/_form.html.erb
create mode 100644 app/views/articles/edit.html.erb
create mode 100644 app/views/articles/index.html.erb
create mode 100644 app/views/articles/new.html.erb
create mode 100644 app/views/articles/show.html.erb
create mode 100644 db/migrate/20130420235045_create_articles.rb
create mode 100644 public/stylesheets/scaffold.css
create mode 100644 test/fixtures/articles.yml
create mode 100644 test/functional/articles_controller_test.rb
create mode 100644 test/unit/article_test.rb
create mode 100644 test/unit/helpers/articles_helper_test.rb
```

When you check the `git status` command now, you see that you have nothing to commit in the `articles` branch:

```
$ git status
```

```
# On branch articles
nothing to commit (working directory clean)
```

The `articles` branch now has an article scaffold, and the `master` branch doesn't. If you switch back to the `master` branch, notice that none of the article scaffold files exist there:

```
$ git checkout master
```

```
Switched to branch 'master'
```

You can modify the project in the `master` branch completely in isolation from the `articles` branch if you want to, but for now let's *merge* the `articles` branch into the `master` branch. You do that using the `git merge` command followed by the branch name you want to merge into the current branch:

```
$ git merge articles
```

```
Updating 88c63c5..33a538e
Fast forward
app/controllers/articles_controller.rb      | 83 ++++++=====
app/helpers/articles_helper.rb              |  2 +
app/models/article.rb                     |  2 +
app/views/articles/_form.html.erb          | 24 ++++++
app/views/articles/edit.html.erb           |  6 ++
app/views/articles/index.html.erb          | 25 ++++++
app/views/articles/new.html.erb            |  5 ++
app/views/articles/show.html.erb           | 15 +////
config/routes.rb                           |  2 +
db/migrate/20130420235045_create_articles.rb | 14 +///
public/stylesheets/scaffold.css           | 60 ++++++=====
test/fixtures/articles.yml                 |  9 +++
test/functional/articles_controller_test.rb | 49 ++++++=====
test/unit/article_test.rb                  |  8 +++
test/unit/helpers/articles_helper_test.rb  |  4 +
15 files changed, 308 insertions(+), 0 deletions(-)
create mode 100644 app/controllers/articles_controller.rb
create mode 100644 app/helpers/articles_helper.rb
create mode 100644 app/models/article.rb
create mode 100644 app/views/articles/_form.html.erb
create mode 100644 app/views/articles/edit.html.erb
create mode 100644 app/views/articles/index.html.erb
create mode 100644 app/views/articles/new.html.erb
create mode 100644 app/views/articles/show.html.erb
create mode 100644 db/migrate/20100420235045_create_articles.rb
```

```
create mode 100644 public/stylesheets/scaffold.css
create mode 100644 test/fixtures/articles.yml
create mode 100644 test/functional/articles_controller_test.rb
create mode 100644 test/unit/article_test.rb
create mode 100644 test/unit/helpers/articles_helper_test.rb
```

The task is complete: you have developed a new feature in a separate branch without affecting the `master` branch; and when you finished, you merged those changes back into `master`.

Remote Repositories and Cloning

As stated previously, Git is a distributed SCM; therefore, your repository is hosted locally on your machine, hidden inside your working copy directory. No one else has access to it.

If you want to set up a repository that you and your team can work on, you first have to create a *remote* repository that all of you can access and clone from. Your remote repository can be hosted on any machine that is available to all developers who need access to the repository and have Git installed. It can be hosted on your local network, online, or with a third-party Git hosting provider like the famous GitHub (<http://github.com>), which hosts Rails, as well as many, many other projects.

We used Git for this book's blog application, and we hosted the repository on GitHub. It's publicly available for you at <http://github.com/adamgamble/blog/>. This means you can clone a copy of the blog repository to your machine and browse the code locally. To do that, you need the Public Clone URL, which is <git://github.com/adamgamble/blog.git>. Let's clone the blog application repository using the `git clone` command:

```
$git clone git://github.com/adamgamble/blog.git
```

```
Initialized empty Git repository in /tmp/blog/.git/
remote: Counting objects: 1085, done.
remote: Compressing objects: 100% (575/575), done.
remote: Total 1085 (delta 539), reused 898 (delta 436)
Receiving objects: 100% (1085/1085), 222.28 KiB | 362 KiB/s, done.
Resolving deltas: 100% (539/539), done.
```

Now you have a local copy of the blog application repository cloned to your machine. You can change files and even commit them to your own local repository, but what you *cannot* do is share those commits with others. In order to *push* your changes, you need write access to the remote repository, which you don't have.

If you want to try that, sign up for a free account on GitHub and create a repository of your own there. You then have two URLs: a *public* URL that everyone can see, and your *clone* URL, which gives you full access to this remote repository.

The concept is simple: after you clone your own repository using your own URL, you can work normally in your working copy, commit changes, and add and remove files. Whenever you want to share those commits with the rest of the world, you push them to the remote repository on GitHub using the `git push` command. If you have teammates pushing changes to the same repository, you can retrieve those changes by using the `git pull` command.

To sum up, you create a remote repository to allow more than one developer to work on the same repository. Although all developers on the team have their own copies, they still need to push their copies to the remote repository to allow the rest to pull from it and stay in sync.

When you sign up for a free account on GitHub, the repositories you create are publicly available for everyone to clone from. If you want your repositories to be private, so only you and your teammates can access them, you can either upgrade your account with GitHub or host them on your own server with your own setup.

Learning More

Git is a great tool and has a lot of commands; however, this appendix has covered only the basic features and commands. We highly encourage you to read more. You can see a list of the most-used Git commands using the `git help` command:

```
$ git help
```

```
usage: git [--version] [--exec-path[=GIT_EXEC_PATH]] [--html-path]
[-p|--paginate|--no-pager] [--bare] [--git-dir=GIT_DIR]
[--work-tree=GIT_WORK_TREE] [--help] COMMAND [ARGS]
```

The most commonly used git commands are:

add	Add file contents to the index
bisect	Find by binary search the change that introduced a bug
branch	List, create, or delete branches
checkout	Checkout a branch or paths to the working tree
clone	Clone a repository into a new directory
commit	Record changes to the repository
diff	Show changes between commits, commit and working tree, etc
fetch	Download objects and refs from another repository
grep	Print lines matching a pattern
init	Create an empty git repository or reinitialize an existing one
log	Show commit logs
merge	Join two or more development histories together
mv	Move or rename a file, a directory, or a symlink
pull	Fetch from and merge with another repository or a local branch
push	Update remote refs along with associated objects
rebase	Forward-port local commits to the updated upstream head
reset	Reset current HEAD to the specified state
rm	Remove files from the working tree and from the index
show	Show various types of objects
status	Show the working tree status
tag	Create, list, delete, or verify a tag object signed with GPG

See '`git help COMMAND`' for more information on a specific command.

To learn more about a specific command, you can use `git help COMMAND`, which shows that command's documentation and how to use the command.

Other SCM Systems

Although Git is the most talked-about SCM nowadays, you may either be required to use a different SCM or may want to investigate the alternatives. Here's a list of other SCMs you could choose:

- *Mercurial*: Just like Git, Mercurial is a distributed SCM. Mercurial is often compared with Git because of their similarities; feel free to try it if you want to explore another option. You can find out more about Mercurial from its official web site: <http://mercurial.selenic.com/>.
- *SVN (Subversion)*: This was the most prominent SCM for a while, but it has since been overtaken by Git. You can find out more about Subversion from its official web site: <http://subversion.tigris.org/>.
- *CVS (Concurrent Versions System)*: This was one of the earliest SCM systems (initial release in 1990). It's still popular, although because of some limitations, such as sparse Unicode support and expensive branching operations, developers have begun moving toward other version control systems like Subversion and Git. You can find out more about CVS from its official web site: www.nongnu.org/cvs/.

Online Resources

After the beta launch of GitHub, Git received huge interest from developers, including the Rails core team—they decided to switch from Subversion to Git and host the official Rails repository on GitHub. This attention to Git encouraged more developers to try it, and a number of tutorials and blog posts began to appear in the community.

The following are some resources you can visit to dig deeper and learn more about Git:

- <http://book.git-scm.com/>: The Git Community Book. As the name implies, this book is written by the Git community to the Git community. It's a cumulative resource to help newcomers to Git get started and quickly find what they're looking for.
- <http://peepcode.com/products/git>: A 60-minute screencast by PeepCode, available with an in-depth PDF explanation on how Git works under the hood.

Index

Symbols

`$()` function, 193

A

Action Controller, 10

Action Pack, 151

controller generator

 command, 152

 empty stub, 152

 form partial, 153–154

 new and edit templates, 153

 new user, 155

 syntax, 151

 user_params method, 155

controllers and templates

 apply filters, 172

 article form, 167

 articles index page, 166

 authentication, 170

 controller callbacks, 170

finishing touches

 action view helpers, 174

 article, 178

 body field, 176

 custom helpers, 180

 edit controls, 177

 escape HTML, 175

 layout style, 182

 style sheet, 183

nested resources

 application layout, 158

 article partial, 158

 article show page, 160

 comment-named routes, 156

 comments, 155

 controller, 157

 delete link, 160

generated named routes, 156

 updated article, 157

sessions and login/logout logic

 logging in, 163

 log out, 164

 resource, 161

 session, 161

 stateless, 161

Action Pack. *See* View and controller

Action View, 10

Active Record, 3, 10, 61, 81

 callbacks

 after_create method, 114

 before_save method, 114

 email_article_author

 methods, 115

 hashing algorithms, 117

 models, 114

 user model, 115

 conventions, 63

 CRUD

 all records, 73

 create method, 70

 creating new records, 67

 deleting with conditions, 78

 find method, 71

 instance, 76

 modeling process, 67

 new constructor, 68

 single record using first, 72

 single record using ID, 71

 updating records, 75

 using delete, 77

 using destroy, 76

 with conditions, 75

 errors collection, 79–80

 finding

 array condition syntax, 102

 association proxies, 104

Active Record (*cont.*)
 default scope, 105
 finder methods, 104
 log file, 103
 named scope, 106
 SQL fragment, 102
 where method, 101
 fresh Article object, 79
 full_messages method, 79
 methods
 basics, 81
 command, 82
 domain logic, 82
 encapsulate, 82
 fat models, 83
 long_title methods, 82
 ORM, 61
 recalling specific validations, 78
 size, 80
 SQL, 62
 updated models, 119
 validations

- account model, 108
- custom validation methods, 112
- validations (*see* Built-in validations)

Active Record. *See* Associations

Ajax, 193
 Apple Developer Tools (Xcode), 14
 Application programming
 interface (API), 283

Associations
 articles and comments tables, 84
 declaration, 84
 foreign key reference, 83
 foreign-key references, 84
 many-to-many relationship
 articles and categories, 96
 seeding data, 98
 new associated object, 92
 one-to-many associations, 90
 one-to-one associations, 85
 options
 default order, 94
 dependencies, 95
 primary key, 83
 relationship, 83
 rich many-to-many
 associations, 99
 user and profile models
 command, 86
 common has_one options, 90
 console session, 88
 create_profile method, 89
 has_one and belongs_to
 relationship, 89

migration, 86
 one-to-one association, 87
 user_id attribute, 88
 users and article models, 90

■ B

Blog application, 242
 Built-in validations
 confirmation, 111
 default options, 109
 format options, 111
 length/size, 110
 method, 112
 unique, 110
 value has been entered, 109

■ C

Concurrent Versions System
 (CVS), 296
 Console
 advantage, 64
 Article model, 64
 column_names class
 method, 65
 irb, 64
 Ruby code, 64
 subclassing and superclass, 65
 Controller walk-through

- action controller, 135
- article form, 138
- articles controller, 144
- blog application, 130
- edit and update actions, 146
- error messages, 145
- form helpers, 140
- GET, 134
- layouts, 136
- named routes, 129
- partials
 collection, 149
 local variables, 149
 object partials, 149
 staying DRY, 148
- redirects, 134–135
- rendering responses, 134
- request parameters, 143
- RESTful routes and
 resources, 129
- routing basics, 128
- scaffold generator, 131
- templates, 135
- validation errors, 142
- views, 146

Crash course, RUBY class, 65
 Create, read, update, and delete (CRUD)
 creating new records, 67
 deleting with conditions, 78
 find method, 71
 instance, 76
 modeling process, 67
 new constructor, 68
 reading records
 all records, 73
 single record using
 first, 72
 single record using ID, 71
 with conditions, 75
 updating records, 75
 using delete, 77
 using destroy, 76

D

Databases 101, 273
 Active Record, 279
 database, 273
 SQL, 273
 SQL commands, 274
 asterisk (*) character, 275
 crux relational databases, 278
 DELETE statement, 276
 foreign key, 277
 INSERT statement, 275
 JOIN query, 278
 primary key, 276
 SELECT statement, 274, 277
 UPDATE statement, 276
 WHERE clause, 275
 table creation, 273
 primary key, 274
 SELECT command, 274
 type, 274

Debian-based variant. *See* Ubuntu Linux
 Deploying rails applications, 267

Capistrano deployment, 272

Git Repository, 269

 database, 271

 Gemfile.lock, 270

 gems, 270

 Heroku creation, 269

 initial commit, 269

 output, 271

 remote server, 269

Heroku account, 267

 Toolbelt, 268

PAAS, 267

Document Object Model (DOM), 193

Don't repeat yourself (DRY), 6

E

E-mail
 Action Mailer, 203
 application settings, 205
 mail server, 203
 server connection, 204
 receiving (*see* Receiving e-mail)
 sending (*see* Sending e-mail)
 Embedded Ruby (ERb), 124

F

File merging, 285
 Foreign key, 277
 Functional testing controllers
 articles controller test, 229
 full test suite, 241
 helper method
 article creation, 236
 destroy article, 239
 index action, 230
 missing test, 241
 new action, 235
 show action, 233
 test_helper file, 230
 overview, 229

G, H

Gemfile.lock, 270
 Git, 285
 adding and committing, 289
 branch, 291
 checkout command, 291
 clone command, 294
 git help command, 295
 GitHub, 294
 git push command, 294
 installing Git, 286
 .gitignore file, 289
 global parameters, 287
 initialize local repository, 288
 Linux, 287
 Mac OS X Leopard, 287
 Windows, 286
 merge command, 293
 Online Resources, 296
 remote repository, 294
 SCM (*see* Source Control Management (SCM))

I

InsertAfter method, 195
 Integration testing

Integration testing (*cont.*)
blog application, 242
overview, 242
singleton methods, 248
story-based testing, 246
Internationalization, 251
bilingual blog
before action, 263
calling language_selector, 262
language selector, 265
language_selector helper
method, 262
language selector link, 265
locale configuration, 264
blog application-Brazilian Portuguese
default locale, 260
localized interface, 261
translation files, 259
blog application-i18n
ArticlesController, 255
default locale file, 257
translation, 255
updated article Partial, 257
updated comment model, 254
Rails
Brazilian Portuguese locale
files, 253
error messages, 253
locales, 251
rails console, 252
translate method, 252
translation missing, 254
Internet Relay Chat (IRC) Channel, 282

■ J, K

JavaScript and CSS, 191
\$() function, 193
Ajax, 193
asset concatenation and
compression, 191
Asset Pipeline, 192
secondary languages, 191
CSS Selectors, 193
DOM, 193
Grand Entrance, 197
insertAfter method, 195
jQuery, 193
js.erb template creation, 194
new comment element, 195
new comment link, 194, 196
Rails, 193
template show, 194
Turbolinks gem, 192
using Ajax, 197

controller comments updated, 198
create method, 198
delete records, 200
error message, 199
JavaScript request, 198
js.erb template creation, 198–199
remote, 197

JOIN query, 278

■ L

Lightweight methodologies, 4
Localization. *See* Internationalization

■ M, N

Mercurial, 296
Migration, 33
Model enhancement. *See* Active Record
Model-view-controller (MVC) pattern
characteristics, 7
controller layer, 8–9
cycle, 7
model layers, 8
spaghetti code, 7
view layer, 8, 10

■ O

Object-oriented (OO), 55
Object-relational gap., 63
Object-relational mapping (ORM), 61

■ P, Q

Platform as a service (PaaS), 267
Primary key, 274, 276
Principle of least surprise (POLS), 5

■ R

Rails, 1
Rails application. *See also* Action pack
action, 23
Active Record (*see* Active Record)
additional fields, 38
advantages, 2
agile principles, 4
Article model, 32
blog application
command line, 28
config/database.yml File, 29
default SQLite database, 30
Rails directory structure, 28–29
users, 27

browser issues, 2
 built-in web server
 commands, 20
 welcome page, 21
 console
 advantage, 64
 Article model, 64
 column_names class method, 65
 irb, 64
 Ruby code, 64
 subclassing and superclass, 65
 controller, 22
 controller generator, 35
 convention over configuration, 5
 creating table
 20130406203049_create_articles.rb, 33
 db:migrate Rake task, 34
 instance method, 34
 migration, 33
 rollback task, 34
 timestamps method, 34
 DRY principle, 6
 fix, 2
 generated files, 41
 history, 3
 installation
 RubyGems, 13
 steps, 13
 less software, 5
 libraries, 10
 Linux
 apt-get update command, 17
 Rails, 18
 Ruby, 18
 SQLite, 19
 Mac OS X 10.8 Mountain Lion
 command line tools, 14
 Homebrew, 14
 install Rails, 15
 RVM, 15
 Xcode, 14
 modular, 10
 MVC pattern
 characteristics, 7
 cycle, 7
 layers of, 8
 spaghetti code, 7
 new project, 19
 no silver bullets, 11
 open source tool, 6
 opinions, 6
 project database, 30
 Rails 4 release, 3
 Ruby, 3
 scaffolding
 Article model, 36
 generator, 36
 http://localhost:3000/articles/new, 37
 speed, 37
 templates
 app/views, 24
 “Hello World” application, 25
 result, 24
 validations, 40
 web-based software, 1
 web framework, 2
 Windows
 Rails, 16
 Ruby installation, 15
 SQLite, 17
 Rails community, 281
 APIs, 283
 Blogs and Podcasts, 282
 communication channels, 281
 Mailing list, 281
 Twitter, 281
 Web site, 281
 Guides, 282
 IRC Channel, 282
 Mailing Lists, 281
 Core mailing list, 281
 Security mailing list, 282
 Talk mailing list, 281
 source and issue tracking, 283
 Receiving e-mail
 POP/IMAP, 216
 Rails process, 215
 Representational State Transfer
 (REST), 126
 Ruby installation
 Linux, 18
 Windows, 15
 Ruby language, 45
 blocks and iterators, 52
 classes, 56
 class variables, 51
 constants, 51
 control-flow statements, 53
 data types
 array, 49
 container objects, 49
 hash objects, 50
 numbers, 48
 shovel operator, 50
 strings, 47
 symbols, 49
 documentation, 59
 instance variables, 51
 interactive interpreter, 45
 local variables, 51

■ INDEX

Ruby language (*cont.*)

- methods, 54
- objects, 55
- operators, 52
- overview, 45
- style, 59
- variables, 51

Ruby Version Manager (RVM), 15

■ S

Scaffolding, 36

Sending e-mail

- add attachments, 213
- comments, 214
- handling basic e-mail
 - article page, 209
 - message delivered, 211
 - notify_friend_article route, 208
 - partial render, 209
 - updated notifier mailer, 210

HTML formatting, 211

mailer instance variables, 207

Notifier class, 206

Simple Mail Transfer Protocol (SMTP), 203

Singleton methods, 248

Source Control Management (SCM)

- CVS, 296
- file merging, 285
- Mercurial, 296
- SCM workflow, 286
- SVN, 296
- versioning, 285

SQLite

- Linux, 19
- Windows, 17

Story-based testing, 246

String interpolation, 47

Structured Query Language (SQL), 273

Subversion (SVN), 296

■ T

Test-driven development (TDD), 229

Testing applications. *See also* Unit testing

- full test suite, 250
- functional testing controllers
 - articles controller test, 229
 - full test suite, 241
 - overview, 229
- functional testing controllers (*see* Testing helper method)
- integration testing

blog application, 242

overview, 242

story-based testing, 246

overview, 219

Rails handles testing

fixtures, 220

points, 219

rails command, 220

unit directory, 220

updated fixture, 220

refactoring existing code, 219

Testing helper method, 230

article creation, 236

destroy article, 239

index action, 230

missing test, 241

new case, 235

POST method, 237

show action, 233

Toolbelt, 268

Turbolinks gem, 192

■ U

Ubuntu Linux

- apt-get update command, 17
- Rails, 18
- Ruby, 18
- SQLite, 19

Unit testing

- Article model
 - assertions, 223
 - create fixtures, 222
 - destroy test, 226
 - find test, 224
 - overview, 222
 - test creation, 222
 - update test, 225
- article test, 221
- assertions, 221
- command prompt, 221
- elements, 221
- overview, 221
- test method, 221
- validations, 227

■ V

Validation testing, 227

Versioning, 285

View and controller, 121. *See also* Controller walk-through

Action pack components, 121

Action View library, 123

controllers, 122

Embedded Ruby, [124](#)
helpers, [125](#)
method, [123](#)
request cycle, [127](#)
RESTful resources, [126](#)
routing, [125](#)
working definition, [123](#)

■ W, X

WHERE clause, [275](#)

■ Y, Z

YAML, [30](#)

Beginning Rails 4



**Adam Gamble
Cloves Carneiro, Jr.
Rida Al Barazi**

Apress®

Beginning Rails 4

Copyright © 2013 by Adam Gamble, Cloves Carneiro, Jr., and Rida Al Barazi

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed. Exempted from this legal reservation are brief excerpts in connection with reviews or scholarly analysis or material supplied specifically for the purpose of being entered and executed on a computer system, for exclusive use by the purchaser of the work. Duplication of this publication or parts thereof is permitted only under the provisions of the Copyright Law of the Publisher's location, in its current version, and permission for use must always be obtained from Springer. Permissions for use may be obtained through RightsLink at the Copyright Clearance Center. Violations are liable to prosecution under the respective Copyright Law.

ISBN-13 (pbk): 978-1-4302-6034-9

ISBN-13 (electronic): 978-1-4302-6035-6

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

President and Publisher: Paul Manning

Lead Editor: Ben Renow-Clarke

Technical Reviewer: Josh Adams

Editorial Board: Steve Anglin, Mark Beckner, Ewan Buckingham, Gary Cornell, Louise Corrigan, Jonathan Gennick, Jonathan Hassell, Robert Hutchinson, Michelle Lowman, James Markham, Matthew Moodie, Jeff Olson, Jeffrey Pepper, Douglas Pundick, Ben Renow-Clarke, Dominic Shakeshaft, Gwenan Spearing, Matt Wade, Steve Weiss, Tom Welsh

Coordinating Editor: Mark Powers

Copy Editor: Mary Bearden

Compositor: SPi Global

Indexer: SPi Global

Artist: SPi Global

Cover Designer: Anna Ishchenko

Distributed to the book trade worldwide by Springer Science+Business Media New York, 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a Delaware corporation.

For information on translations, please e-mail rights@apress.com, or visit www.apress.com.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales-eBook Licensing web page at www.apress.com/bulk-sales.

Any source code or other supplementary materials referenced by the author in this text is available to readers at www.apress.com/9781430260349. For detailed information about how to locate your book's source code, go to www.apress.com/source-code.

To my dad, Gary, who always invested in and encouraged my love of technology.

Contents

About the Authors.....	xvii
About the Technical Reviewer	xix
Introduction	xxi
■ Chapter 1: Introducing the Rails Framework	1
The Rise and Rise of the Web Application	1
The Web Isn't Perfect	2
The Good Web Framework	2
Enter Rails	3
Rails Is Ruby	3
Rails Encourages Agility	4
Rails Is Opinionated Software.....	6
Rails Is Open Source.....	6
The MVC Pattern.....	7
The MVC Cycle.....	7
The Layers of MVC.....	8
The Libraries That Make Up Rails.....	10
Rails Is Modular.....	10
Rails Is No Silver Bullet	10
Summary.....	11

■ Chapter 2: Getting Started	13
An Overview of Rails Installation.....	13
Installing on Mac OS X 10.8 Mountain Lion.....	14
Installing the Apple Developer Tools (Xcode).....	14
Installing Command Line Tools	14
Installing Homebrew.....	14
Installing RVM.....	15
Installing Rails	15
Installing on Windows	15
Installing Ruby	15
Installing Rails	16
Installing SQLite.....	17
Installing on Linux	17
Installing Ruby	18
Installing Rails	18
Installing SQLite.....	19
Creating Your First Rails Application	19
Starting the Built-In Web Server	20
Generating a Controller.....	22
Creating an Action	23
Creating a Template.....	23
Summary.....	25
■ Chapter 3: Getting Something Running	27
An Overview of the Project.....	27
Creating the Blog Application	27
Creating the Project Databases	30
Creating the Article Model	32
Creating a Database Table	33
Generating a Controller.....	35
Up and Running with Scaffolding	36

Adding More Fields.....	38
Adding Validations	40
Generated Files.....	41
Summary	43
■ Chapter 4: Introduction to the Ruby Language.....	45
Instant Interaction	45
Ruby Data Types	46
Strings	47
Numbers	48
Symbols.....	49
Arrays and Hashes.....	49
Language Basics	51
Variables	51
Operators.....	52
Blocks and Iterators.....	52
Control Structures	53
Methods.....	54
Classes and Objects	55
Objects.....	55
Classes	56
Ruby Documentation	59
Summary	59
■ Chapter 5: Working with a Database: Active Record.....	61
Introducing Active Record: Object-Relational Mapping on Rails	61
What About SQL?	62
Active Record Conventions	63
Introducing the Console	64
Active Record Basics: CRUD	67
Creating New Records.....	67
Reading (Finding) Records.....	71

Updating Records	75
Deleting Records	76
When Good Models Go Bad	78
Summary	80
■ Chapter 6: Advanced Active Record: Enhancing Your Models.....	81
Adding Methods	81
Using Associations	83
Declaring Associations	84
Creating One-to-One Associations.....	85
Creating One-to-Many Associations	90
Applying Association Options	94
Creating Many-to-Many Associations.....	96
Creating Rich Many-to-Many Associations.....	99
Advanced Finding	101
Using the <i>where</i> Method	101
Using an SQL Fragment.....	102
Using an Array Condition Syntax.....	102
Using Association Proxies.....	104
Other Finder Methods	104
Default Scope	105
Named Scope	106
Applying Validations	108
Using Built-In Validations.....	109
Building Custom Validation Methods	112
Making Callbacks	114
Updating the User Model	115
Reviewing the Updated Models.....	119
Summary	120

■ Chapter 7: Action Pack: Working with the View and the Controller	121
Action Pack Components	121
Action Controller	122
Action View	123
Embedded Ruby.....	124
Helpers	125
Routing	125
RESTful Resources	126
Action Pack Request Cycle	127
A Controller Walk-Through	128
Setting Up Routes	128
Revisiting the Scaffold Generator	131
Rendering Responses.....	134
Redirecting	134
Understanding Templates	135
Working with Layouts	136
Looking at the Article Form	138
Using Form Helpers	140
Processing Request Parameters	143
Revisiting the Controller	144
Displaying Error Messages in Templates.....	145
The <i>edit</i> and <i>update</i> Actions	146
Revisiting the Views	146
Staying DRY with Partials	148
Summary	150
■ Chapter 8: Advanced Action Pack	151
Generating a Controller	151
Nested Resources	155
Sessions and the Login/Logout Logic.....	161
Lying in State.....	161
Using the Session	161

Session as a Resource	161
Logging In a User.....	163
Logging Out a User	164
Improving Controllers and Templates.....	166
Cleaning Up the Articles Index Page.....	166
Adding Categories to the Article Form.....	167
Using Controller Callbacks.....	170
Requiring Authentication with Filters	170
Applying Filters to Controllers	172
Adding Finishing Touches.....	174
Using Action View Helpers.....	174
Escaping HTML in Templates.....	175
Formatting the Body Field	176
Adding Edit Controls	177
Making Sure Articles Have Owners	178
Adding Custom Helpers	180
Giving It Some Style.....	182
Summary.....	189
■ Chapter 9: JavaScript and CSS.....	191
Asset Concatenation and Compression.....	191
Secondary Languages	191
Asset Pipeline Locations.....	192
Turbolinks.....	192
Let's Build Something!	193
Ajax and Rails	193
jQuery and DOM.....	193
Moving to Practice.....	194
Summary.....	201

■ Chapter 10: Sending and Receiving E-Mail	203
Setting Up Action Mailer.....	203
Configuring Mail Server Settings.....	203
Configuring Application Settings	205
Sending E-Mail	206
Handling Basic E-Mail	208
Sending HTML E-Mail	211
Adding Attachments	213
Letting Authors Know About Comments	214
Receiving E-Mail	215
Using a Rails Process	215
Reading E-Mail Using POP or IMAP	216
Summary.....	217
■ Chapter 11: Testing Your Application	219
How Rails Handles Testing	219
Unit Testing Your Rails Application	221
Testing the Article Model	222
Testing Validations.....	227
Functional Testing Your Controllers	229
Testing the Articles Controller.....	229
Creating a Test Helper	230
Running the Full Test Suite	241
Integration Testing.....	242
Integration Testing the Blog Application	242
Story-Based Testing.....	246
Running the Full Test Suite.....	250
Summary.....	250

■ Chapter 12: Internationalization.....	251
Internationalization Logic in Rails	251
Setting Up i18n in the Blog Application	254
Localizing the Blog Application to Brazilian Portuguese	259
Bilingual Blog	262
Summary.....	266
■ Chapter 13: Deploying Your Rails Applications	267
Set Up an Account with Heroku.....	267
Preparing Your Git Repository.....	269
That's it!	272
Summary.....	272
■ Appendix A: Databases 101	273
Examining a Database Table	273
Working with Tables	274
Selecting Data	274
Inserting Data	275
Updating Data.....	276
Deleting Data.....	276
Understanding Relationships	277
SQL and Active Record	279
■ Appendix B: The Rails Community.....	281
Beginning Rails 4 Channels.....	281
Rails Mailing Lists	281
Rails IRC Channel	282
Rails Blogs and Podcasts	282
Rails Guides.....	282
Rails APIs.....	283
Rails Source and Issue Tracking.....	283

Appendix C: Git	285
What Is Source Control Management?	285
How Does It Work?	285
Git.....	286
Installing Git.....	286
Setting Global Parameters.....	287
Initializing a Repository	288
Ignoring Files	289
Adding and Committing	289
Branching and Merging	291
Remote Repositories and Cloning.....	294
Learning More	295
Other SCM Systems.....	296
Online Resources	296
Index.....	297

About the Authors



Adam Gamble is a professional web developer currently working for Isotope 11 in Birmingham, AL. He has over 10 years experience building web applications for everything from startups to multiple Fortune 500 companies. His passion for technology has enabled him to turn a hobby into a career that he loves. Adam was born in Birmingham, Alabama and currently live there with his wife Monica.



Cloves Carneiro Jr is a software engineer with over 15 years of experience creating web applications for companies in many fields, including startups and telecommunication and financial companies. He has been using Ruby on Rails since its early days and has been a full-time Rails developer for six years. He currently helps building and scaling services at LivingSocial. Born in Brazil and having lived in many parts of the world, Cloves now lives in South Florida with his wife, Jane, and children, Noah and Sofia. He also maintains a personal web site at <http://www.clovescarneirojr.com/>.

■ ABOUT THE AUTHORS



Rida Al Barazi is a passionate software engineer experienced in building smart web applications for startups. He has been designing and building for the web since 2002. He started working with Rails in 2005 and spoke at different web and Rails conferences in North America, Europe, and the Middle East. Rida was raised in Kuwait, grew up in Syria, started his career in Dubai and currently lives in Toronto, Canada, with his wife, Norah. In his free time he enjoys music, concerts, movies, traveling and meeting new people. Rida's contact information can be found on his website, www.rida.me.

About the Technical Reviewer



Josh Adams is a developer and architect with over eleven years of professional experience building production-quality software and managing projects. Josh is Isotope Eleven's CTO and lead architect, and is responsible for overseeing architectural decisions and translating customer requirements into working software. Josh graduated from the University of Alabama at Birmingham (UAB) with Bachelor of Science degrees in both Mathematics and Philosophy. When he's not working, Josh enjoys spending time with his family. He'd like to thank his wife Kristen, and his children, Matthew and Gracie, for putting up with him.